

6.033 – Operating Systems: Structure
Lecture 6
Katrina LaCurts, lacurts@mit.edu

1. Virtual Machines

- How to run multiple OSes on one machine?
- Constraint: compatibility. Don't want to change existing kernel code.
- We'll run multiple virtual machines (VMs) on a single CPU. Kernel equivalent is the "virtual machine monitor" (VMM)
- Can run VMM as user-mode app inside host OS, or run VMM on hardware in kernel mode with guest OSes in user mode. We'll talk about second, but the issues are the same.
- Role of VMM:
 - Allocate resources
 - Dispatch events
 - Deal with instructions from guest OS that require interaction with the physical hardware
- Attempt 1: emulate every single instruction
 - Problem: Slow
- Attempt 2: guest OSes run instructions directly on CPU
 - Problem: dealing with privileged instructions (can't run in kernel mode; then we'd be back to our original problem)
- VMM will deal with handling privileged instructions

2. VMM Implementation

- Trap and emulate
 - Guest OS in user mode
 - Privileged instructions cause an exception; VMM intercepts these and emulates
 - If VMM can't emulate, send exception back up to guest OS
- Problems:
 - How to do the emulate
 - How to deal with instructions that don't trigger an interrupt but that the VMM still needs to intercept

3. Virtualizing memory

- VMM needs to translate guest OS addresses into physical memory addresses. Three layers: guest virtual, guest physical, host physical
- Approach 1: Shadow pages
 - Guest OS loads PTR; causes interrupt. VMM intercepts
 - VMM locates guest OS's page table. Combines guest OS's table with its own table, constructing a third table mapping guest virtual to host physical
 - VMM loads host physical addr of this new page table into the hardware PTR
 - If guest OS modifies its page table, no interrupt thrown. To force an interrupt, VMM marks guest OS's page table as read-only memory

- Approach 2
 - Modern hardware has support for virtualization
 - Physical hardware (effectively) knows about both levels of tables: will do lookup in the guest OS's page table and then the VMM's page table

4. Virtualizing U/K bit

- Problem with basic trap-and-emulate: U/K bit involved in some instructions that don't cause exception (e.g., reading U/K bit, writing it to U)
- Few solutions:
 - Para-virtualization: modify guest OS. Hard to do, and goes against our compatibility goal
 - Binary translation: VMM analyzes code from guest OS and replaces problematic instructions
 - Hardware support: some architectures have virtualization support built in. Have special VMM operating mode in addition to the U/K bit
- Hardware support is arguably the best. Makes VMM's job easier.

5. Monolithic kernels

- VMs protect OSes from each other's faults, protect physical machine from OS faults. Why so many bugs, though?
- The Linux kernel is, effectively, one large C program. Careful software engineering, but very little modularity within the kernel itself.
- Bugs come about because of its complexity
- Kernel bugs = entire system failure (recall the in-class demo)
- Even worse: adversary can exploit these bugs

6. Microkernels: alternative to monolithic kernels

- Put subsystems -- file servers, device drivers, etc. -- in user programs. More modular.
- There will still be bugs but:
 - Fewer, because of decreased complexity
 - A single bug is less likely to crash the entire system
- Why isn't Linux a microkernel, then?
 - High communication cost between modules
 - Not clear that moving programs to userspace is worth it
 - Hard to balance dependencies (e.g., sharing memory across modules)
 - Redesign is tough!
 - Spend a year of developer time rewriting the kernel or adding new features?
 - Microkernels can make it more difficult to change interfaces
- Some parts of Linux do have microkernel design aspects

7. Summary

- Cool things we do with VMs: run different OSes on a single

- machine, move VMs from one physical machine to another
- Microkernels and VMs solve orthogonal problems
 - Microkernels: split up monolithic designs
 - VMs: let us run many instances of an existing OS. They are, in some sense, a partial solution to monolithic kernels (at least we can run these kernels safely). But their goal is to run multiple OSes on a single piece of hardware, not to target monolithic OSes specifically.
 - VMs most commonly implemented with hardware support (a special VMM mode in addition to U/K bit)