

6.033: Networking: P2P Networks + CDNs

Lecture 13

Katrina LaCurts, lacurts@mit.edu

0. Introduction

- "New" technologies on the Internet. How do they work? Are they overcoming any problems in the existing architecture? Do they invalidate any of our assumptions? Do they provide opportunities?
- Today: file-sharing, VoIP, and video-streaming
- Commonalities: all deal with P2P networks, or related constructs (CDNs).

1. File-sharing: getting a file from one person (machine) to another

- Can use client/server
 - client requests file, server responds with the data
 - HTTP, FTP work this way
- Downsides: single point of failure, expensive, doesn't scale
- Could use CDNs
 - Buy multiple servers, put them near clients to decrease latency
 - No single point of failure, scales better
 - See the next recitation for more discussion

2. Peer-to-peer (P2P) networks for file-sharing

- Distribute the architecture to the extreme
- Once a client downloads (part of) the file from the server, that client can upload (part of) the file to others. Put clients to work!
- In theory: infinitely scalable
- P2P networks create overlays on top of the underlying Internet (so do CDNs)
- Problem: what if users aren't willing to upload?

3. BitTorrent: how to incentivize peers to upload

- Basics of original BitTorrent (BT) protocol:
 - Create a .torrent file, which contains meta-information about the file (file name, length, info about pieces that comprise the file, URL of tracker)
 - Have a tracker. A server that knows the identity of all the peers involved in your file transfer.
 - To download:
 - Peer contacts tracker
 - Tracker responds with list of other peers involved in transfer
 - Peer connects to these other peers, begins to transfer blocks (see below)
 - Some peers are seeders: already have the entire file (maybe servers that host the file, or just nice peers who are sticking around)
- In the actual download, peers request blocks: pieces of pieces
 - Details/terminology doesn't matter. Just know that blocks are

- small (~16KB) chunks of the file.
- Request blocks in a random order (more or less)
- What incentivizes users to upload (UL) rather than just download(DL)ing?
 - High-level: users aren't allowed to DL from a user unless they're also ULing to that user
 - So peers want mutual interest: A has to have blocks that B needs, and vice versa.
 - Protocol is divided into rounds. In round n, some number of peers upload blocks to Peer X. In round n+1, Peer X will send blocks to the peers that uploaded the most in round n. (Typically, to the top four peers.)
 - How do peers get started? Each peer reserves some (small) amount of bandwidth to give away freely
- This method of incentivizing peers is part of what allowed P2P file-sharing to take off.
- Lingering problem: tracker is central point of failure
- Most BT clients today are "trackerless", and use Distributed Hash Tables (DHTs) instead.

4. VoIP: Voice over IP

- Talking specifically about Skype, a proprietary system
- Skype used to use a P2P network for two things: to improve performance, allow certain connections to work at all.
- Recall the first networking lecture. Internet bred NATs: Network Address Translators.
 - Consider client A behind a NAT, who wants to initiate a connection to server S. A's IP is private (can't route to it);

S's and N's are public.

A --- N ---- S


- A sends a packet: [to:S from:A]
- N rewrites the header: [to:S from:N]
 - and stores some state
- S receives it, sends response back to N: [to:N from:S]
- N uses stored state to figure out that this packet is really meant for A
 - N will keep track of the port(s) that A is communicating on. Communication via those ports is then meant for A.
- Now imagine two clients, both behind NATs

A --- N1 ---- N2 --- S

- Now A doesn't even know S's IP (private IPs aren't routable). It also doesn't know N2's IP; it has no way to get that.
- For Skype: means that A and S can't call each other
- Skype provides a directory service, so assume we can get N2's public IP. When N2 gets packet destined for S, it has no idea

what to do with it.

- (See slides for example)
- Skype will employ an additional node -- a "supernode" -- P, with a public IP, and route A and S's calls through P:



- P keeps a bunch of state to get this to work, and A and S must both be registered users of Skype. A and S will connect to P as part of starting up their Skype client (so private IP users initiate connections to public IPs)
 - In reality, there is not one supernode, but a network of supernodes. A, S are both connected to nodes in that network, and the overlay network routes data between them.
- Seems like this will affect performance, so Skype only let you be a supernode if your memory/CPU is sufficient (and you have a public IP)
- Good idea?
 - A/S might not want their (encrypted) call routed through someone else
 - P might not want to pay to transit traffic for A and S
- Today: Microsoft owns all of the supernodes, making this less of a P2P network and more of a hierarchy
 - Skype also claims that its P2P system improves quality by allowing for more optimal routing

<https://support.skype.com/en/faq/FA10983/what-are-p2p-communications>

6. Video-streaming (briefly)

- Can we just use BitTorrent to stream (live) video?
 - streaming requires getting blocks (roughly) in order
 - also requires certain amount of bandwidth at all times
- Probably not
 - BT works because peers can acquire blocks in any order
 - Moreover, most BT peers are on residential links, which have underwhelming upload bandwidth.
- What's good for streaming? CDNs!
 - Thursday's recitation: what CDNs bring to the table that P2P networks don't
 - Also think about whether you want to reconsider CDNs for file-sharing