

6.033: Security - Bitcoin  
Lecture 24  
Katrina LaCurts, lacurts@mit.edu

```
*****
* Disclaimer: This is part of the security section in 6.033. Only *
* use the information you learn in this portion of the class to *
* secure your own systems, not to attack others. *
*****
```

0. Today

- Bitcoin! Deals with anonymity, which we haven't considered yet.

1. Distributed Public Logs

- Most of this lecture is about the blockchain, the data structure that underlies Bitcoin.
- The blockchain provides a distributed public log (or ledger) that multiple users can read to/write from.
- Basics: users append their transactions to the log, everyone in the system has a copy of the log, totally decentralized so no part of the system is more powerful than any other.
- Challenge: anonymity; Solution: users identified only by their public keys.
- Challenge: prevent adversaries from editing transactions (e.g., causing me to give Howie \$100 instead of \$10); Solution: users sign their transactions.
- Challenge: prevent adversaries from duplicating, removing, reordering transactions; Solution: hash chains.
- Each block includes the hash of the block before it in the chain.

2. Consensus

- With multiple users writing to the log, how do we know which log is "correct"? How do we confirm a transaction?
  - Alternate motivation: Eve wants to "double-spend" a coin
- Idea: Raft? No; we don't want a leader
- Idea: Broadcast transactions, wait for majority of network to confirm them before appending to log. No: thwarted by Sybil Attacks.
  - Eve creates a lot of identities, confirms anything she wants
- Idea: If a user receives two (or more) copies of the log, and they conflict in some way, only keep (and forward) the one that is longest. A transaction is confirmed when it appears in the log with N blocks after it.
- Not bad, but Eve can still thwart this with a Sybil Attack.
- Underlying issue: it's \*easy\* to valid blocks; calculating hashes is easy.

3. Proofs of Work (one way to handle this issue; NOT the only way)

- Idea: make it computationally expensive to validate a block

- Bitcoin's mechanism:
  - A user receives the log with an unvalidated transaction pending. Checks that that transaction is legit (e.g., if the transaction is A transferring coin X to B, need to check via prior log entries that A indeed owns coin X). Then begins to solve puzzle:
    - $t$  = pending block
    - Find nonce such that  $H(t|nonce) < \text{target}$ , where  $H$  is the hash function, target is set by the system (can change)
    - In practice, target is tuned so that nonce takes ~10min to find
  - Once a user finds a valid nonce, they publish it as part of the block, and broadcast the solution
- V can't mount Sybil attack now; would need to own a lot of computational resources, not just a lot of identities

#### 4. Vocabulary

- The log is the blockchain
- In the world of Bitcoin, validating blocks is the process of mining
  - There is incentive to do this; users get paid (in Bitcoin) for successfully mined blocks
- If two users validate the same block at (roughly) the same time, we say the blockchain forks

#### 5. Discussion

- Performance
  - It takes time and energy (power) to validate blocks
    - In reality, each block contains multiple transactions
  - It takes time to get a transaction confirmed
  - Downloading the entire blockchain at the start of using Bitcoin can take a very long time
  - Proofs of work waste computation, and normal users are unlikely to mine successfully
- Lots of other uses for blockchains, not just digital currency
- Are users truly anonymous?