

6.033 Spring 2019

Lecture #5

- **Threads**
- **Condition Variables**
- **Preemption**




operating systems enforce modularity on a single machine using **virtualization**

in order to enforce modularity + build an effective operating system

- | | | |
|---|---|--|
| 1. programs shouldn't be able to refer to (and corrupt) each others' memory | → | virtual memory |
| 2. programs should be able to communicate | → | bounded buffers
(virtualize communication links) |
| 3. programs should be able to share a CPU without one program halting the progress of the others | → | assume one program
per CPU |

operating systems enforce modularity on a single machine using **virtualization**

in order to enforce modularity + build an effective operating system

- 1. programs shouldn't be able to refer to (and corrupt) each others' **memory**  **virtual memory**
- 2. programs should be able to **communicate**  **bounded buffers**
(virtualize communication links)
- 3. programs should be able to **share a CPU** without one program halting the progress of the others  **threads**
(virtualize processors)

today's goal: use **threads** to allow multiple programs to share a CPU

thread: a virtual processor

thread API:

suspend() : save state of current thread
to memory

resume() : restore state from memory

```
send(bb, message):  
    acquire(bb.lock)  
    while bb.in - bb.out == N:  
        release(bb.lock)  
        acquire(bb.lock)  
    bb.buf[bb.in mod N] <- message  
    bb.in <- bb.in + 1  
    release(bb.lock)  
    return
```

```
send(bb, message):  
    acquire(bb.lock)  
    while bb.in - bb.out == N:  
        release(bb.lock)  
        yield()  
        acquire(bb.lock)  
    bb.buf[bb.in mod N] <- message  
    bb.in <- bb.in + 1  
    release(bb.lock)  
    return
```

yield():

```
// Suspend the running thread  
// Choose a new thread to run  
// Resume the new thread
```

```
yield():  
    acquire(t_lock)
```

```
    // Suspend the running thread  
    // Choose a new thread to run  
    // Resume the new thread
```

```
    release(t_lock)
```



```
yield():  
    acquire(t_lock)  
  
    id = id of current thread  
    threads[id].state = RUNNABLE  
    threads[id].sp = SP  
    threads[id].ptr = PTR  
  
    // Choose a new thread to run  
    // Resume the new thread  
  
    release(t_lock)
```

Suspend
current thread

```
yield():  
    acquire(t_lock)  
  
    id = cpus[CPU].thread  
    threads[id].state = RUNNABLE  
    threads[id].sp = SP  
    threads[id].ptr = PTR  
  
    // Choose a new thread to run  
    // Resume the new thread  
  
    release(t_lock)
```

Suspend
current thread

```
yield():
```

```
    acquire(t_lock)
```

```
    id = cpus[CPU].thread
```

```
    threads[id].state = RUNNABLE
```

```
    threads[id].sp = SP
```

```
    threads[id].ptr = PTR
```

Suspend
current thread

```
    do:
```

```
        id = (id + 1) mod N
```

```
    while threads[id].state != RUNNABLE
```

Choose new
thread

```
    // Resume the new thread
```

```
    release(t_lock)
```

```
yield():  
    acquire(t_lock)
```

```
    id = cpus[CPU].thread  
    threads[id].state = RUNNABLE  
    threads[id].sp = SP  
    threads[id].ptr = PTR
```

Suspend
current thread

```
    do:  
        id = (id + 1) mod N  
    while threads[id].state != RUNNABLE
```

Choose new
thread

```
    SP = threads[id].sp  
    PTR = threads[id].ptr  
    threads[id].state = RUNNING  
    cpus[CPU].thread = id
```

Resume new
thread

```
    release(t_lock)
```

```
send(bb, message):  
    acquire(bb.lock)  
    while bb.in - bb.out == N:  
        release(bb.lock)  
        yield()  
        acquire(bb.lock)  
    bb.buf[bb.in mod N] <- message  
    bb.in <- bb.in + 1  
    release(bb.lock)  
    return
```

condition variables: let threads wait for events, and get notified when they occur

condition variable API:

`wait(cv)`: yield processor and wait to be notified of `cv`

`notify(cv)`: notify waiting threads of `cv`

```
send(bb, message):  
    acquire(bb.lock)  
    while bb.in - bb.out == N:  
        release(bb.lock)  
        wait(bb.not_full)  
        acquire(bb.lock)  
    bb.buf[bb.in mod N] <- message  
    bb.in <- bb.in + 1  
    release(bb.lock)  
    notify(bb.not_empty)  
    return
```

(threads in receive() will
wait on `bb.not_empty` and
notify of `bb.not_full`)

problem: lost notify

condition variable API:

`wait(cv, lock)`: yield processor, release lock, wait to be notified of cv

`notify(cv)`: notify waiting threads of cv


```
send(bb, message):  
    acquire(bb.lock)  
    while bb.in - bb.out == N:  
        wait(bb.not_full, bb.lock)  
    bb.buf[bb.in mod N] <- message  
    bb.in <- bb.in + 1  
    release(bb.lock)  
    notify(bb.not_empty)  
    return
```

```
wait(cv, lock):  
    acquire(t_lock)
```

```
release(t_lock)
```

```
wait(cv, lock):  
    acquire(t_lock)  
    release(lock)
```

```
release(t_lock)  
acquire(lock)
```

```
wait(cv, lock):  
    acquire(t_lock)  
    release(lock)  
    id = cpus[CPU].thread  
    threads[id].cv = cv  
    threads[id].state = WAITING  
  
release(t_lock)  
acquire(lock)
```

```
wait(cv, lock):  
    acquire(t_lock)  
    release(lock)  
    id = cpus[CPU].thread  
    threads[id].cv = cv  
    threads[id].state = WAITING  
    yield_wait() ←  
    release(t_lock)  
    acquire(lock)
```

will be different
than `yield()`

```
wait(cv, lock):  
    acquire(t_lock)  
    release(lock)  
    id = cpus[CPU].thread  
    threads[id].cv = cv  
    threads[id].state = WAITING  
    yield_wait() ← will be different  
    release(t_lock)      than yield()  
    acquire(lock)
```

```
notify(cv):  
    acquire(t_lock)  
    for id = 0 to N-1:  
        if threads[id].cv == cv &&  
            threads[id].state == WAITING:  
                threads[id].state = RUNNABLE  
    release(t_lock)
```

```
yield_wait(): // called by wait()  
    acquire(t_lock)
```

```
    id = cpus[CPU].thread  
threads[id].state = RUNNABLE  
threads[id].sp = SP  
threads[id].ptr = PTR
```

```
do:  
    id = (id + 1) mod N  
while threads[id].state != RUNNABLE
```

```
SP = threads[id].sp  
PTR = threads[id].ptr  
threads[id].state = RUNNING  
cpus[CPU].thread = id
```

```
release(t_lock)
```

problem: wait() holds **t_lock**

```
yield_wait(): // called by wait()
```

```
id = cpus[CPU].thread  
threads[id].state = RUNNABLE  
threads[id].sp = SP  
threads[id].ptr = PTR
```

```
do:  
    id = (id + 1) mod N  
while threads[id].state != RUNNABLE
```

```
SP = threads[id].sp  
PTR = threads[id].ptr  
threads[id].state = RUNNING  
cpus[CPU].thread = id
```

problem: current thread's state shouldn't be RUNNABLE


```
yield_wait(): // called by wait()
```

```
id = cpus[CPU].thread
```

```
threads[id].sp = SP
```

```
threads[id].ptr = PTR
```

```
do:
```

```
    id = (id + 1) mod N
```

```
while threads[id].state != RUNNABLE
```

```
SP = threads[id].sp
```

```
PTR = threads[id].ptr
```

```
threads[id].state = RUNNING
```

```
cpus[CPU].thread = id
```

problem: deadlock (`wait()` holds `t_lock`)

`yield_wait(): // called by wait()`

```
id = cpus[CPU].thread
threads[id].sp = SP
threads[id].ptr = PTR
```

do:

```
id = (id + 1) mod N
```

```
release(t_lock)
```

```
acquire(t_lock)
```

```
while threads[id].state != RUNNABLE
```

```
SP = threads[id].sp
```

```
PTR = threads[id].ptr
```

```
threads[id].state = RUNNING
```

```
cpus[CPU].thread = id
```

problem: stack corruption

`yield_wait(): // called by wait()`

```
id = cpus[CPU].thread
threads[id].sp = SP
threads[id].ptr = PTR
SP = cpus[CPU].stack
```

do:

```
id = (id + 1) mod N
```

```
release(t_lock)
```

```
acquire(t_lock)
```

```
while threads[id].state != RUNNABLE
```

```
SP = threads[id].sp
```

```
PTR = threads[id].ptr
```

```
threads[id].state = RUNNING
```

```
cpus[CPU].thread = id
```

preemption: forcibly interrupt threads

```
timer_interrupt():  
    push PC  
    push registers  
    yield()  
    pop registers  
    pop PC
```

problem: what if timer interrupt occurs while running
`yield()` or `yield_wait()`?




preemption: forcibly interrupt threads

```
timer_interrupt():  
    push PC  
    push registers  
    yield()  
    pop registers  
    pop PC
```

solution: hardware mechanism to disable interrupts

operating systems enforce modularity on a single machine using **virtualization**

in order to enforce modularity + build an effective operating system

1. programs shouldn't be able to refer to (and corrupt) each others' **memory**  **virtual memory**
2. programs should be able to **communicate**  **bounded buffers**
(virtualize communication links)
3. programs should be able to **share a CPU** without one program halting the progress of the others  **threads**
(virtualize processors)

- **Threads** virtualize a processor so that we can share it among programs. **yield()** allows the kernel to suspend the current thread and resume another.
- **Condition Variables** provide a more efficient API for threads, where they **wait** for an event and are **notified** when it occurs. `wait()` requires a new version of `yield()`, **yield_wait()**.
- **Preemption** forces a thread to be interrupted so that we don't have to rely on programmers correctly using `yield()`. Requires a special **interrupt** and hardware support to disable other interrupts.