

6.033 Spring 2019

Lecture #12

- **In-network resource management**
 - **Queue management schemes**
 - **Traffic differentiation**

Internet of Problems

How do we **route** (and address) scalably, while dealing with issues of policy and economy?



BGP

How do we **transport** data scalably, while dealing with varying application demands?



TCP,
in-network
resource management

How do we **adapt** new applications and technologies to an inflexible architecture?

problem: TCP reacts to drops, and packets aren't dropped until queues are full

Queue Management


given a queue, when do we drop packets?

1. droptail

drop packets only when the queue is full. simple, but leads to high delays and synchronizes flows.

2. RED

drop packets before the queue is full



Queue Management

given a queue, when do we drop packets?

1. droptail

drop packets only when the queue is full. simple, but leads to high delays and synchronizes flows.

2. RED

drop packets before the queue is full: with increasing probability as the queue grows. prevents queue lengths from oscillating, decreases delay, flows don't synchronize

Queue Management

given a queue, when do we drop (or mark) packets?

1. droptail

drop packets only when the queue is full. simple, but leads to high delays and synchronizes flows.

2. RED (drops) / ECN (marks)

drop (or mark) packets before the queue is full: with increasing probability as the queue grows. prevents queue lengths from oscillating, decreases delay, flows don't synchronize, but complex and hard to pick parameters

what if we want to give **latency
guarantees to certain types of
traffic?**

(or at least try to prioritize latency-sensitive traffic)

Delay-based Scheduling

how could we give latency guarantees for some traffic?

1. **priority queueing**

put latency-sensitive traffic in its own queue and serve that queue first. does not prevent the latency-sensitive traffic from “starving out” the other traffic (in other queues).

**what if we want to allocate different
amounts of **bandwidth** to different
types of traffic?**

Bandwidth-based Scheduling

how can we allocate a specific amount of network bandwidth to some traffic?

1. round-robin

can't handle variable packet sizes (and in its most basic form doesn't allow us to weight traffic differently)

2. weighted round-robin

can set weights and deal with variable packet sizes

Weighted Round Robin

in each round:

for each queue q :

$q.\text{norm} = q.\text{weight} / q.\text{mean_packet_size}$

$\text{min} = \text{min of } q.\text{norm}'\text{s over all flows}$

for each queue q :

$q.\text{n_packets} = q.\text{norm} / \text{min}$

send $q.\text{n_packets}$ from queue q

Bandwidth-based Scheduling

how can we allocate a specific amount of network bandwidth to some traffic?

1. round-robin

can't handle variable packet sizes (and in its most basic form doesn't allow us to weight traffic differently)

2. weighted round-robin

can set weights and deal with variable packet sizes, but needs to know mean packet sizes

3. deficit round-robin

Deficit Round Robin

```
in each round:  
  for each queue q:  
    q.credit += q.quantum  
    while q.credit >= size of next packet p:  
      q.credit -= size of p  
      send p
```

Bandwidth-based Scheduling

how can we allocate a specific amount of network bandwidth to some traffic?

1. round-robin

can't handle variable packet sizes (and in its most basic form doesn't allow us to weight traffic differently)

2. weighted round-robin

can set weights and deal with variable packet sizes, but needs to know mean packet sizes

3. deficit round-robin

doesn't need mean packet sizes. near-perfect fairness and low packet processing overhead

Delay-based Scheduling

how could we give latency guarantees for some traffic?

1. **priority queueing**

put latency-sensitive traffic in its own queue and serve that queue first. does not prevent the latency-sensitive traffic from “starving out” the other traffic (in other queues).



**can solve this problem by doing
something similar to bandwidth-based
scheduling across the two queues**

In-network Resource Management

Queue Management

switches can signal congestion before queues are full

DropTail
RED
ECN

Delay-based Scheduling

switches can prioritize latency-sensitive traffic

Priority Queueing

Bandwidth-based Scheduling

switches can enforce (weighted) fairness among different types of traffic

Round-robin
Weighted Round-robin
Deficit Round-robin

in-network resource management: a good idea?

- Active **queue management schemes**, such as **RED** or **ECN**, drop or mark packets before a queue is full, in hopes of getting TCP senders to react earlier to congestion. They are difficult to get to work on the Internet-at-large, but the ideas can be useful in other types of networks.
- **Traffic differentiation** requires a scheduling discipline, such as **weighted round robin** or **deficit round robin**. The goal of these schemes is to give weighted fairness in the face of variable packet sizes while having low processing overhead
- Both of these are examples of **in-network resource management**