

6.033 Spring 2019

Lecture #15

- **When replication fails us**
 - **Atomicity via shadow copies**
 - **Isolation**
 - **Transactions**

high-level goal: build reliable systems from unreliable components

this is difficult because reasoning about failures is difficult. we need some abstractions that will let us simplify.

atomicity

an action is atomic if it **happens completely or not at all**. if we can guarantee atomicity, it will be much easier to reason about failures

```
transfer (bank, account_a, account_b, amount):  
    bank[account_a] = bank[account_a] - amount  
    bank[account_b] = bank[account_b] + amount ← crash! 💥
```

problem: `account_a` lost `amount` dollars, but
`account_b` didn't gain `amount` dollars

```
transfer (bank, account_a, account_b, amount):  
    bank[account_a] = bank[account_a] - amount  
    bank[account_b] = bank[account_b] + amount ← crash! 💥
```

solution: make this action atomic. ensure that the system completes both steps or neither step.

quest for atomicity: attempt 1

```
transfer (bank_file, account_a, account_b, amount):  
    bank = read_accounts(bank_file)  
    bank[account_a] = bank[account_a] - amount  
    bank[account_b] = bank[account_b] + amount ← crash! 💥  
    write_accounts(bank_file)
```

quest for atomicity: attempt 1

```
transfer (bank_file, account_a, account_b, amount):
```

```
    bank = read_accounts(bank_file)
```

```
    bank[account_a] = bank[account_a] - amount
```

```
    bank[account_b] = bank[account_b] + amount
```

```
    write_accounts(bank_file) ← crash! 💥
```

problem: a crash during `write_accounts`
leaves `bank_file` in an intermediate state

quest for atomicity: attempt 2

(shadow copies)

```
transfer (bank_file, account_a, account_b, amount):  
    bank = read_accounts(bank_file)  
    bank[account_a] = bank[account_a] - amount  
    bank[account_b] = bank[account_b] + amount ← crash! 💥  
    write_accounts(tmp_file)  
    rename(tmp_file, bank_file)
```


quest for atomicity: attempt 2

(shadow copies)

```
transfer (bank_file, account_a, account_b, amount):  
    bank = read_accounts(bank_file)  
    bank[account_a] = bank[account_a] - amount  
    bank[account_b] = bank[account_b] + amount  
    write_accounts(tmp_file)  
    rename(tmp_file, bank_file) ← crash! 💥
```

quest for atomicity: attempt 2

(shadow copies)

```
transfer (bank_file, account_a, account_b, amount):  
    bank = read_accounts(bank_file)  
    bank[account_a] = bank[account_a] - amount  
    bank[account_b] = bank[account_b] + amount  
    write_accounts(tmp_file)  
    rename(tmp_file, bank_file) ← crash! 💥
```

problem: a crash during rename potentially leaves **bank_file** in an intermediate state

quest for atomicity: attempt 2

(shadow copies)

```
transfer (bank_file, account_a, account_b, amount):  
    bank = read_accounts(bank_file)  
    bank[account_a] = bank[account_a] - amount  
    bank[account_b] = bank[account_b] + amount  
    write_accounts(tmp_file)  
    rename(tmp_file, bank_file) ← crash! 💥
```

solution: make rename atomic

quest for atomicity: making rename atomic

directory entries

filename “**bank_file**” -> inode **1**

filename “**tmp_file**” -> inode **2**

inode 1: // old data
data blocks: [...]
refcount: 1

inode 2: // new data
data blocks: [...]
refcount: 1

```
rename(tmp_file, orig_file):  
    // point bank_file's dirent at inode 2  
    // delete tmp_file's dirent  
    // remove refcount on inode 1
```

quest for atomicity: making rename atomic

directory entries

filename “**bank_file**” -> inode **1**

filename “**tmp_file**” -> inode **2**

inode **1**: // old data
data blocks: [...]
refcount: 1

inode **2**: // new data
data blocks: [...]
refcount: 1

```
rename(tmp_file, orig_file):  
    tmp_inode = lookup(tmp_file)    // = 2  
    orig_inode = lookup(orig_file) // = 1  
  
    // point bank_file's dirent at inode 2  
    // delete tmp_file's dirent  
    // remove refcount on inode 1
```

quest for atomicity: making rename atomic

directory entries

filename “**bank_file**” -> inode **2**

filename “**tmp_file**” -> inode **2**

inode **1**: // old data
data blocks: [...]
refcount: 1

inode **2**: // new data
data blocks: [...]
refcount: 1

```
rename(tmp_file, orig_file):  
    tmp_inode = lookup(tmp_file)    // = 2  
    orig_inode = lookup(orig_file) // = 1  
  
    orig_file dirent = tmp_inode  
    // delete tmp_file's dirent  
    // remove refcount on inode 1
```

quest for atomicity: making rename atomic

directory entries

filename “**bank_file**” -> inode **2**

inode **1**: // old data
data blocks: [...]
refcount: 1

inode **2**: // new data
data blocks: [...]
refcount: 1

```
rename(tmp_file, orig_file):  
    tmp_inode = lookup(tmp_file)    // = 2  
    orig_inode = lookup(orig_file)  // = 1  
  
    orig_file dirent = tmp_inode  
    remove tmp_file dirent  
    // remove refcount on inode 1
```

quest for atomicity: making rename atomic

directory entries

filename “**bank_file**” -> inode **2**

inode **1**: // old data
data blocks: [...]
refcount: **0**

inode **2**: // new data
data blocks: [...]
refcount: 1

```
rename(tmp_file, orig_file):  
    tmp_inode = lookup(tmp_file)    // = 2  
    orig_inode = lookup(orig_file)  // = 1  
  
    orig_file dirent = tmp_inode  
    remove tmp_file dirent  
    decref(orig_inode)
```


quest for atomicity: making rename atomic


directory entries

filename “**bank_file**” -> inode **1**

filename “**tmp_file**” -> inode **2**

inode **1**: // old data
data blocks: [...]
refcount: 1

inode **2**: // new data
data blocks: [...]
refcount: 1

```
rename(tmp_file, orig_file):  
    tmp_inode = lookup(tmp_file)    // = 2  
    orig_inode = lookup(orig_file)  // = 1  
                                     ← crash!   
    orig_file dirent = tmp_inode    rename didn't happen  
    remove tmp_file dirent  
    decref(orig_inode)
```

quest for atomicity: making rename atomic

directory entries

filename “**bank_file**” -> inode **2**


filename “**tmp_file**” -> inode **2**

inode **1**: // old data
data blocks: [...]
refcount: 1

inode **2**: // new data
data blocks: [...]
refcount: 1

```
rename(tmp_file, orig_file):  
    tmp_inode = lookup(tmp_file)    // = 2  
    orig_inode = lookup(orig_file)  // = 1
```

```
    orig_file dirent = tmp_inode  
    remove tmp_file dirent  
    decref(orig_inode)
```

crash! 
rename happened,
but refcounts are wrong

quest for atomicity: making rename atomic

directory entries

filename “**bank_file**” -> inode **?**

filename “**tmp_file**” -> inode **2**

inode **1**: // old data
data blocks: [...]
refcount: 1

inode **2**: // new data
data blocks: [...]
refcount: 1

```
rename(tmp_file, orig_file):  
    tmp_inode = lookup(tmp_file)    // = 2  
    orig_inode = lookup(orig_file)  // = 1
```

orig_file dirent = **tmp_inode** ← **crash!** 

remove **tmp_file** dirent crash *during* this line seems bad..
decref(**orig_inode**) but is okay because single-sector writes
are themselves atomic

interlude

we're trying to make a sequence of actions atomic using shadow copies: write to a temporary file, and then rename it to the original.

rename itself must be atomic, and we've *almost* got that working — thanks in part to atomic single-sector writes — but our refcounts aren't quite correct.

quest for atomicity: making rename atomic

directory entries

filename “**bank_file**” -> inode **2**


filename “**tmp_file**” -> inode **2**

inode **1**: // old data
data blocks: [...]
refcount: 1

inode **2**: // new data
data blocks: [...]
refcount: 1

```
rename(tmp_file, orig_file):  
    tmp_inode = lookup(tmp_file)    // = 2  
    orig_inode = lookup(orig_file)  // = 1
```

```
    orig_file dirent = tmp_inode  
    remove tmp_file dirent  
    decref(orig_inode)
```

crash! 
rename happened,
but refcounts are wrong

quest for atomicity: making rename atomic

directory entries

filename “**bank_file**” -> inode **2**

filename “**tmp_file**” -> inode **2**

inode **1**: // old data

data blocks: [..]

refcount: 1

inode **2**: // new data

data blocks: [..]

refcount: 1

```
rename(tmp_file, orig_file):
```

```
    tmp_inode = lookup(tmp_file)    // = 2
```

```
    orig_inode = lookup(orig_file) // = 1
```

```
    incref(tmp_inode)
```

```
    orig_file dirent = tmp_inode
```

```
    decref(orig_inode)
```

```
    remove tmp_file dirent
```

```
    decref(tmp_inode)
```

problem: this is a mess,
and is still incorrect

solution: **recover** from failure

(clean things up)

```
recover(disk):  
    for inode in disk.inodes:  
        inode.refcount = find_all_refs(disk.root_dir, inode)  
    if exists("tmp_file"):  
        unlink("tmp_file")
```

quest for atomicity: attempt 2

(shadow copies)

```
transfer (bank_file, account_a, account_b, amount):  
    bank = read_accounts(bank_file)  
    bank[account_a] = bank[account_a] - amount  
    bank[account_b] = bank[account_b] + amount  
    write_accounts(tmp_file)  
    rename(tmp_file, bank_file)
```


atomicity

(first abstraction)

not quite solved; shadow copies perform poorly even for a single user and a single file, and we haven't even talked about concurrency

isolation

(second abstraction)

if we guarantee isolation, then two actions A1 and A2 will appear to have run **serially** even if they were executed concurrently (i.e., A1 before A2, or vice versa)

transactions: provide atomicity and isolation

Transaction 1

```
begin  
transfer(A, B, 20)  
withdraw(B, 10)  
end
```

Transaction 2

```
begin  
transfer(B, C, 5)  
deposit(A, 5)  
end
```

atomicity: each transaction will appear to have run to completion, or not at all

isolation: when multiple transactions are run concurrently, it will appear as if they were run sequentially (serially)

**atomicity and isolation — and thus,
transactions — make it easier to reason
about failures (and concurrency)**

```
transfer (bank_file, account_a, account_b, amount):  
    acquire(lock)  
    bank = read_accounts(bank_file)  
    bank[account_a] = bank[account_a] - amount  
    bank[account_b] = bank[account_b] + amount  
    write_accounts("tmp_file")  
    rename("tmp_file", bank_file)  
    release(lock)
```

**couldn't we just put locks around
everything?**

(isn't that what locks are *for*?)

```
transfer (bank_file, account_a, account_b, amount):  
    acquire(lock)  
    bank = read_accounts(bank_file)  
    bank[account_a] = bank[account_a] - amount  
    bank[account_b] = bank[account_b] + amount  
    write_accounts("tmp_file")  
    rename("tmp_file", bank_file)  
    release(lock)
```

this particular strategy will perform poorly
(would force a single transfer at a time)

**locks sometimes require global reasoning,
which is messy**

eventually, we'll incorporate locks, but in a systematic way

goal: to implement **transactions**,
which provide atomicity and isolation,
while not hindering performance

atomicity → **shadow copies.** work, but perform
poorly and don't allow for concurrency

?

isolation → (coarse-grained locks perform poorly,
finer-grained locks are difficult to
reason about)

eventually, we also want transaction-based systems to
be **distributed**: to run across multiple machines

- **Transactions** provide **atomicity** and **isolation**, both of which make it easier for us to reason about failures because we don't have to deal with intermediate states.
- **Shadow copies** are one way to achieve atomicity. They work, but perform poorly: require copying an entire file even for small changes, and don't allow for concurrency.