


6.033 Spring 2019

Lecture #17

- **Isolation**
 - **Conflict serializability**
 - **Conflict graphs**
 - **Two-phase locking**

goal: build reliable systems from unreliable components
the abstraction that makes that easier is

transactions, which provide **atomicity** and **isolation**, while not hindering **performance**



eventually, we also want transaction-based systems to be **distributed**: to run across multiple machines

goal: build reliable systems from unreliable components
the abstraction that makes that easier is

transactions, which provide **atomicity** and **isolation**, while not hindering **performance**

atomicity —————→ **shadow copies** (simple, poor performance) or **logs** (better performance, a bit more complex)

isolation —————→ **two-phase locking**

eventually, we also want transaction-based systems to be **distributed**: to run across multiple machines

goal: run transactions T1, T2, ..., TN concurrently, and have it “appear” as if they ran sequentially

T1

```
begin  
read(x)  
tmp = read(y)  
write(y, tmp+10)  
commit
```

T2

```
begin  
write(x, 20)  
write(y, 30)  
commit
```

naive approach: actually run them sequentially, via (perhaps) a single global lock

goal: run transactions T1, T2, ..., TN concurrently, and have it “appear” as if they ran sequentially

↖ what does this even mean?

T1

```
begin  
read(x)  
tmp = read(y)  
write(y, tmp+10)  
commit
```

T2

```
begin  
write(x, 20)  
write(y, 30)  
commit
```

T1

begin

read(x)

tmp = read(y)

write(y, tmp+10)

commit

T2

begin

write(x, 20)

write(y, 30)

commit

possible sequential schedules

T1 -> **T2**: x=20, y=30

T2 -> **T1**: x=20, y=40

T2: write(x, 20)

T1: read(x)

T2: write(y, 30)

T1: tmp = read(y)

T1: write(y, tmp+10)

at end:

x=20, y=40

T1: read(x)

T2: write(x, 20)

T1: tmp = read(y)

T2: write(y, 30)

T1: write(y, tmp+10)

at end:

x=20, y=10

(assume x, y initialized to zero)

T1

```
begin
read(x)
tmp = read(y)
write(y, tmp+10)
commit
```

T2

```
begin
write(x, 20)
write(y, 30)
commit
```

possible sequential schedules

T1 -> **T2**: x=20, y=30

T2 -> **T1**: x=20, y=40

```
T2: write(x, 20)
T1: read(x)
T2: write(y, 30)
T1: tmp = read(y)
T1: write(y, tmp+10)
```

at end:
x=20, y=40

```
T1: read(x)
T2: write(x, 20)
T1: tmp = read(y)
T2: write(y, 30)
T1: write(y, tmp+10)
```

~~at end:
x=20, y=10
(assume x, y initialized to zero)~~

T1

begin

read(x)

tmp = read(y)

write(y, tmp+10)

commit

T2

begin

write(x, 20)

write(y, 30)

commit

possible sequential schedules

T1 -> **T2**: x=20, y=30

T2 -> **T1**: x=20, y=40

T2: write(x, 20)

T1: read(x)

T2: write(y, 30)

T1: tmp = read(y)

T1: write(y, tmp+10)

at end:

x=20, y=40

T1: read(x) // x=0

T2: write(x, 20)

T2: write(y, 30)

T1: tmp = read(y) // y=30

T1: write(y, tmp+10)

at end:

x=20, y=40

In the second schedule, **T1** reads x=0 and y=30; those two reads together aren't possible in a sequential schedule.
is that okay?

it depends.

there are many ways for multiple transactions to “appear” to have been run in sequence; we say there are different notions of **serializability**. what type of serializability you want depends on what your application needs.

conflicts

two operations conflict if they operate on the same object and at least one of them is a write.

T1

begin

T1.1 read(x)

T1.2 tmp = read(y)

T1.3 write(y, tmp+10)

commit

T2

begin

T2.1 write(x, 20)

T2.2 write(y, 30)

commit

conflicts

T1.1 read(x)	and	T2.1 write(x, 20)
T1.2 tmp = read(y)	and	T2.2 write(y, 30)
T1.3 write(y, tmp+10)	and	T2.2 write(y, 30)

conflicts

two operations conflict if they operate on the same object and at least one of them is a write.

in any schedule, two conflicting operations A and B will have an order: either A is executed before B, or B is executed before A. we'll call this the **order** of the conflict (in that schedule).

T1

begin

T1.1 read(x)

T1.2 tmp = read(y)

T1.3 write(y, tmp+10)

commit

T2

begin

T2.1 write(x, 20)

T2.2 write(y, 30)

commit

conflicts

T1.1 read(x) and **T2.1** write(x, 20)

T1.2 tmp = read(y) and **T2.2** write(y, 30)

T1.3 write(y, tmp+10) and **T2.2** write(y, 30)

T1

begin

T1.1 read(x)

T1.2 tmp = read(y)

T1.3 write(y, tmp+10)

commit

T2

begin

T2.1 write(x, 20)

T2.2 write(y, 30)

commit

conflicts

T1.1 read(x)	->	T2.1 write(x, 20)
T1.2 tmp = read(y)	->	T2.2 write(y, 30)
T1.3 write(y, tmp+10)	->	T2.2 write(y, 30)

if we execute **T1** before **T2**, within any conflict, **T1**'s operation will occur first

T1

begin

T1.1 read(x)

T1.2 tmp = read(y)

T1.3 write(y, tmp+10)

commit

T2

begin

T2.1 write(x, 20)

T2.2 write(y, 30)

commit

conflicts

T1.1 read(x)	<-	T2.1 write(x, 20)
T1.2 tmp = read(y)	<-	T2.2 write(y, 30)
T1.3 write(y, tmp+10)	<-	T2.2 write(y, 30)

if we execute **T2** before **T1**, within any conflict, **T2**'s operation will occur first

conflicts

two operations conflict if they operate on the same object and at least one of them is a write.

conflict serializability

a schedule is **conflict serializable** if the order of all of its conflicts is the same as the order of the conflicts in some sequential schedule.

conflicts

T1.1, T2.1

T1.2, T2.2

T1.3, T2.2

a schedule is **conflict serializable** if the order of all of its conflicts is the same as the order of the conflicts in some sequential schedule.

(here, that means we will see one transaction's — T1's or T2's — operation occurring first in each conflict)

T2.1: write(x, 20)

T1.1: read(x)

T2.2: write(y, 30)

T1.2: tmp = read(y)

T1.3: write(y, tmp+10)

T2.1 -> T1.1

T2.2 -> T1.2

T2.2 -> T1.3

T1.1: read(x)

T2.1: write(x, 20)

T2.2: write(y, 30)

T1.2: tmp = read(y)

T1.3: write(y, tmp+10)

T1.1 -> T2.1

T2.2 -> T1.2

T2.2 -> T1.3

conflicts

T1.1, T2.1

T1.2, T2.2

T1.3, T2.2

a schedule is **conflict serializable** if the order of all of its conflicts is the same as the order of the conflicts in some sequential schedule.

(here, that means we will see one transaction's — **T1**'s or **T2**'s — operation occurring first in each conflict)

T2.1: write(x, 20)

T1.1: read(x)

T2.2: write(y, 30)

T1.2: tmp = read(y)

T1.3: write(y, tmp+10)

T2.1 -> T1.1

T2.2 -> T1.2

T2.2 -> T1.3

~~T1.1: read(x)~~

~~T2.1: write(x, 20)~~

~~T2.2: write(y, 30)~~

~~T1.2: tmp = read(y)~~

~~T1.3: write(y, tmp+10)~~

~~T1.1 -> T2.1~~

~~T2.2 -> T1.2~~

~~T2.2 -> T1.3~~

conflict graph

edge from T_i to T_j iff T_i and T_j have a conflict between them and the first step in the conflict occurs in T_i

T2: write(x, 20)
T1: read(x)
T2: write(y, 30)
T1: tmp = read(y)
T1: write(y, tmp+10)

T2.1 -> **T1.1**
T2.2 -> **T1.2**
T2.2 -> **T1.3**

T1: read(x)
T2: write(x, 20)
T2: write(y, 30)
T1: tmp = read(y)
T1: write(y, tmp+10)

T1.1 -> **T2.1**
T2.2 -> **T1.2**
T2.2 -> **T1.3**

conflict graph

edge from T_i to T_j iff T_i and T_j have a conflict between them and the first step in the conflict occurs in T_i

T2: write(x, 20)
T1: read(x)
T2: write(y, 30)
T1: tmp = read(y)
T1: write(y, tmp+10)

T2 \longrightarrow **T1**

T1: read(x)
T2: write(x, 20)
T2: write(y, 30)
T1: tmp = read(y)
T1: write(y, tmp+10)

T2 \Longleftrightarrow **T1**

a schedule is conflict serializable iff it has an acyclic conflict graph

problem: how do we generate schedules that are conflict serializable? generate all possible schedules and check their conflict graphs?

solution: two-phase locking (2PL)


1. each shared variable has a lock
2. before **any** operation on a variable, the transaction must acquire the corresponding lock
3. after a transaction releases a lock, it may **not** acquire any other locks

we will usually release locks after commit or abort, which is technically *strict* two-phase locking

2PL produces a conflict-serializable schedule

(equivalently, 2PL produces a conflict graph without a cycle)


proof: suppose not. then a cycle exists in the conflict graph



2PL produces a conflict-serializable schedule

(equivalently, 2PL produces a conflict graph without a cycle)

proof: suppose not. then a cycle exists in the conflict graph




to cause the conflict, each pair of conflicting **transactions** must have some **shared variable** that they conflict on

2PL produces a conflict-serializable schedule

(equivalently, 2PL produces a conflict graph without a cycle)

proof: suppose not. then a cycle exists in the conflict graph



to cause the conflict, each pair of conflicting **transactions** must have some **shared variable** that they conflict on

in the schedule, each pair of **transactions** needs to acquire a lock on their **shared variable**

T₁ acquires **x₁.lock**

T₂ acquires **x₁.lock**

T₂ acquires **x₂.lock**

T₃ acquires **x₂.lock**

...


T_k acquires **x_k.lock**

T₁ acquires **x_k.lock**

2PL produces a conflict-serializable schedule

(equivalently, 2PL produces a conflict graph without a cycle)

proof: suppose not. then a cycle exists in the conflict graph



to cause the conflict, each pair of conflicting **transactions** must have some **shared variable** that they conflict on

in the schedule, each pair of **transactions** needs to acquire a lock on their **shared variable**

in order for the schedule to progress,
T₁ must have released its lock on **x₁**
before **T₂** acquired it

T₁ acquires **x₁.lock**

T₁ releases **x₁.lock**

T₂ acquires **x₁.lock**

T₂ acquires **x₂.lock**

T₃ acquires **x₂.lock**

...


T_k acquires **x_k.lock**

T₁ acquires **x_k.lock**

2PL produces a conflict-serializable schedule

(equivalently, 2PL produces a conflict graph without a cycle)

proof: suppose not. then a cycle exists in the conflict graph



to cause the conflict, each pair of conflicting **transactions** must have some **shared variable** that they conflict on

in the schedule, each pair of **transactions** needs to acquire a lock on their **shared variable**

in order for the schedule to progress,
 T_1 must have released its lock on X_1
before T_2 acquired it

contradiction: this is not a valid 2PL schedule

T_1 acquires $X_1.lock$

T_1 releases $X_1.lock$

T_2 acquires $X_1.lock$

T_2 acquires $X_2.lock$

T_3 acquires $X_2.lock$

...

T_k acquires $X_k.lock$

T_1 acquires $X_k.lock$

T1

```
acquire(x.lock)
read(x)
acquire(y.lock)
read(y)
release(y.lock)
release(x.lock)
```

T2

```
acquire(y.lock)
read(y)
acquire(x.lock)
read(x)
release(x.lock)
release(y.lock)
```

problem: 2PL can result in deadlock

T1

```
acquire(x.lock)
read(x)
acquire(y.lock)
read(y)
release(y.lock)
release(x.lock)
```

T2

```
acquire(y.lock)
read(y)
acquire(x.lock)
read(x)
release(x.lock)
release(y.lock)
```

“solution”: global ordering on locks

T1

```
acquire(x.lock)
read(x)
acquire(y.lock)
read(y)
release(y.lock)
release(x.lock)
```

T2

```
acquire(y.lock)
read(y)
acquire(x.lock)
read(x)
release(x.lock)
release(y.lock)
```

better solution: take advantage of atomicity and abort one of the transactions!

performance improvement: allow concurrent reads with reader- and writer-locks

T1

```
acquire(x.reader_lock)
read(x)
acquire(y.writer_lock)
write(y)
release(y.writer_lock)
release(x.reader_lock)
```

T2

```
acquire(x.reader_lock)
read(x)
acquire(y.writer_lock)
write(y)
release(y.writer_lock)
release(x.reader_lock)
```

multiple transactions can hold reader locks for the same variable at once. a transaction can only hold a writer lock for a variable if there are *no* other locks held for that variable

- Different types of **serializability** allow us to specify precisely what we want when we run transactions in parallel. **Conflict-serializability** is common in practice.
- **Two-phase locking** allows us to generate conflict serializable schedules. We can improve its performance by allowing concurrent reads via reader- and writer-locks.