

# 6.033 Spring 2019

## Lecture #19

- **Distributed transactions**
  - **Availability**
  - **Replicated State Machines**

**goal:** build reliable systems from unreliable components  
the abstraction that makes that easier is

**transactions**, which provide **atomicity** and **isolation**, while not hindering **performance**

<b>atomicity</b>	→	<b>shadow copies</b> (simple, poor performance) or <b>logs</b> (better performance, a bit more complex)
<b>isolation</b>	→	<b>two-phase locking</b>

we also want transaction-based systems to be **distributed** — to run across multiple machines — and to remain **available** even through failures

**C<sub>1</sub>** **write<sub>1</sub>(X)**

**S<sub>1</sub>**

**C<sub>2</sub>** **write<sub>2</sub>(X)**

**S<sub>2</sub>**

(replica of S<sub>1</sub>)

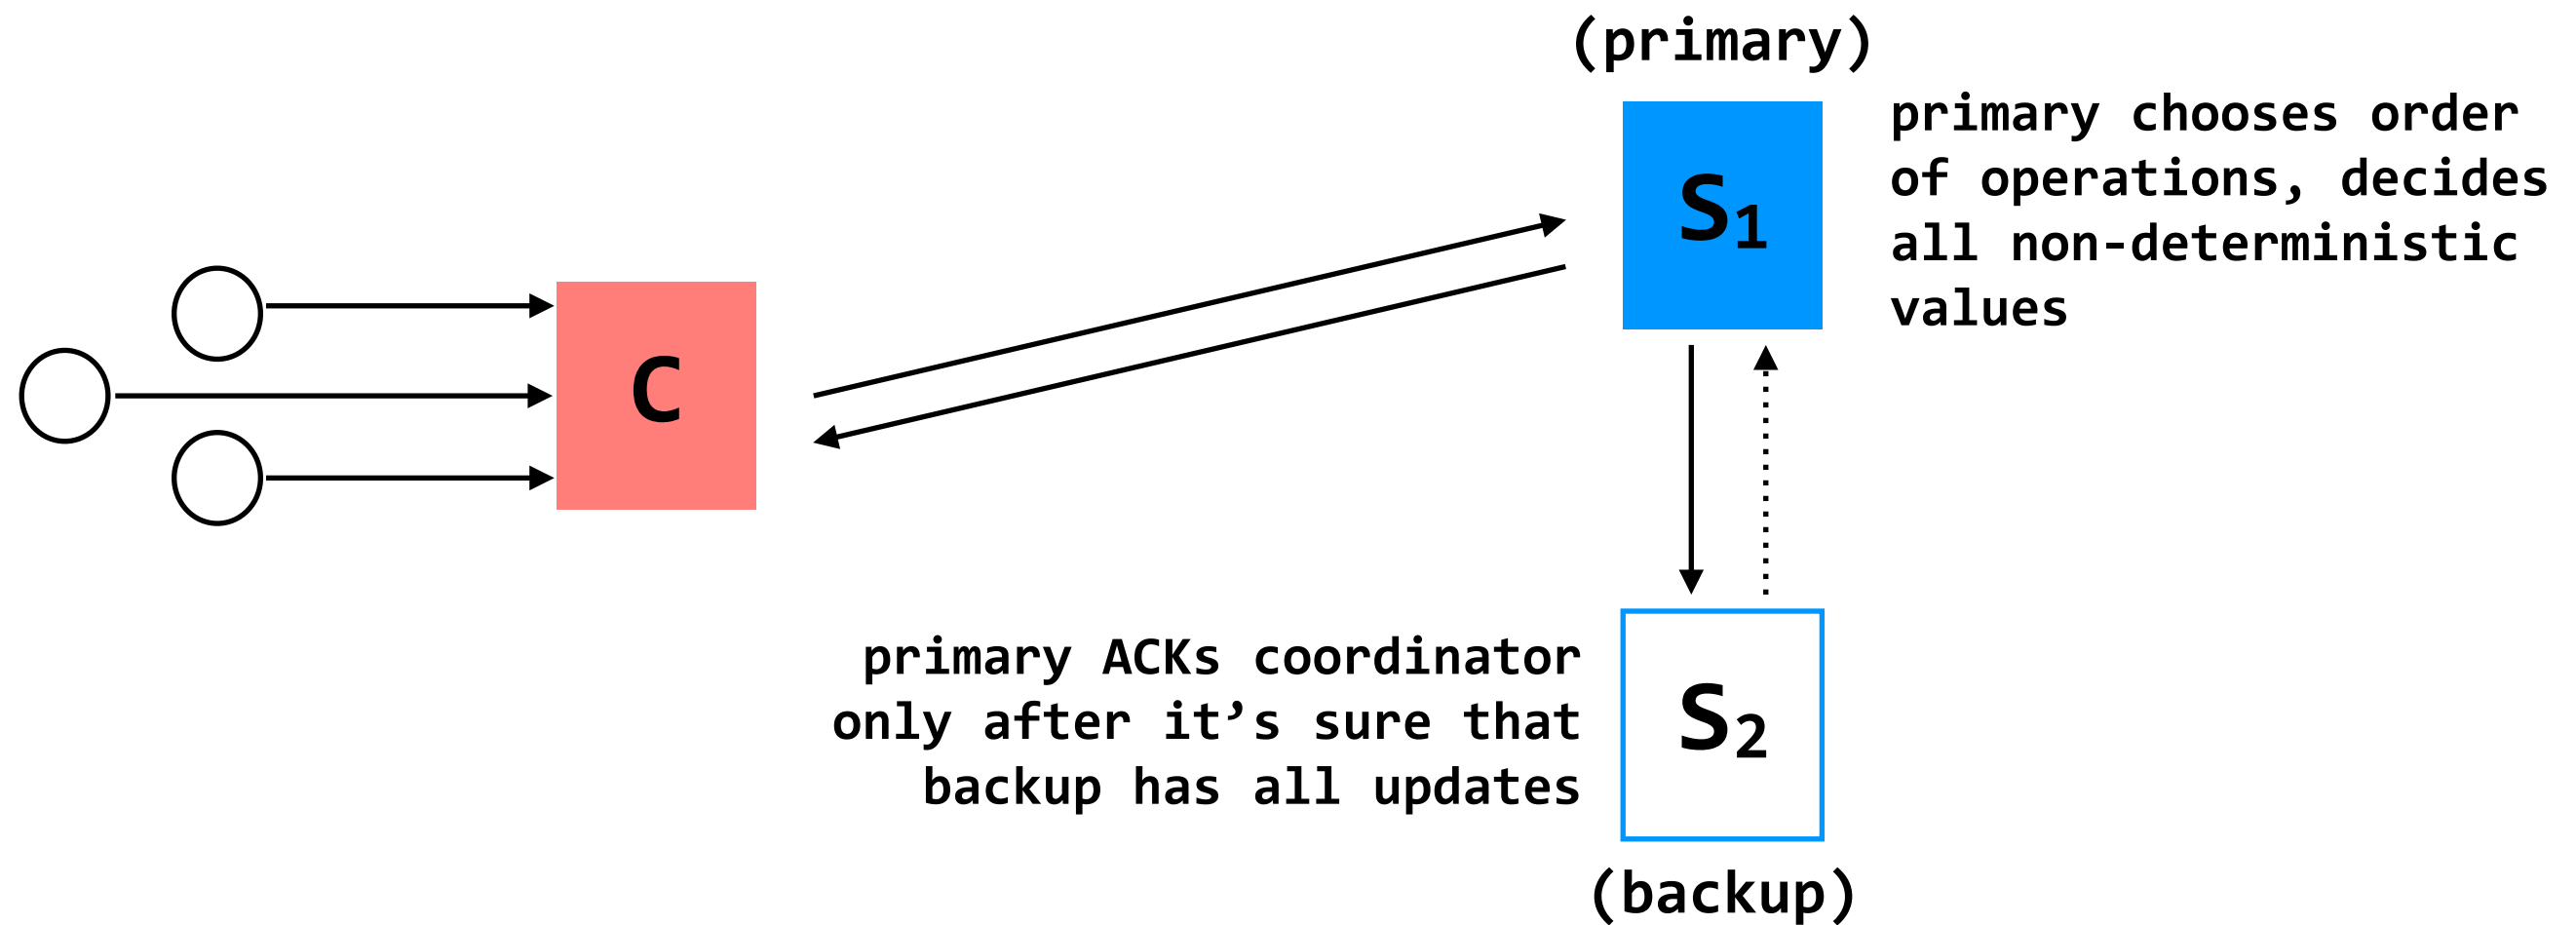
**C<sub>1</sub>**

**S<sub>1</sub>** **write<sub>1</sub>(X)**  
**write<sub>2</sub>(X)**

**C<sub>2</sub>**

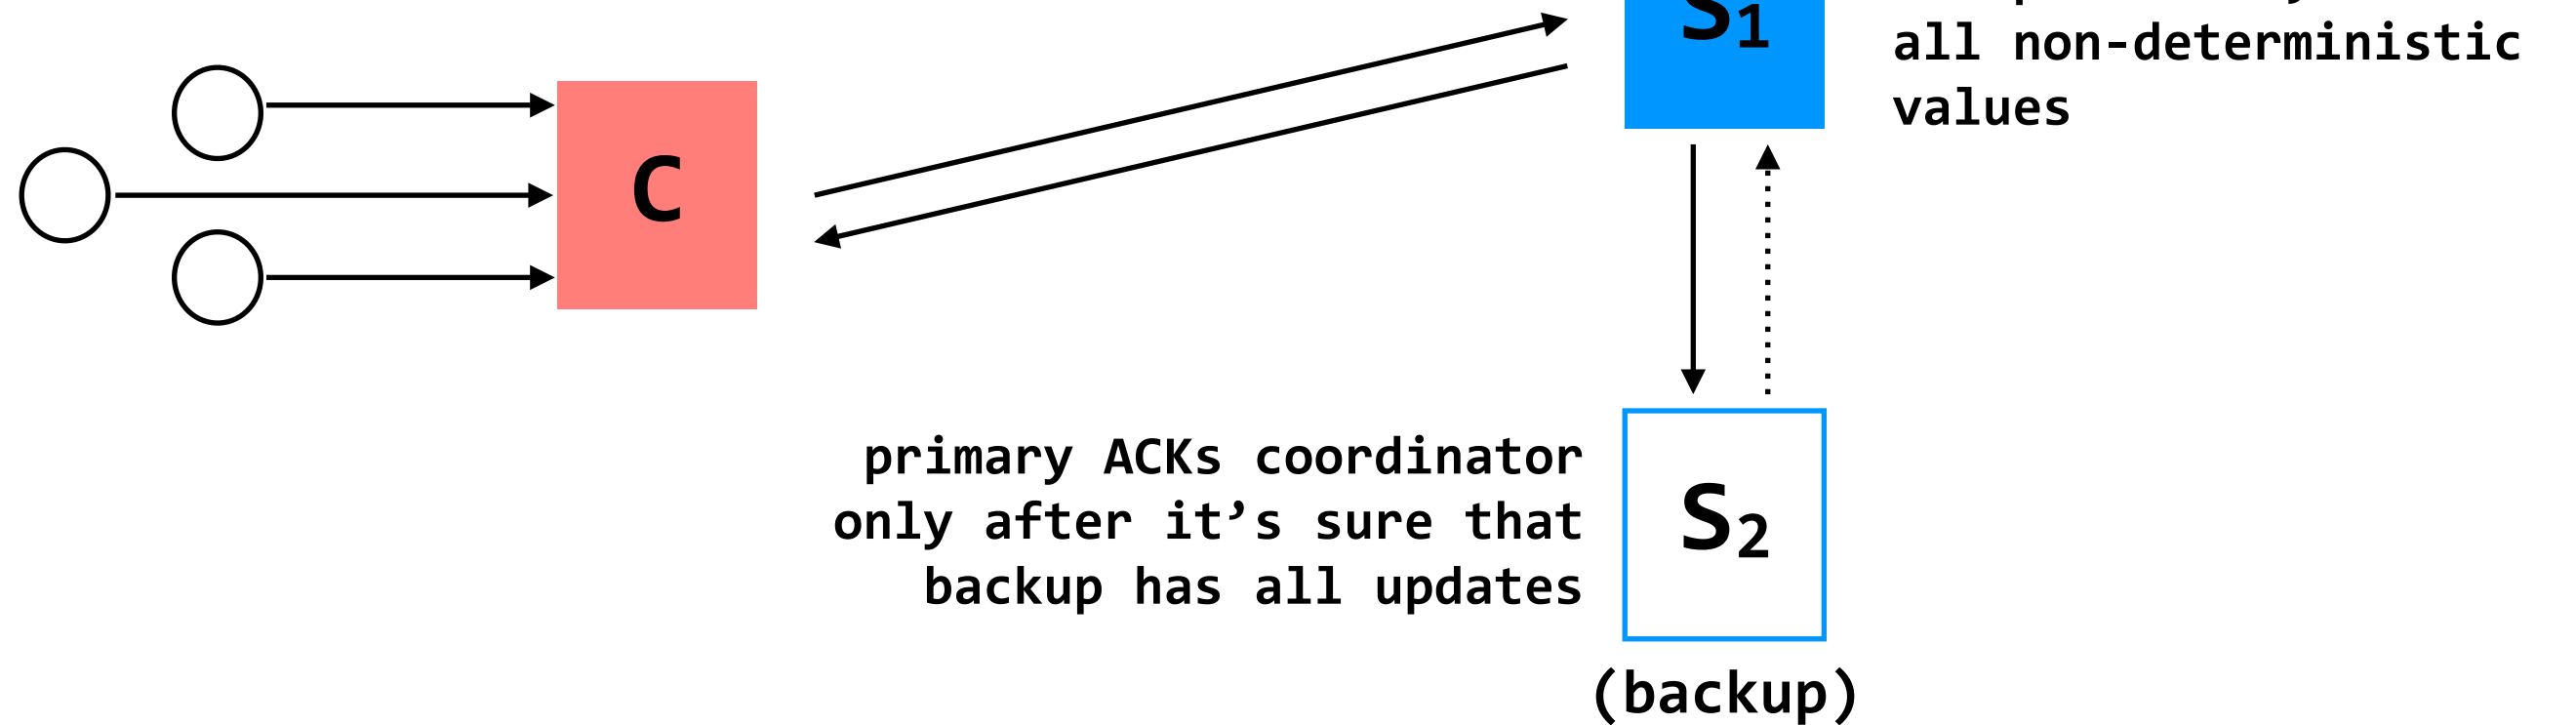
**S<sub>2</sub>** **write<sub>2</sub>(X)**  
**write<sub>1</sub>(X)**  
(replica of S<sub>1</sub>)

**problem:** replica servers can become inconsistent




**attempt:** coordinators communicate with primary servers, who communicate with backup servers

if primary fails, **C** switches to backup  
(**C** knows how to contact backup servers)




**attempt:** coordinators communicate with primary servers, who communicate with backup servers

if primary fails, **C** switches to backup  
(**C** knows how to contact backup servers)



**attempt:** coordinators communicate with primary servers, who communicate with backup servers


if primary fails, **C** switches to backup  
(**C** knows how to contact backup servers)



**attempt:** coordinators communicate with primary servers, who communicate with backup servers




# multiple coordinators + the network = problems




**attempt:** coordinators communicate with primary servers, who communicate with backup servers

# multiple coordinators + the network = problems




**attempt:** coordinators communicate with primary servers, who communicate with backup servers

# multiple coordinators + the network = problems




..... network partition .....



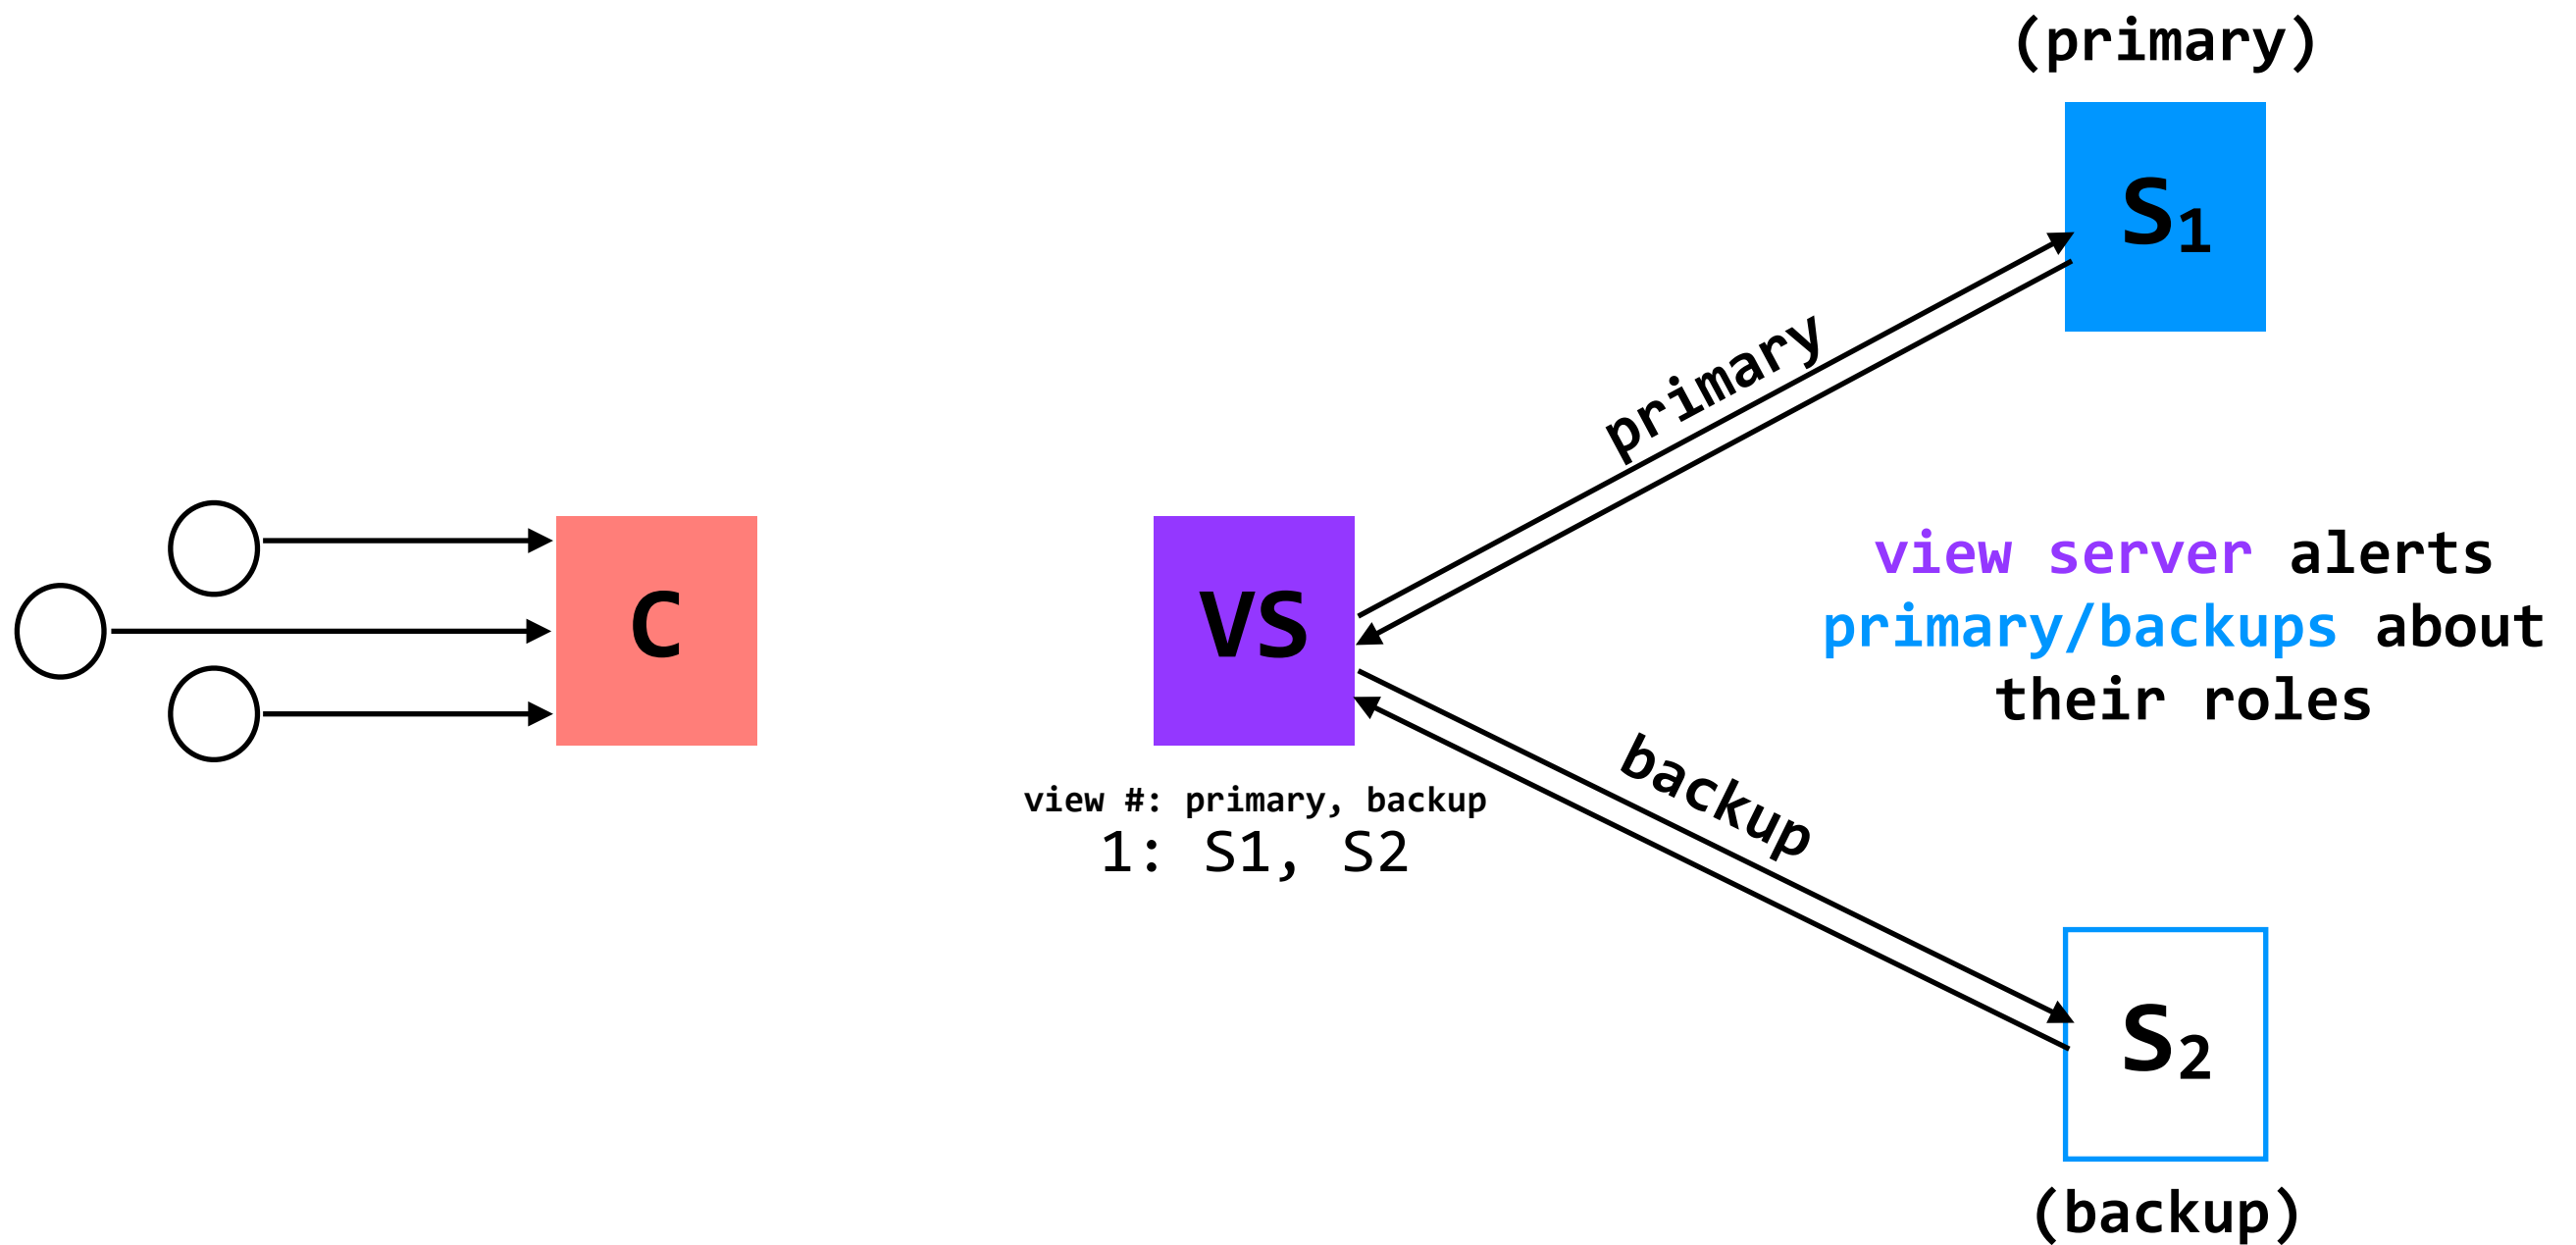
$C_1$  and  $C_2$  are using different primaries;

$S_1$  and  $S_2$  are no longer consistent

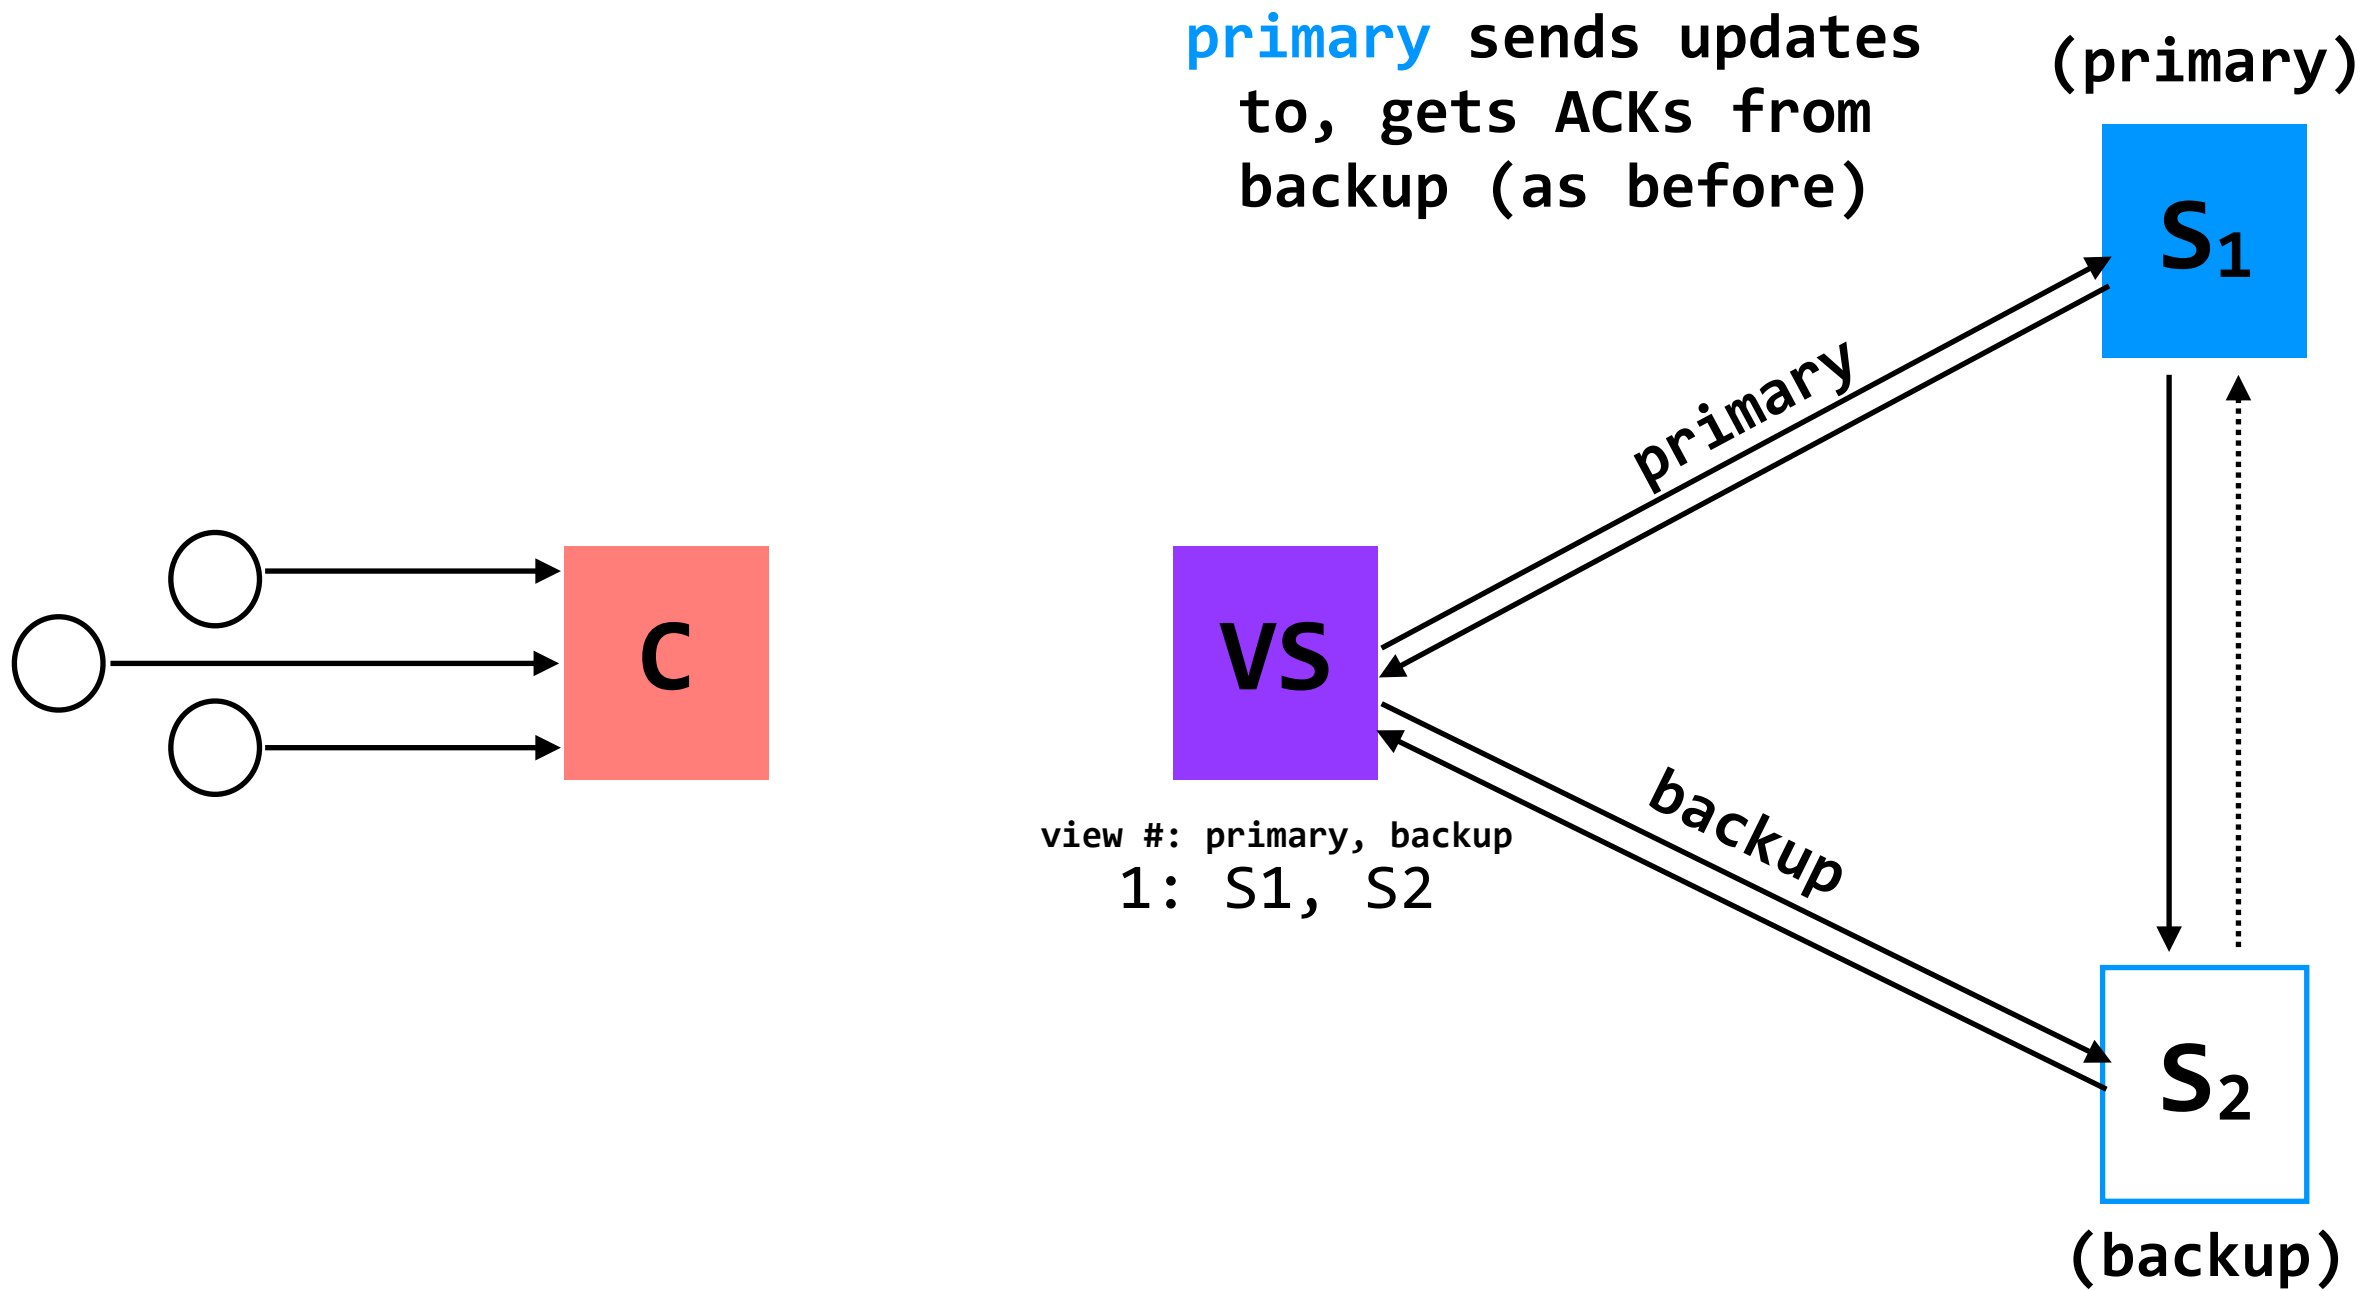
**attempt:** coordinators communicate with primary servers, who communicate with backup servers



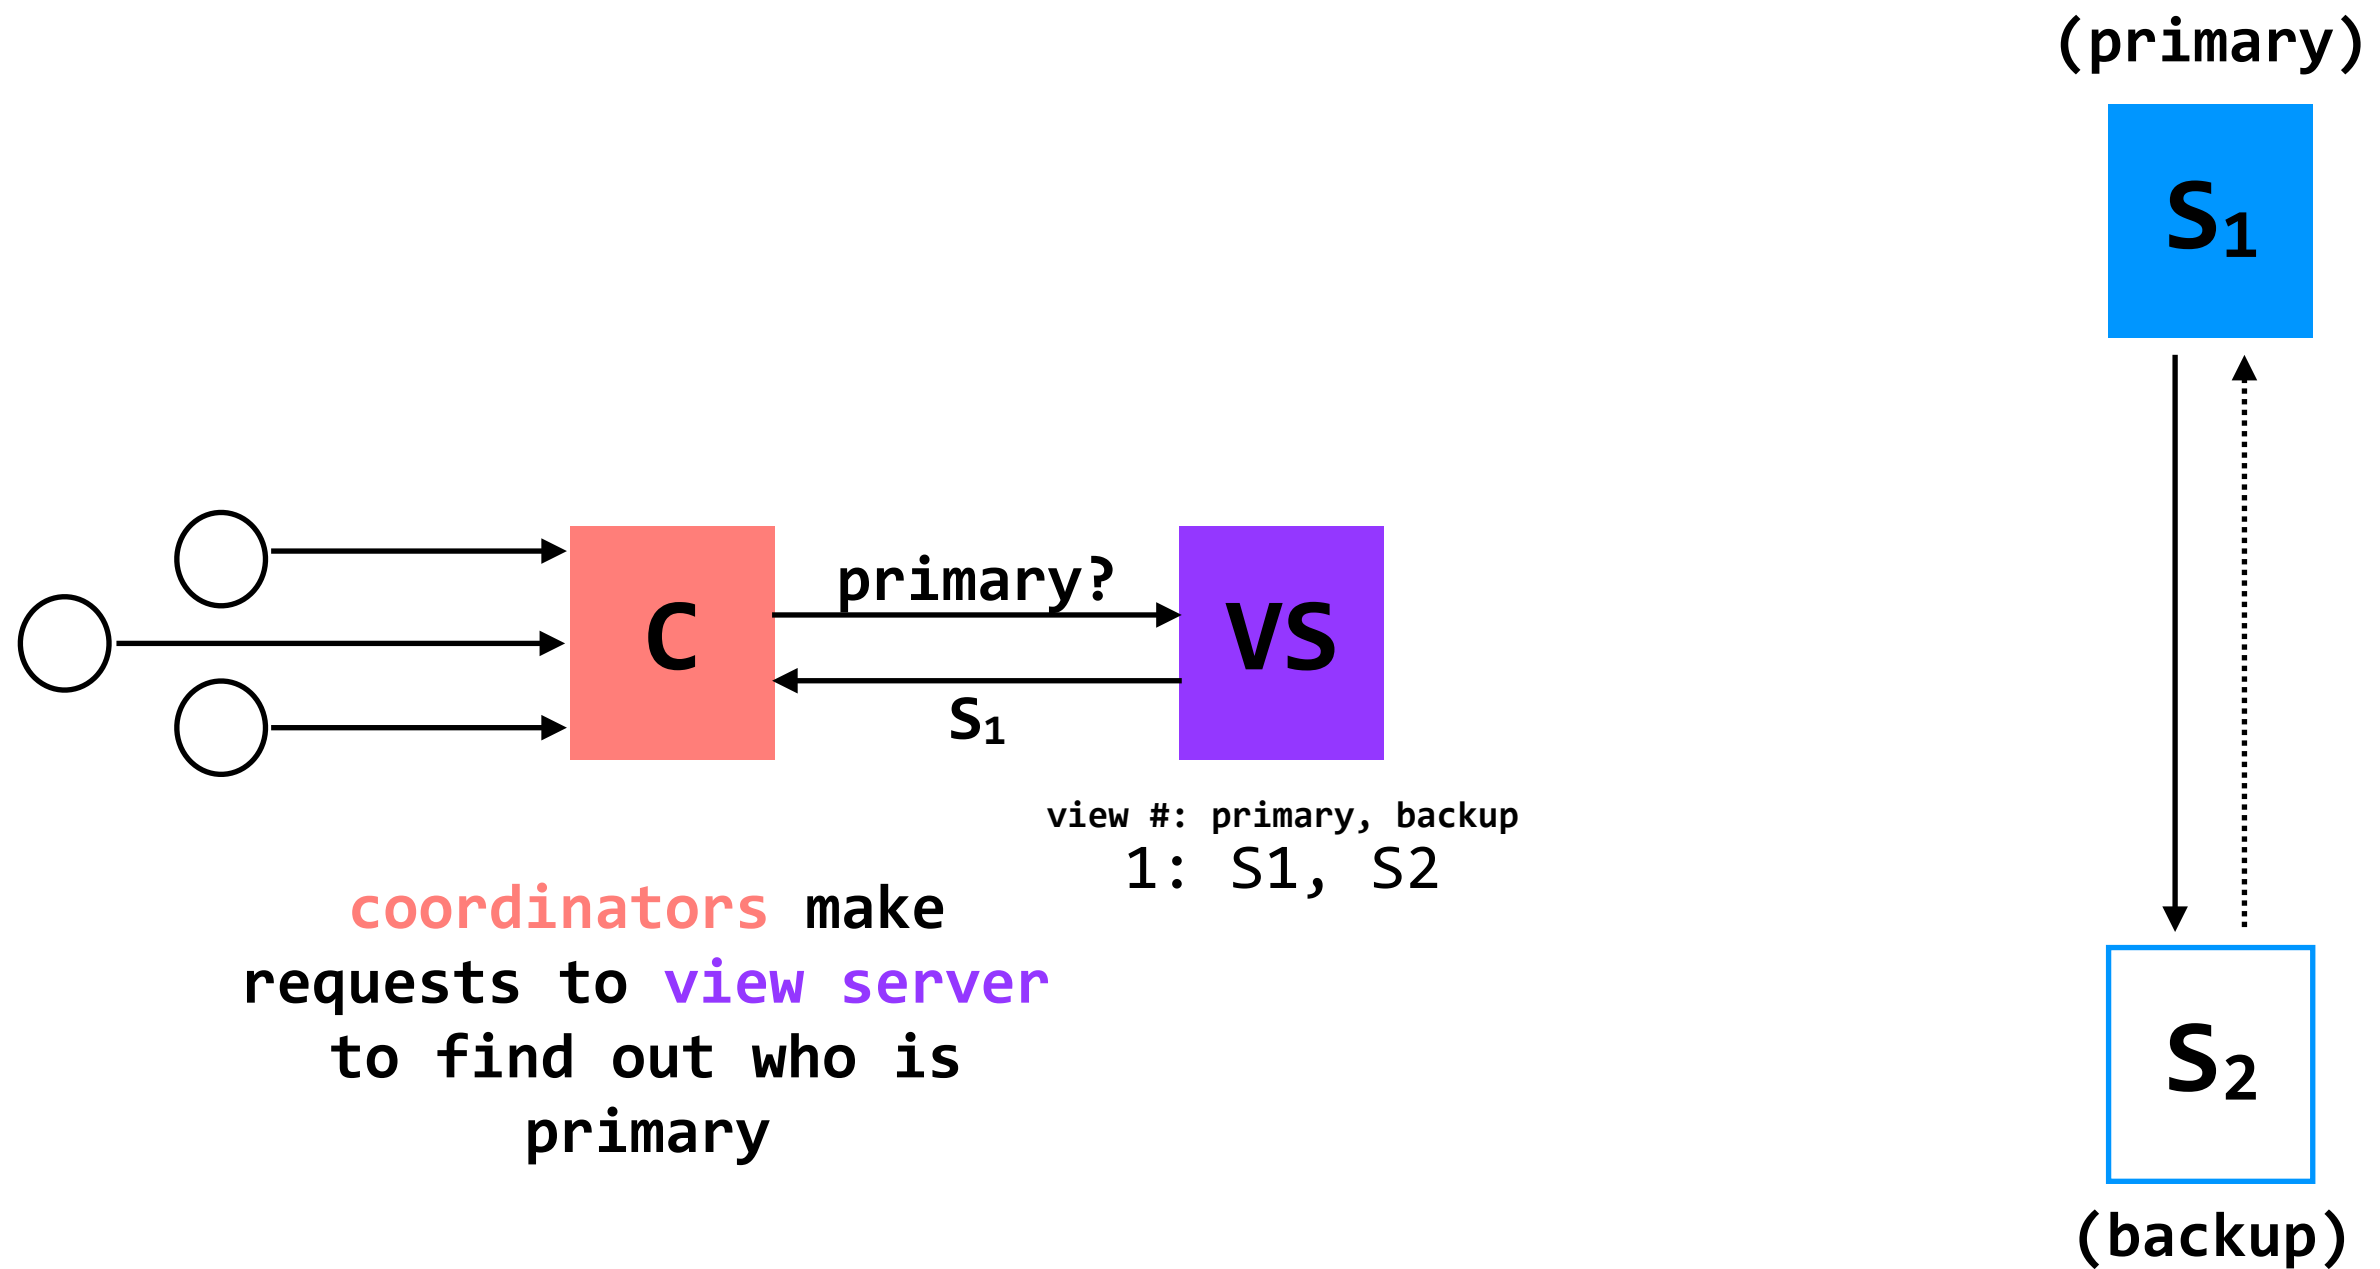
use a **view server**, which determines which replica is the primary



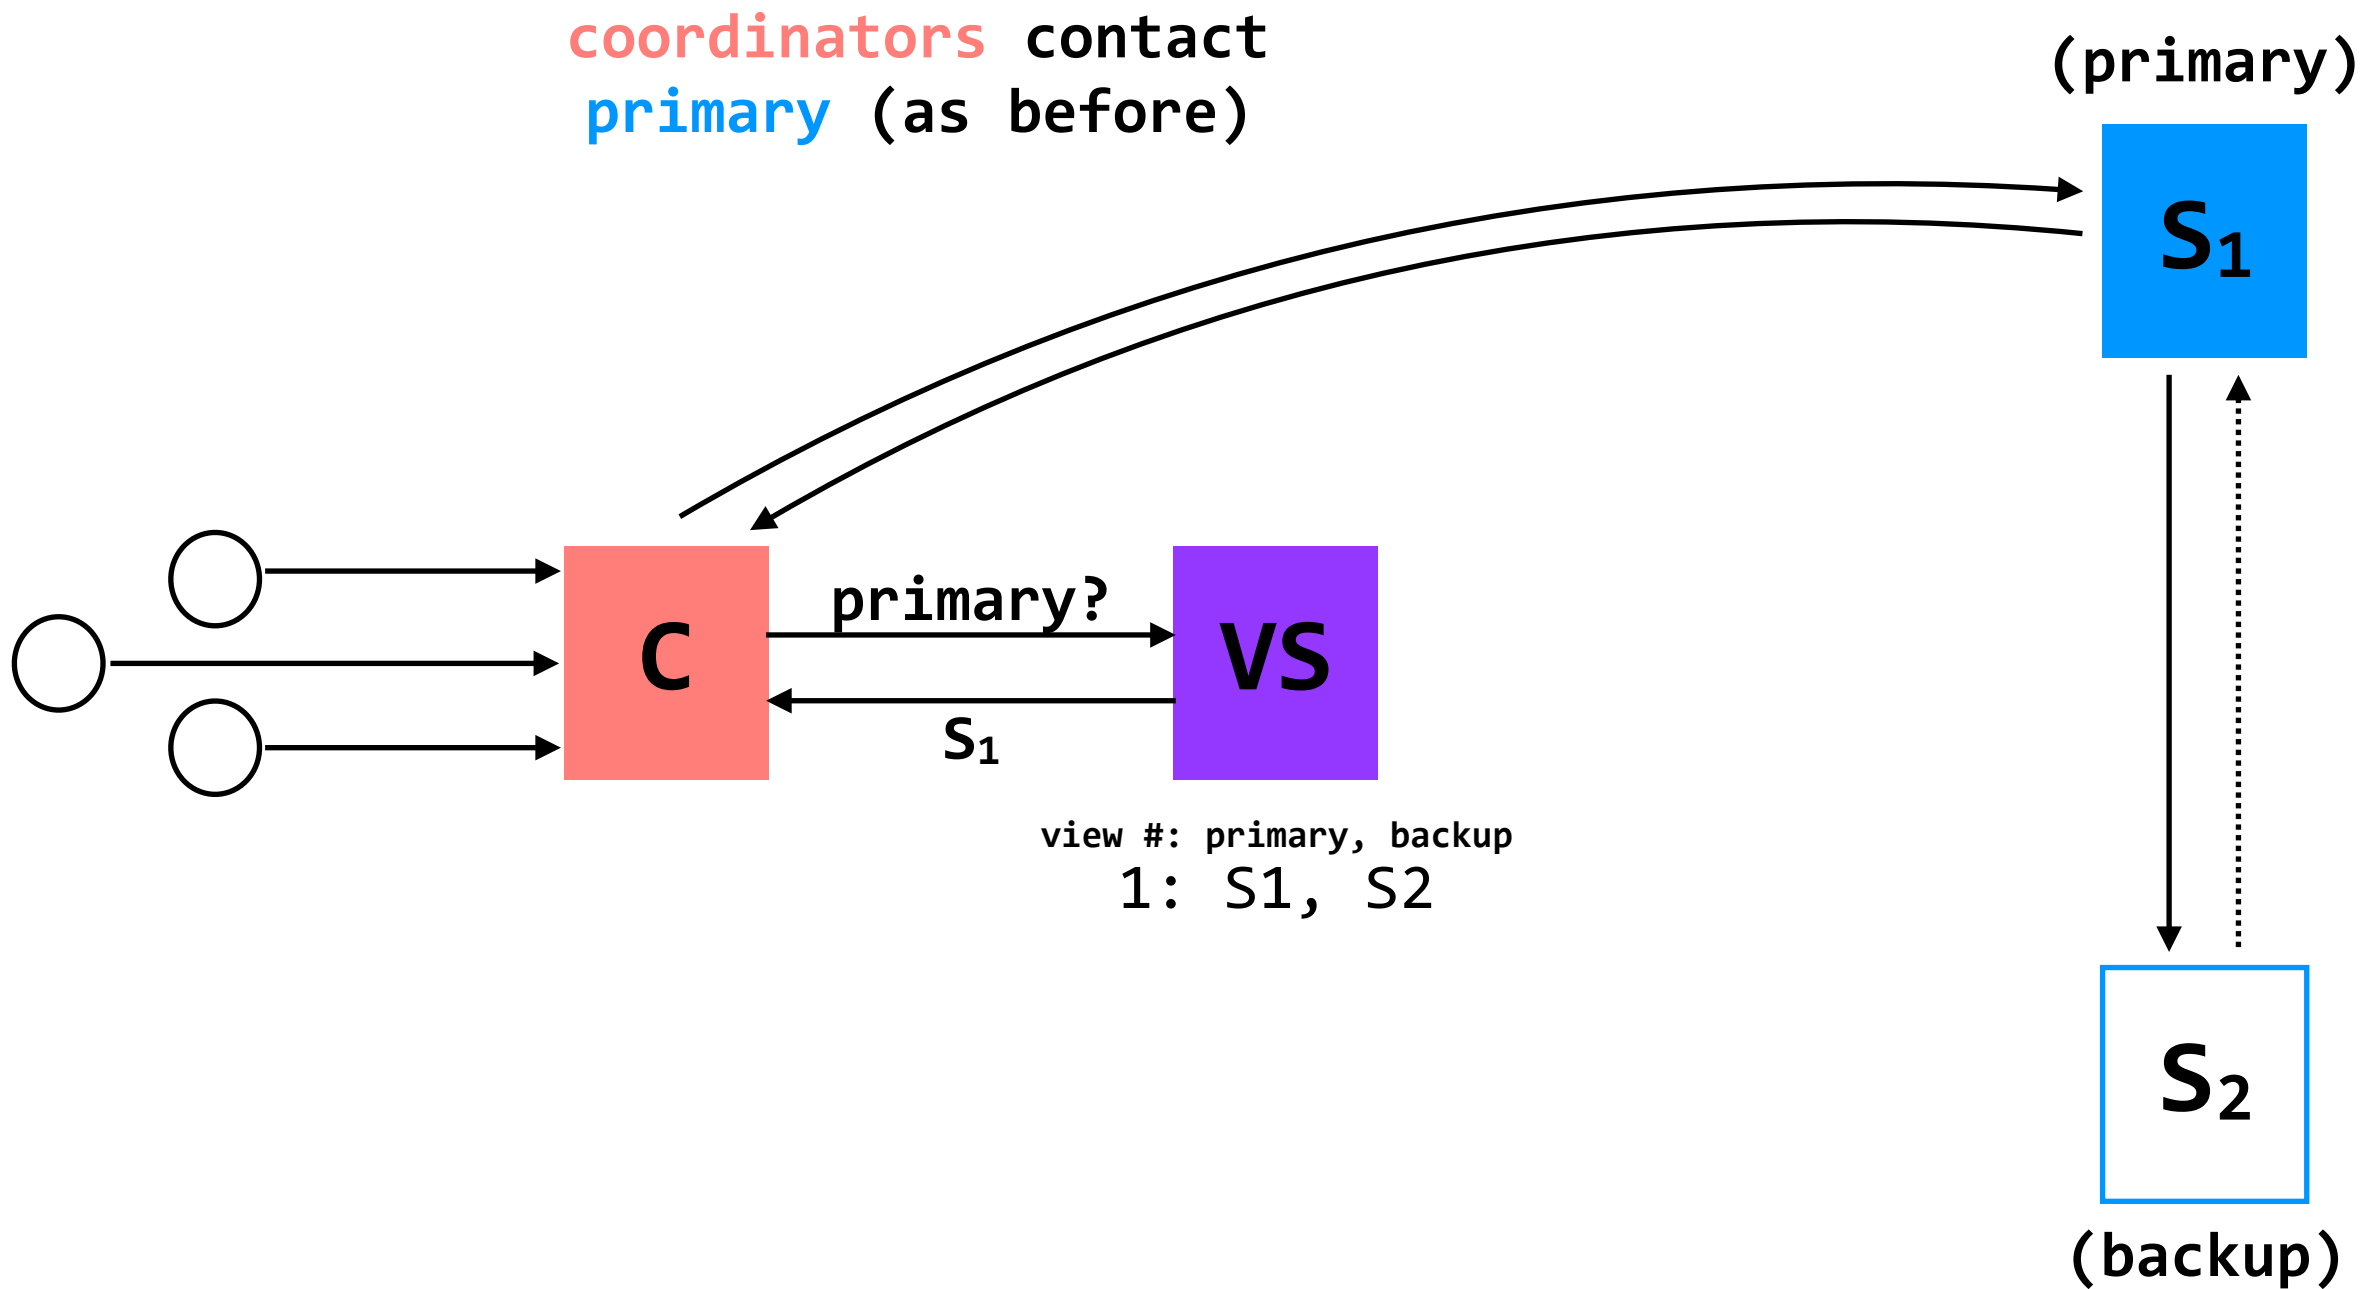
use a **view server**, which determines which replica is the primary



use a **view server**, which determines which replica is the primary

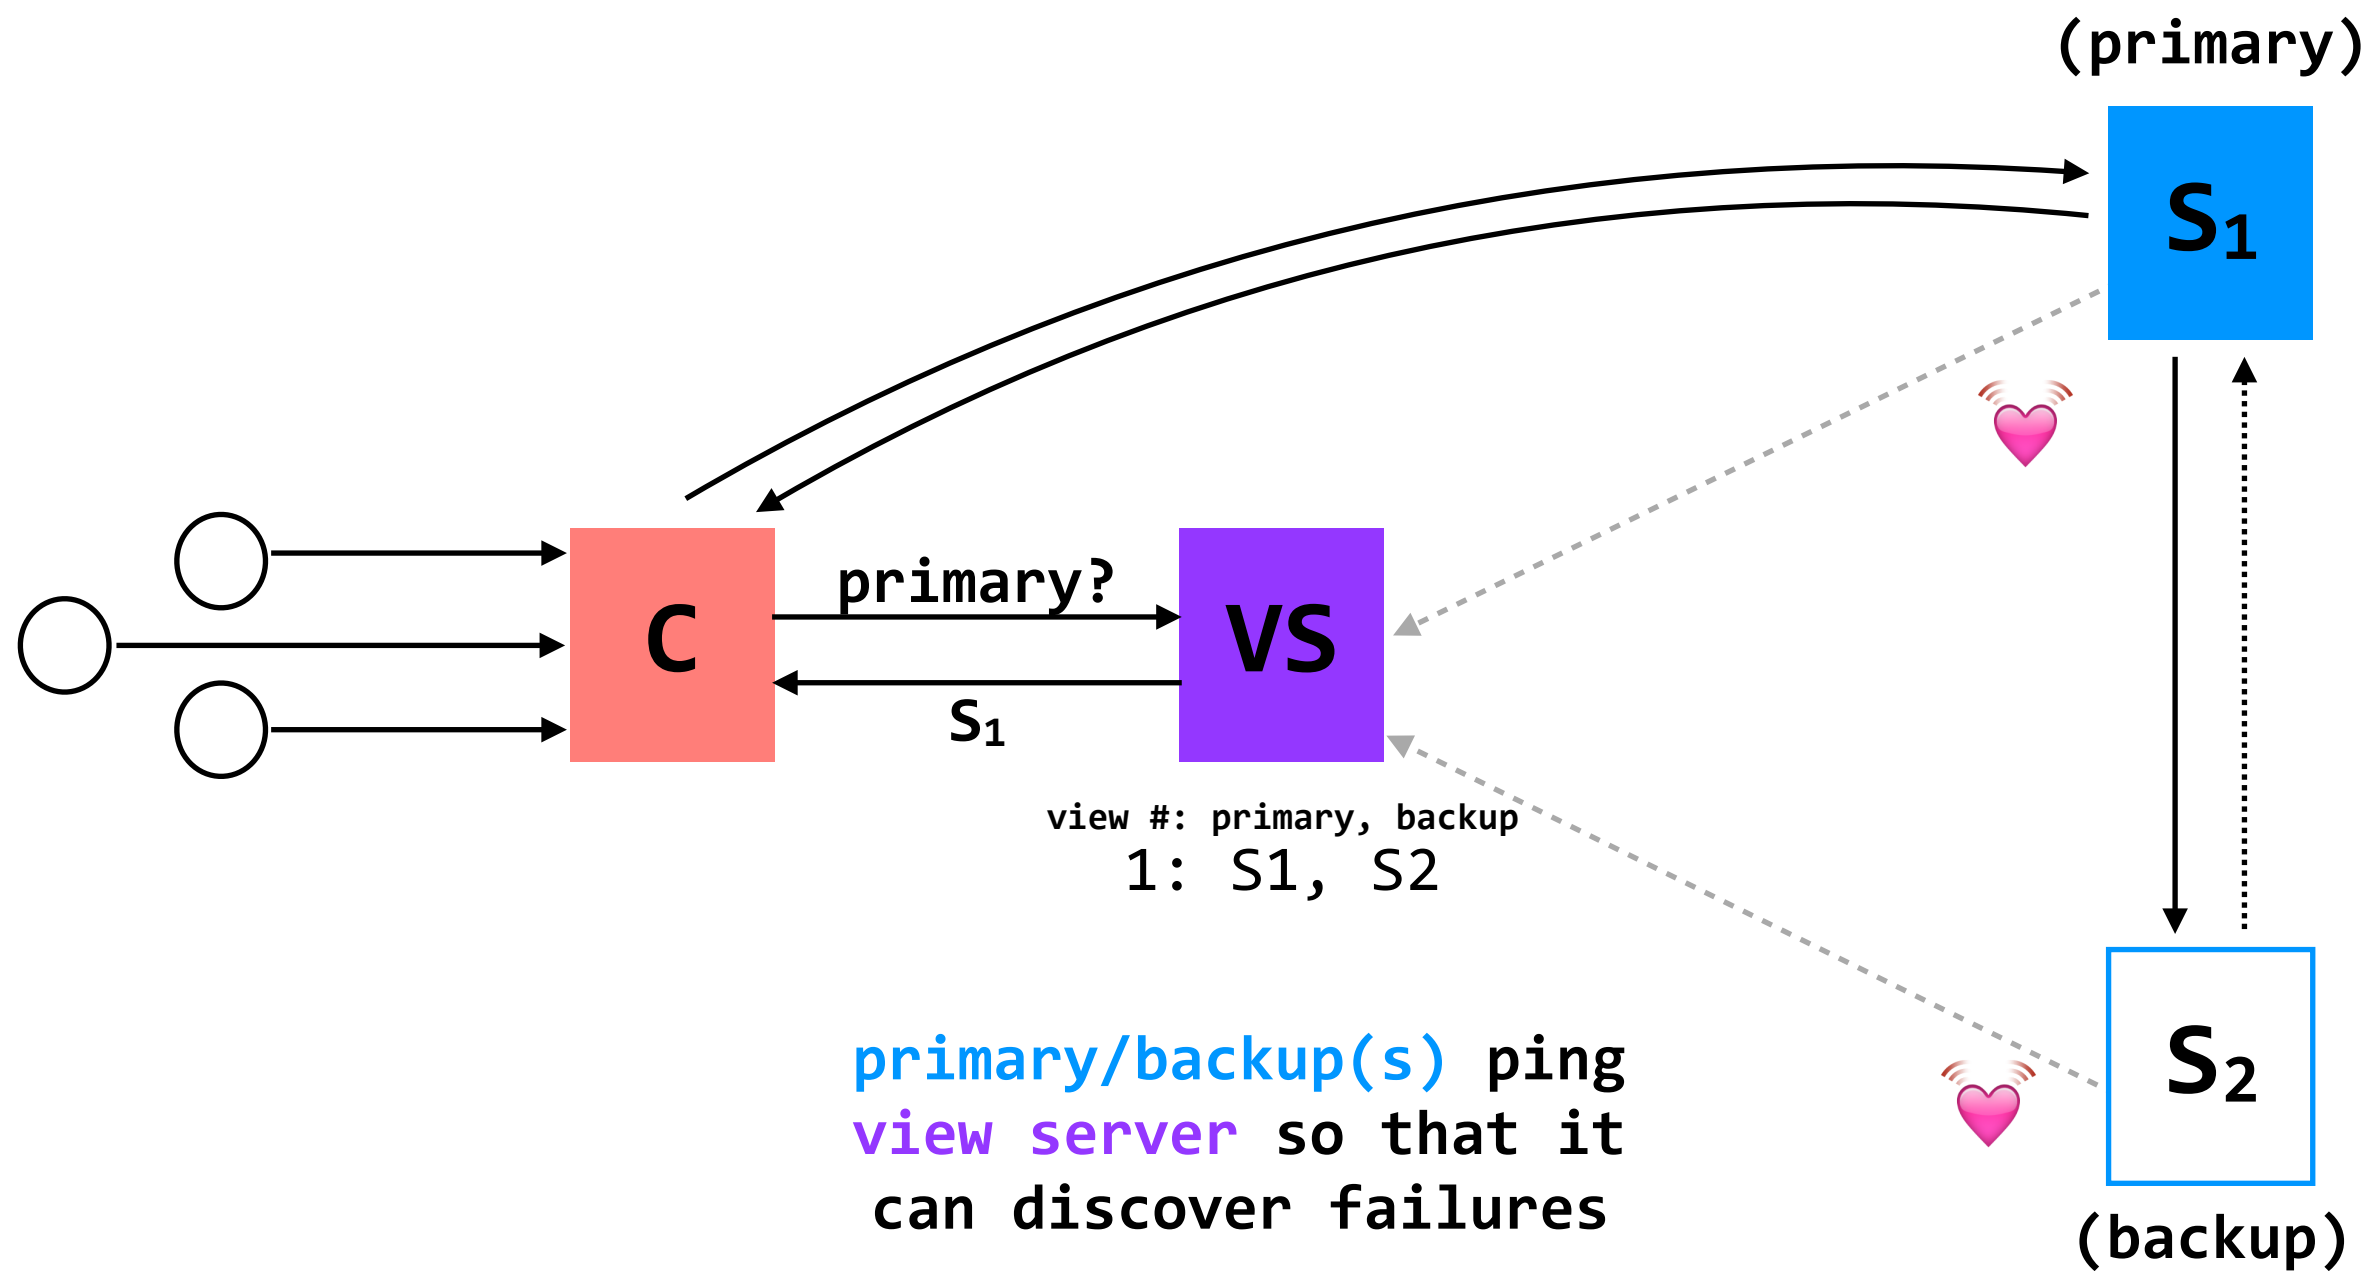


use a **view server**, which determines which replica is the primary



use a **view server**, which determines which replica is the primary







use a **view server**, which determines which replica is the primary


# handling primary failure

(dead)



lack of pings indicates  
to **VS** that **S<sub>1</sub>** is down



view #: primary, backup  
1: S<sub>1</sub>, S<sub>2</sub>



(backup)

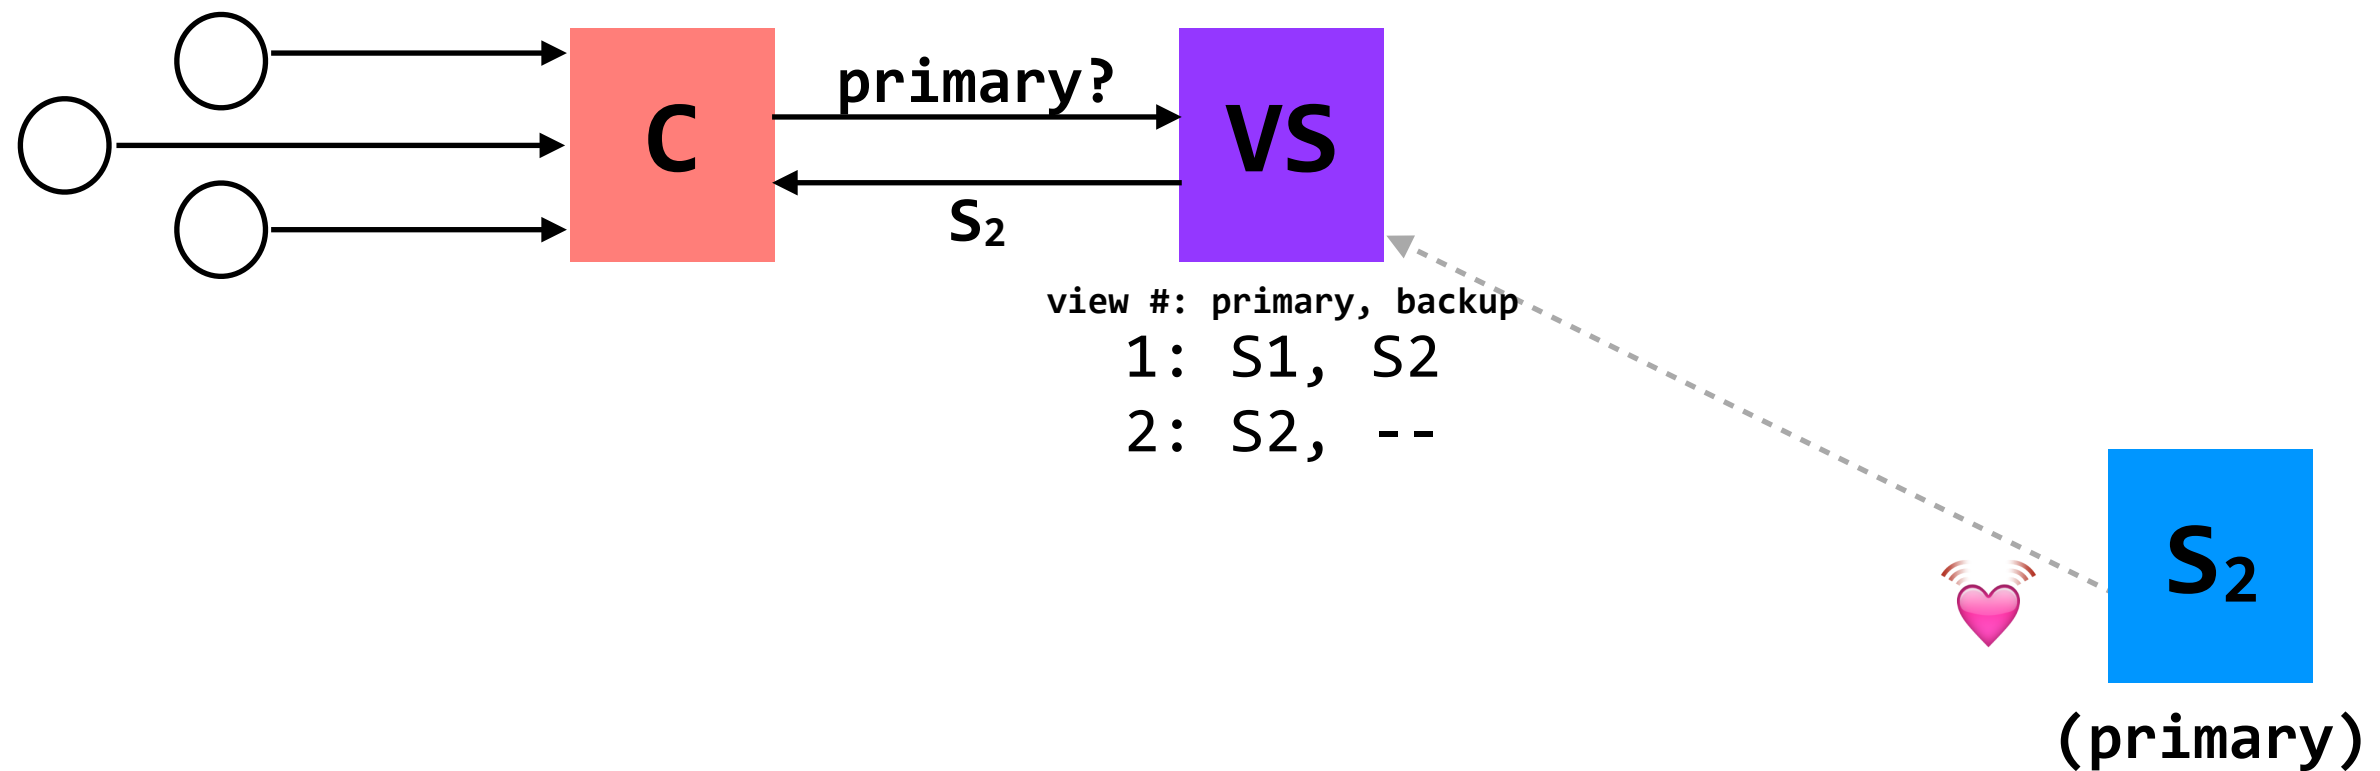
# handling primary failure

(dead)




# handling primary failure

(dead)

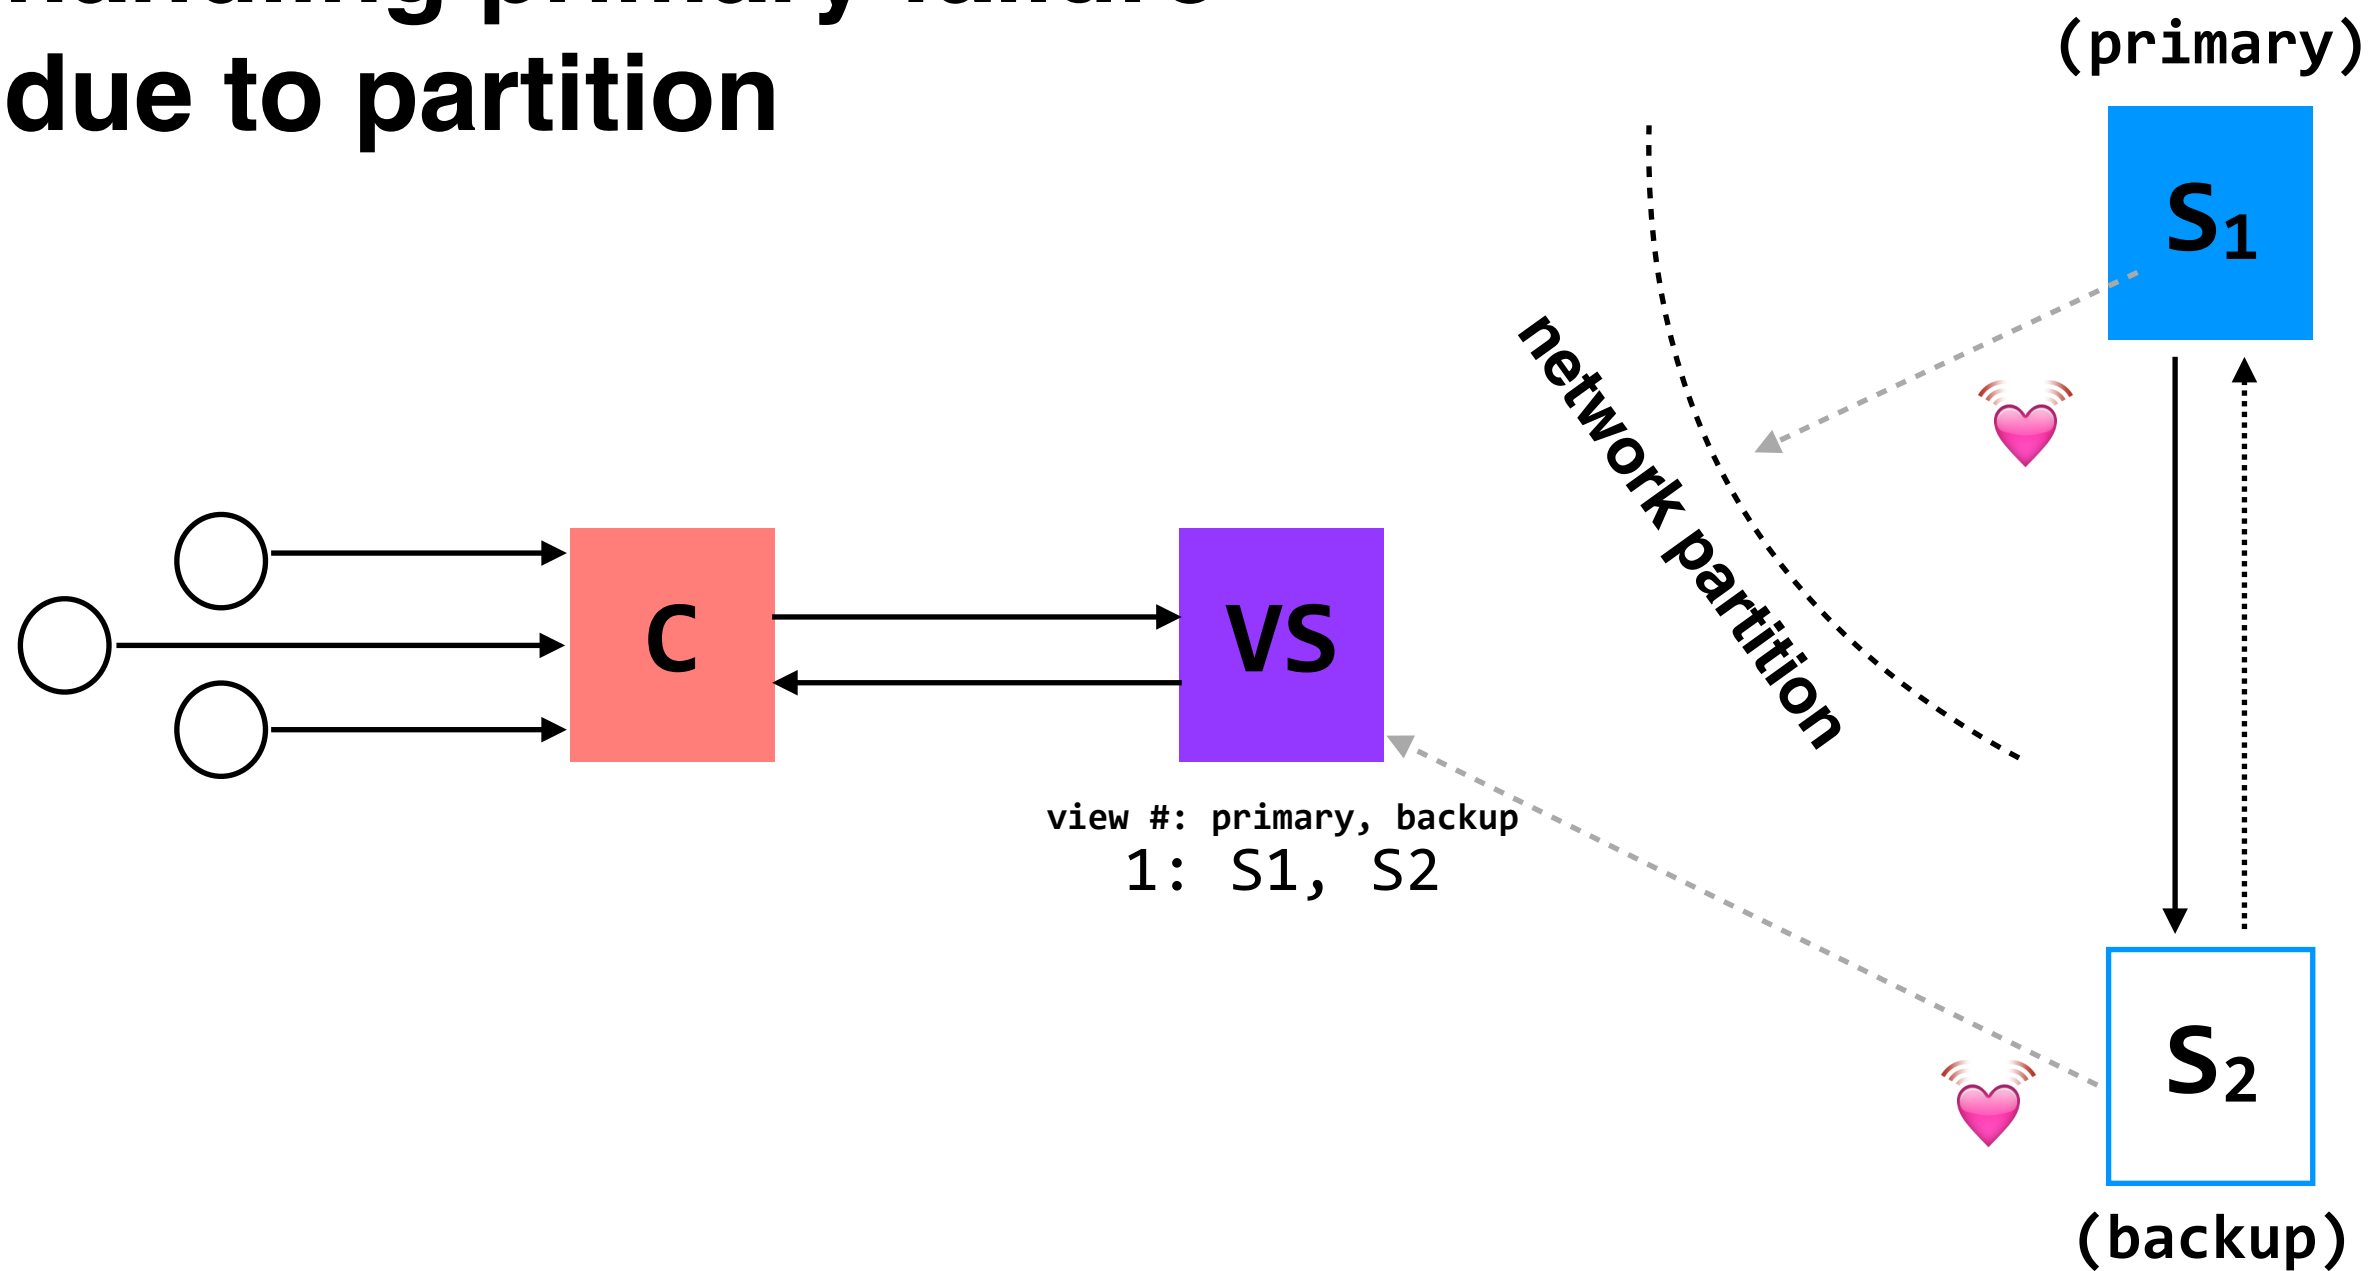


# handling primary failure

(dead)

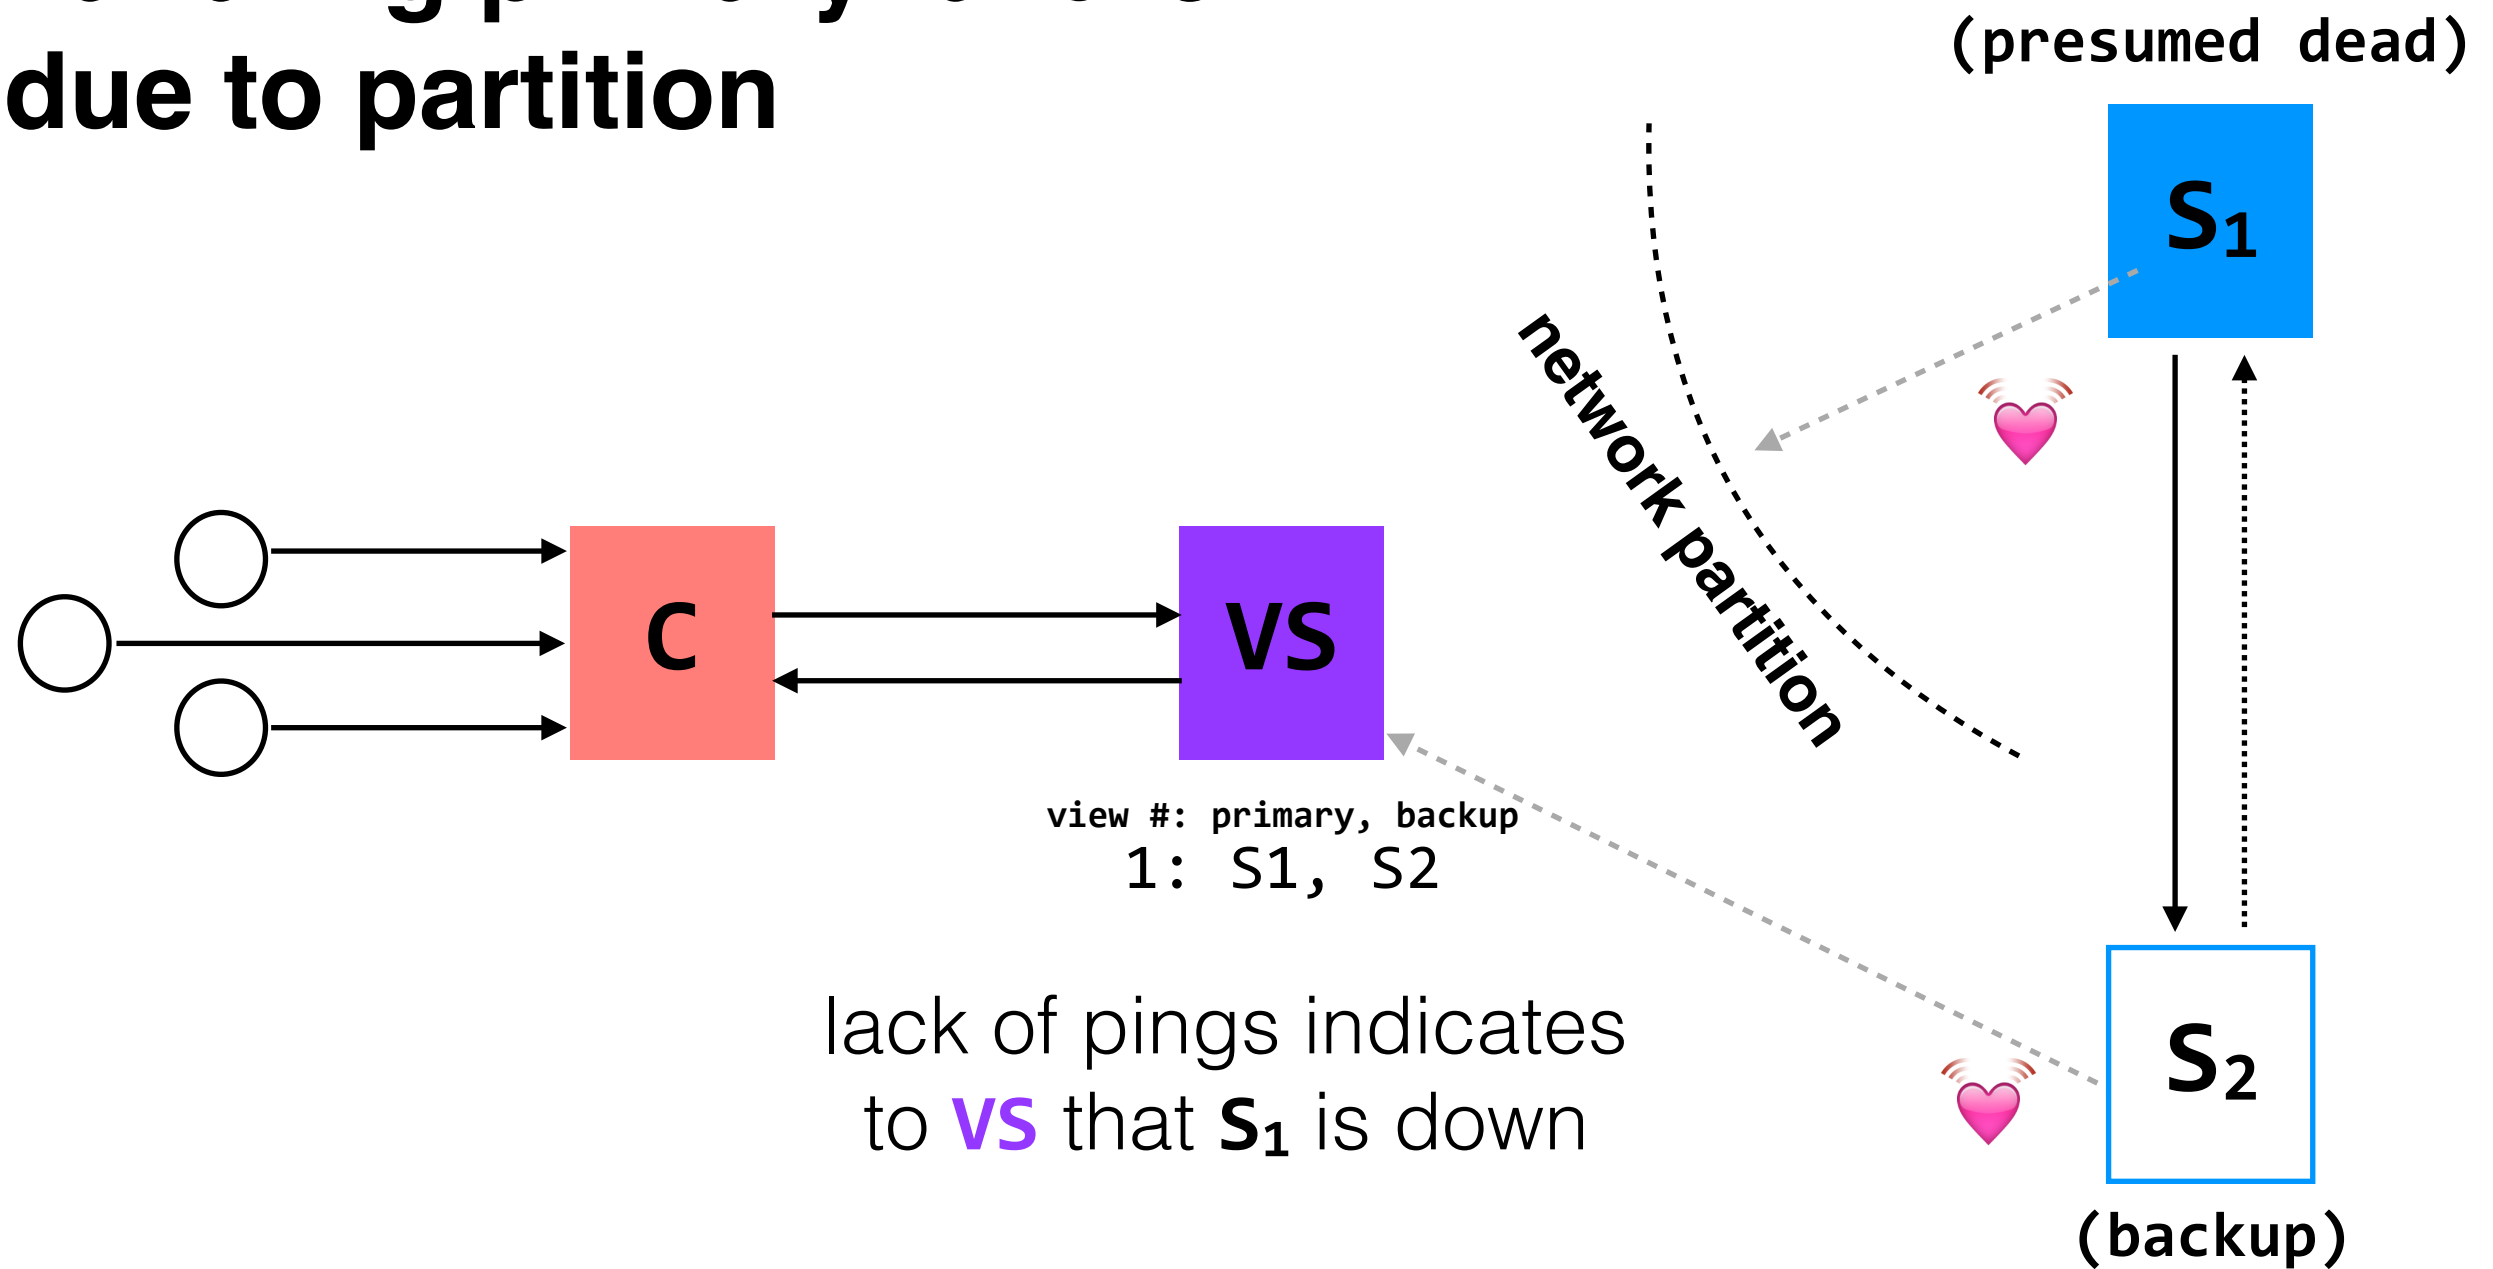


# handling primary failure due to partition

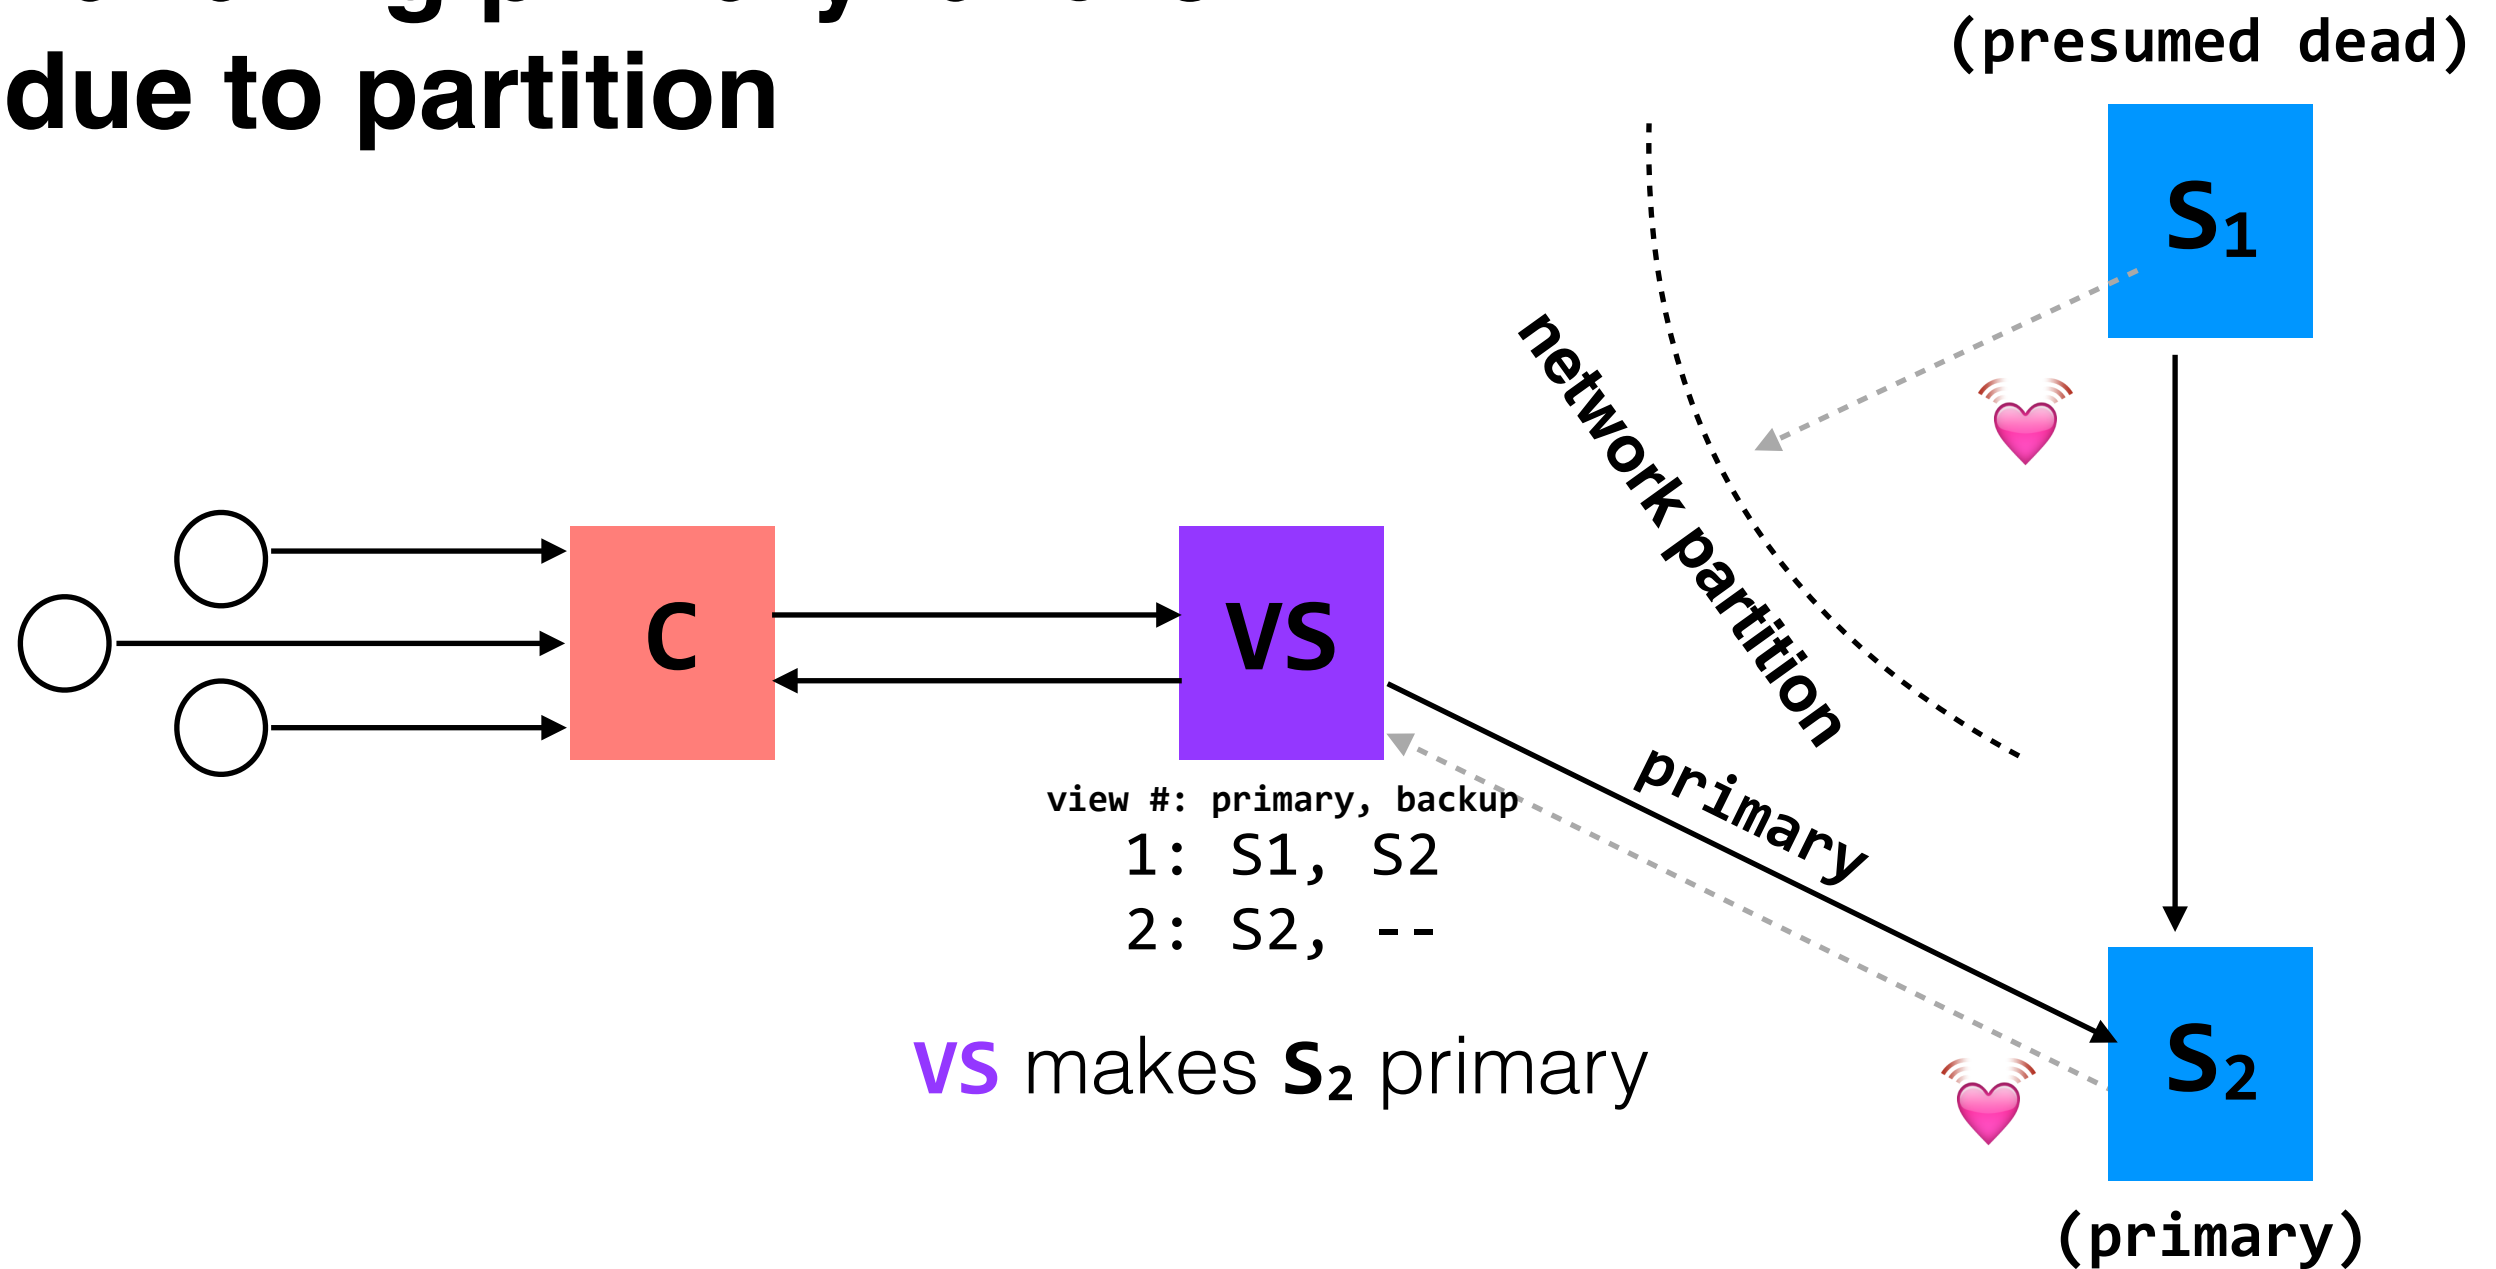


suppose a partition keeps **S<sub>1</sub>** from communicating with the  
**view server**

# handling primary failure due to partition

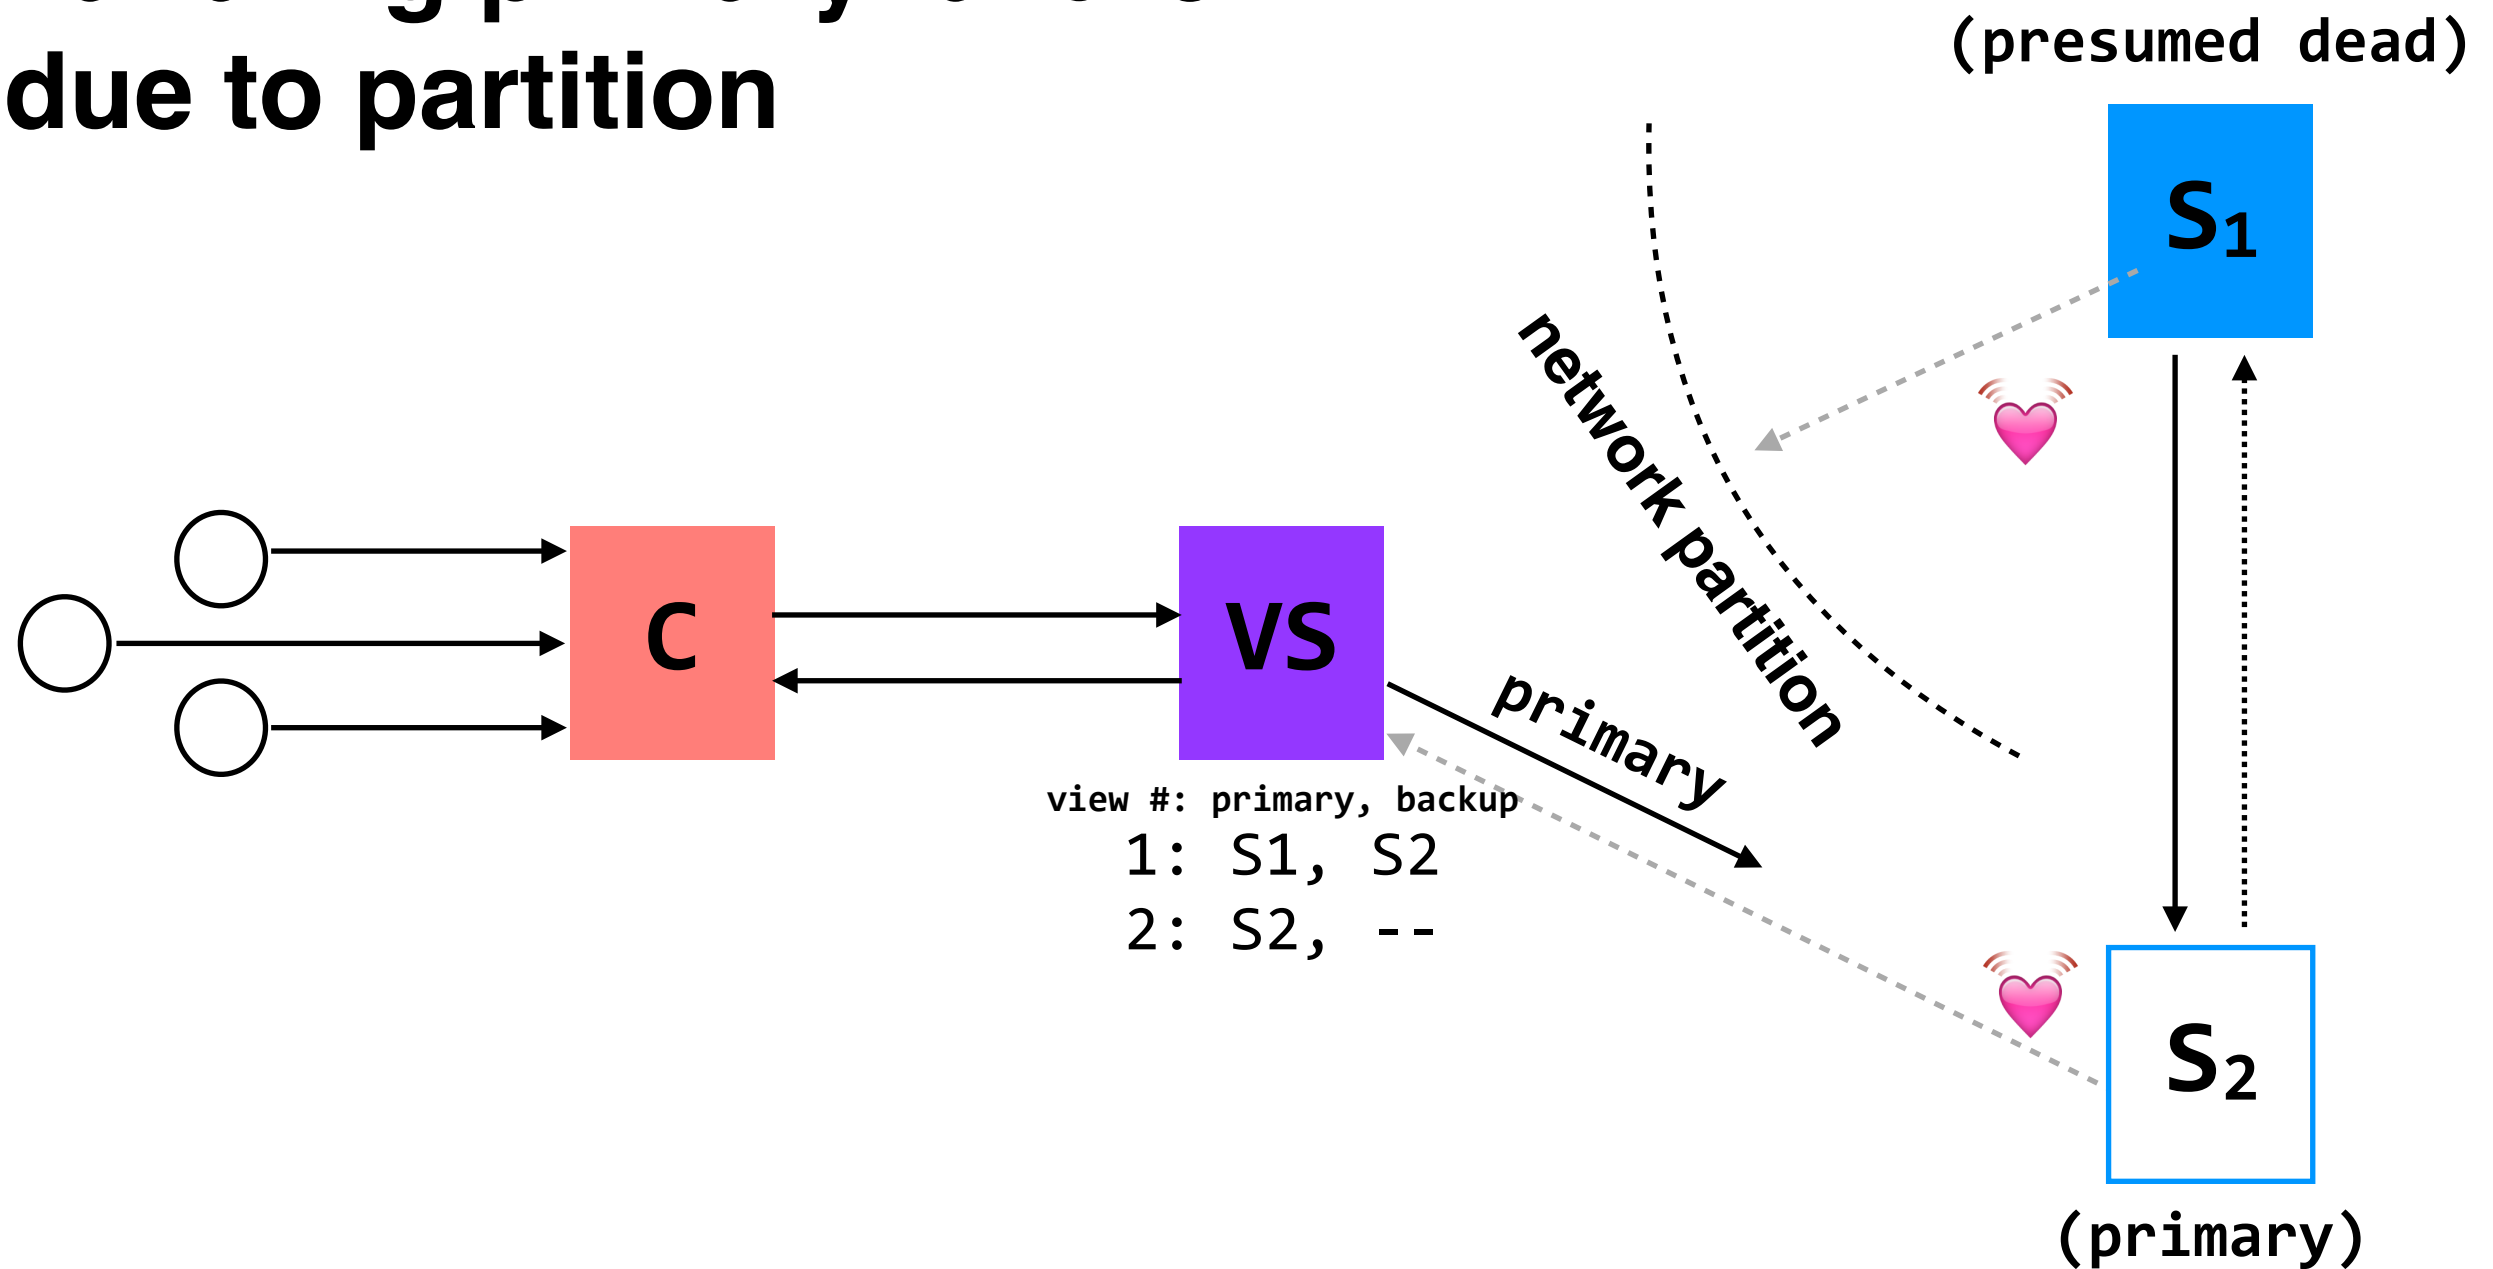


# handling primary failure due to partition



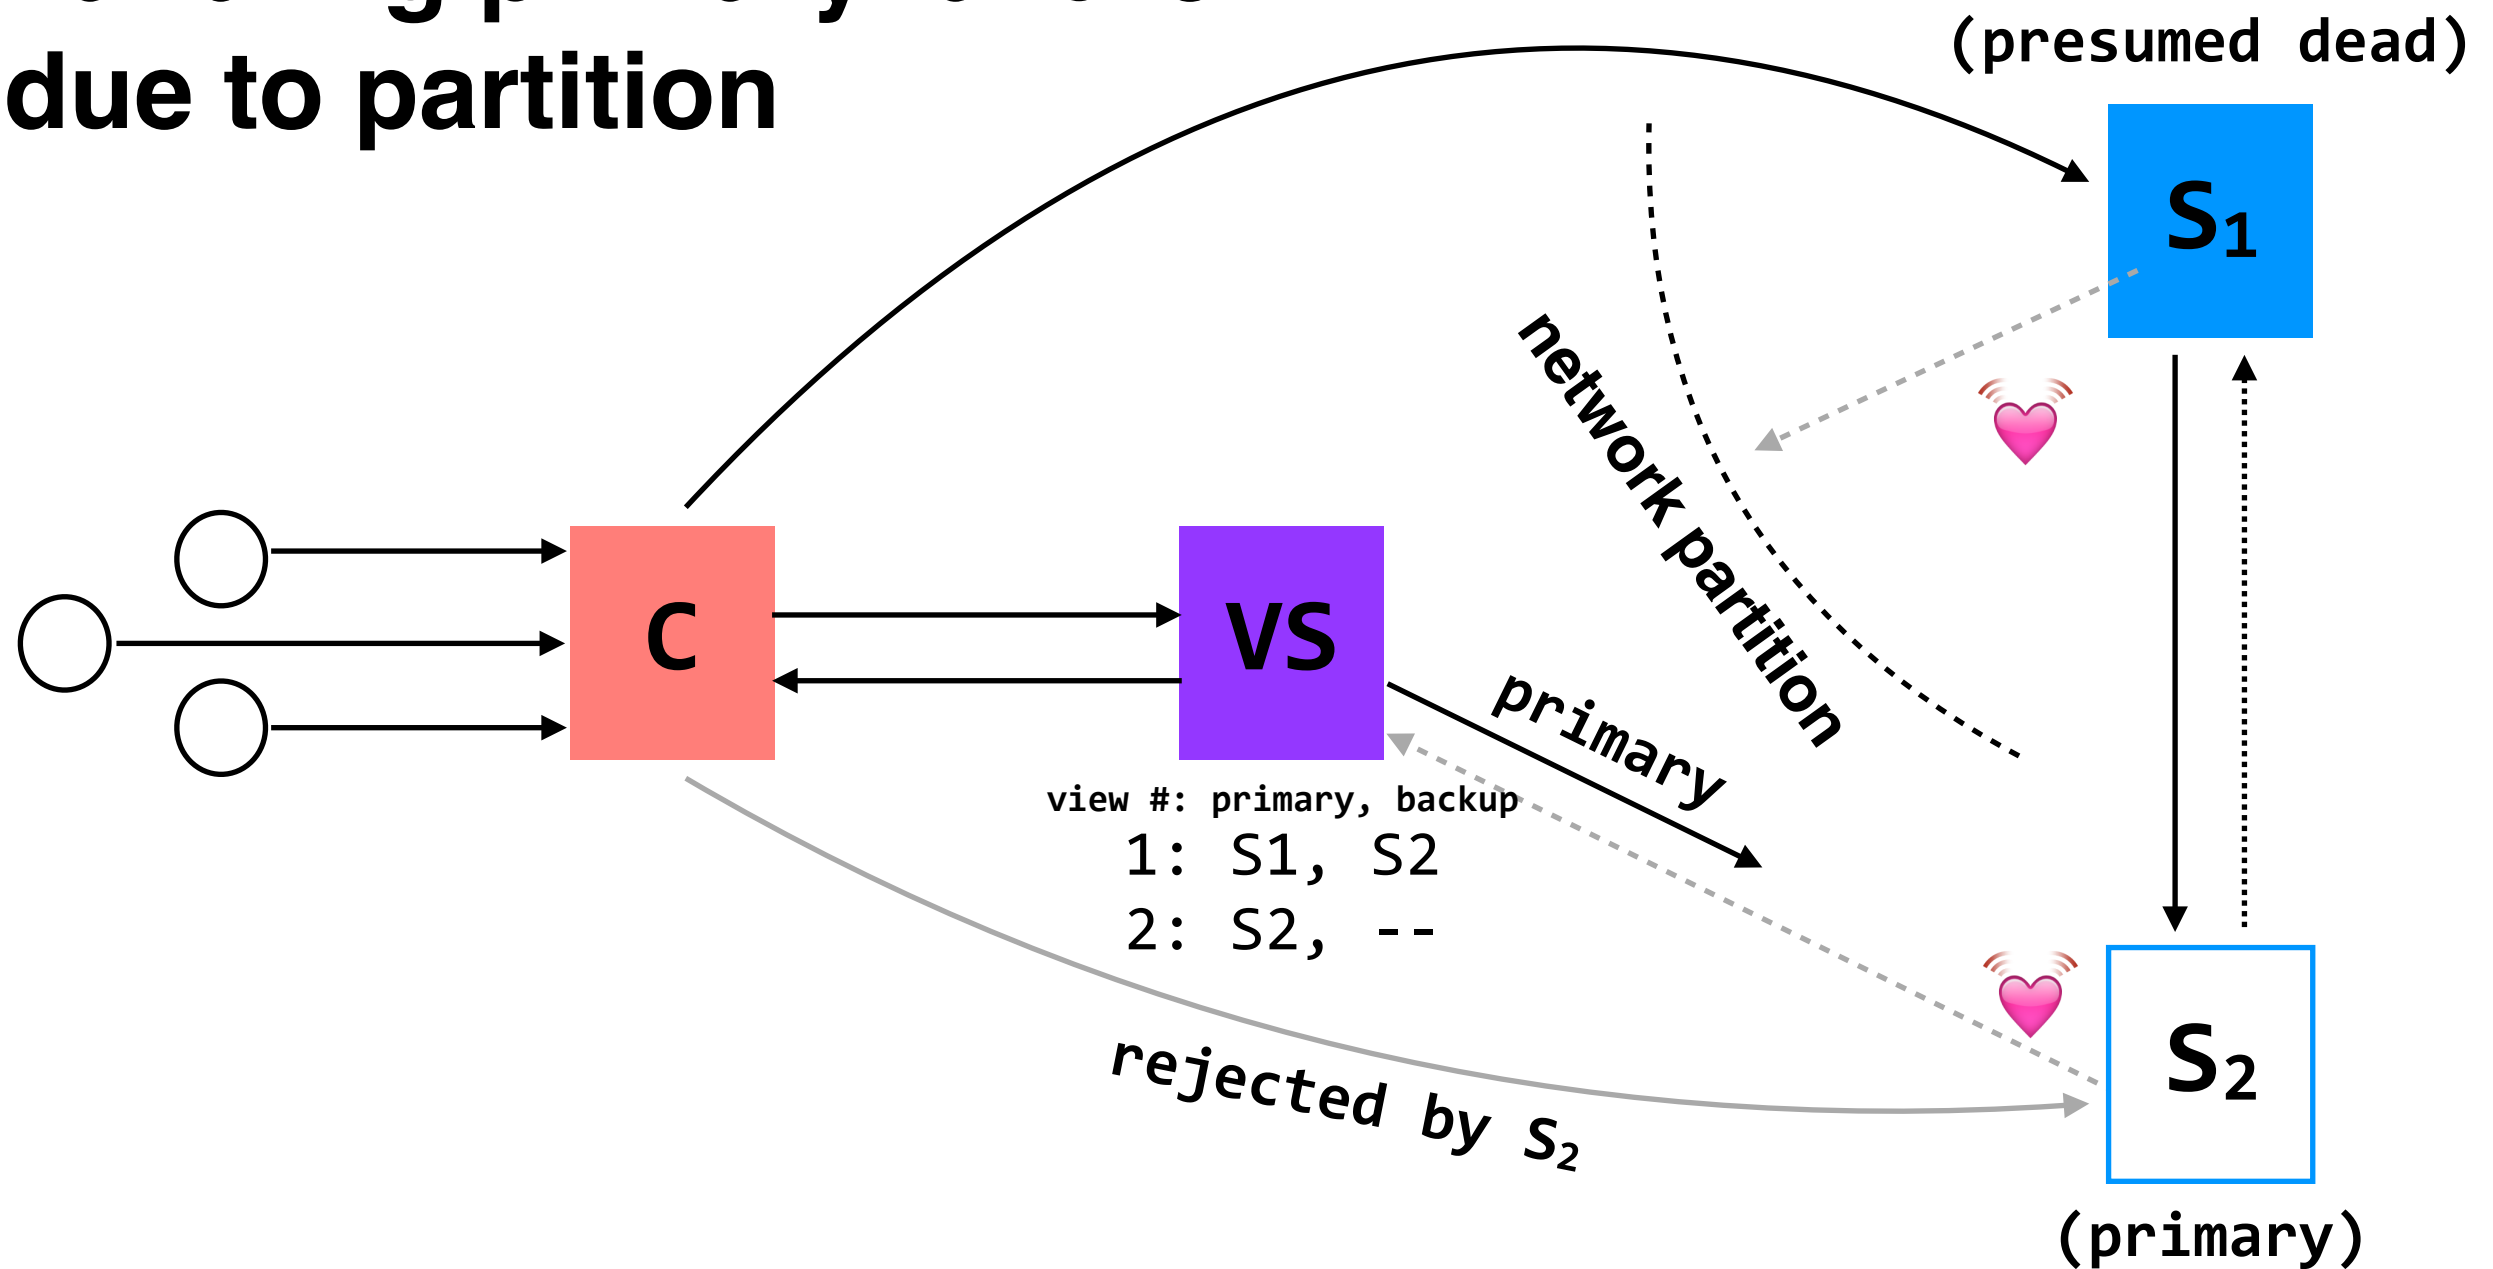


# handling primary failure due to partition



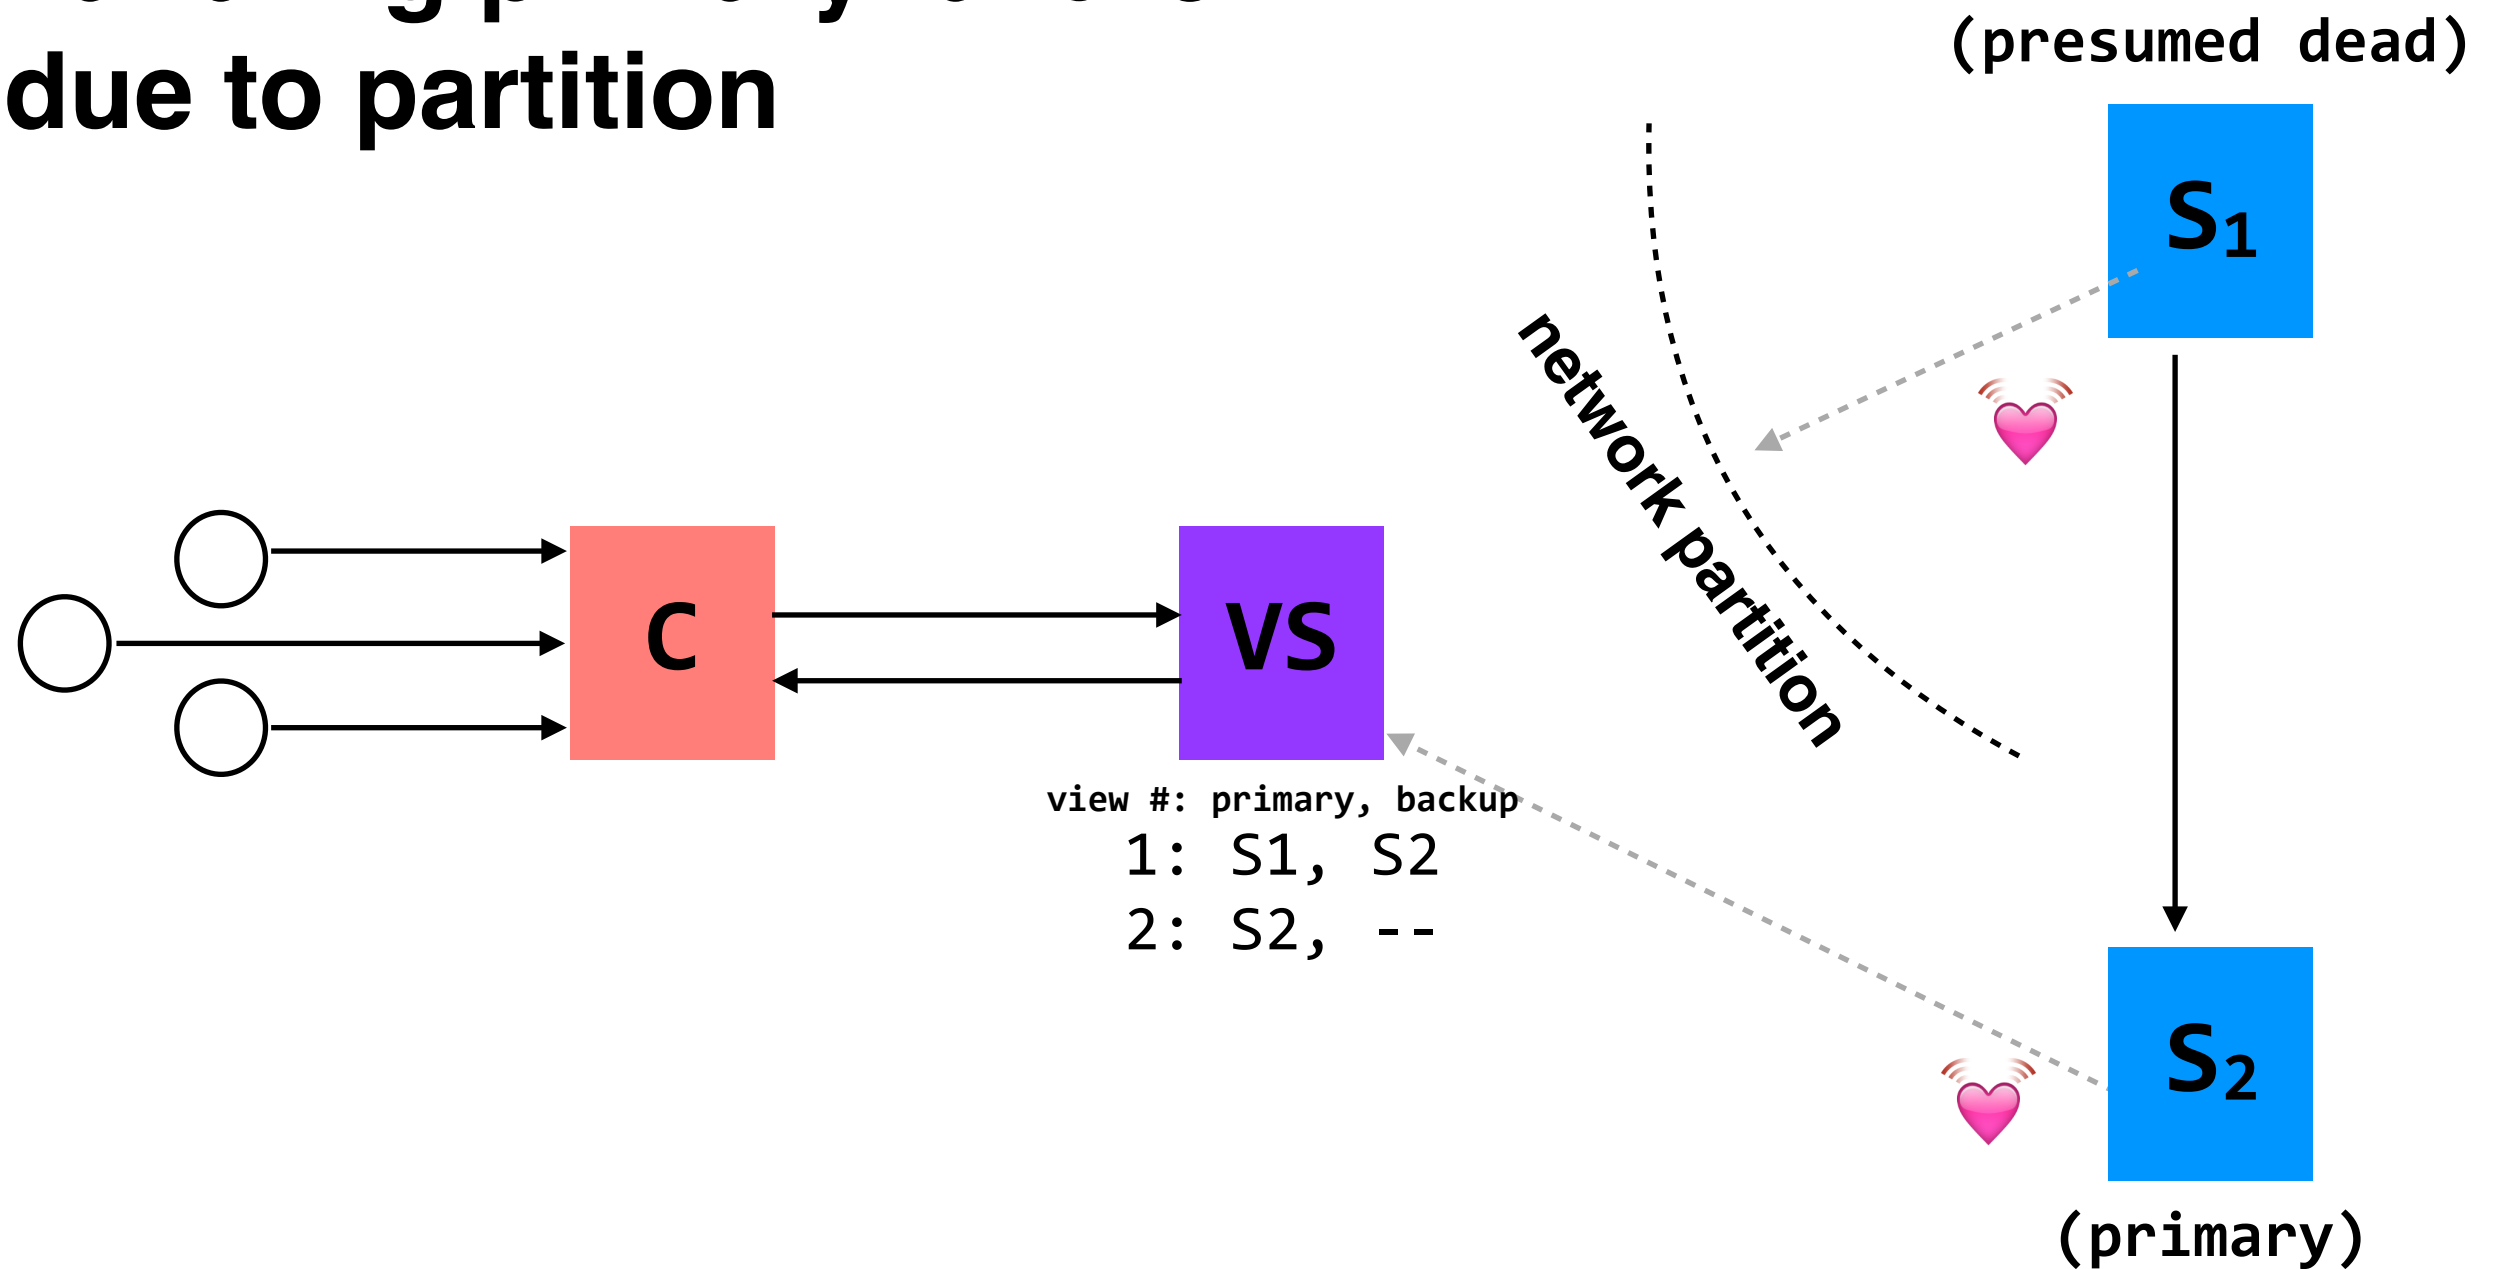
**question:** what happens before S<sub>2</sub> knows  
it's the primary?

# handling primary failure due to partition



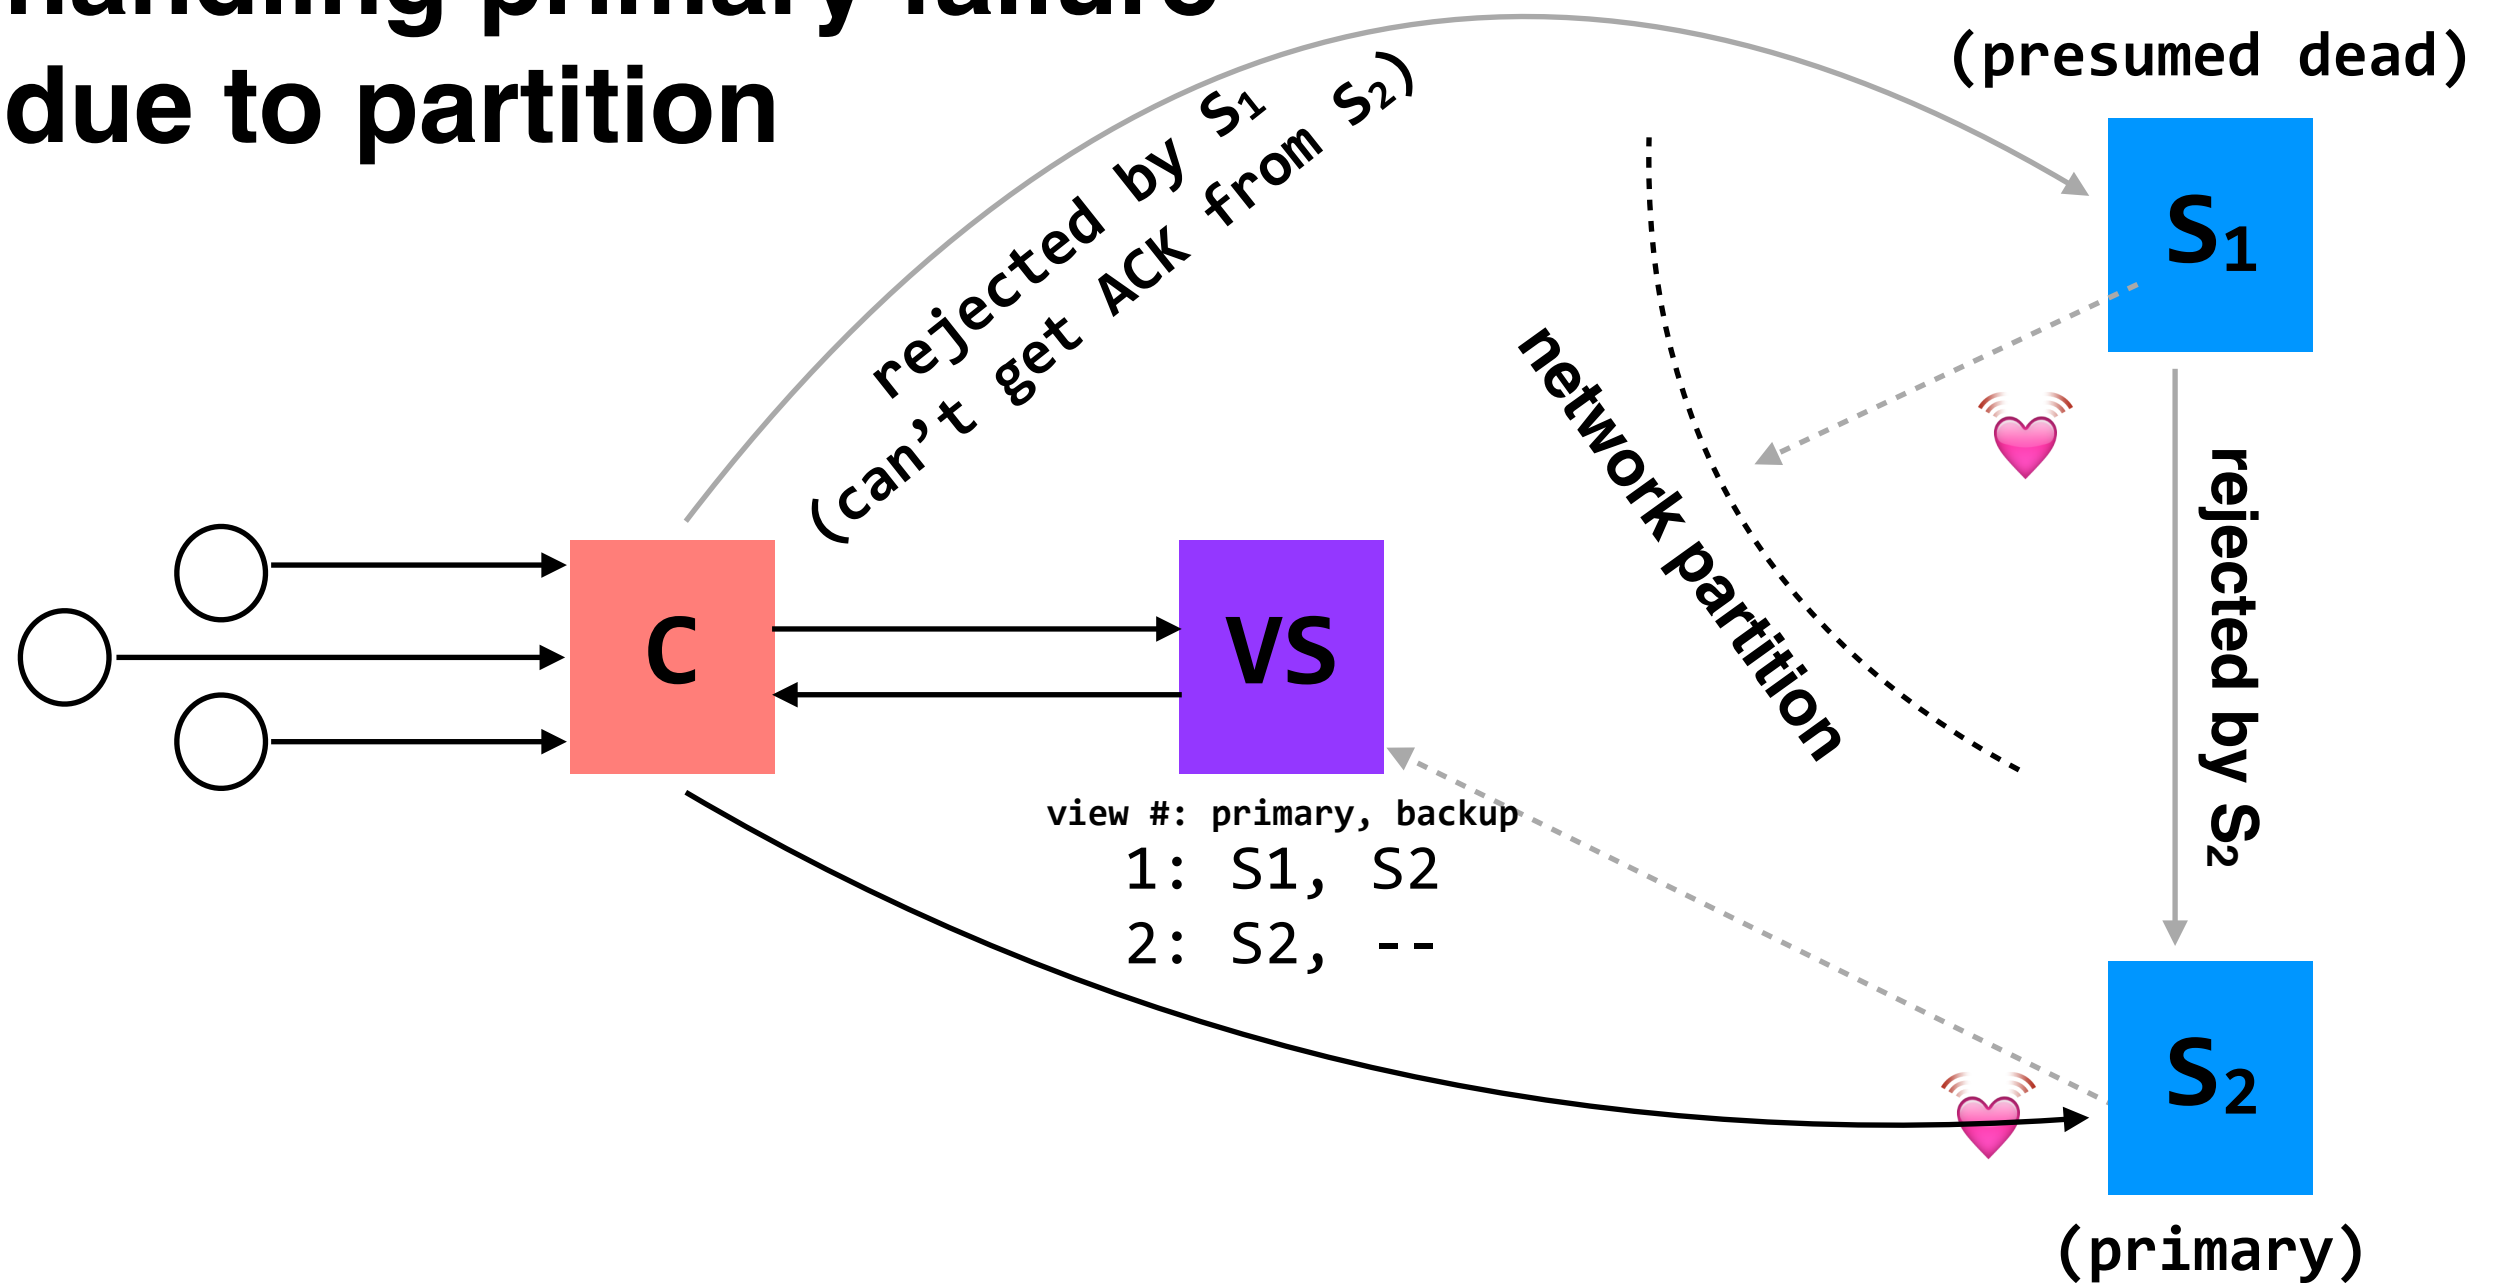
**S<sub>2</sub> will act as backup**  
(accept updates from S<sub>1</sub>, reject coordinator requests)

# handling primary failure due to partition

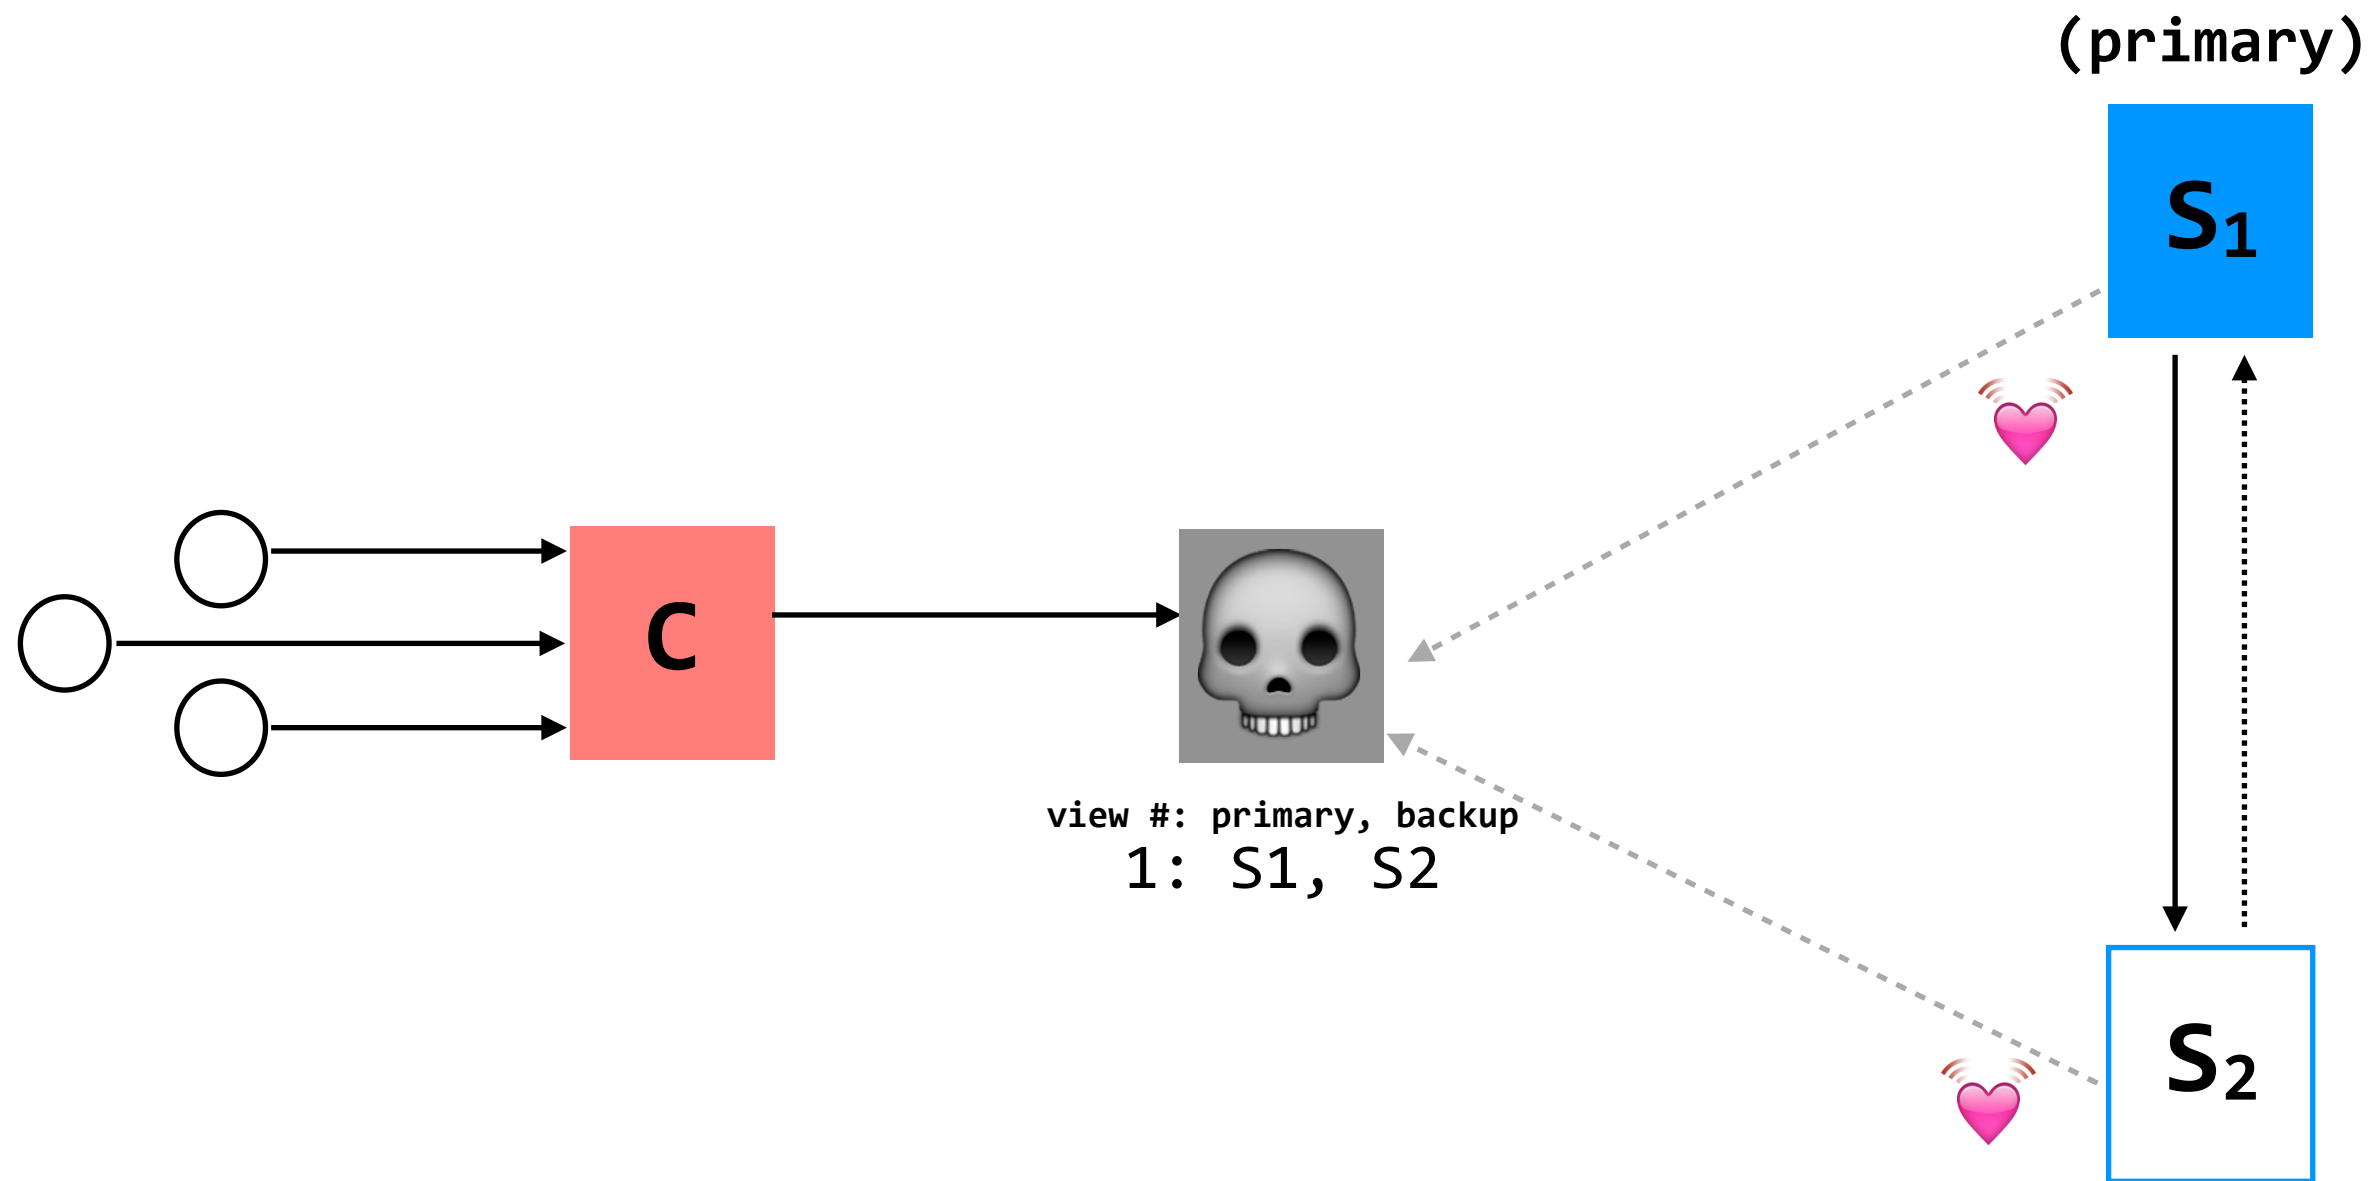


**question:** what happens after  $S_2$  knows it's the primary, but  $S_1$  also thinks it is?

# handling primary failure due to partition



**S<sub>1</sub> won't be able to act as primary**  
(can't accept client requests because it won't get ACKs from S<sub>2</sub>)



**problem:** what if view server fails?

**go to recitation tomorrow and find out!**

- **Replicated state machines (RSMs)** provide **single-copy consistency**: operations complete as if there is a single copy of the data, though internally there are replicas.
- RSMs use a **primary-backup** mechanism for replication. The **view server** ensures that only one replica acts as the primary. It can also recruit new backups after servers fail.
- To extend this model to handle view-server failures, we need a mechanism to provide **distributed consensus**; see tomorrow's recitation (on Raft).