

REMINDER

course evaluations are online

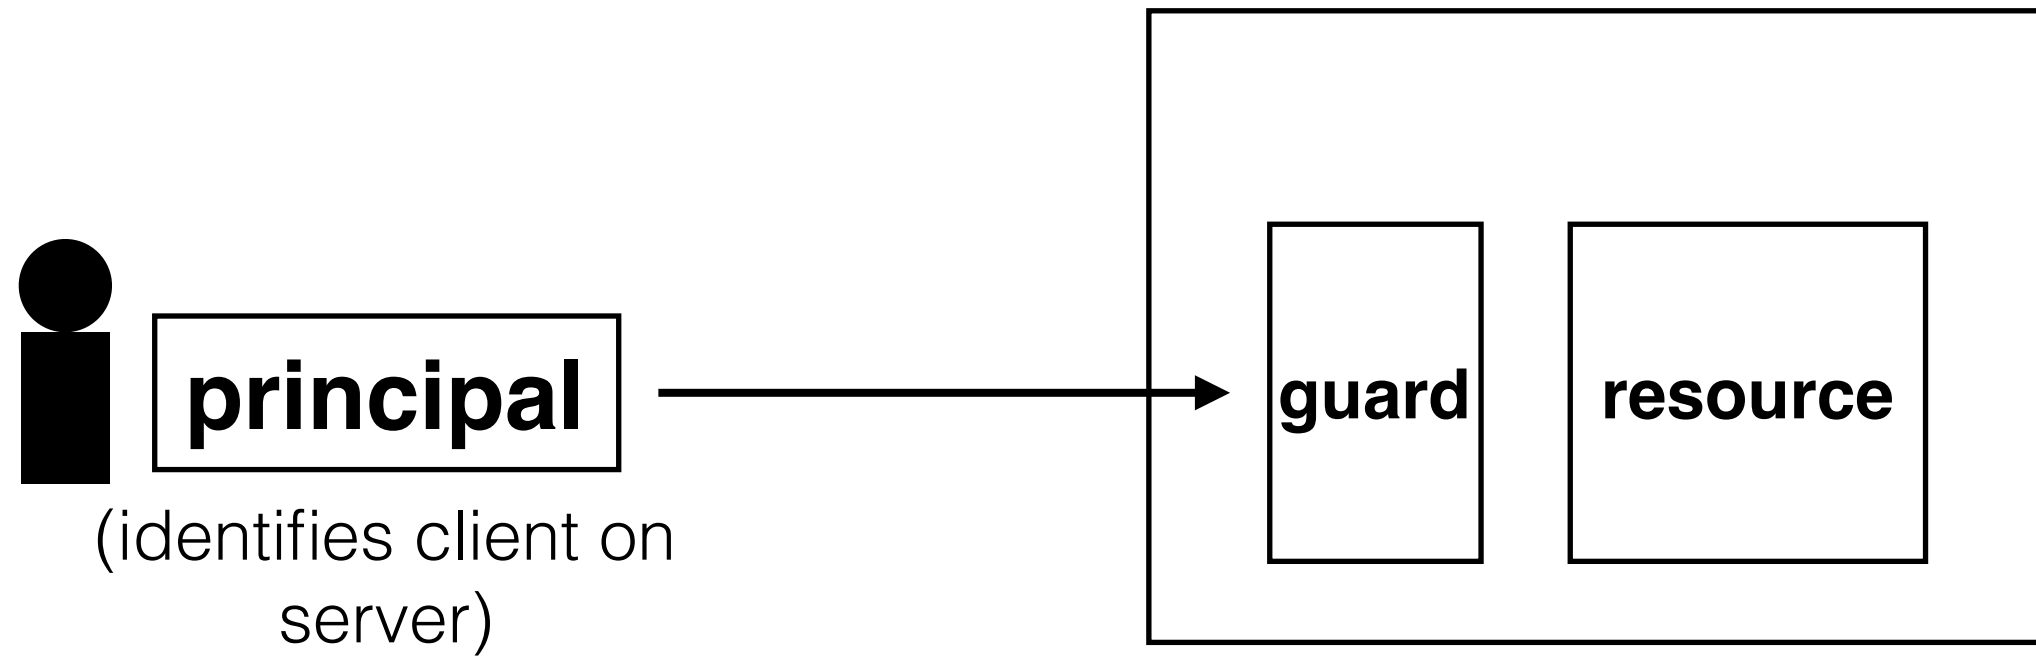
<http://web.mit.edu/subjectevaluation/>

**please fill them out — they provide
extremely valuable feedback to all
instructors**

6.033 Spring 2019

Lecture #24

- **Bitcoin and blockchains (and anonymity)**



**suppose we want a distributed (and decentralized)
public log that multiple users can read from/write to**

one motivation: digital currency

**challenge: what if I don't want other users to be able
to tie my transactions back to me personally?**

e.g., I don't want everyone in the world to know that Katrina
LaCurts just spent money buying 10,000 tiny hands

**suppose we want a distributed (and decentralized)
public log that multiple users can read from/write to**

one motivation: digital currency

users will be anonymous

identified only by their public keys

**challenge: how do we prevent an adversary from
editing our transactions?**

**suppose we want a distributed (and decentralized)
public log that multiple users can read from/write to**

one motivation: digital currency

users will be anonymous

identified only by their public keys

users will sign their transactions

prevents adversary from tampering with log data

**challenge: how do we prevent an adversary from
reordering/removing/duplicating transactions?**

**suppose we want a distributed (and decentralized)
public log that multiple users can read from/write to**

one motivation: digital currency

users will be anonymous

identified only by their public keys


users will sign their transactions

prevents adversary from tampering with log data

**users will include a hash of the previous
transaction as part of their signed data**


prevents adversary from re-ordering

suppose we want a **distributed** (and decentralized)
public log that multiple users can read from/write to
one motivation: digital currency




suppose we want a **distributed** (and decentralized)
public log that multiple users can read from/write to

one motivation: digital currency



suppose we want a **distributed** (and decentralized) public log that multiple users can read from/write to

one motivation: digital currency




in the world of digital currency,
 PK_V might do this to try to spend
the same “coin” in multiple
places

which log is correct?

we need **consensus**


note that a similar situation could occur
for completely innocuous reasons, not
just because of an attack from PK_V

goal: the system should come to a **consensus on what the correct log is, and ignore all other logs**




idea: PK_A broadcasts their transaction to the network, waits for a majority of users to confirm it

goal: the system should come to a **consensus on what the correct log is, and ignore all other logs**




idea: PK_A broadcasts their transaction to the network, waits for a majority of users to confirm it

goal: the system should come to a **consensus on what the correct log is, and ignore all other logs**




idea: PK_V broadcasts their transaction to the network, waits for a majority of users to confirm it

goal: the system should come to a **consensus on what the correct log is, and ignore all other logs**




problem: PK_v can create many identities (“sybils”) on the network, verify any transaction they want

goal: the system should come to a **consensus** on what the correct log is, and ignore all other logs



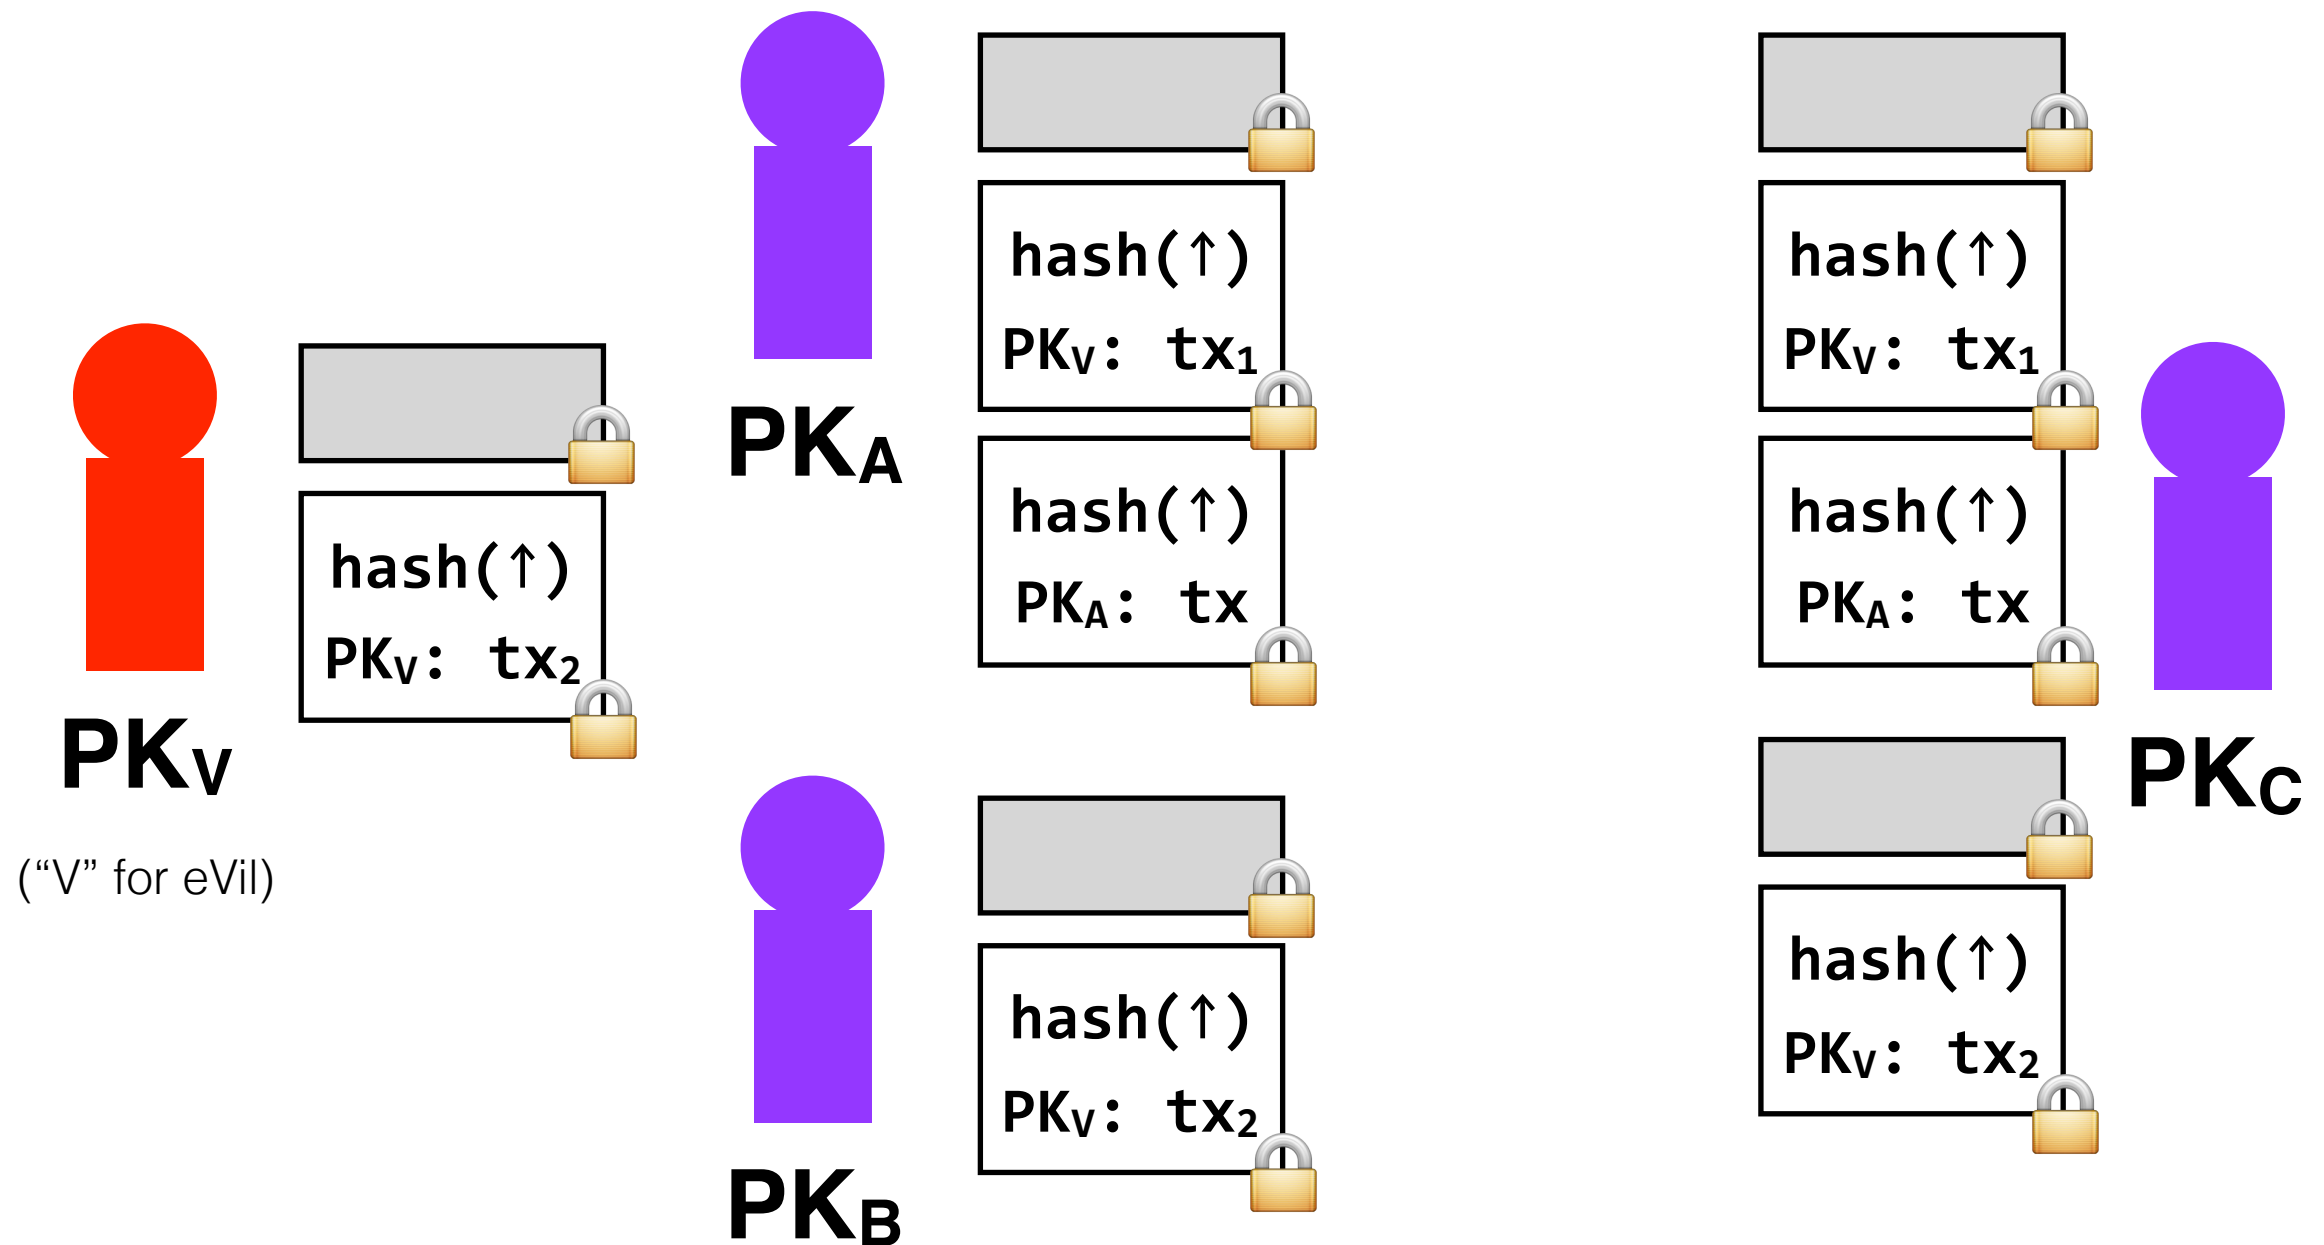
idea: if users receive different logs, they only append to the one that is longest, and don't distribute any others

goal: the system should come to a **consensus** on what the correct log is, and ignore all other logs



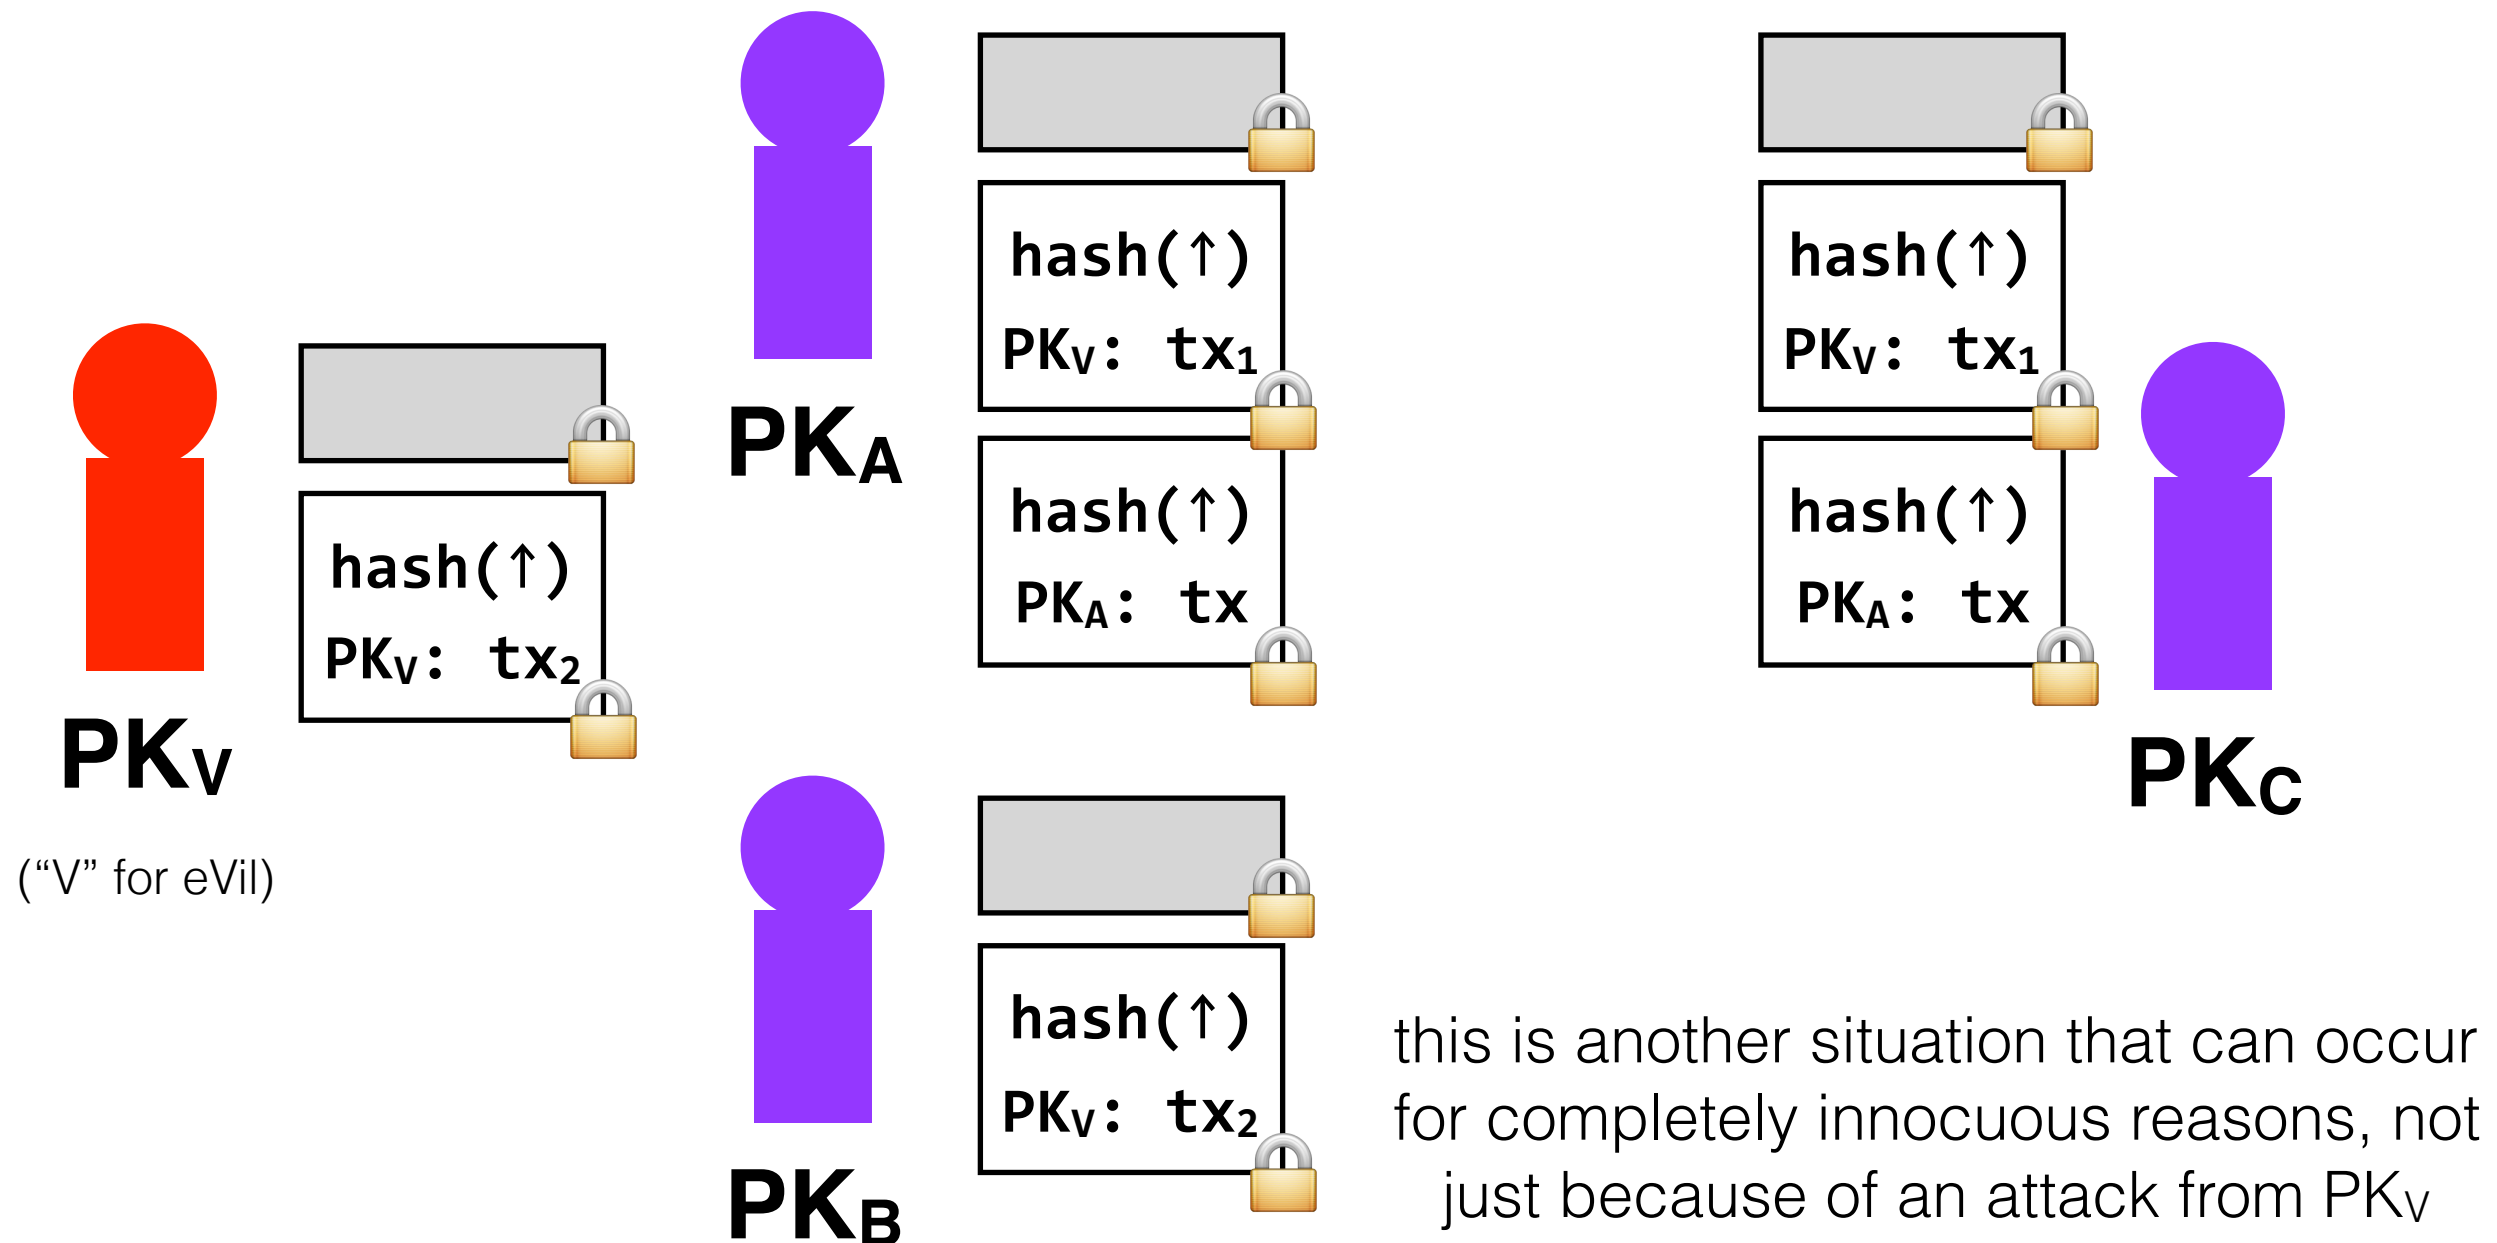
idea: if users receive different logs, they only append to the one that is longest, and don't distribute any others

goal: the system should come to a **consensus** on what the correct log is, and ignore all other logs




idea: if users receive different logs, they only append to the one that is longest, and don't distribute any others

goal: the system should come to a **consensus on what the correct log is, and ignore all other logs**



idea: if users receive different logs, they only append to the one that is longest, and don't distribute any others


goal: the system should come to a **consensus** on what the correct log is, and ignore all other logs



problem: at what point can a user decide that their transaction is confirmed?

e.g., how does PK_A know that their transaction is part of the “correct” log?


goal: the system should come to a **consensus on what the correct log is, and ignore all other logs**



if $N = 2$, PK_A 's transaction is confirmed at this point


idea: transaction t is confirmed when it's followed by N blocks in the log

goal: the system should come to a **consensus** on what the correct log is, and ignore all other logs




idea: transaction t is confirmed when it's followed by N blocks in the log

goal: the system should come to a **consensus on what the correct log is, and ignore all other logs**



problem: PK_V can create sybil and validate as many blocks as is necessary; it's *easy* to validate blocks


goal: the system should make it difficult to **validate blocks**
(validate = add a block to the log)



solution: proofs of work

goal: the system should make it difficult to **validate** blocks

(validate = add a block to the log)




now it is **difficult** for **PK_V** to validate all of the
necessary blocks, even with Sybils

(requires a lot of compute power, not just a lot of identities)

goal: the system should make it difficult to **validate** blocks

(validate = add a block to the log)

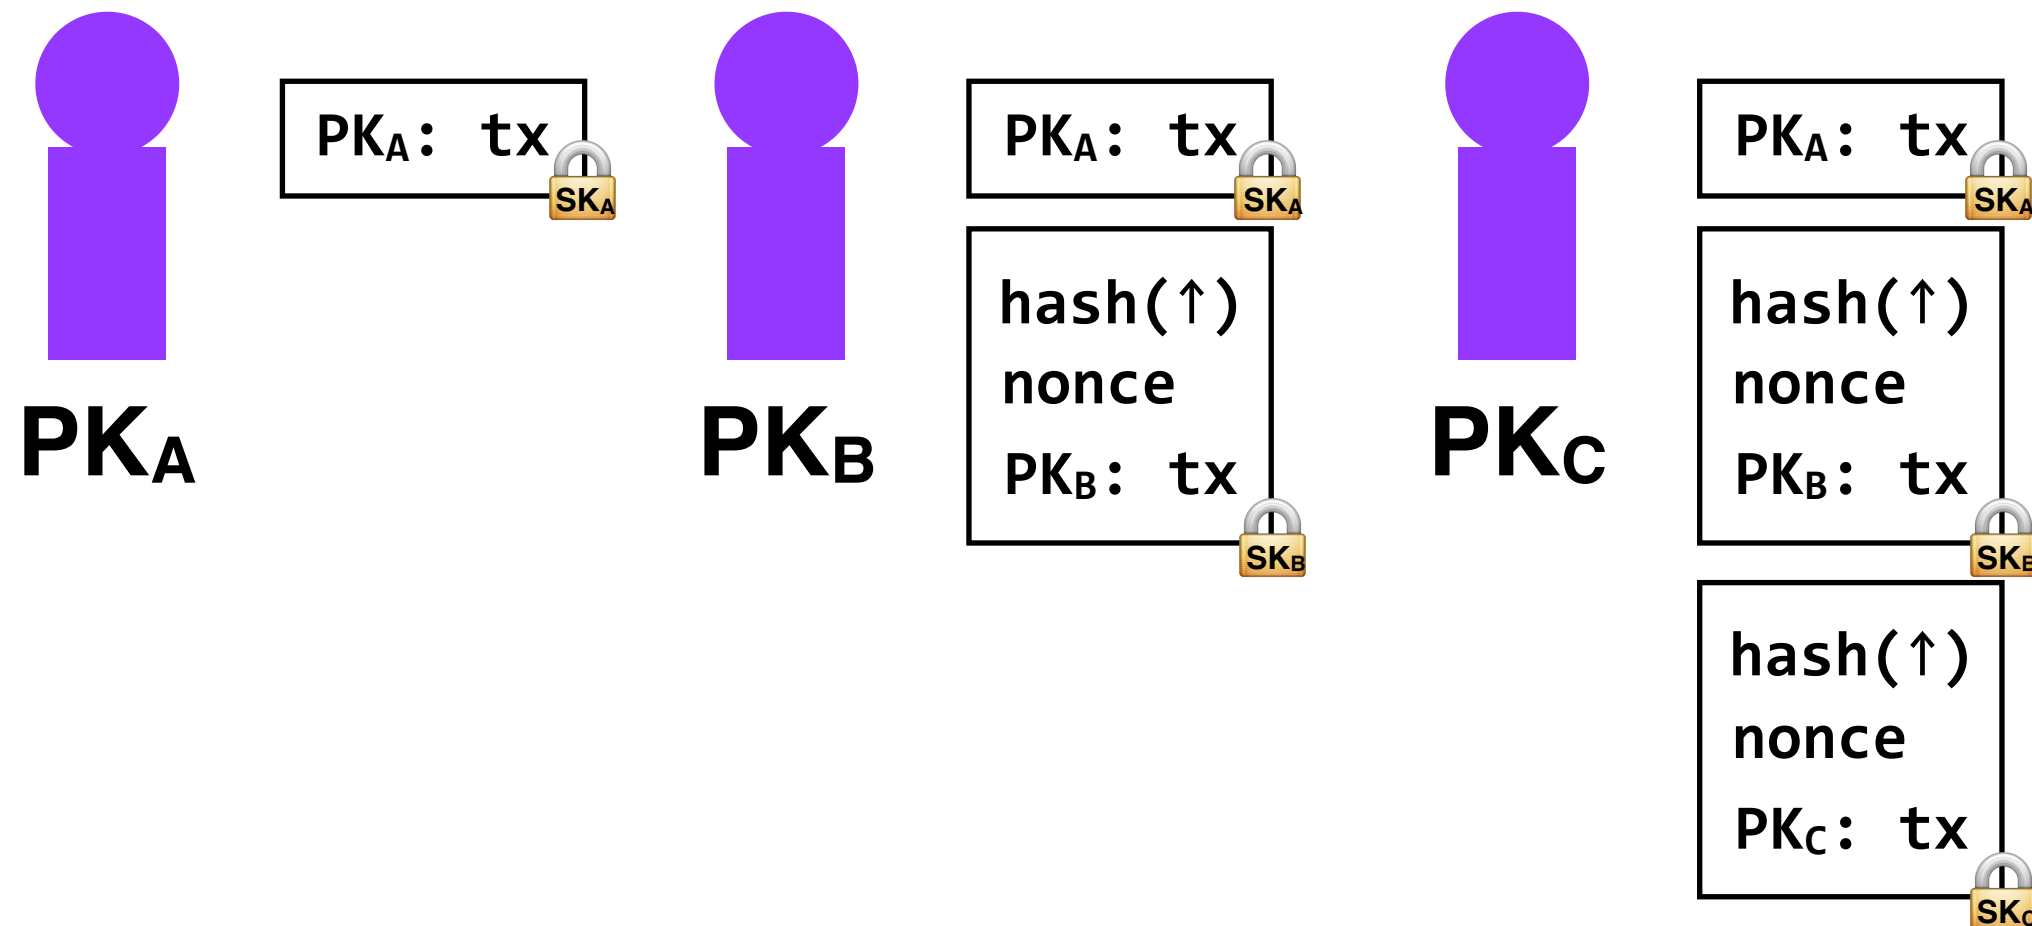


now it is **difficult** for **PK_V** to validate all of the **necessary blocks, even with Sybils**

(requires a lot of compute power, not just a lot of identities)

a **distributed** public log that multiple users can read from/write to

one motivation: digital currency

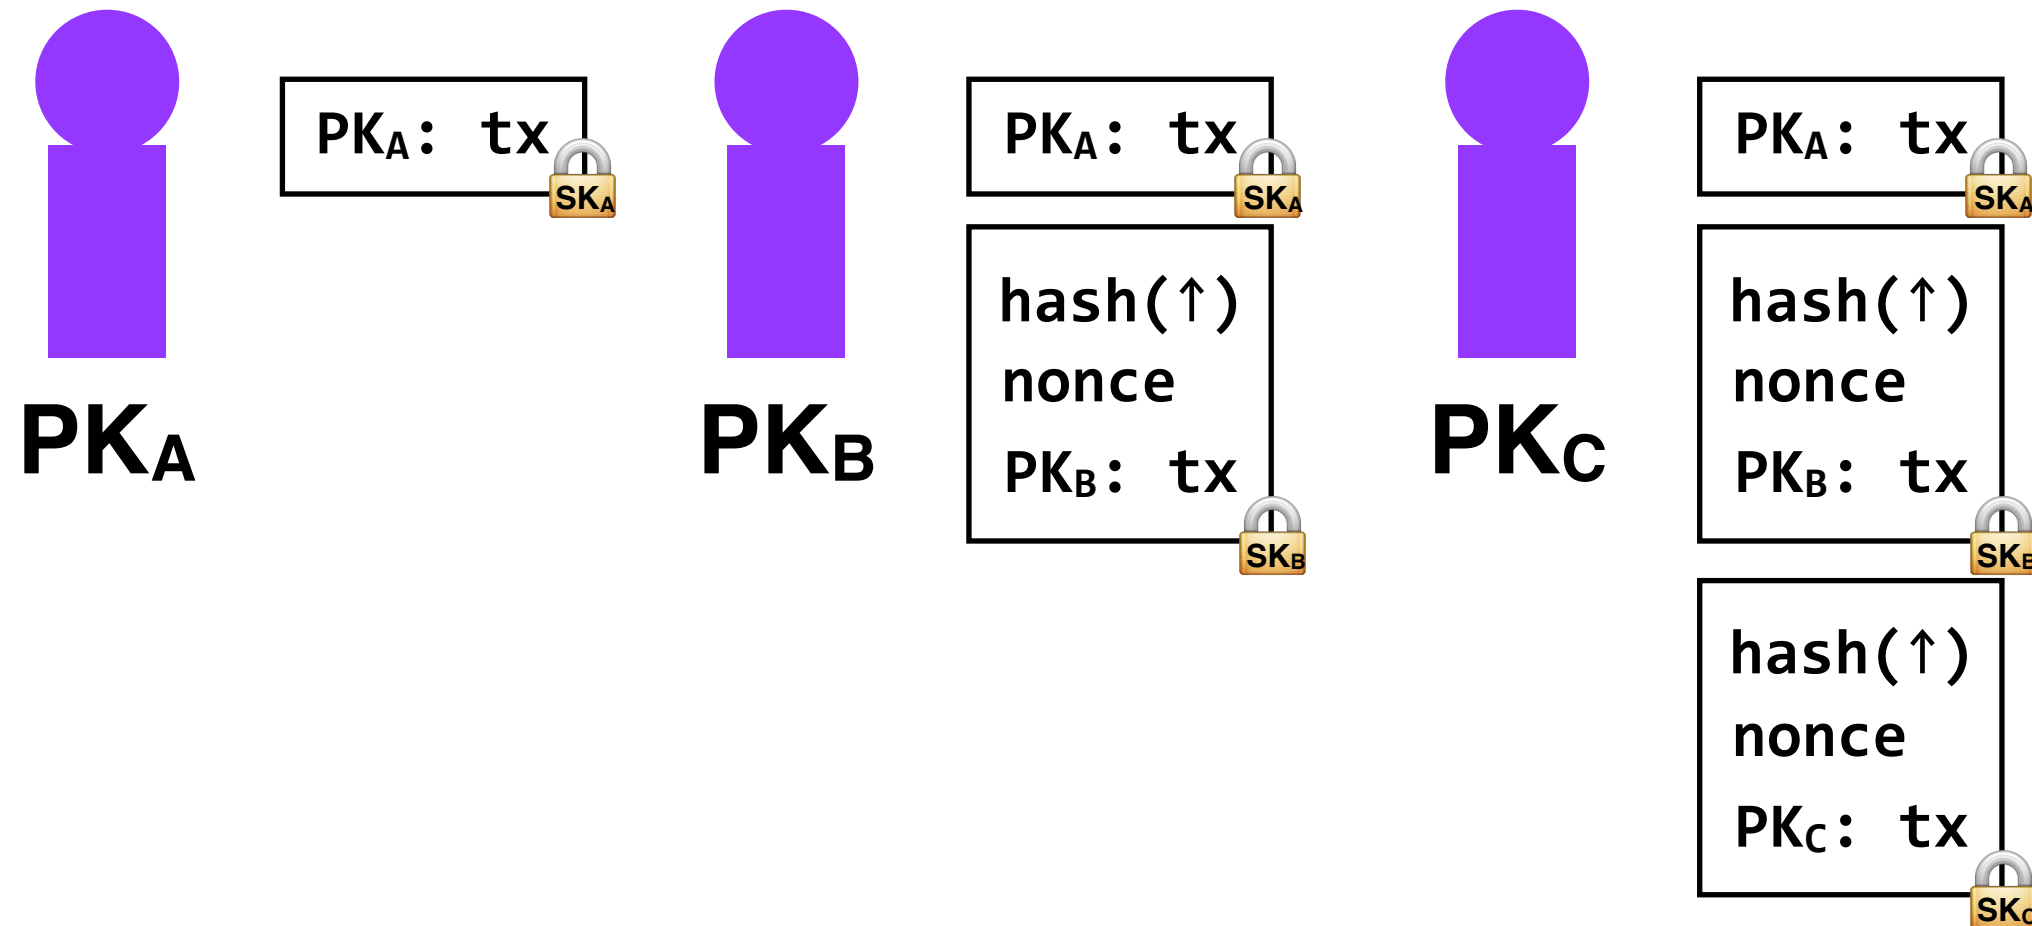


this log is the **blockchain**

(typically a block will contain multiple transactions, not one as we've been showing)

a **distributed** public log that multiple users can read from/write to

one motivation: digital currency




the process of validating transactions is known as **mining**

users receive payment (bitcoins) for validating blocks

a **distributed** public log that multiple users can read from/write to

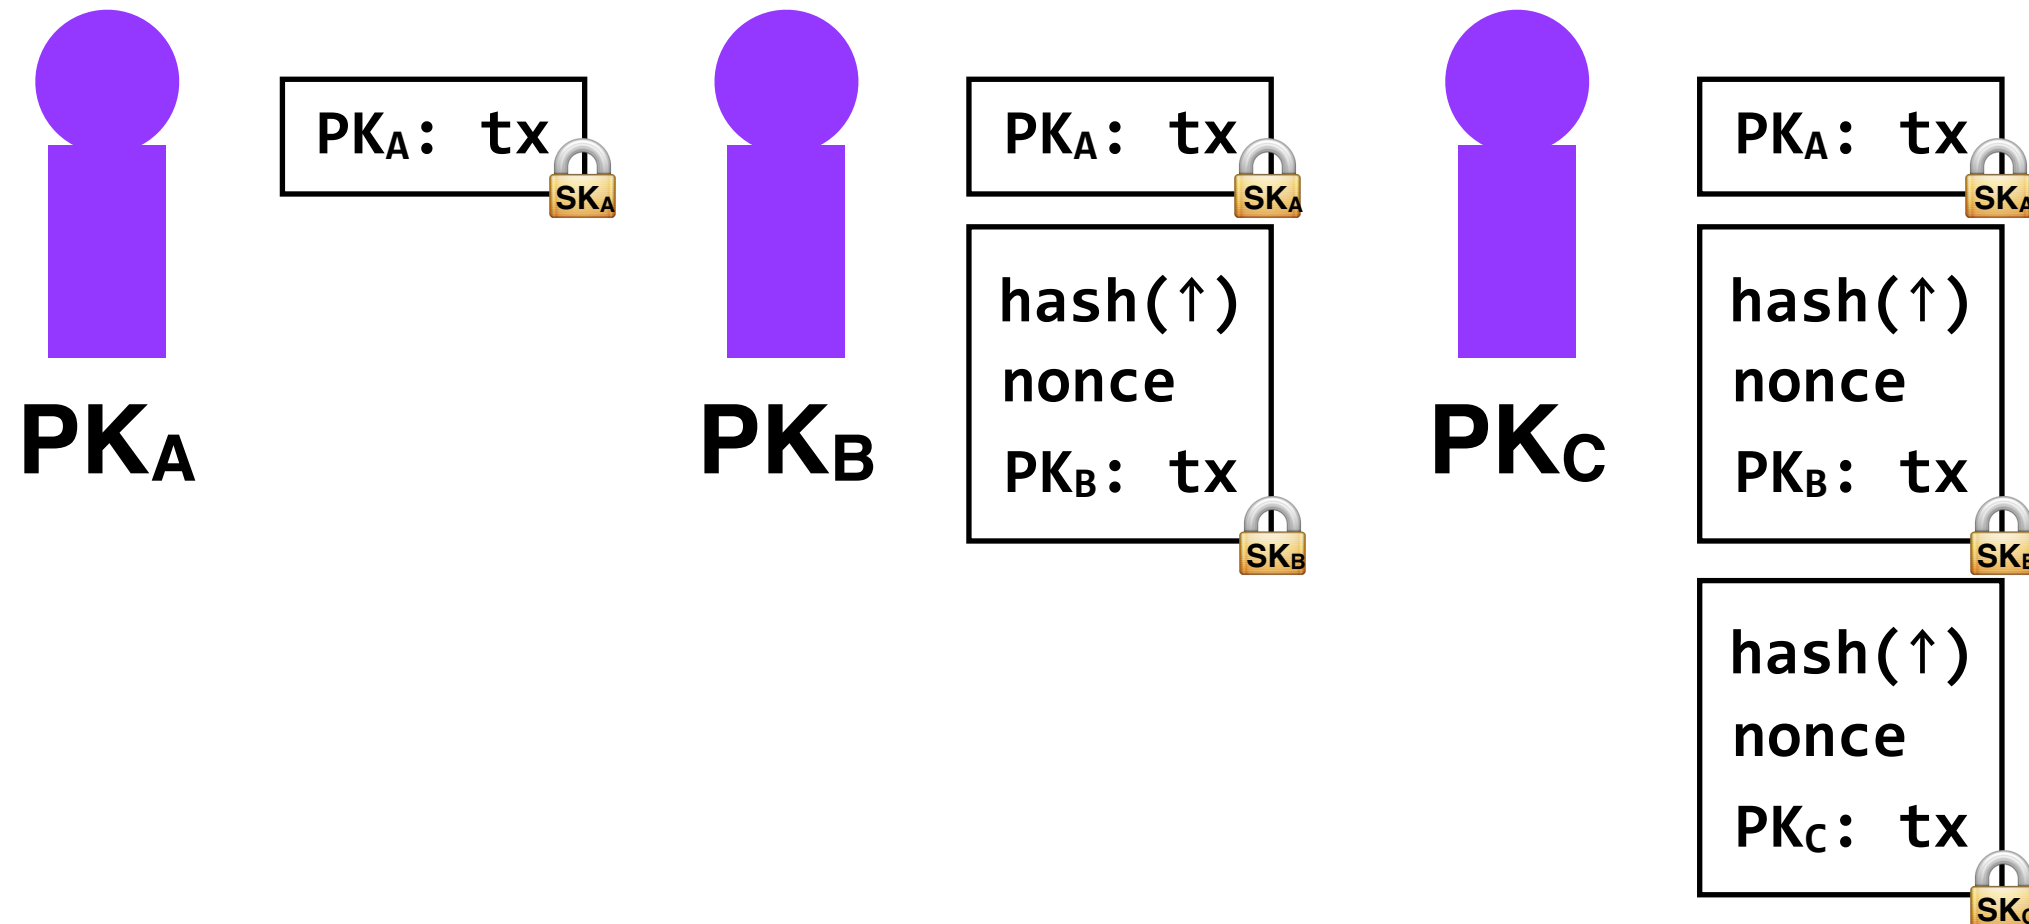
one motivation: digital currency



**if multiple users validate the same block at
(roughly) the same time, the blockchain will **fork****
the system resolves this problem since miners will only work on the longest fork

a **distributed** public log that multiple users can read from/write to

one motivation: digital currency



a block is confirmed after 5* valid blocks follow it
in the blockchain

* 5 is common, but this is a tunable parameter

- **Bitcoin** is a decentralized digital currency. Being decentralized means that there is no bank; in Bitcoin, everyone is the bank.
- Bitcoin provides a distributed public log called the **blockchain** that can be used for purposes other than digital currency. It uses **proofs-of-work** to prevent Sybil Attacks, since strong identities won't work.
- In theory, users of Bitcoin are **anonymous**; in practice, it's not clear how true that is.