

# **CSE 120**

# **Principles of Operating Systems**

**Spring 2018**

**Lecture 13: File System Implementation**

Geoffrey M. Voelker

# File Systems


---

- First we'll discuss properties of physical disks
  - ♦ Structure
  - ♦ Performance
- Then how file systems support users and programs
  - ♦ Files
  - ♦ Directories
  - ♦ Sharing
  - ♦ Protection
- End with how file systems are implemented
  - ♦ File System Data Structures
  - ♦ File Buffer Cache
  - ♦ Read Ahead

# File System Layout

---

- We start with an empty disk



- Goal for the file system is to manage the disk space to implement the file and directory abstractions that are so convenient for programs and users

# Key Questions

---

- How do we keep track of blocks used by a file?
- Where do we store metadata information?
- How do we (really) do path name translation?
- How do we implement common file operations?
- How can we cache data to improve performance?
- Our discussion will be Unix-oriented
  - ♦ Other file systems face same challenges, with analogous approaches and data structures for solving them

# File System Layout

---

How do file systems use the disk to store files?

- File systems define a block size (e.g., 4KB)
  - ♦ Disk space is allocated in granularity of blocks
- A “Master Block” determines location of root directory
  - ♦ Always at a well-known disk location
  - ♦ Often replicated across disk for reliability
- A free map determines which blocks are free, allocated
  - ♦ Usually a bitmap, one bit per block on the disk
  - ♦ Also stored on disk, cached in memory for performance
- Remaining disk blocks used to store files (and dirs)
  - ♦ There are many ways to do this

# File System Layout

---


How do file systems use the disk to store files?

- File systems define a block size (e.g., 4KB)
  - ♦ Disk space is allocated in granularity of blocks
- A “Master Block” determines location of root directory
  - ♦ Always at a well-known disk location
  - ♦ Often replicated across disk for reliability
- A free map determines which blocks are free, allocated
  - ♦ Usually a bitmap, one bit per block on the disk
  - ♦ Also stored on disk, cached in memory for performance
- Remaining disk blocks used to store files (and dirs)
  - ♦ There are many ways to do this

# File System Layout

---

- Partition it into fixed-size file system blocks



- Typically 4KB in size
  - ♦ Block size set when file system is formatted
- Independent of disk physical sector size
  - ♦ Sector 512 bytes, file system will use 8 sectors/block

# File System Layout

---


- Files span multiple disk blocks
  - ♦  $2\text{MB file uses } 2 \times 1024 \times 1024 / 4096 = 512 \text{ blocks (4KB block size)}$
- A small file still uses an entire block
  - ♦ A file of size 4001 bytes uses one block
  - ♦ What kind of fragmentation is this, internal or external?
- Challenge: How do we keep track of all of the blocks used by one file?



# Contiguous Layout

---


- Can layout file blocks contiguously



# Contiguous Layout

---


- Can layout file blocks contiguously



# Contiguous Layout

---


- Can layout file blocks contiguously



# Contiguous Layout

---

- Simple to keep track of where a file's blocks are




- Directory stores a pointer to the first block
  - ♦ All others are a simple offset from the first
  - ♦ Makes random access also straightforward
- Enables fast sequential access to disk for reads/writes
- But there are multiple disadvantages

# Contiguous Layout

---

- Difficult to grow a file once it has been written




- If the blue or orange files need to grow, we're stuck

# Contiguous Layout

---

- As files are created and deleted, gaps will occur




- If we need to store a file using 8 blocks, we're stuck
  - ♦ What kind of fragmentation is this, internal or external?
  - ♦ What would be one method for rearranging?

# Linked Layout


---

- Need to store a file with 8 blocks into the “gaps”



# Linked Layout

- Another option is to link each block to the next
  - ♦ Essentially a linked list on disk for each file




- Directory still just stores pointer to first block of file
- Fragmentation no longer a problem, can fill in gaps
- Random access now expensive
  - ♦ Need to traverse pointers to access a random block
  - ♦ Potentially many disk reads just to get to desired block



# Indexed Layout

- Indexed layouts use a special block (index block) ■ to store pointers to the data blocks ■



- Directory points to the index block
- Still solves fragmentation problem (can fill in gaps)
- Also solves random access problem
  - After reading the index block, know the locations of all blocks
- For large files, need multiple index blocks

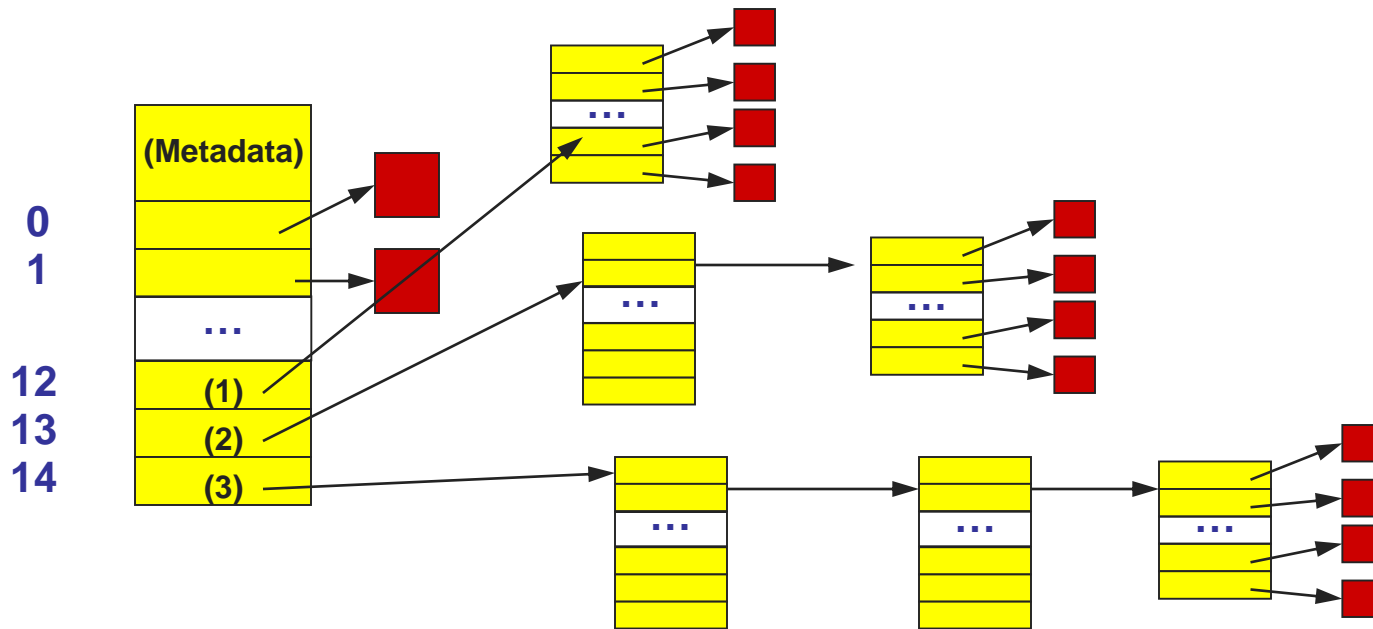
# Disk Layout Summary

---

- Files span multiple disk blocks
- How do you find all of the blocks for a file?
  1. **Contiguous allocation**
    - » Like memory
    - » Fast, simplifies directory access
    - » Inflexible, causes fragmentation, needs compaction
  2. **Linked structure**
    - » Each block points to the next, directory points to the first
    - » Good for sequential access, bad for all others
  3. **Indexed structure (indirection, hierarchy)**
    - » An “index block” contains pointers to many other blocks
    - » Handles random better, still good for sequential
    - » May need multiple index blocks (linked together)

# Unix Inodes

- Unix inodes implement an indexed structure for files
- Each inode contains 15 block pointers
  - ♦ First 12 are direct blocks (e.g., 4 KB blocks)
  - ♦ Then single, double, and triple indirect



# File Metadata

---

- Unix inodes also store all of the metadata for a file
- File size
  - ♦ In bytes (actual file size) and blocks (data blocks allocated)
- User & group of file owner
- Protection bits
  - ♦ User/group/other, read/write/execute
- Reference count
  - ♦ How many directory entries point to this inode
- Timestamps
  - ♦ Created, modified, last accessed, any change
- “ls -l” reads this info from the inode (syscall: stat())


# File System Layout


---



How do file systems use the disk to store files?

- File systems define a block size (e.g., 4KB)
  - ♦ Disk space is allocated in granularity of blocks
- A “Master Block” determines location of root directory
  - ♦ Always at a well-known disk location
  - ♦ Often replicated across disk for reliability
- A free map determines which blocks are free, allocated
  - ♦ Usually a bitmap, one bit per block on the disk
  - ♦ Also stored on disk, cached in memory for performance
- Remaining disk blocks used to store files (and dirs)
  - ♦ There are many ways to do this

# Master Block (Superblock)

- “/” is the directory that is the root of the file system
- How do we find the inode  for “/”?



- The master block  stores a pointer to the inode of “/”
  - ♦ The inode for “/” has pointers to all of the blocks  storing the directory entries for “/”
- It is the basis for translating all path names
- It is at a fixed, pre-defined location on disk
  - ♦ Replicated deterministically across the FS for redundancy

# File System Layout


---

How do file systems use the disk to store files?

- File systems define a block size (e.g., 4KB)
  - ♦ Disk space is allocated in granularity of blocks
- A “Master Block” determines location of root directory
  - ♦ Always at a well-known disk location
  - ♦ Often replicated across disk for reliability
- A free map determines which blocks are free, allocated
  - ♦ Usually a bitmap, one bit per block on the disk
  - ♦ Also stored on disk, cached in memory for performance
- Remaining disk blocks used to store files (and dirs)
  - ♦ There are many ways to do this

# Block Allocation

- The file system needs to keep track of which blocks have been allocated and which are free




- Free map blocks B store a bitmap, one bit per block
  - ♦ Bit is set → block is allocated
  - ♦ One bitmap for data blocks
  - ♦ Another for inode blocks



# Types of Blocks on Disk

- Four basic kinds of blocks on disk
  - ♦ Only data blocks store file data and directory data



- Master block M
- Bitmap blocks B
- Inode blocks
- Data blocks 0 1 2

# Unix Inodes != Directories

---

- Unix inodes are **not** directories
  - ♦ Inodes describe where on the disk the blocks of a file are
- Directories are files, so inodes also describe where the blocks for directories are placed on the disk
  - ♦ Every directory and file on disk has an associated inode for it

# Path Name Translation (v2)

---

- Directory entries map file names to inodes
  - ♦ To open “/one”, use master block to find inode for “/” on disk
  - ♦ Open “/”, look for entry for “one”
  - ♦ This entry gives the disk block number for the inode for “one”
  - ♦ Read the inode for “one” into memory
  - ♦ The inode says where first data block is on disk
  - ♦ Read that block into memory to access the data in the file

# Symbolic (Soft) Links

---

- It is convenient to be able to create **aliases** in the FS
  - ♦ Have a file be referred to by multiple names
- Soft links are the most familiar form in Unix
  - `% ln -s file alias (ln -s /a/b/c /tmp/softlink)`
  - ♦ Syscall: `symlink()`
- Soft links create aliases via path name translation
  - ♦ Path name translation starts again when hitting a soft link
  - ♦ `/tmp/softlink` → `/a/b/c`
- Implemented simply by **storing the alias as a string in a file** and marking the inode as a soft link
  - ♦ FS reads the path alias from the file and restarts translation

# Hard Links

---

- Hard links are another form of aliasing
  - ◊ *% In file alias* (ln /a/b/c /tmp/hardlink)
  - ◊ Syscall: link()
- Hard links create aliases via inode pointers in dirs
  - ◊ Recall that a directory entry maps a name to an inode
  - ◊ Creating a hard link adds another directory entry mapping the new name to **the same inode** as the old name
  - ◊ It adds a new pointer, or link, to the inode
  - ◊ Reference count in the inode is also incremented
- The “.” and “..” names are hard links to directories
  - ◊ /a/b/c and /a/b/c/. point to the same inode

# Create

---

- Creating a file “new” is relatively straightforward
- **Allocate an inode**
  - ♦ Initialize the metadata (owner, protection, timestamps, ...)
  - ♦ Update inode bitmap
- **Allocate a directory entry** in the directory for the file
  - ♦ Entry maps “new” to the inode allocated for “new”
- When process starts writing to file, **allocate data blocks**
  - ♦ Update inode to point to data blocks allocated
  - ♦ Update data block bitmap
  - ♦ Continue to allocate data blocks on demand
    - » Preallocating blocks in “extents” helps keep blocks contiguous

# Rename

---

- One way to rename a file is to simply create a new one with the new name, copy the contents, and delete the old file
  - ♦ Method used in original version of Unix (test/mv.c in Nachos)
- More efficient to implement in FS
  - % mv old new
  - ♦ Syscall: rename()
- Rename creates a new directory entry with the **new name that points (links) to the same inode** as the old
  - ♦ Then it deletes the entry directory for the old name
  - ♦ Only directories are modified, file and inode stay the same

# Delete

---

- Deleting a file has a few steps
  - ♦ Remove the directory entry for the name being deleted
    - » Hence the syscall name unlink()
  - ♦ Decrement the reference count in the inode
  - ♦ If the file still has links to it, nothing else happens
- If there are no remaining links
  - ♦ Free up the data blocks (update the data block bitmap)
  - ♦ Free up the inode blocks (update the inode bitmap)
  - ♦ Block data is not erased
- If the file is still open in any process, the directory entry is removed but the file blocks are not
  - ♦ Until the last process with the file open finally closes it



# Partitions

---

- What if we want multiple file systems on one disk?



# Partitions

---

- What if we want multiple file systems on one disk?



- Split up physical disk into multiple **partitions**
- Each partition has an entire file system inside of it
  - ♦ Master block, bitmaps, etc.
- How do we link them together into one name space?

# Mounting File Systems

---

- Mounting is the mechanism used to piece together multiple file systems into a single global name space
- One file system is mounted as “/” (root)
- Other file systems attached at **mount points**
  - ♦ An empty directory in file system, e.g., /home
  - ♦ Mounting the “home” file system attaches the root for “home” to /home in the name space
  - ♦ Opening “/home/voelker/file” starts path name translation in the “root” file system, continues in the “home” file system when crossing the mount point
- Mostly invisible to users and processes
  - ♦ Some exceptions (e.g., cannot hard link across file systems)

# File Buffer Cache

---

- Applications exhibit significant locality for reading and writing files
- Idea: Cache file blocks in memory to capture locality
  - ♦ Called the **file buffer cache**
  - ♦ Cache is system wide, used and shared by all processes
  - ♦ Reading from the cache makes a disk perform like memory
  - ♦ Even a small cache can be very effective
- Issues
  - ♦ The file buffer cache competes with VM
    - » **Tradeoff: More physical memory for file cache, less for VM**
  - ♦ Like VM, it has limited size
  - ♦ Need replacement algorithms again (**LRU usually used**)

# Caching Writes

---

- On a write, some applications assume that data makes it through the buffer cache and onto the disk
  - ♦ As a result, writes are often slow even with caching
- OSes typically do write back caching
  - ♦ Maintain a queue of uncommitted blocks
  - ♦ Periodically flush the queue to disk (30 second threshold)
  - ♦ If blocks changed many times in 30 secs, only need one I/O
  - ♦ If blocks deleted before 30 secs (e.g., /tmp), no I/Os needed
- Unreliable, but practical
  - ♦ On a crash, all writes within last 30 secs are lost
  - ♦ Modern OSes do this by default; too slow otherwise
  - ♦ System calls (Unix: fsync) enable apps to force data to disk

# Read Ahead

---

- Many file systems implement “read ahead”
  - ♦ FS predicts that the process will request next block
  - ♦ FS goes ahead and requests it from the disk
  - ♦ Can happen while the process is computing on previous block
    - » Overlap I/O with execution
  - ♦ When the process requests block, it will be in cache
  - ♦ Compliments the disk cache, which also is doing read ahead
- For sequentially accessed files can be a big win
  - ♦ Unless blocks for the file are scattered across the disk
  - ♦ File systems try to prevent that, though (during allocation)

# Next time...

---

- Read Chapters 37, 39, 40