



CS 423

Operating System Design: The Kernel Abstraction

Professor Adam Bates
Fall 2018

Goals for Today



- Learning Objectives:
 - Understand the Kernel/Process Abstraction
 - Gain a working knowledge of Mode/Context switches
- Announcements:
 - C4 template available on Compass! **Due Jan 26**
 - MP0 is available on Compass! **Due Jan 29**
 - HW0 is available on Compass! **Due Jan 29**



Reminder: Please put away devices at the start of class



Process concept

- A process is the OS abstraction for executing a program with limited privileges

Dual-mode operation: user vs. kernel

- Kernel-mode: execute with complete privileges
- User-mode: execute with fewer privileges

Safe control transfer

- How do we switch from one mode to the other?



Process: an **instance** of a program that runs with limited rights on the machine

- Thread: a sequence of instructions within a process
 - Potentially many threads per process (for now, assume 1:1)
- Address space: set of rights of a process
 - Memory that the process can access
 - Other permissions the process has (e.g., which system calls it can make, what files it can access)



How can we permit a process to execute with only limited privileges?

Thought Experiment



How can we implement execution with limited privilege?

- Execute each program instruction in a simulator
- If the instruction is permitted, do the instruction
- Otherwise, stop the process
- Basic model in Javascript and other interpreted languages

Ok... but how do we go faster?

- Run the unprivileged code directly on the CPU!

Thought Experiment



How can we implement execution with limited privilege?

- Execute each program instruction in a simulator
- If the instruction is permitted, do the instruction
- Otherwise, stop the process
- Basic model in Javascript and other interpreted languages

Ok... but how do we go faster?

Thought Experiment



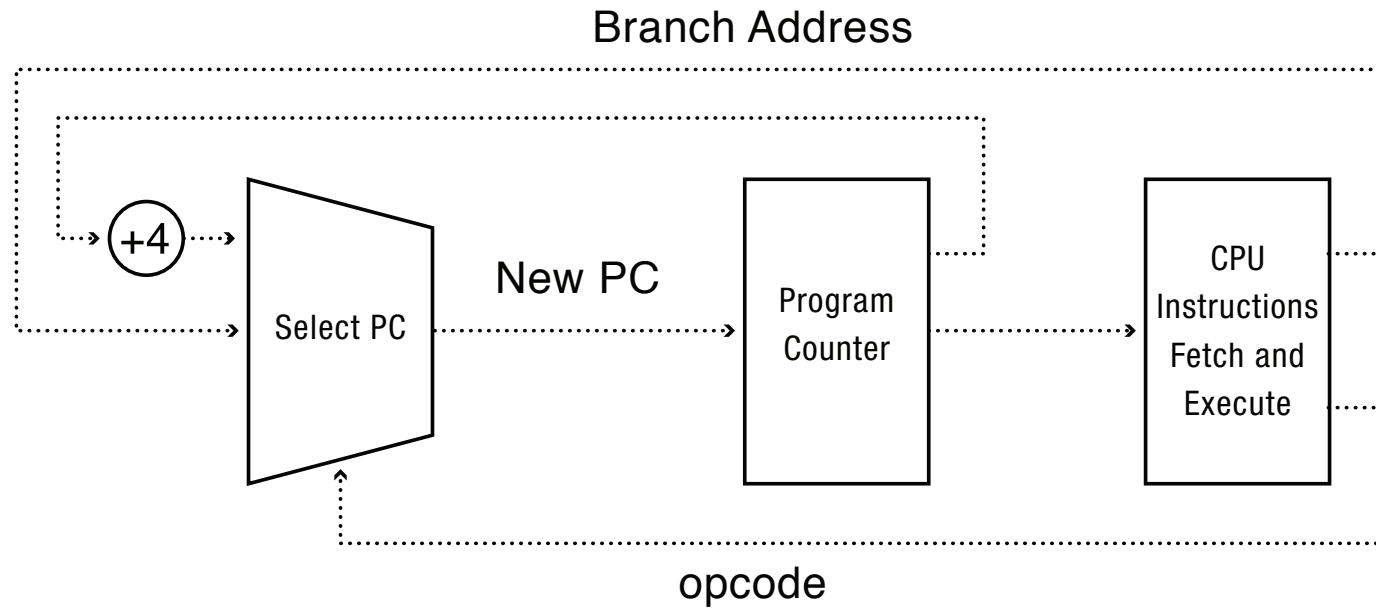
How can we implement execution with limited privilege?

- Execute each program instruction in a simulator
- If the instruction is permitted, do the instruction
- Otherwise, stop the process
- Basic model in Javascript and other interpreted languages

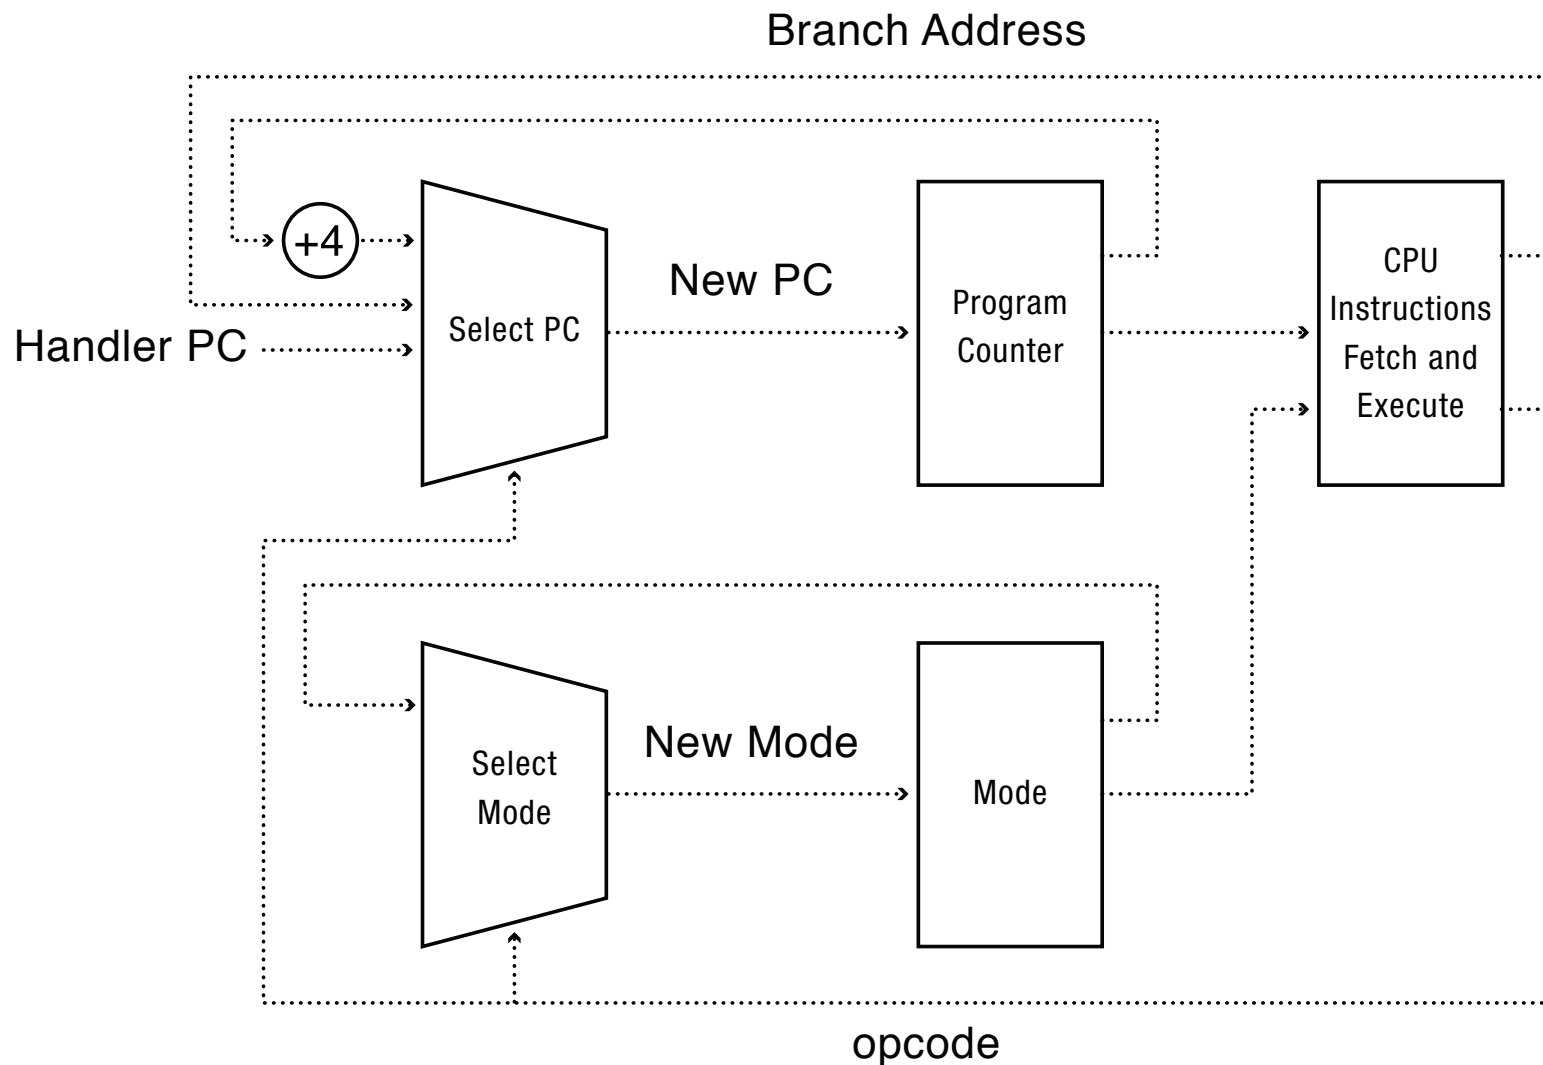
Ok... but how do we go faster?

- Run the unprivileged code directly on the CPU!

A Model of a CPU



A CPU with Dual-Mode Operation





Privileged instructions

- Available to kernel
- Not available to user code

Limits on memory accesses

- To prevent user code from overwriting the kernel

Timer

- To regain control from a user program in a loop

Safe way to switch from user mode to kernel mode,
and vice versa

Privileged Instructions



Examples?

What should happen if a user program attempts to execute a privileged instruction?

User->Kernel Switches



How/when do we switch from user to kernel mode?

1. Interrupts
 - Triggered by timer and I/O devices
2. Exceptions
 - Triggered by unexpected program behavior
 - Or malicious behavior!
3. System calls (aka protected procedure call)
 - Request by program for kernel to do some operation on its behalf
 - Only limited # of very carefully coded entry points



**How does the OS know
when a process is in an
infinite loop?**

Hardware Timer



Hardware device that periodically interrupts the processor

- Returns control to the kernel handler
- Interrupt frequency set by the kernel
Not by user code!
- Interrupts can be temporarily deferred
Not by user code!
Interrupt deferral crucial for implementing mutual exclusion

Kernel->User Switches



How/when do we switch from kernel to user mode?

1. New process/new thread start
 - Jump to first instruction in program/thread
2. Return from interrupt, exception, system call
 - Resume suspended execution (return to PC)
3. Process/thread context switch
 - Resume some other process (return to PC)
4. User-level upcall (UNIX signal)
 - Asynchronous notification to user program

What's a 'real' CPU?

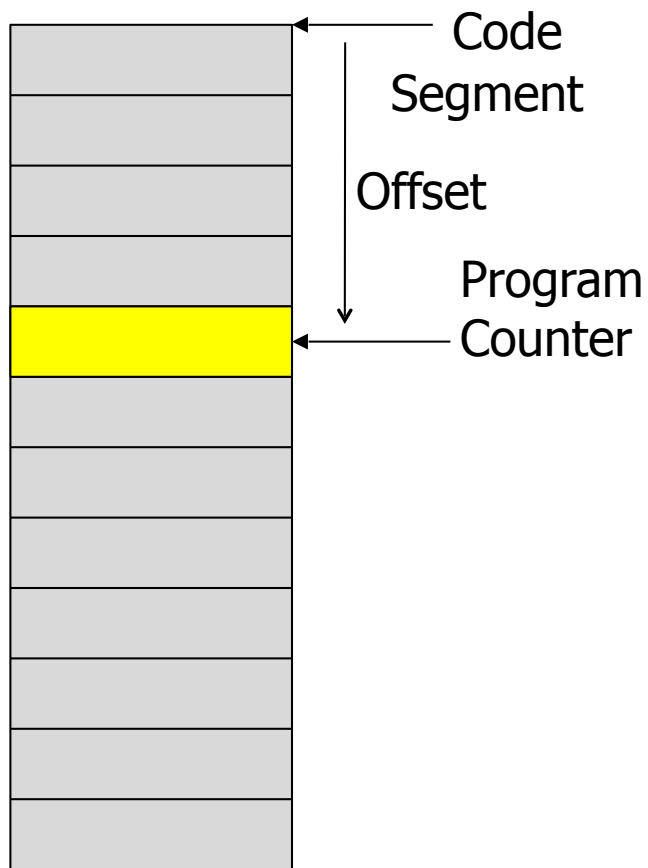


What is the “real” CPU comprised of at any given moment?

What's a 'real' CPU?



What is the "real" CPU comprised of at any given moment?

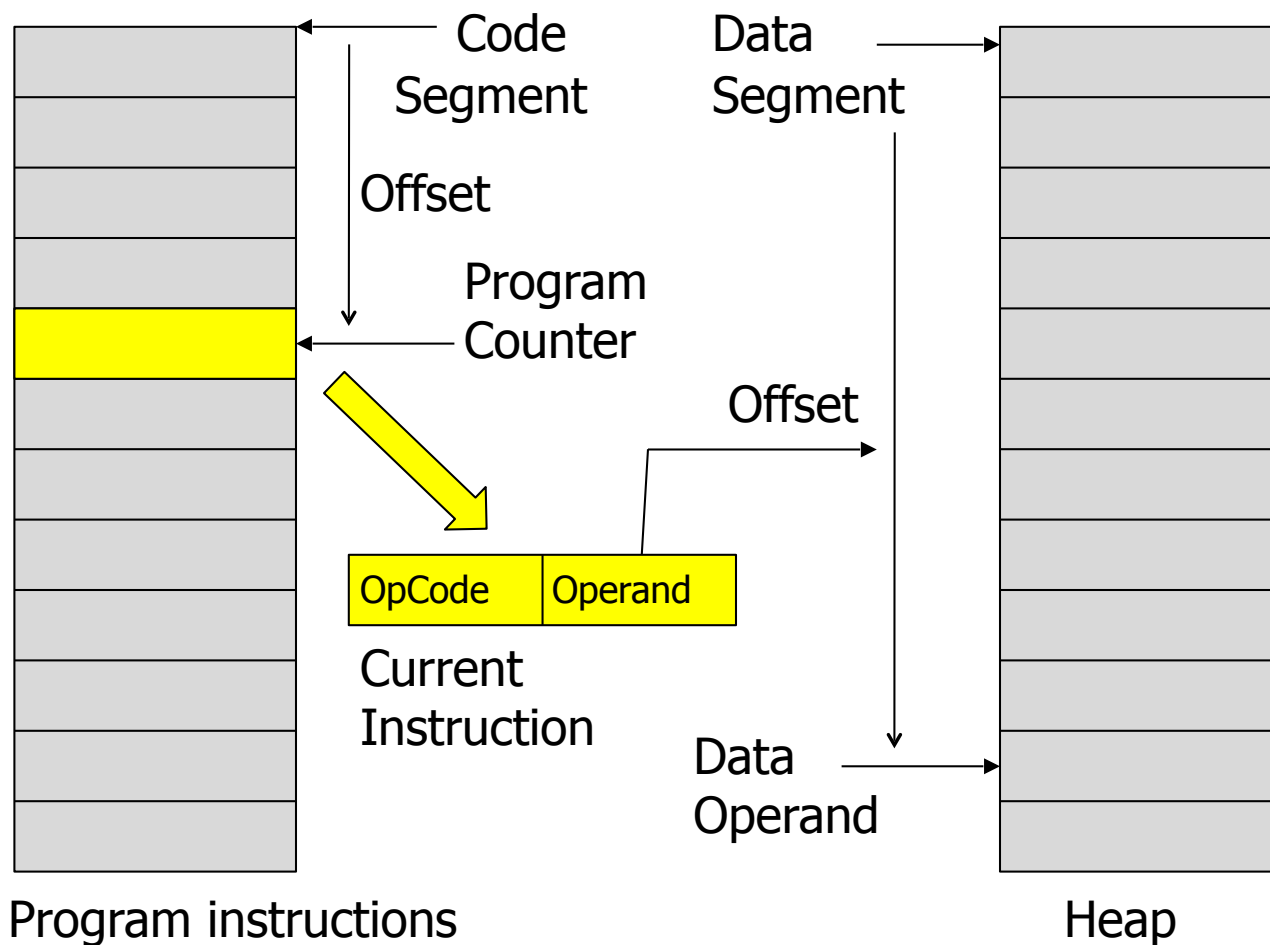


Program instructions

What's a 'real' CPU?




What is the “real” CPU comprised of at any given moment?



What's a 'real' CPU?



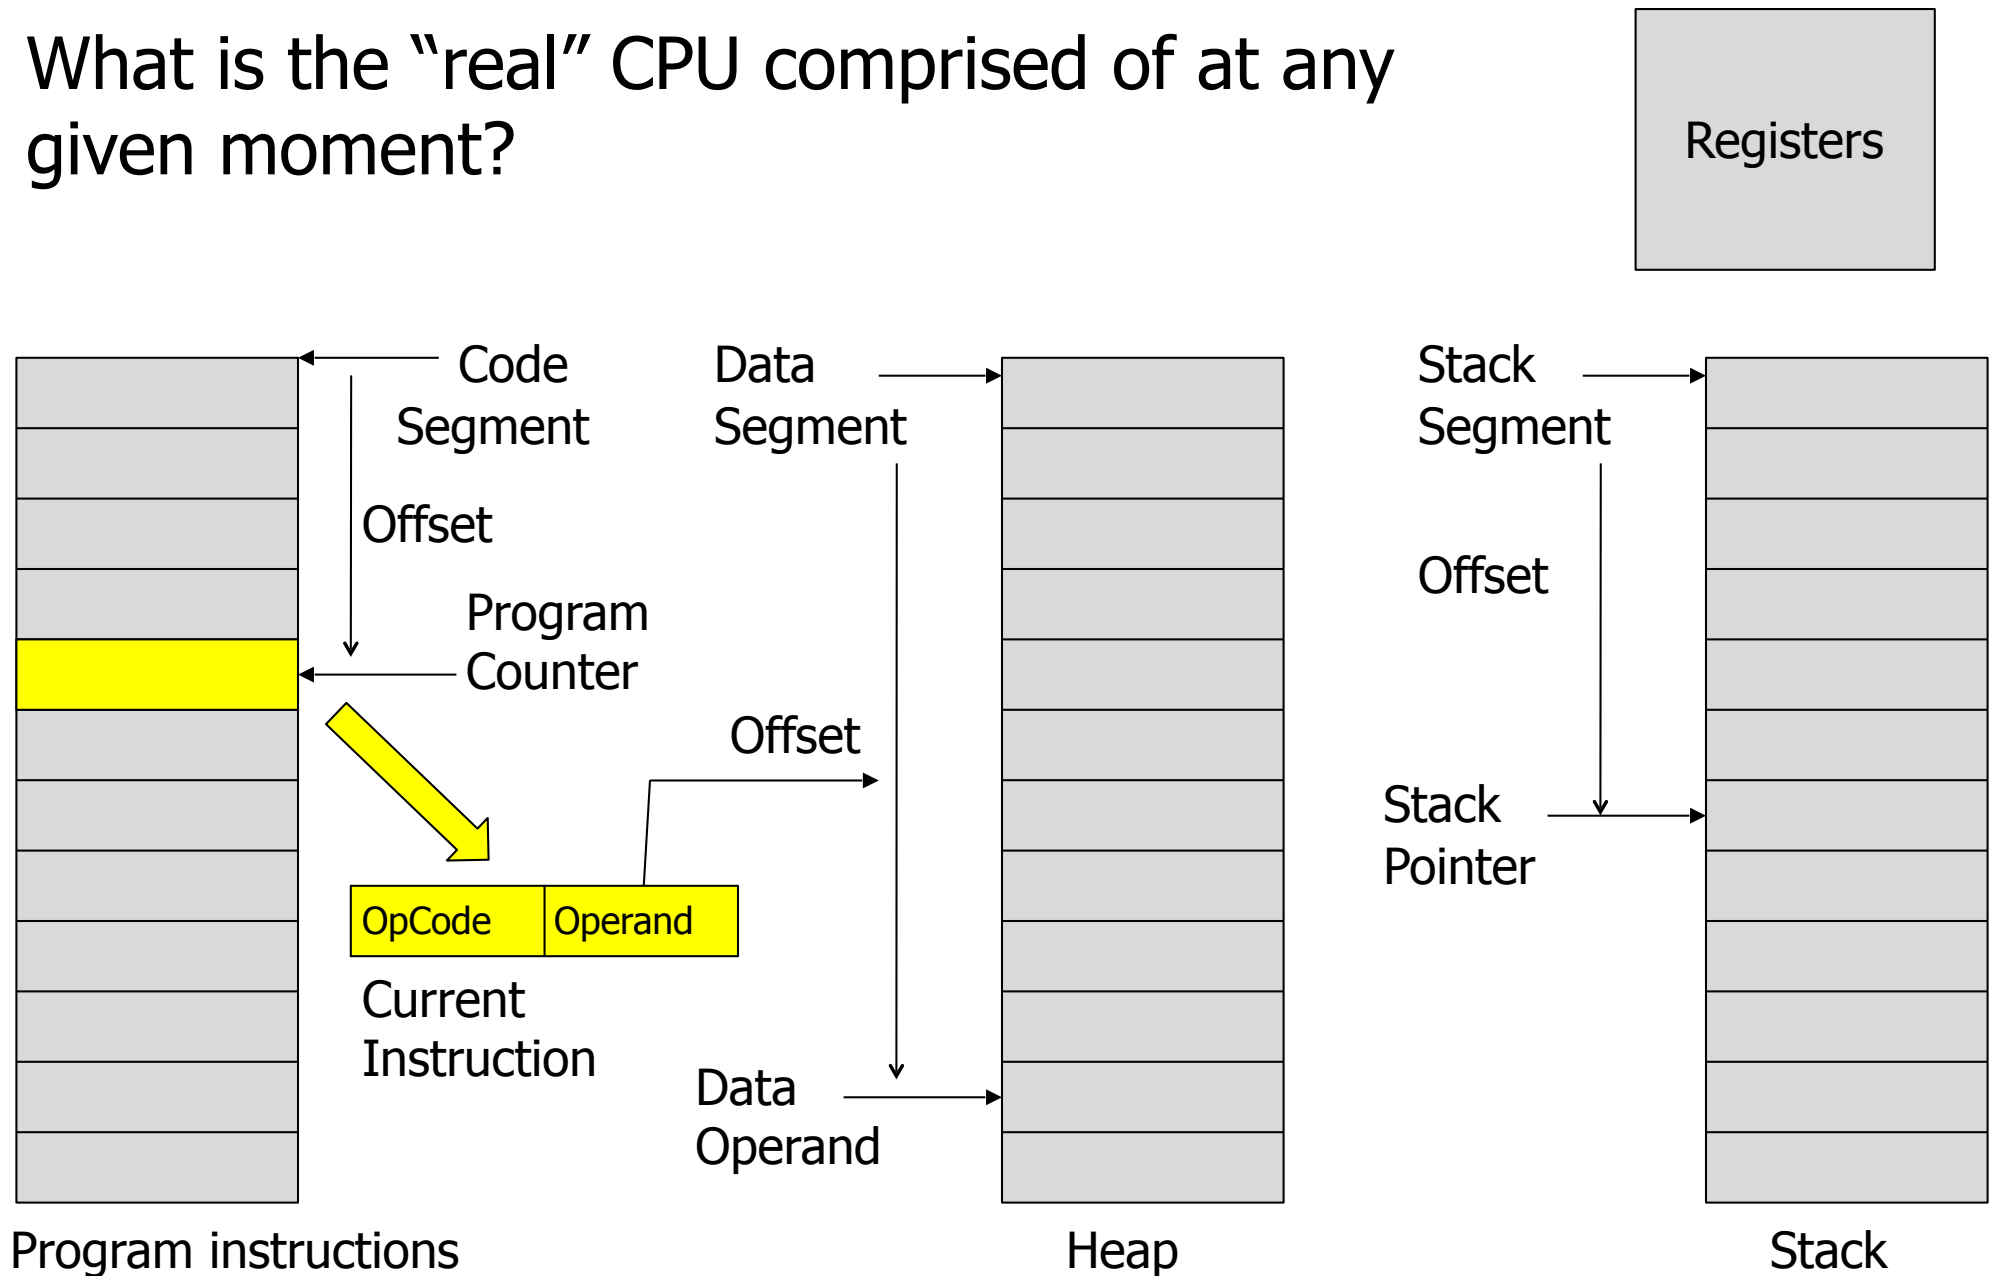
What is the "real" CPU comprised of at any given moment?



What's a 'real' CPU?



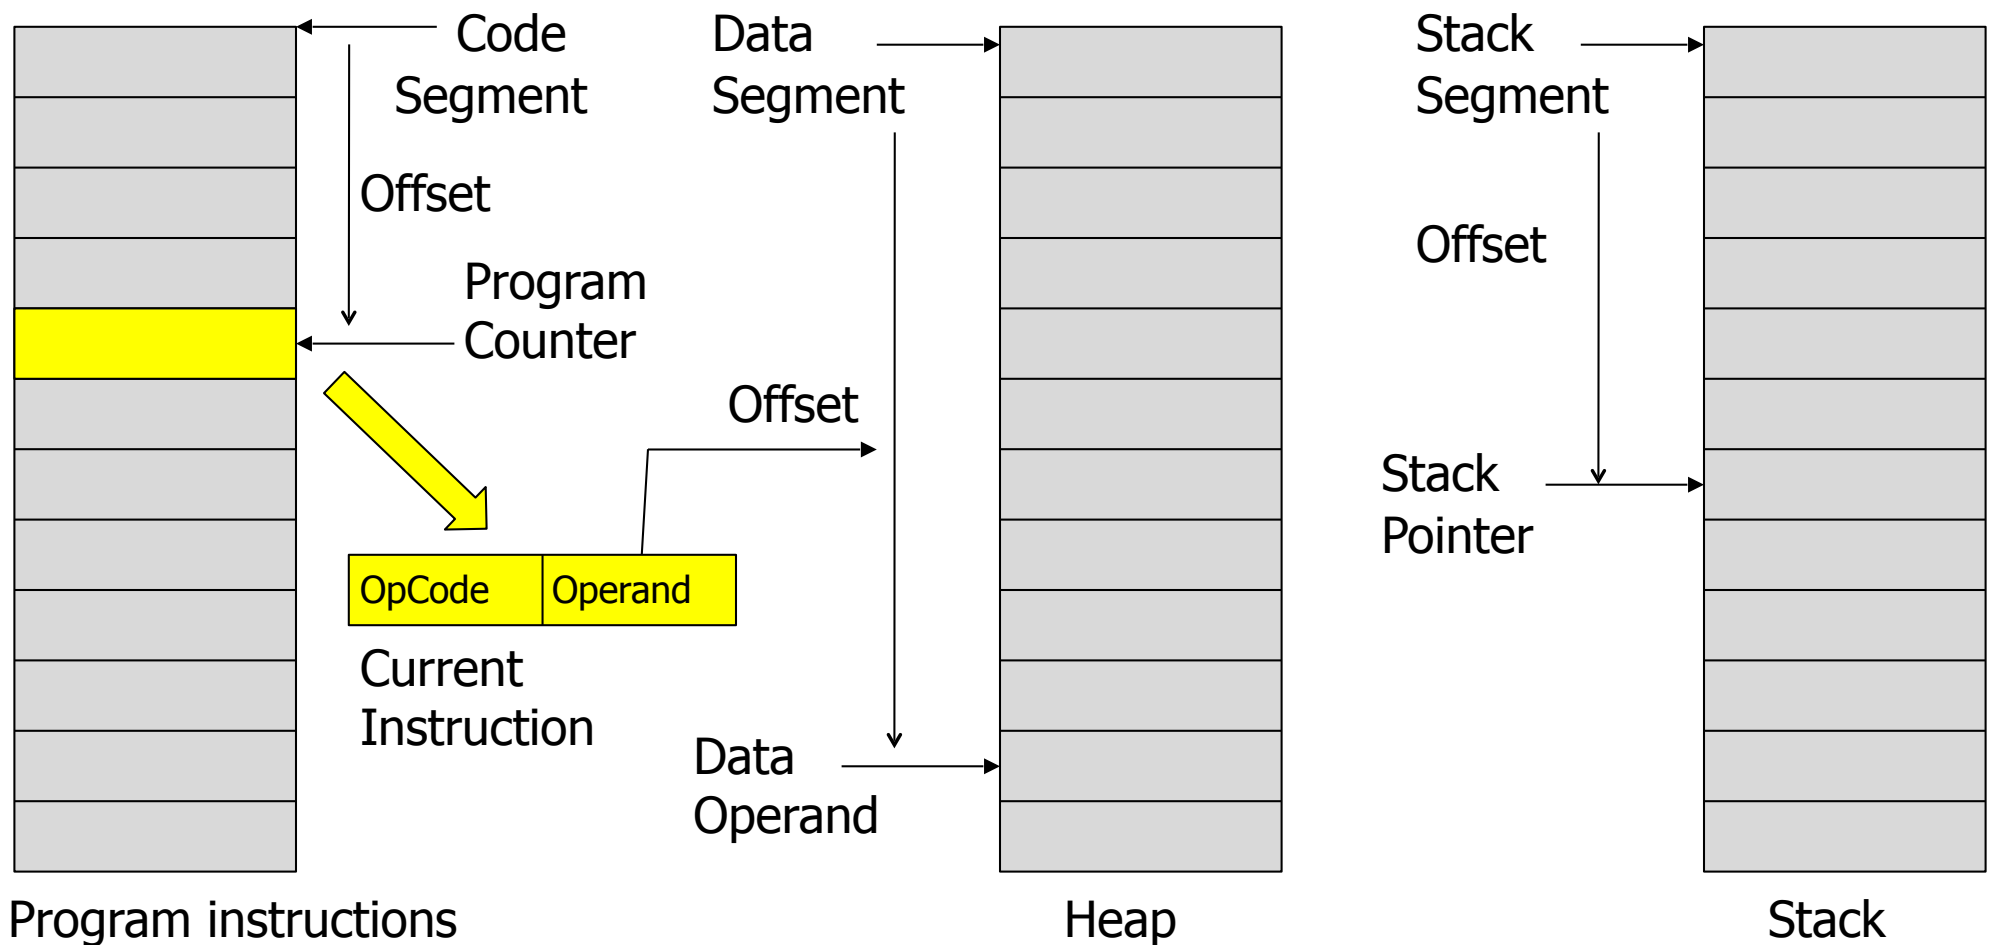

What is the "real" CPU comprised of at any given moment?



What's a 'real' CPU?



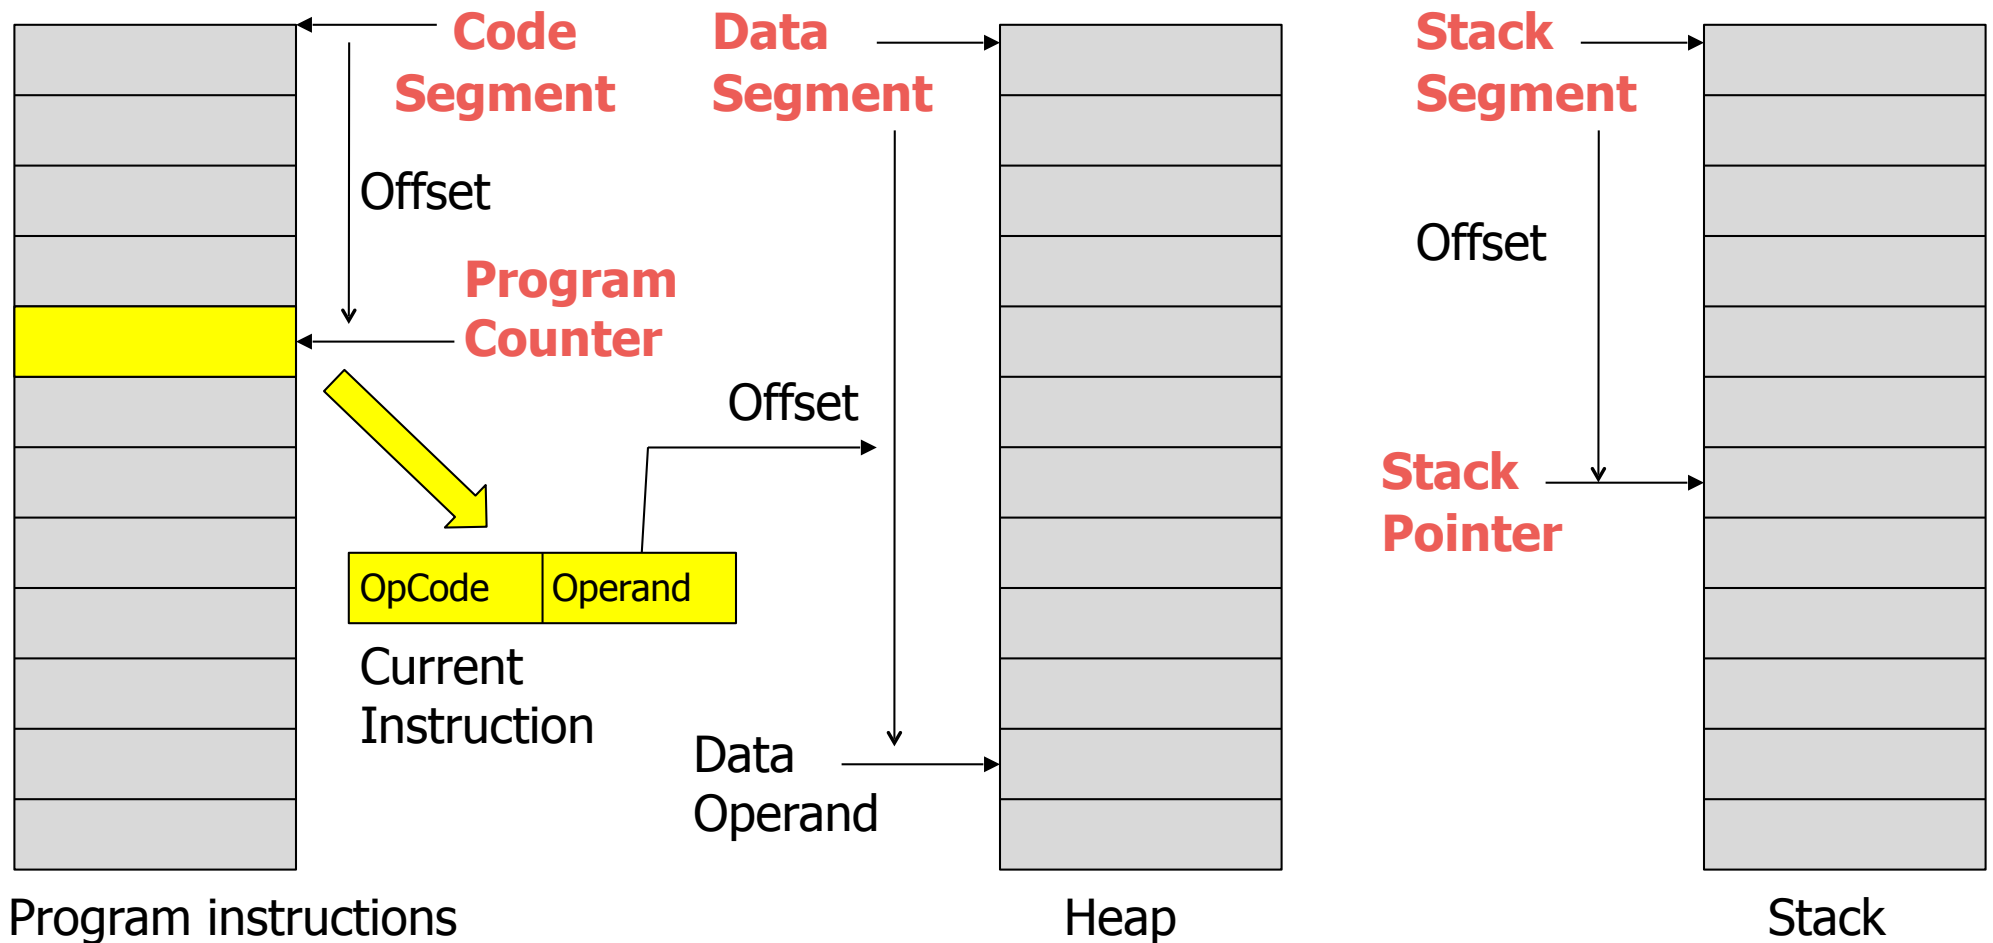

What's the **STATE** of a real CPU?



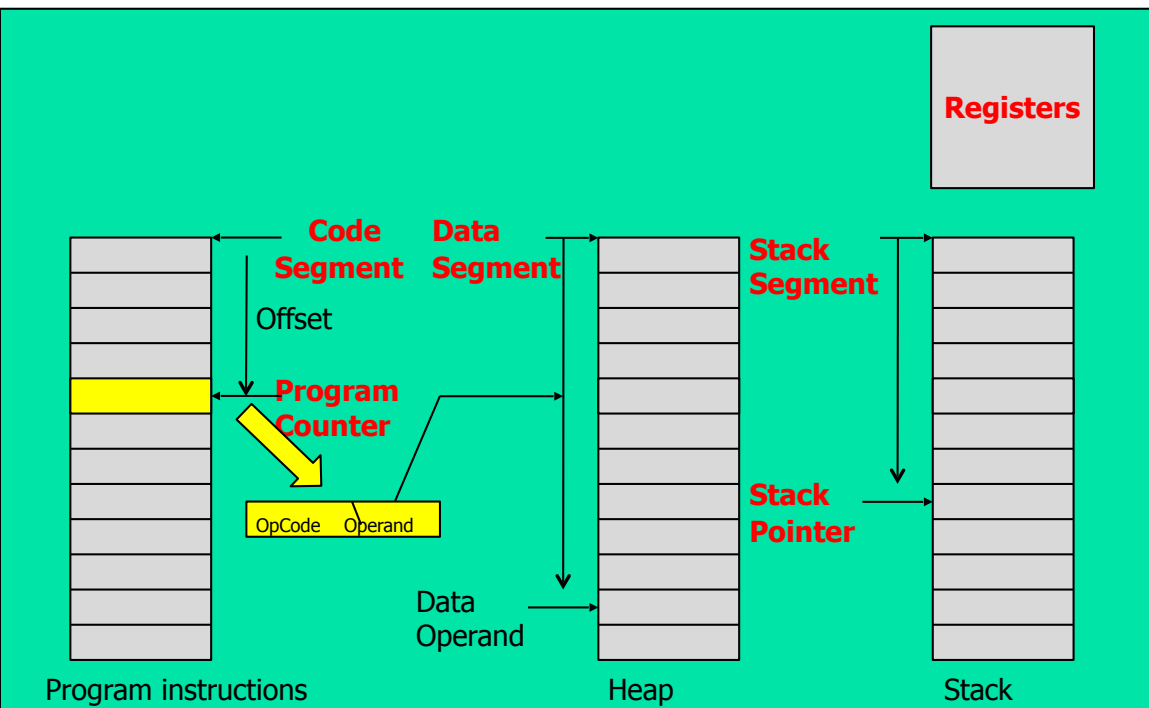
What's a 'real' CPU?



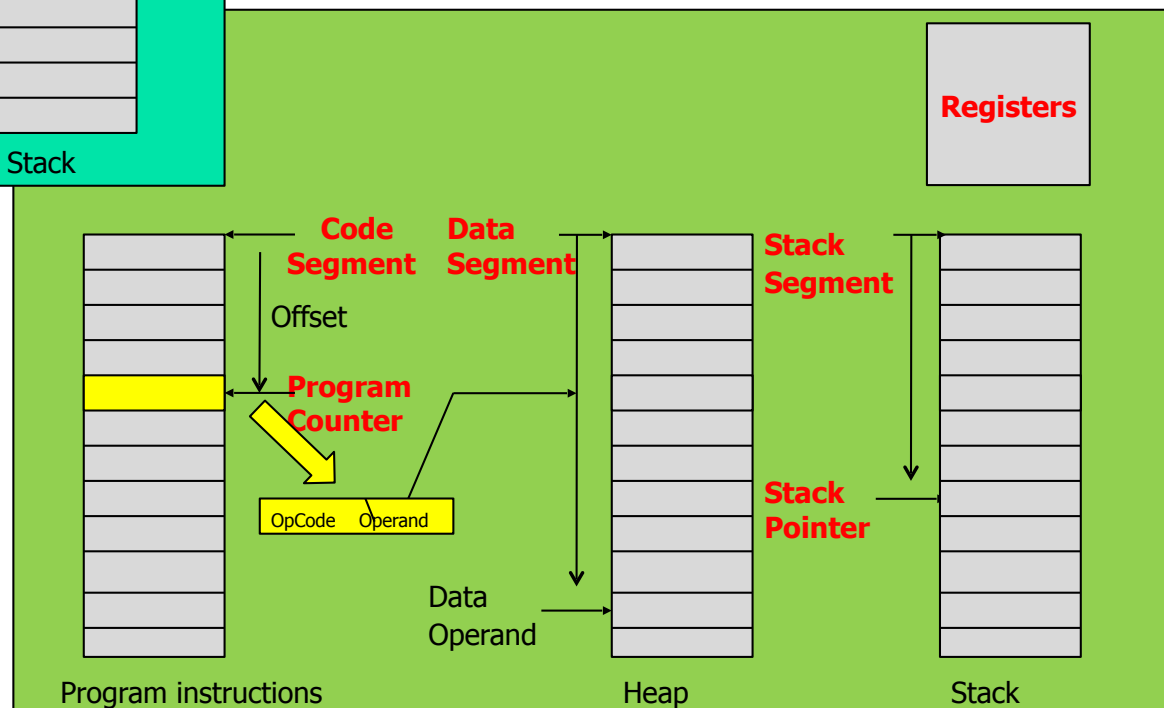

What's the **STATE** of a real CPU?




The Context Switch



**Load State
(Context)**



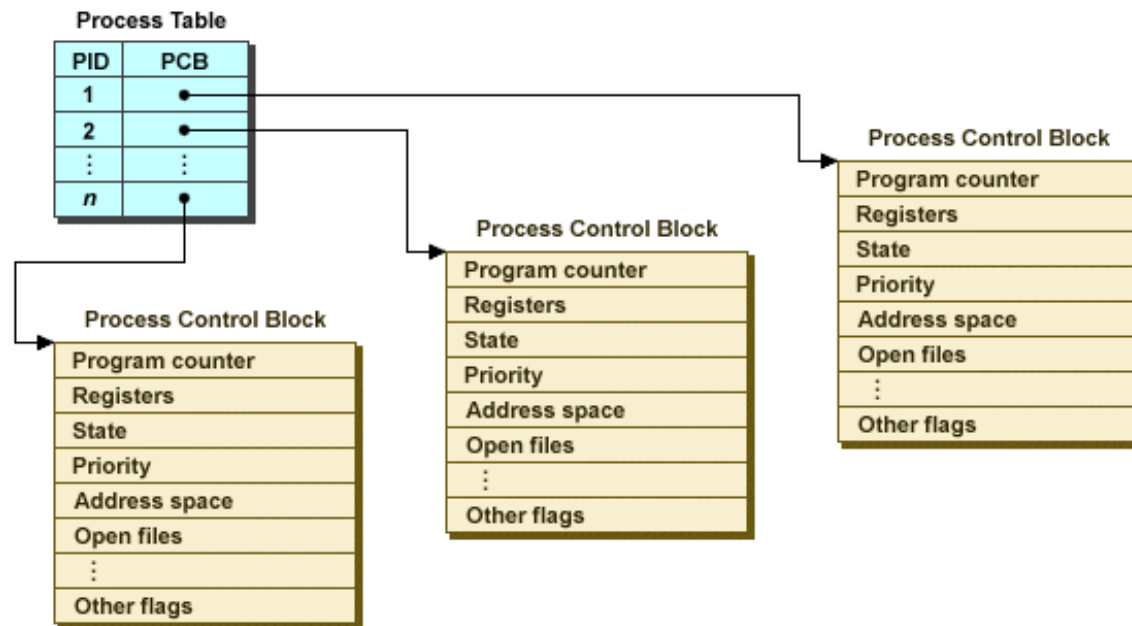
**Save State
(Context)**



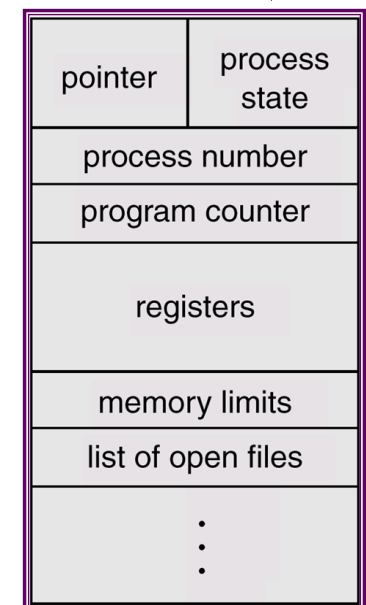
Process Control Block



The state for processes that are not running on the CPU are maintained in the Process Control Block (PCB) data structure

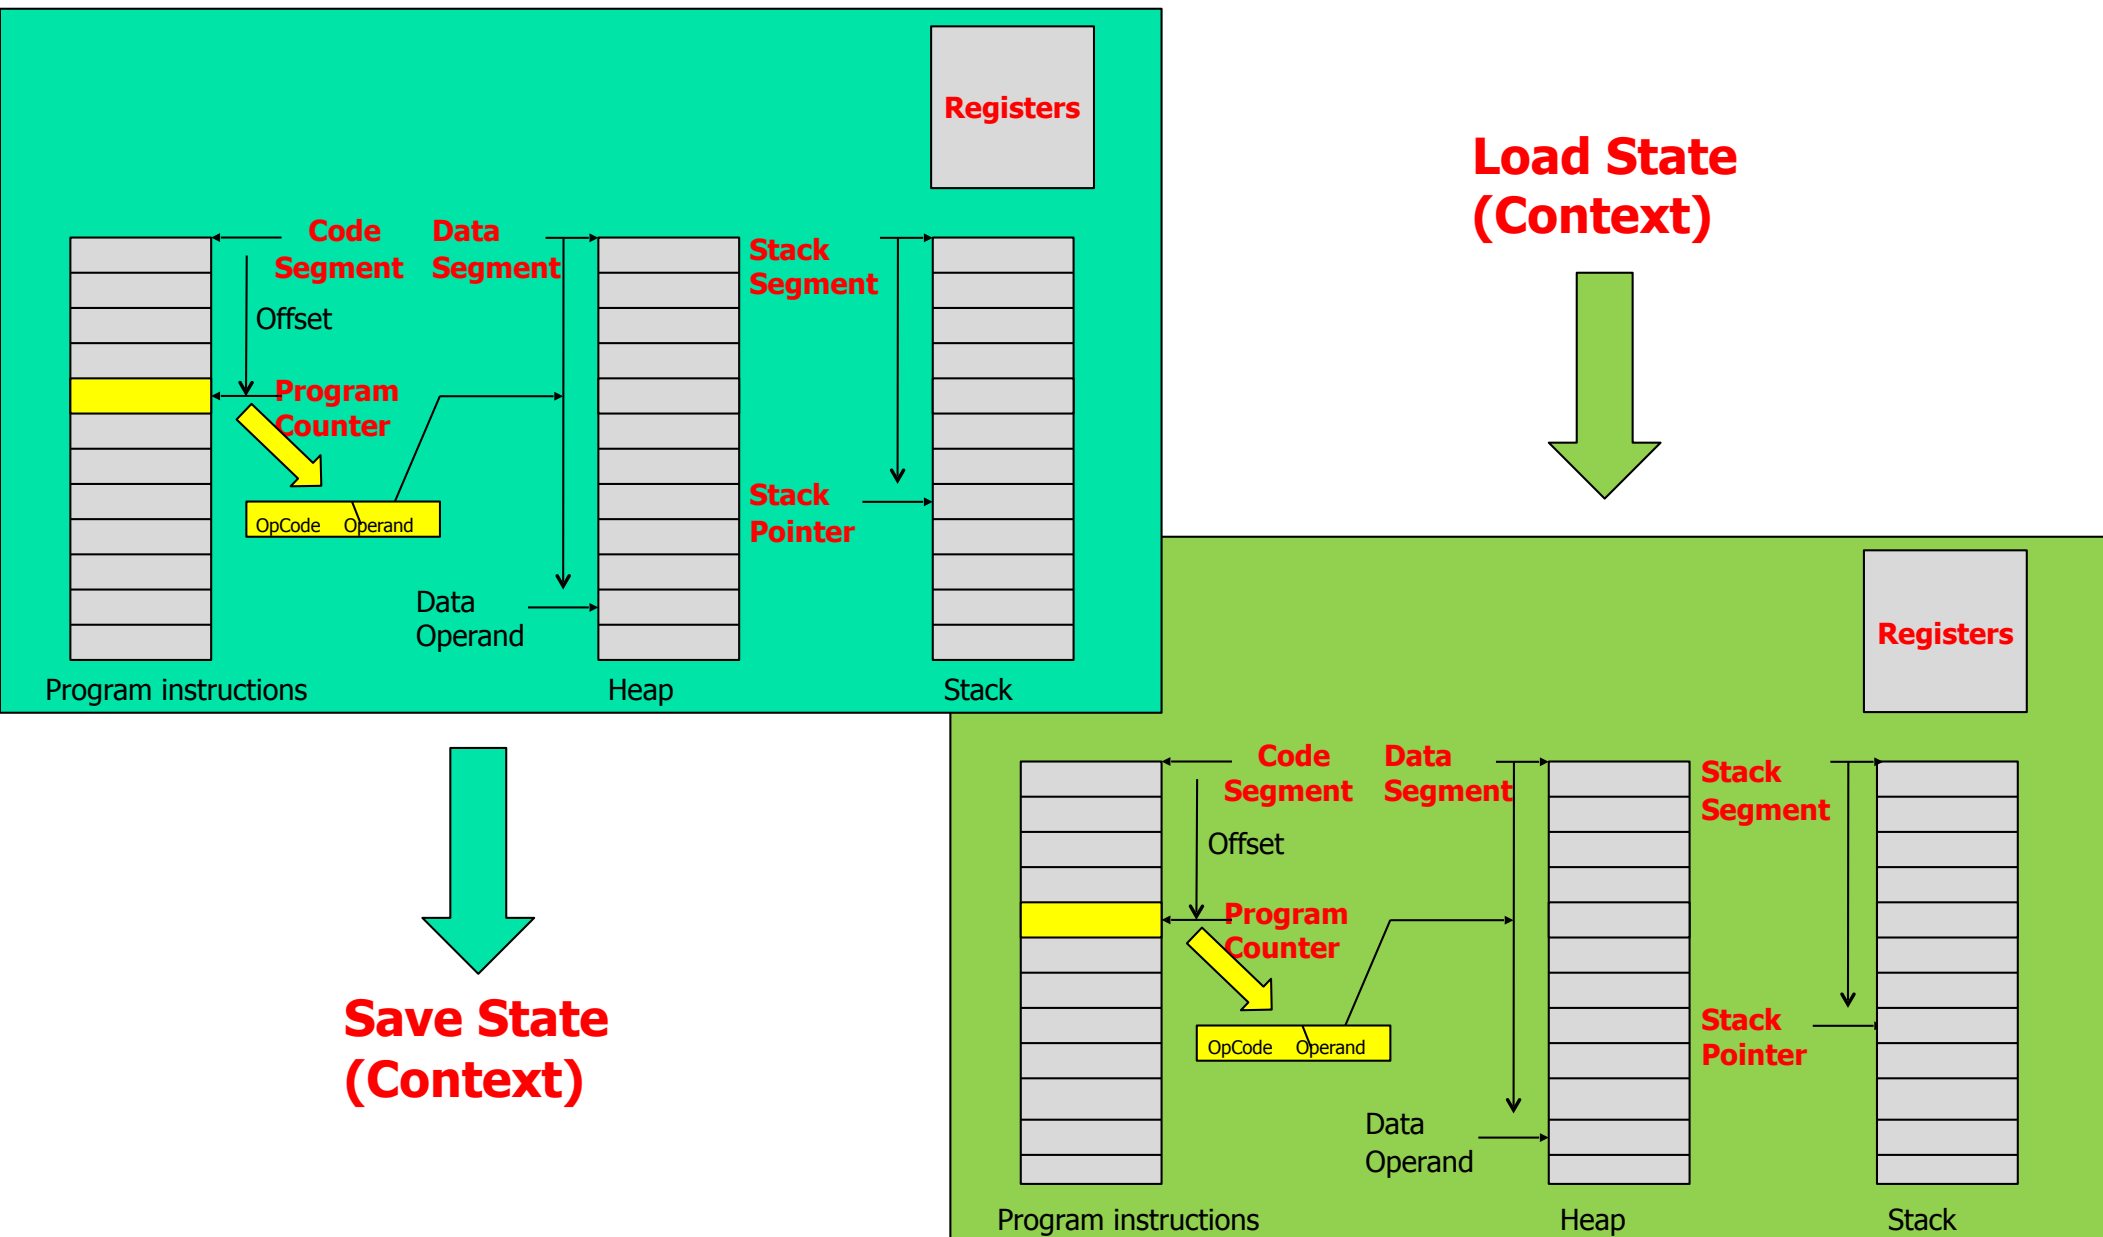


Updated during context switch

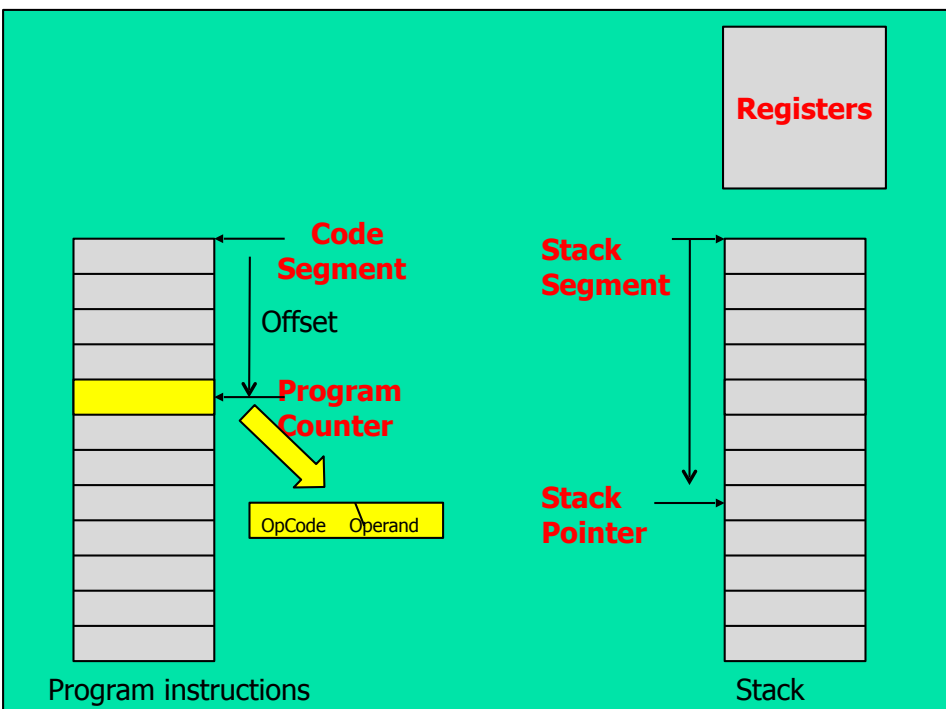


An alternate PCB diagram


The Context Switch



The Context Switch

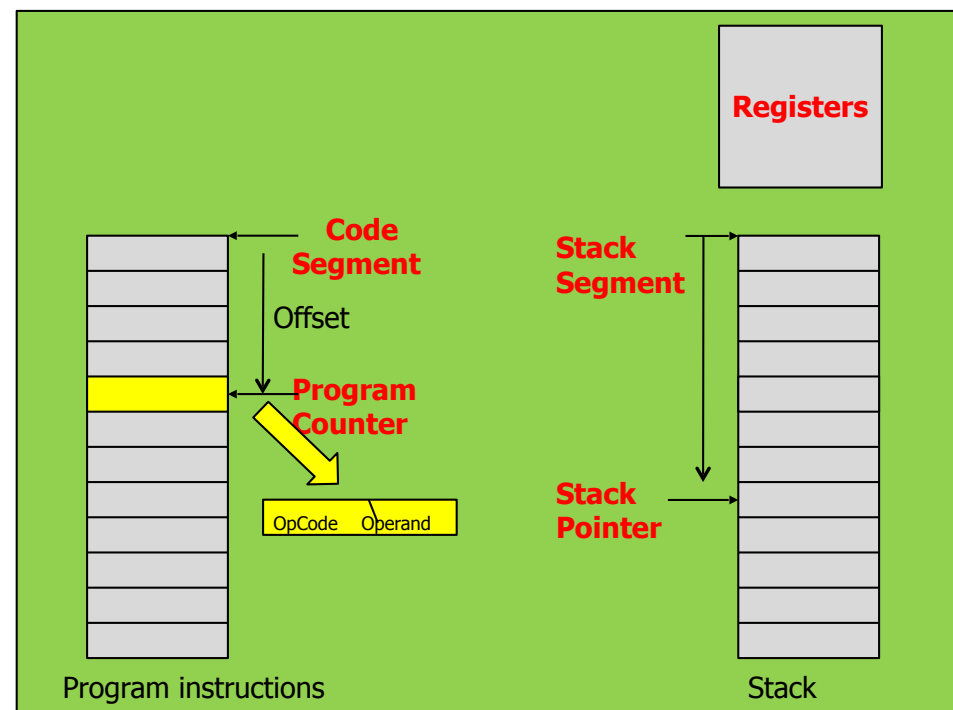



**Save State
(Context)**

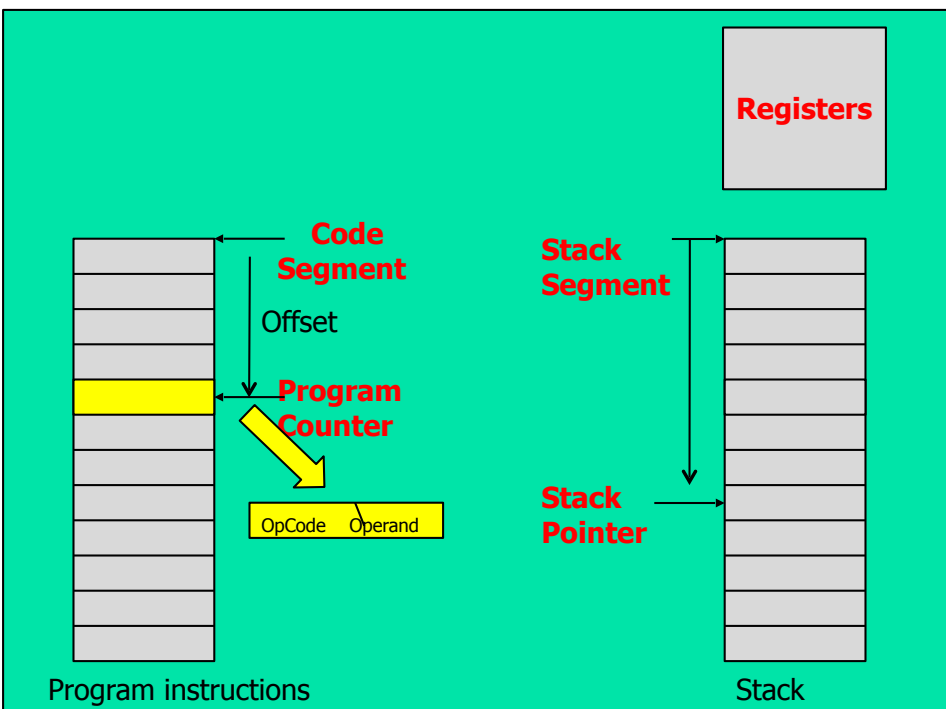


Note: In **thread** context switches, heap is not switched!


**Load State
(Context)**



The Context Switch

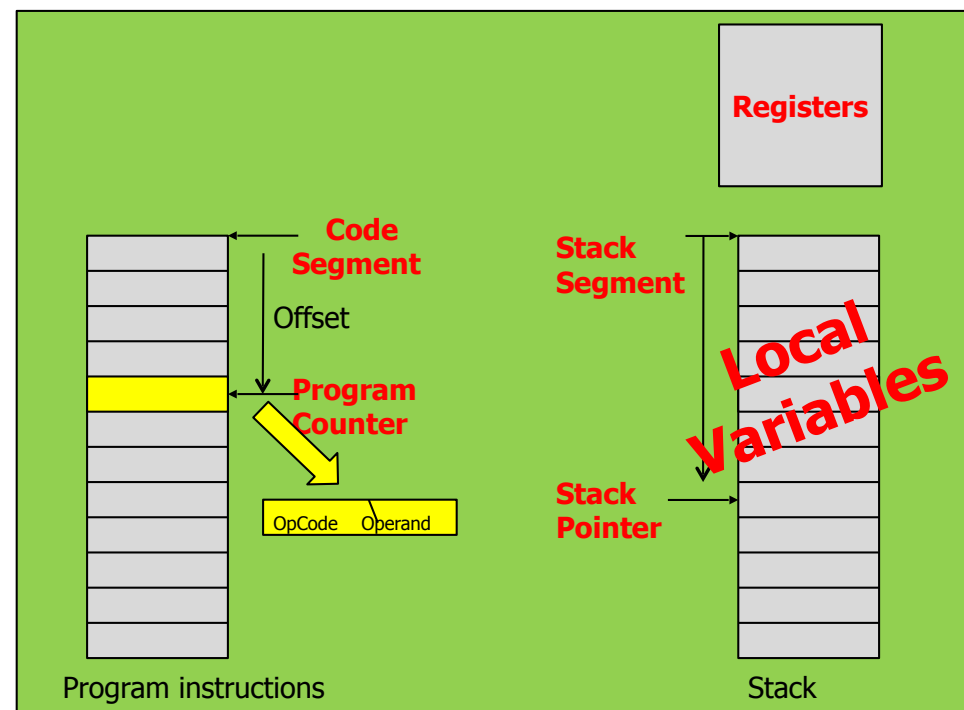


**Save State
(Context)**



Note: In **thread** context switches, heap is not switched!

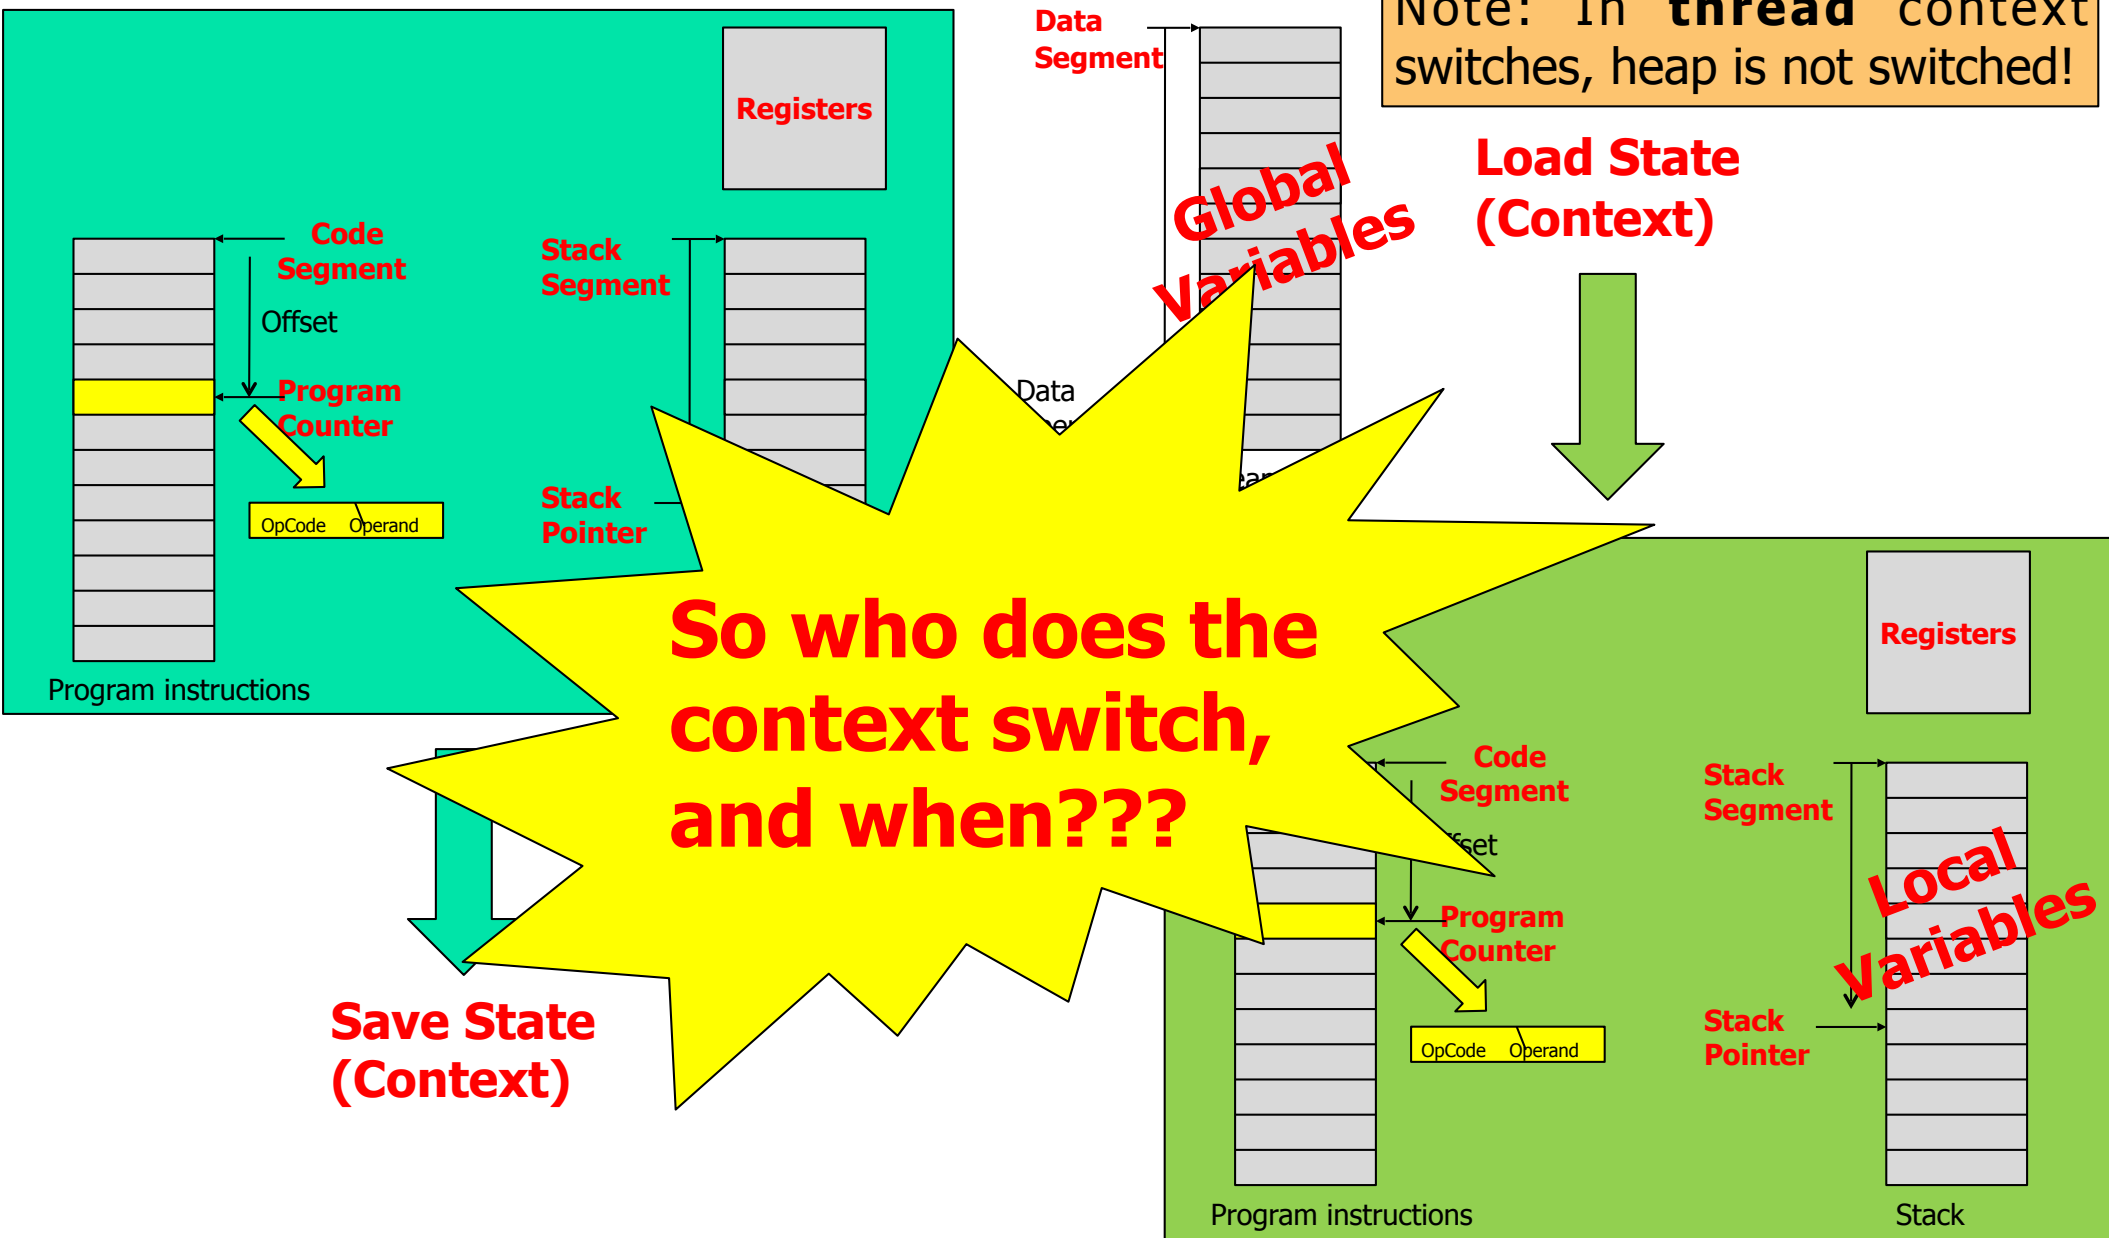
**Load State
(Context)**



Thread Context Switch




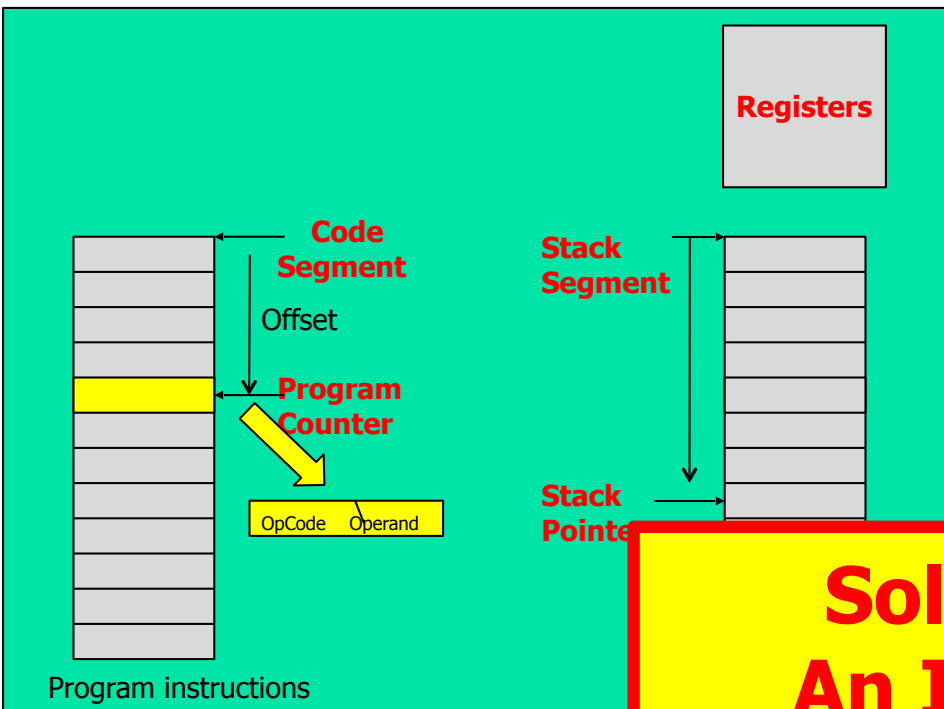
Note: In **thread** context switches, heap is not switched!



Thread Context Switch



Note: In **thread** context switches, heap is not switched!

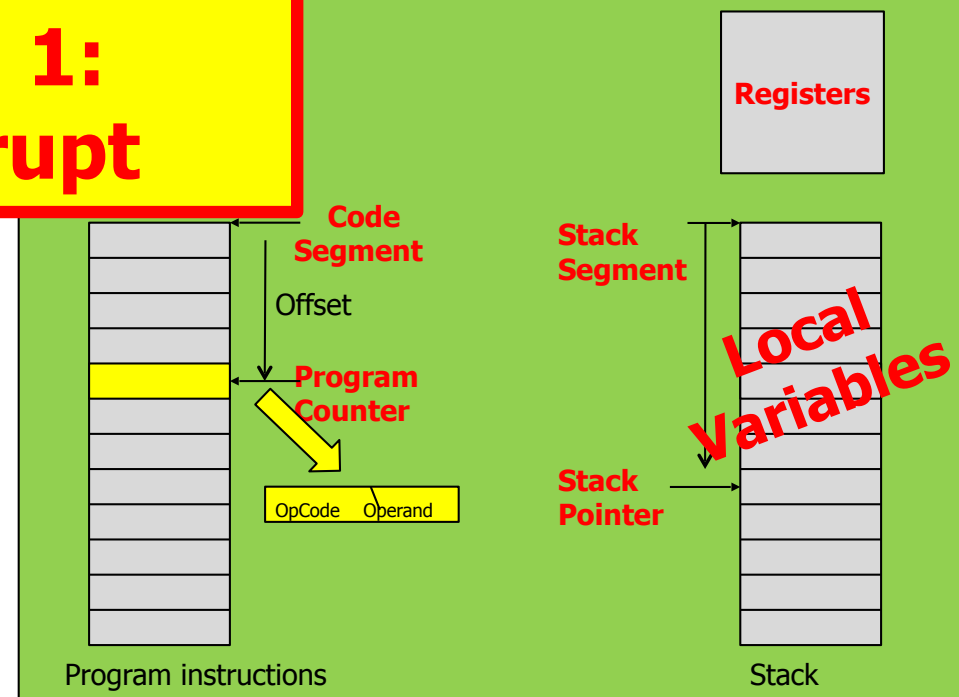



Load State
(Context)



**Solution 1:
An Interrupt**

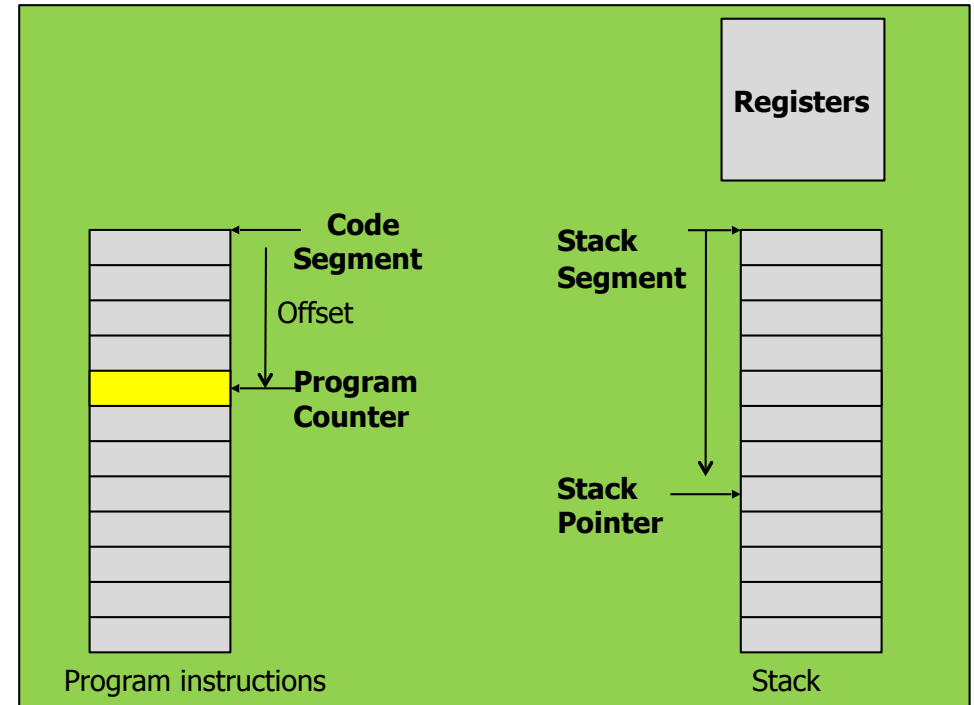
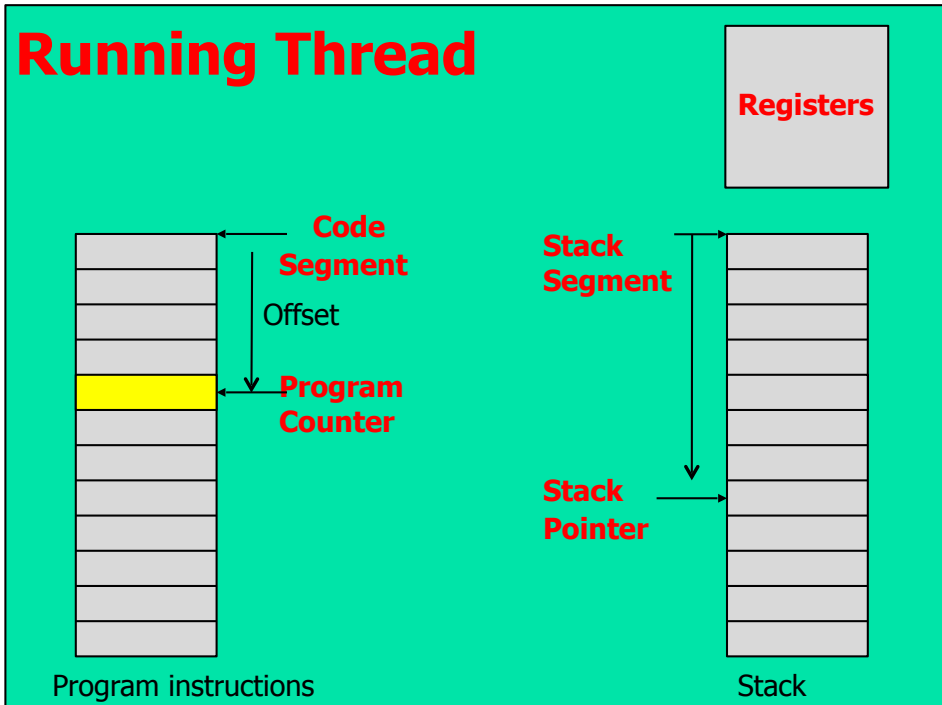
Save State
(Context)



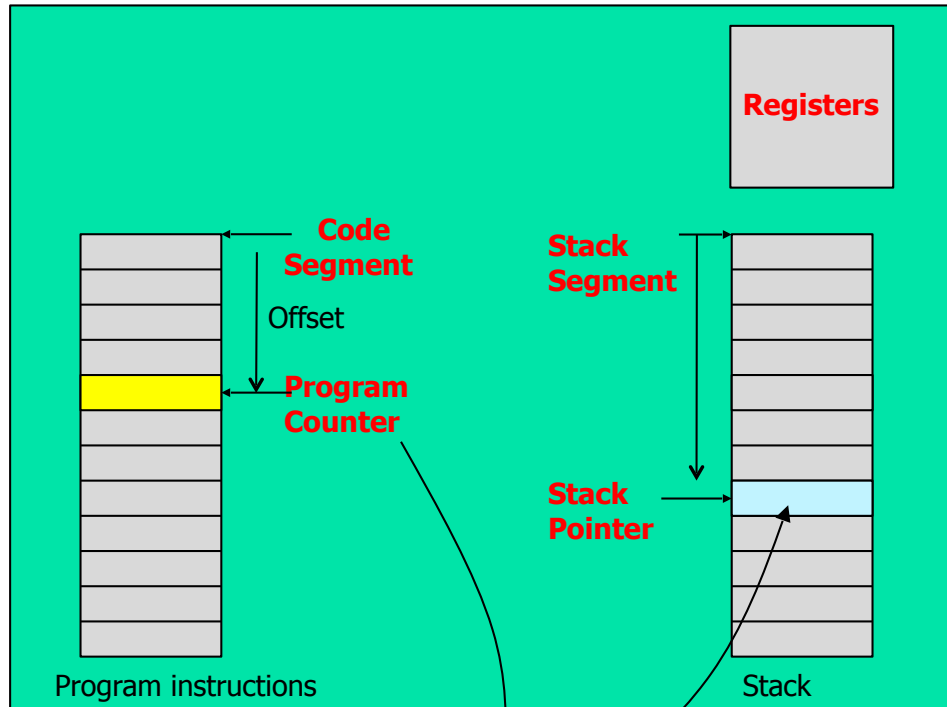
CTX Switch: Interrupt



Running Thread

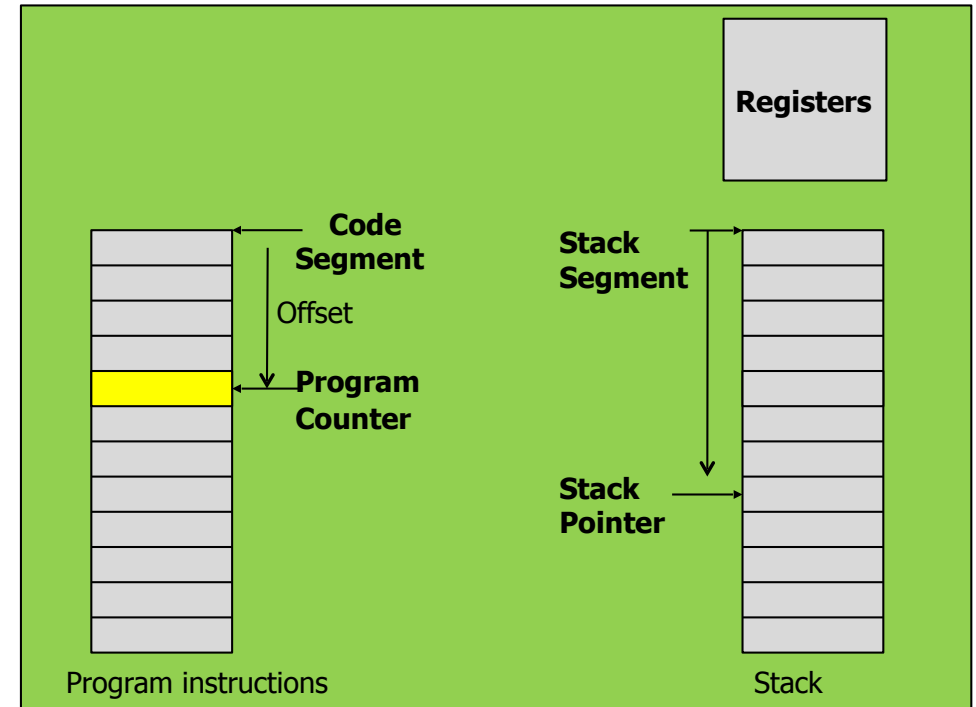


CTX Switch: Interrupt

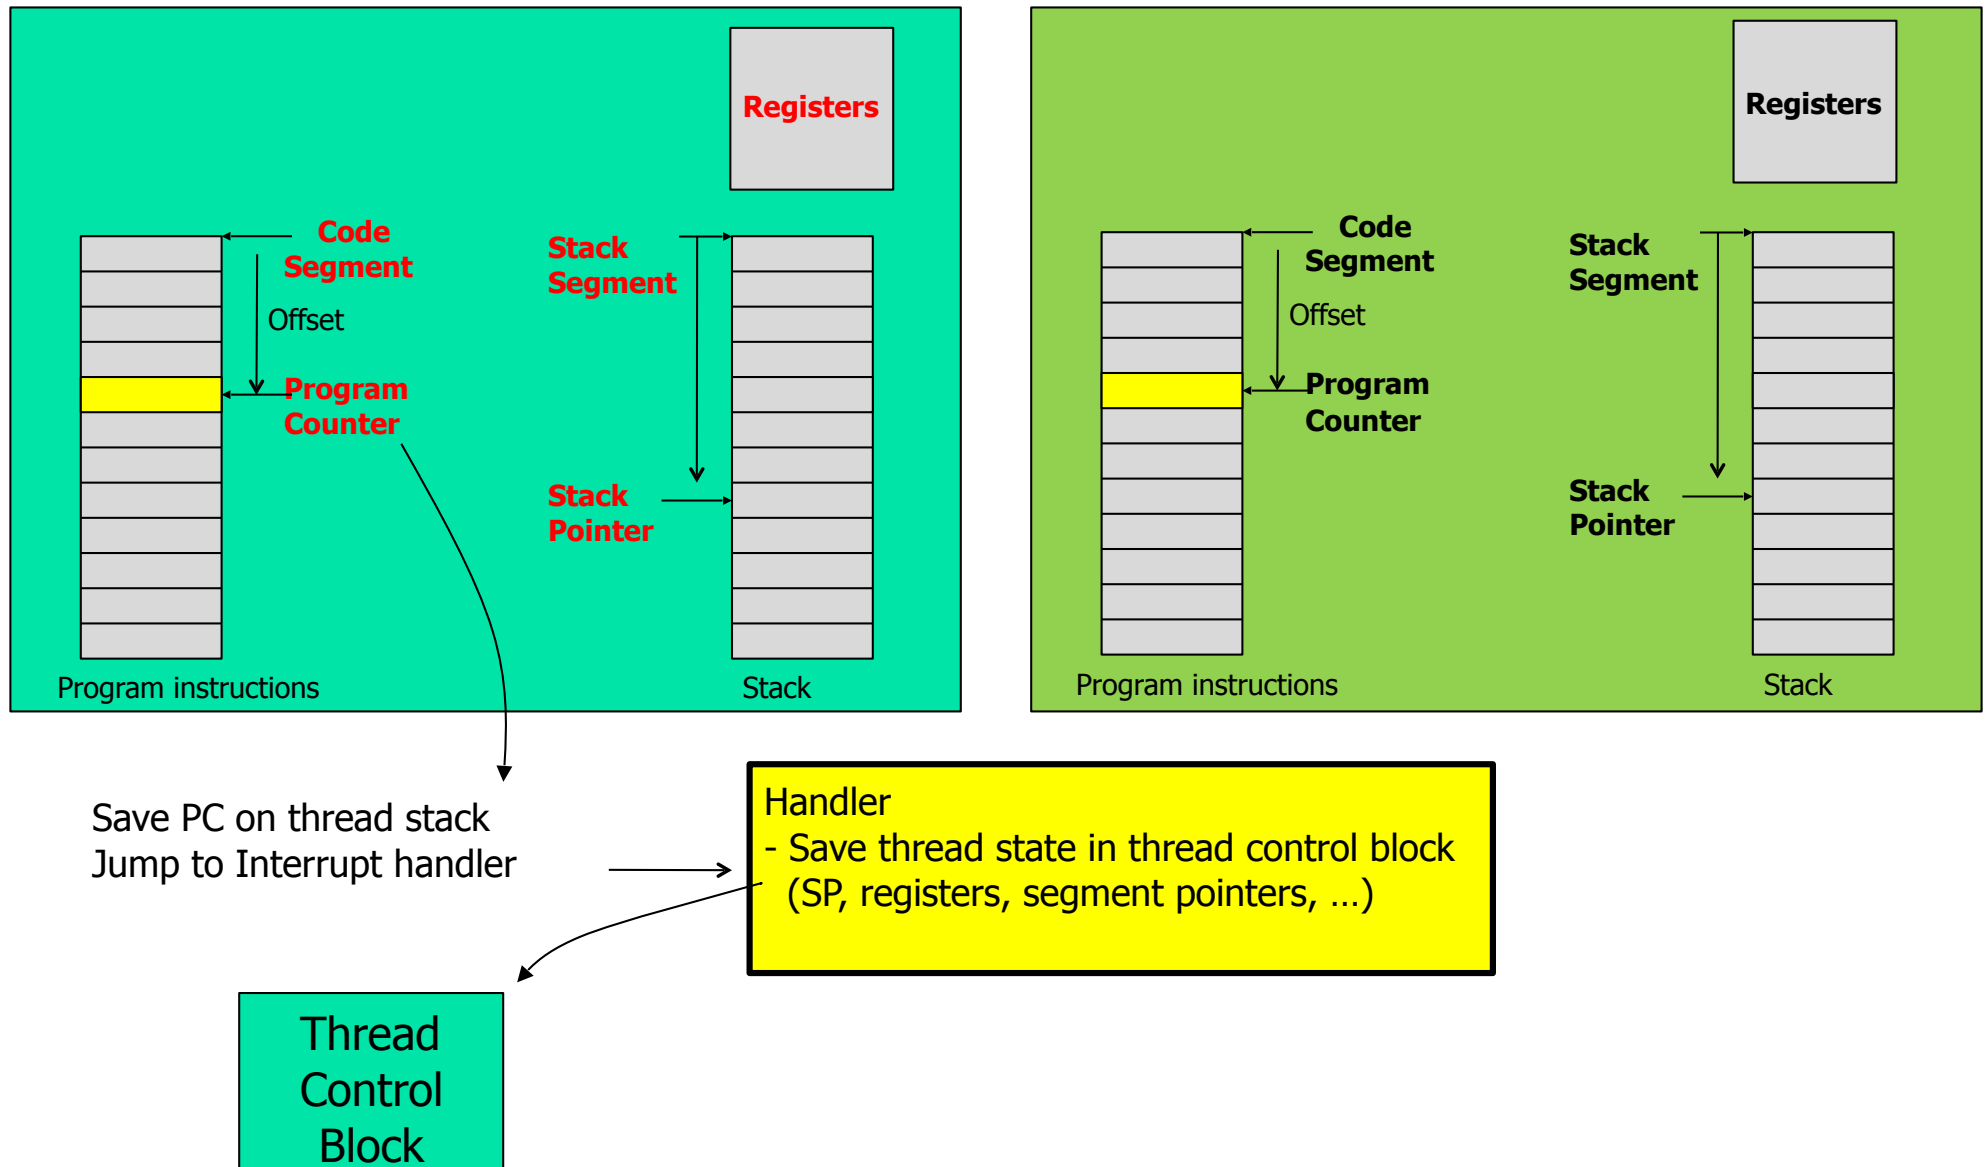


Interrupt

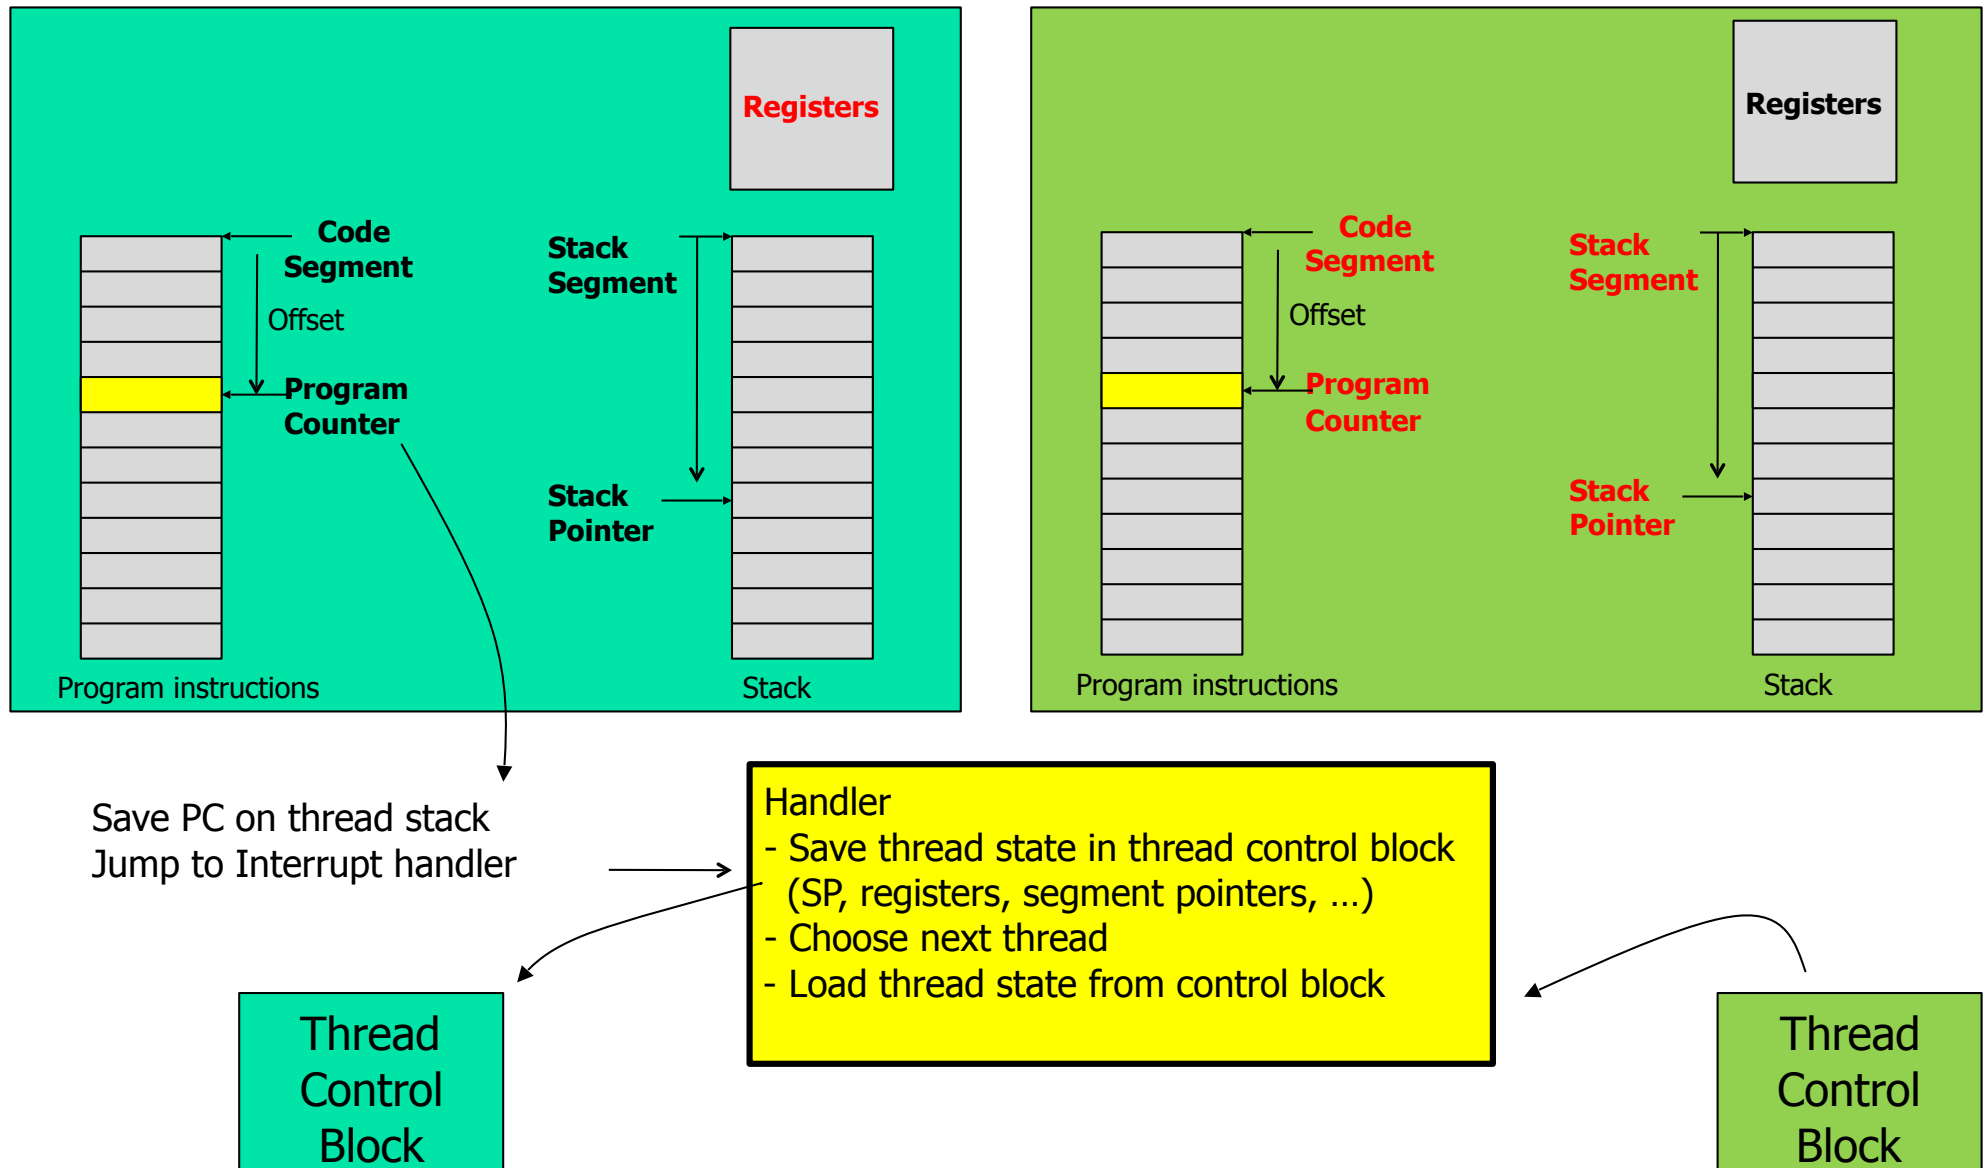
Save PC on thread stack
Jump to Interrupt handler



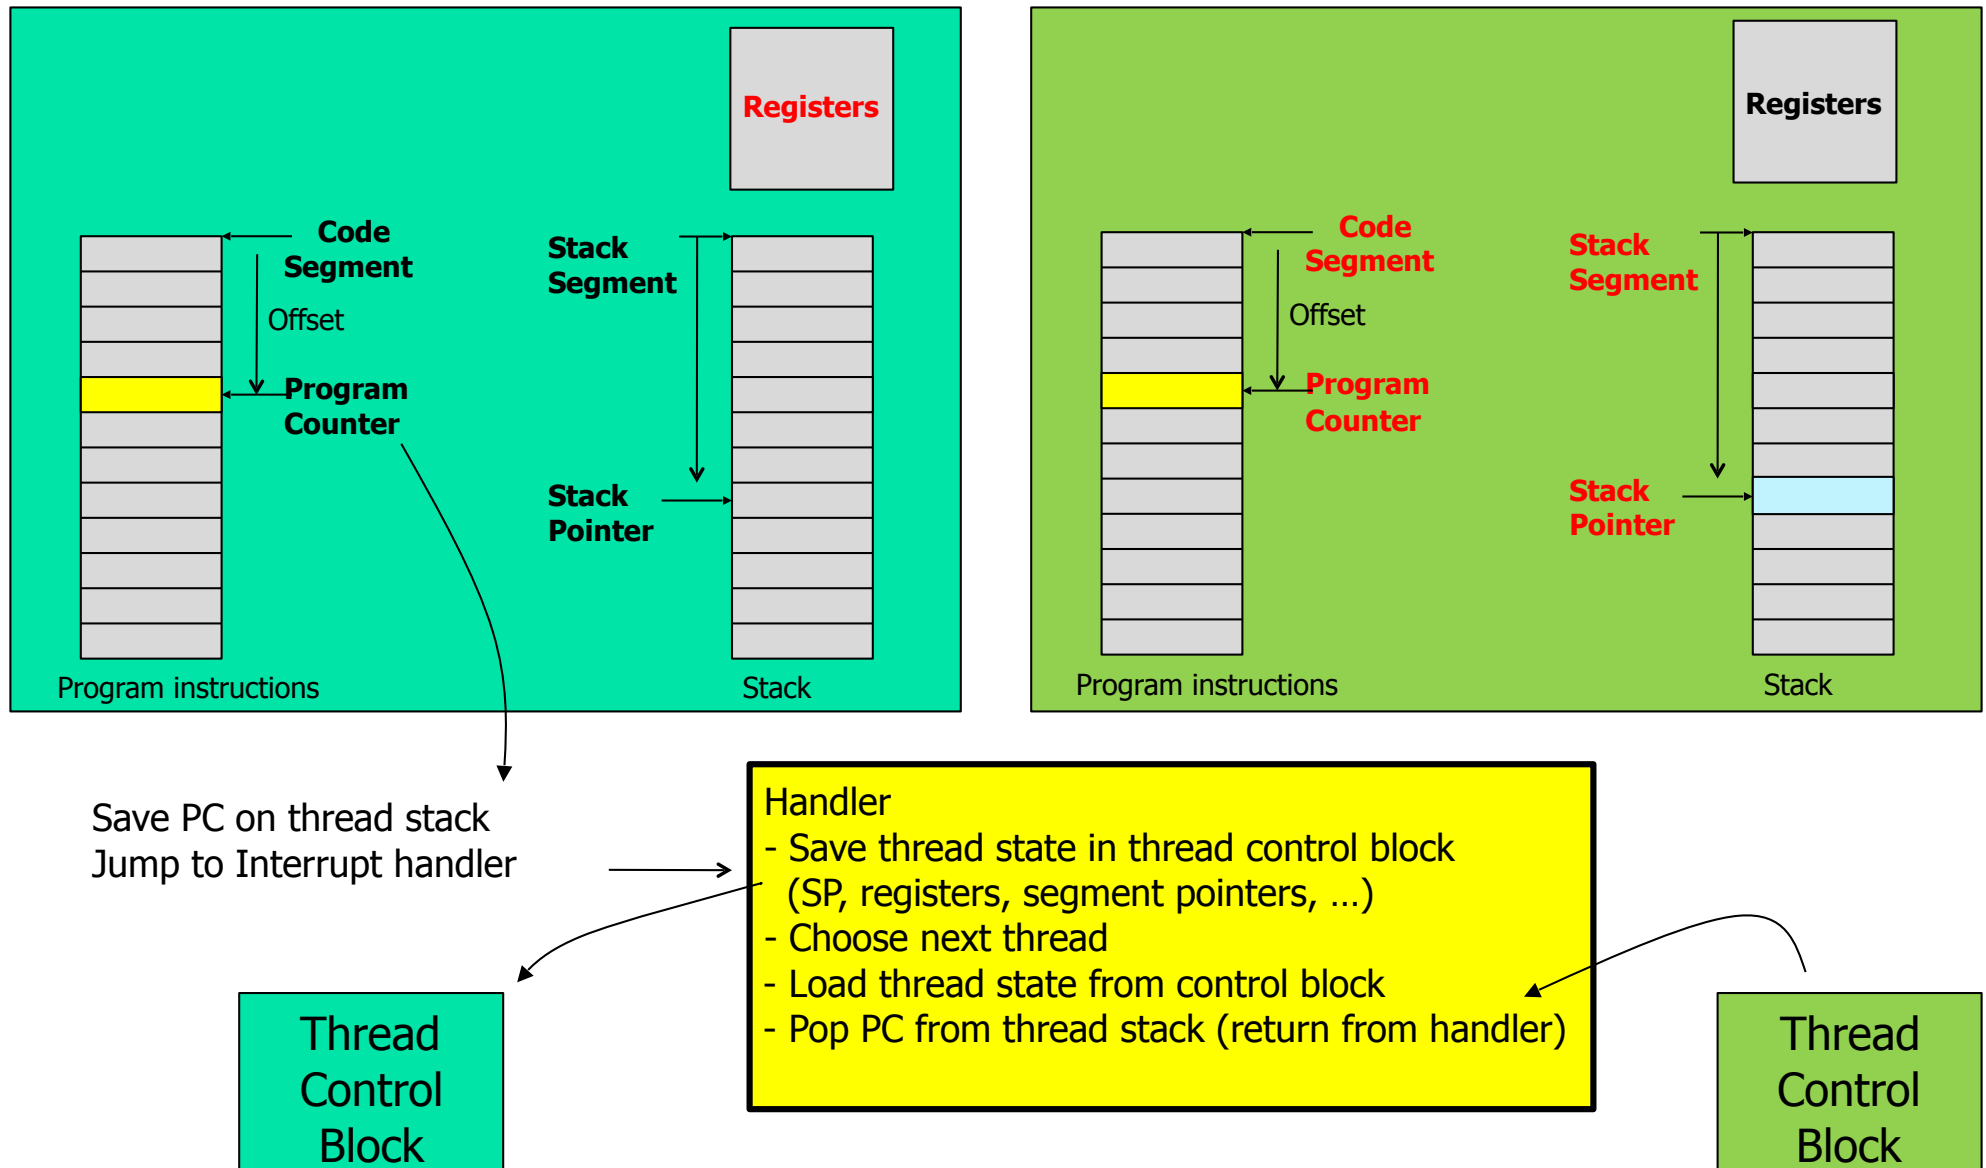
CTX Switch: Interrupt



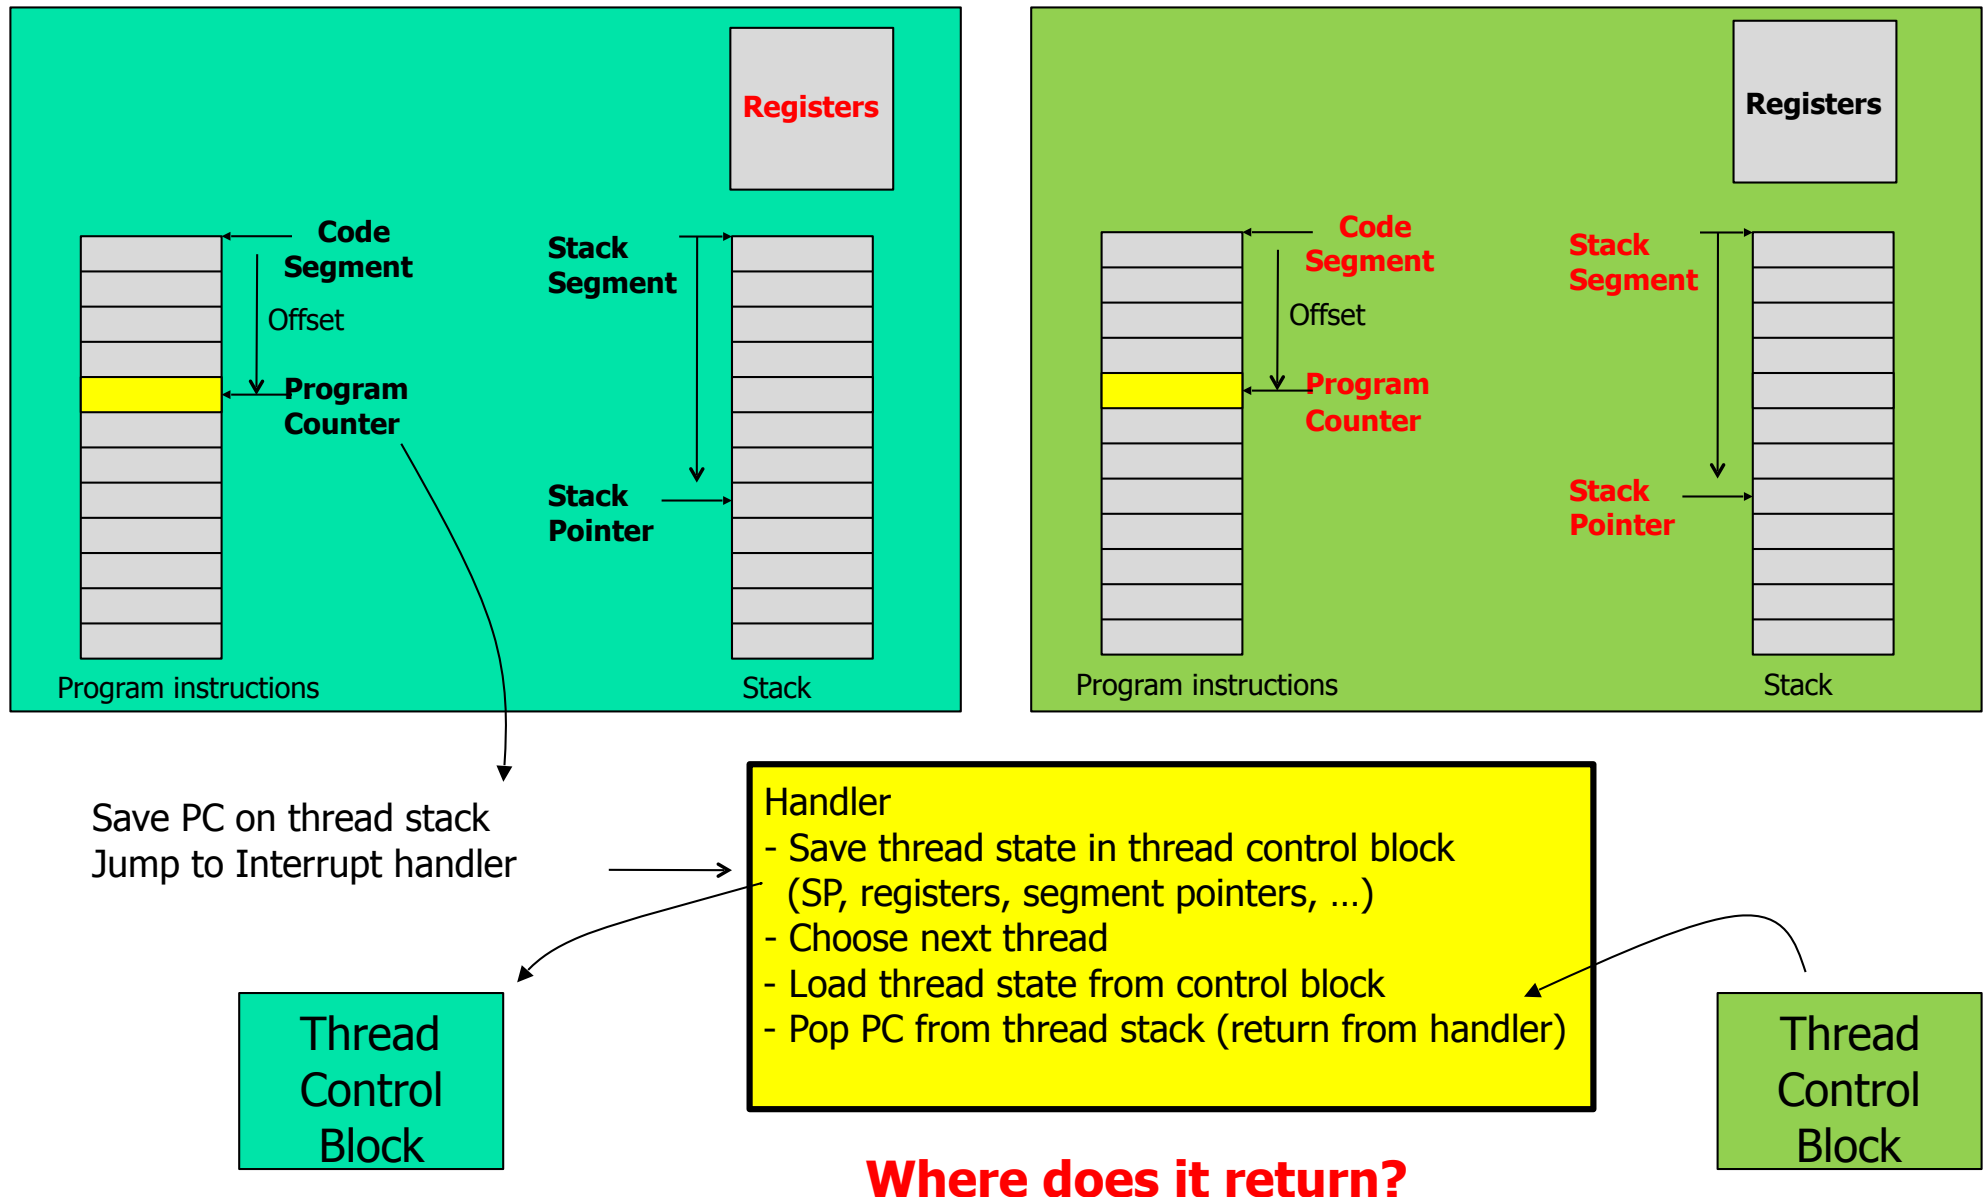
CTX Switch: Interrupt



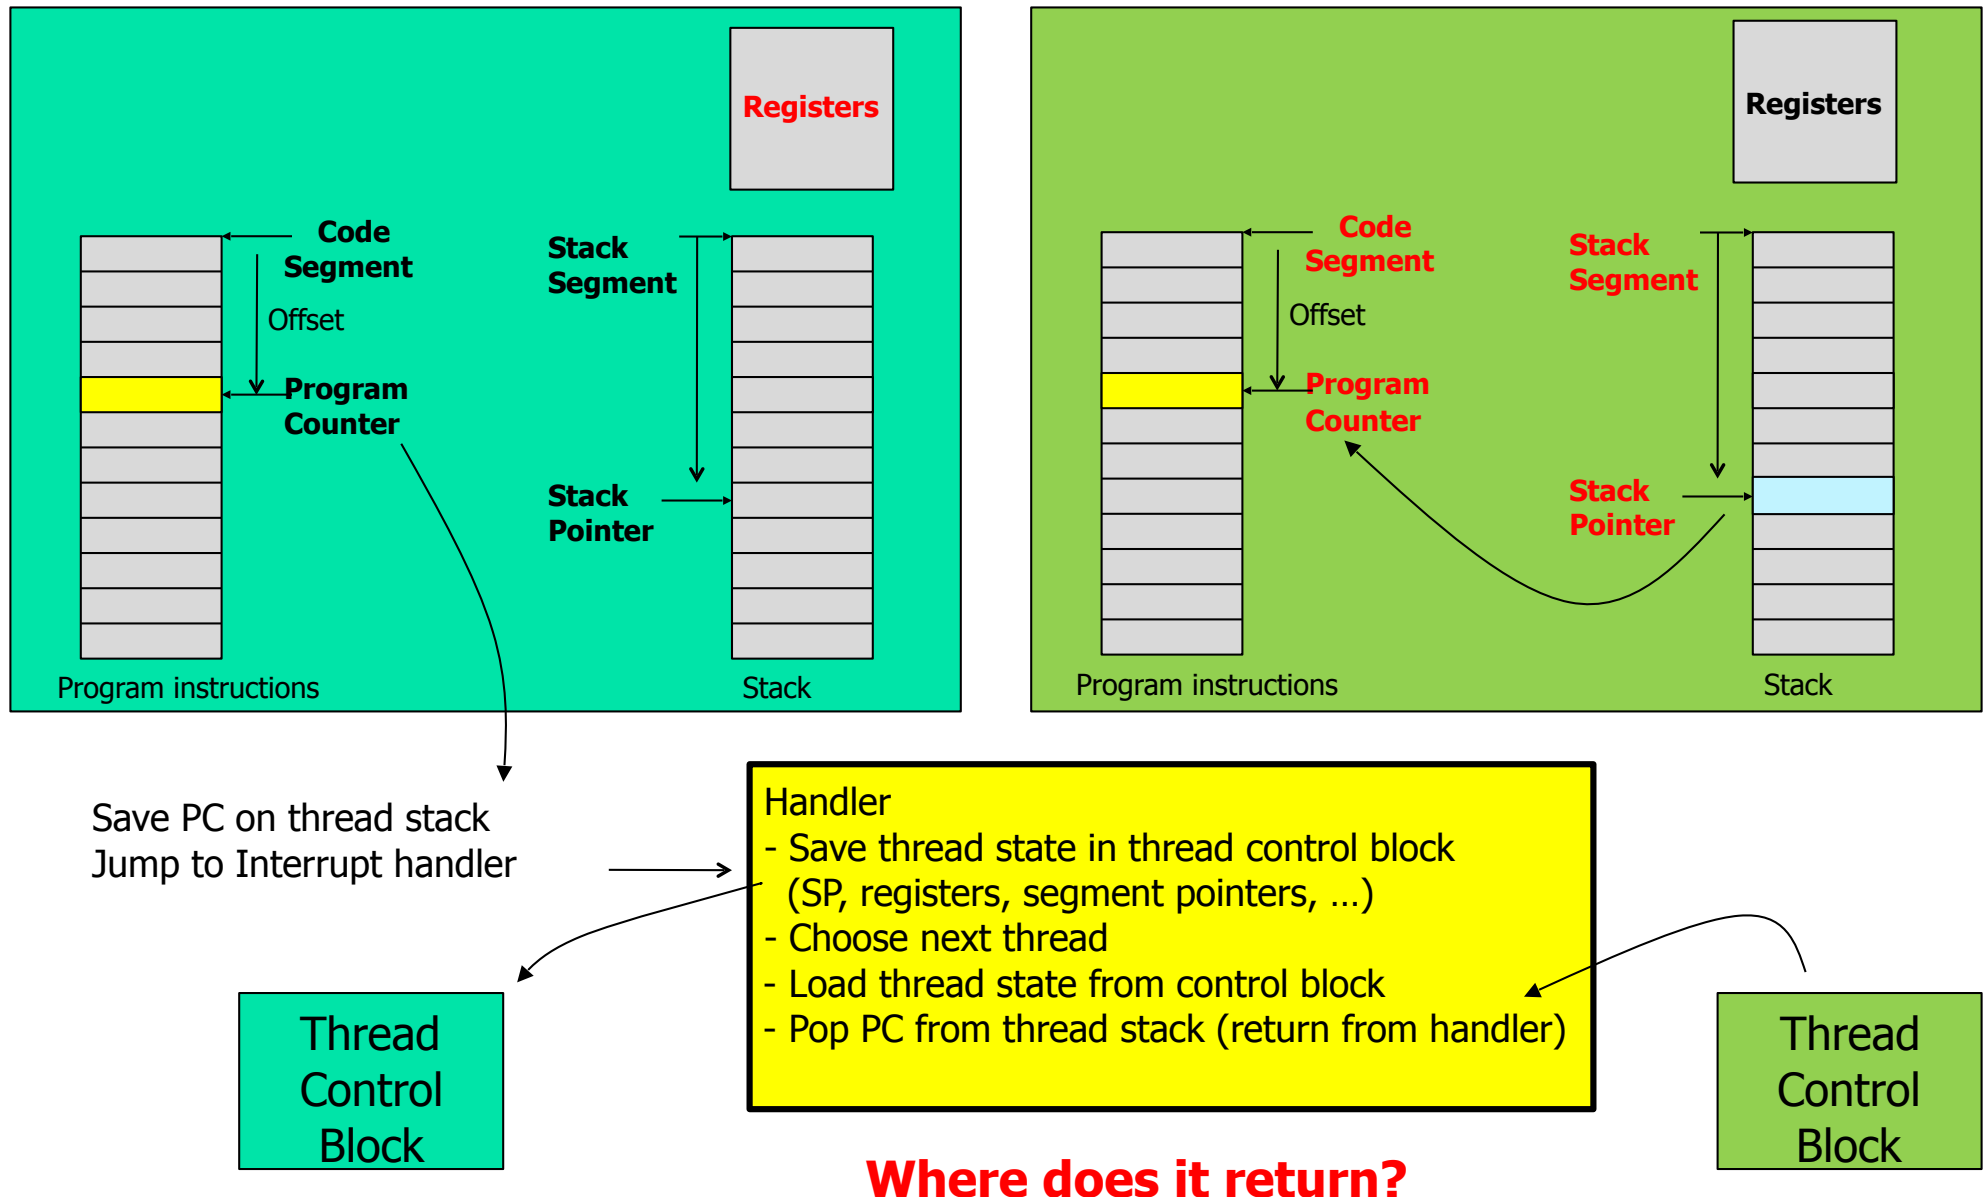
CTX Switch: Interrupt



CTX Switch: Interrupt



CTX Switch: Interrupt





What are some examples of context switches due to interrupts?

- **Clock Interrupt:** Task exceeds its time slice
- **I/O Interrupt:** Waiting processes may be preempted
- **Memory Fault:** CPU attempts encounters a virtual memory address that is not in main memory. OS may resume execution of another process while retrieving the block, then moves process to ready state.