



# CS 423

## Operating System Design: Concurrency (more Threads)

Professor Adam Bates  
Fall 2018

# Goals for Today



- Learning Objectives:
  - Dive yet further into concurrency and threading
- Announcements:
  - **MP1 out on Friday!**



**Reminder:** Please put away devices at the start of class

# HWI Wrap-Up



- Average Score: 88.8%...




- Three questions of medium difficulty (<80%):
  - “Hold-and-wait” situations?
  - which code snippets are wrong?
  - UNIX shell?



Which of the following systems may never exhibit a "hold-and-wait" situation? You may assume that the only blocking that occurs in these systems occurs on mutexes.

- a. Systems that ensure that all resources needed for an application are locked in a single atomic operation that either succeeds and locks all requested resources or fails and locks none.
- b. Systems with only one mutex.
- c. Systems where all mutexes are numbered. A user cannot lock a mutex with a lower number,  $X$ , after they have locked a mutex with a larger number,  $Y > X$ .
- d. Systems of type (a) and (b) only
- e. Systems of type (a), (b), or (c)




# HWI Wrap-Up



Which of the following code snippets are wrong?


Case 1

```
int *p;  
*p=10;
```



Case 2

```
char a[2];  
strcpy (a, "Hi");
```




Case 3

```
int b[10];  
*b=11;
```



Which of the following best describes a UNIX shell?



- a. Part of the UNIX kernel that executes user commands
- b. A process forked off at UNIX initialization to accept inputs from a user
- c. A system call executed by the UNIX startup routine to accept commands from users
- d. A library that implements various UNIX commands
- e. The UNIX keyboard device driver that interprets keyboard input

# Why Concurrency?



- Servers
  - Multiple connections handled simultaneously
- Parallel programs
  - To achieve better performance
- Programs with user interfaces
  - To achieve user responsiveness while doing computation
- Network and disk bound programs
  - To hide network/disk latency

# Definitions



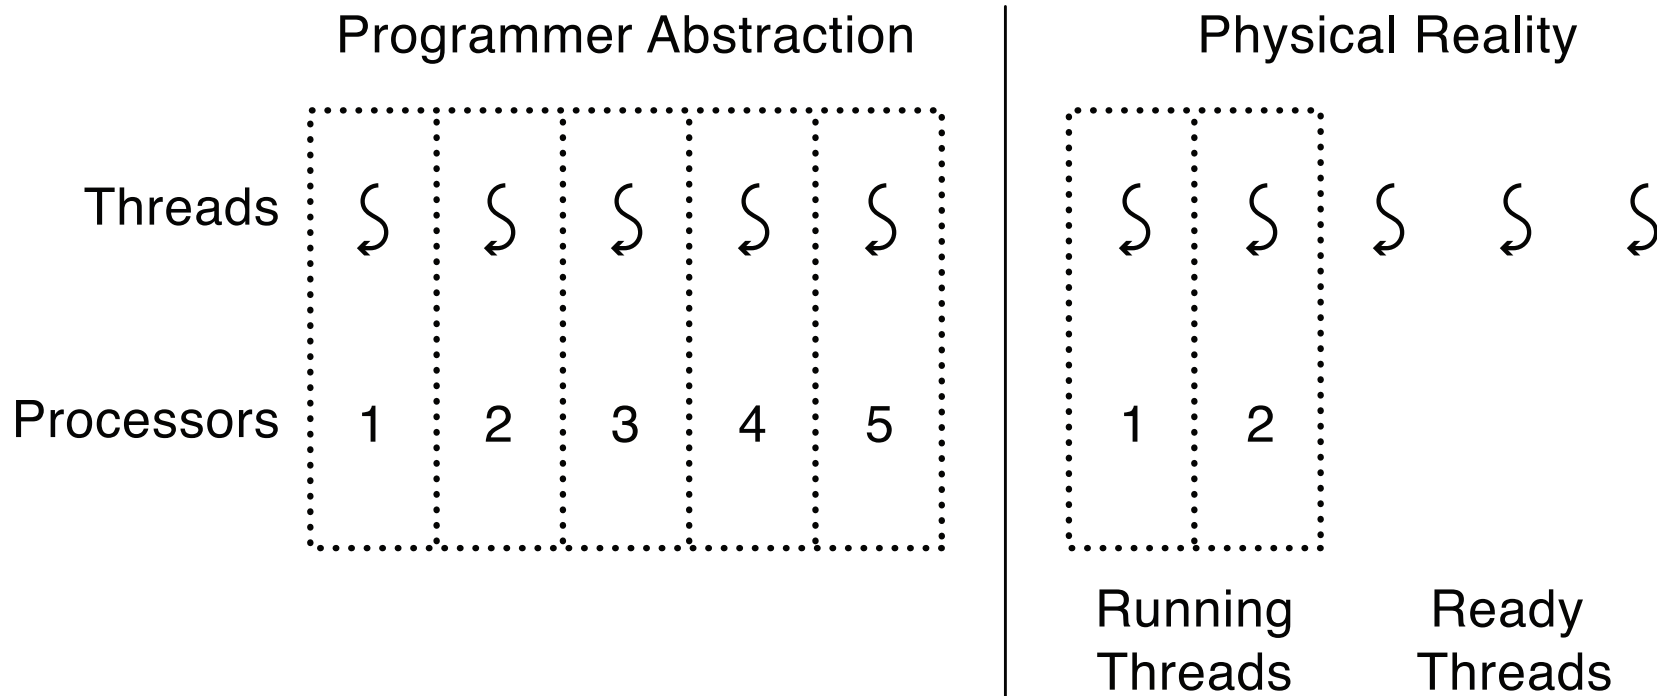
- Thread: A single execution sequence that represents a separately schedulable task.
  - *Single execution sequence*: intuitive and familiar programming model
  - *separately schedulable*: OS can run or suspend a thread at any time.
  - Schedulers operate over threads/tasks, both kernel and user threads.
- *Does the OS protect all threads from one another?*



# The Thread Abstraction



- Infinite number of processors
- Threads execute with variable speed



# Programmer vs. Processor View



## Programmer View

### Programmer's View

.  
.  
.  
x = x + 1;  
y = y + x;  
z = x + 5y;  
.  
.  
.

### Possible Execution #1

.  
.  
.  
x = x + 1;  
y = y + x;  
z = x + 5y;  
.  
.  
.

### Possible Execution #2

.  
.  
.  
x = x + 1;  
.....  
Thread is suspended.  
Other thread(s) run.  
Thread is resumed.  
.....  
y = y + x;  
z = x + 5y;

### Possible Execution #3

.  
.  
.  
x = x + 1;  
y = y + x;  
.....  
Thread is suspended.  
Other thread(s) run.  
Thread is resumed.  
.....  
z = x + 5y;


**Variable Speed: Program must anticipate all of these possible executions**

# Possible Executions




## Processor View


### One Execution



### Another Execution



### Another Execution



**Something to look forward to when we discuss scheduling!**

# Thread Ops



- `thread_create(thread, func, args)`  
Create a new thread to run `func(args)`
- `thread_yield()`  
Relinquish processor voluntarily
- `thread_join(thread)`  
In parent, wait for forked thread to exit, then return
- `thread_exit`  
Quit thread and clean up, wake up joiner if any

# Ex: threadHello



```
#define NTHREADS 10
thread_t threads[NTHREADS];
main() {
    for (i = 0; i < NTHREADS; i++) thread_create(&threads[i], &go, i);
    for (i = 0; i < NTHREADS; i++) {
        exitValue = thread_join(threads[i]);
        printf("Thread %d returned with %ld\n", i, exitValue);
    }
    printf("Main thread done.\n");
}
void go (int n) {
    printf("Hello from thread %d\n", n);
    thread_exit(100 + n);
    // REACHED?
}
```

# Ex: threadHello output



```
bash-3.2$ ./threadHello
```

```
Hello from thread 0  
Hello from thread 1  
Thread 0 returned 100  
Hello from thread 3  
Hello from thread 4  
Thread 1 returned 101  
Hello from thread 5  
Hello from thread 2  
Hello from thread 6  
Hello from thread 8  
Hello from thread 7  
Hello from thread 9  
Thread 2 returned 102  
Thread 3 returned 103  
Thread 4 returned 104  
Thread 5 returned 105  
Thread 6 returned 106  
Thread 7 returned 107  
Thread 8 returned 108  
Thread 9 returned 109  
Main thread done.
```

- Must “thread returned” print in order?
- What is maximum # of threads running when thread 5 prints hello?
- Minimum?

# Fork/Join Concurrency



- Threads can create children, and wait for their completion
- Data only shared before fork/after join
- Examples:
  - Web server: fork a new thread for every new connection
    - As long as the threads are completely independent
  - Merge sort
  - Parallel memory copy

# Ex: bzero

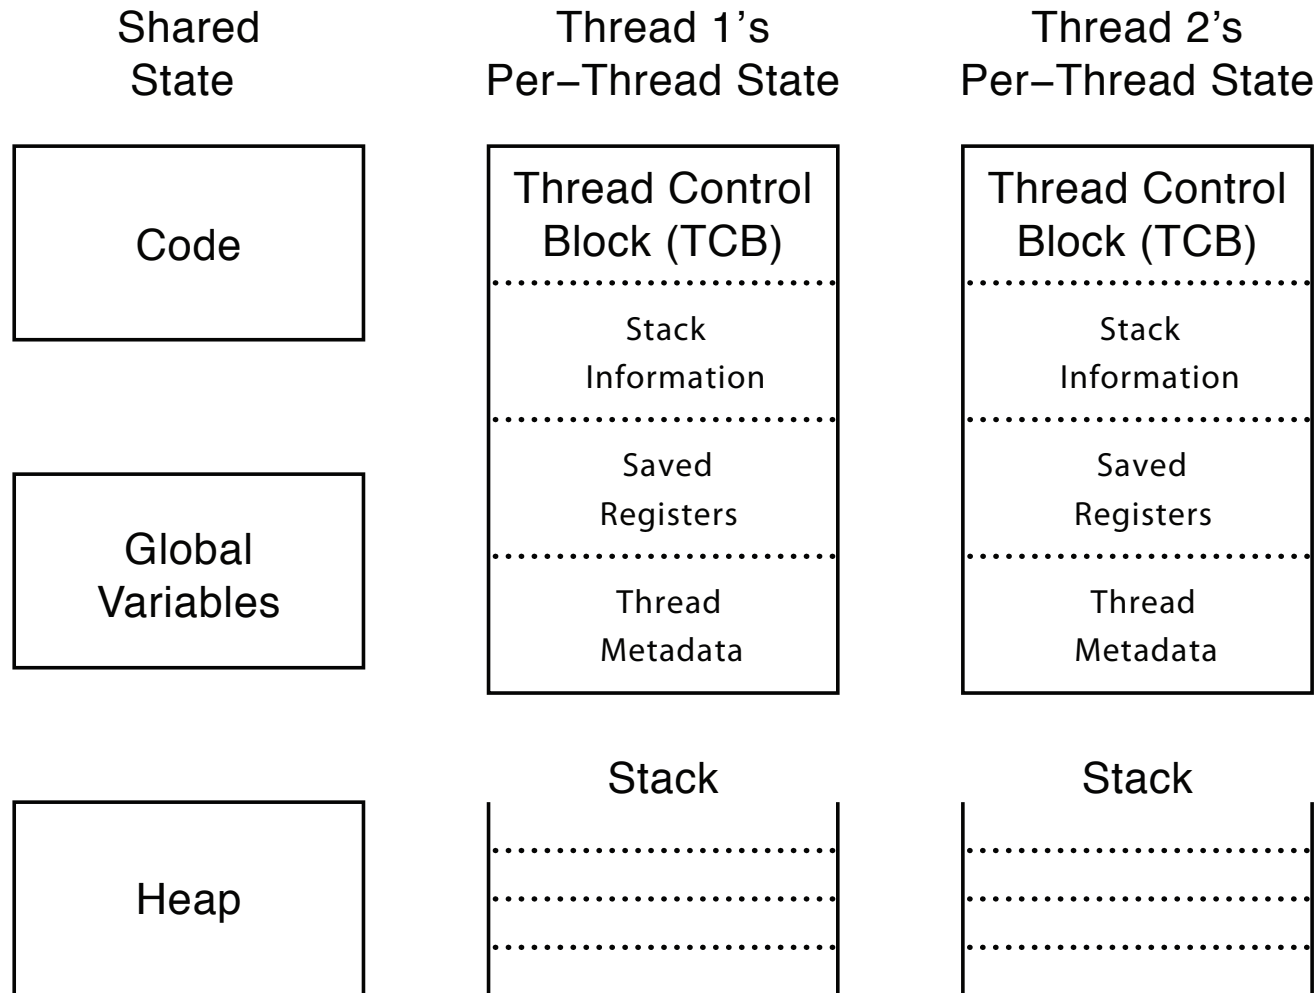


```
void blockzero (unsigned char *p, int length) {
    int i, j;
    thread_t threads[NTHREADS];
    struct bzeroparams params[NTHREADS];

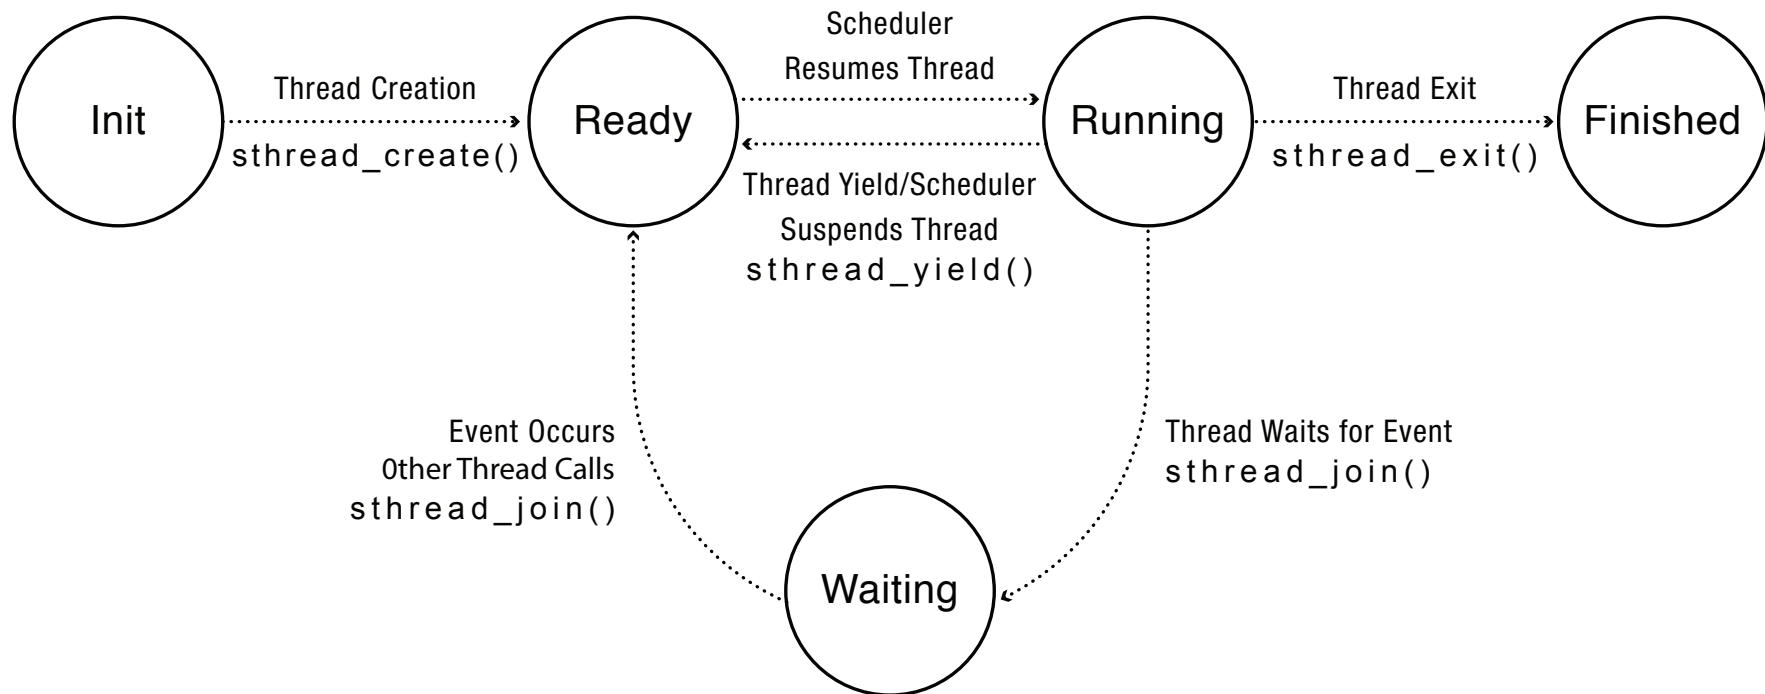
    // For simplicity, assumes length is divisible by NTHREADS.
    for (i = 0, j = 0; i < NTHREADS; i++, j += length/NTHREADS) {
        params[i].buffer = p + i * length/NTHREADS;
        params[i].length = length/NTHREADS;
        thread_create_p(&(threads[i]), &go, &params[i]);
    }
    for (i = 0; i < NTHREADS; i++) {
        thread_join(threads[i]);
    }
}
```



# Thread Data Structures



# Thread Lifecycle




# Thread Implementations



- Kernel threads
  - Thread abstraction only available to kernel
  - To the kernel, a kernel thread and a single threaded user process look quite similar
- Multithreaded processes using kernel threads
  - Kernel thread operations available via syscall
- User-level threads
  - Thread operations without system calls

# Multithreaded OS Kernel



# Implementing Threads



- Thread\_fork(func, args)
  - Allocate thread control block
  - Allocate stack
  - Build stack frame for base of stack (stub)
  - Put func, args on stack
  - Put thread on ready list
  - Will run sometime later (maybe right away!)
- stub(func, args):
  - Call (\*func)(args)
  - If return, call thread\_exit()

# Implementing Threads



- Thread\_Exit
  - Remove thread from the ready list so that it will never run again
  - Free the per-thread state allocated for the thread

# switchframe



How do we switch out thread state? (i.e., ctx switch)

```
# Save caller's register state
# NOTE: %eax, etc. are ephemeral
pushl %ebx
pushl %ebp
pushl %esi
pushl %edi

# Get offset of (struct thread, stack)
mov thread_stack_ofs, %edx
# Save current stack pointer to old thread's stack, if any.
movl SWITCH_CUR(%esp), %eax
movl %esp, (%eax,%edx,1)

# Change stack pointer to new thread's stack
# this also changes currentThread
movl SWITCH_NEXT(%esp), %ecx
movl (%ecx,%edx,1), %esp

# Restore caller's register state.
popl %edi
popl %esi
popl %ebp
popl %ebx
ret
```

# A subtlety



- Thread\_create puts new thread on ready list
- When it first runs, some thread calls switchframe
  - Saves old thread state to stack
  - Restores new thread state from stack
- Set up new thread's stack as if it had saved its state in switchframe
  - "returns" to stub at base of stack to run func



# Ex: Two Threads call Yield



## Thread 1's instructions

"return" from thread\_switch  
into stub  
call go  
call thread\_yield  
choose another thread  
call thread\_switch  
save thread 1 state to TCB  
load thread 2 state

return from thread\_switch  
return from thread\_yield  
call thread\_yield  
choose another thread  
call thread\_switch

## Thread 2's instructions

"return" from thread\_switch  
into stub  
call go  
call thread\_yield  
choose another thread  
call thread\_switch  
save thread 2 state to TCB  
load thread 1 state

## Processor's instructions

"return" from thread\_switch  
into stub  
call go  
call thread\_yield  
choose another thread  
call thread\_switch  
save thread 1 state to TCB  
load thread 2 state  
"return" from thread\_switch  
into stub  
call go  
call thread\_yield  
choose another thread  
call thread\_switch  
save thread 2 state to TCB  
load thread 1 state  
return from thread\_switch  
return from thread\_yield  
call thread\_yield  
choose another thread  
call thread\_switch



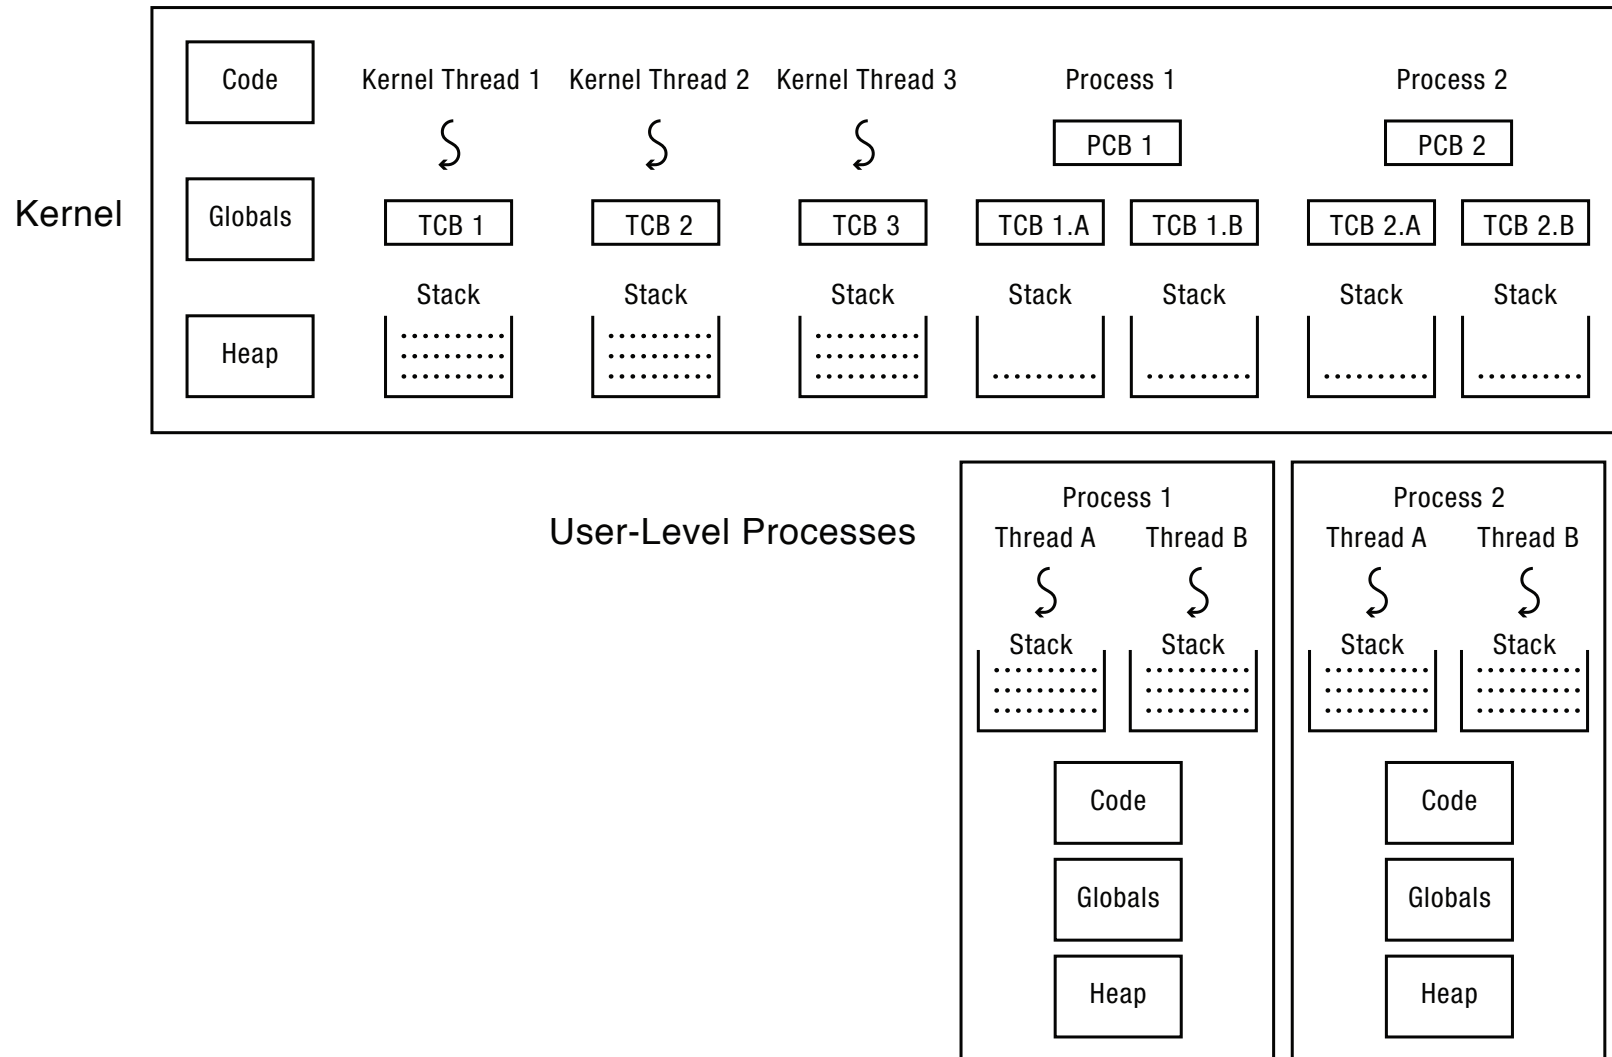
## Take 1:

- User thread = kernel thread (Linux, MacOS)
  - System calls for thread fork, join, exit (and lock, unlock,...)
  - Kernel does context switch
  - Simple, but a lot of transitions between user and kernel mode

# Multi-threaded User Processes



## Take 1:





## Take 2:

- Green threads (early Java)
  - User-level library, within a single-threaded process
  - Library does thread context switch
  - Preemption via upcall/UNIX signal on timer interrupt
  - Use multiple processes for parallelism
  - Shared memory region mapped into each process



## Take 3:


- Scheduler activations (Windows 8):
  - Kernel allocates processors to user-level library
  - Thread library implements context switch
  - Thread library decides what thread to run next
- Upcall whenever kernel needs a user-level scheduling decision:
  - Process assigned a new processor
  - Processor removed from process
  - System call blocks in kernel

# Multi-threaded User Processes



## Take 3: (What's old is new again)

Multiple user threads  
on a kernel thread



M:N model multiplexes N user-level threads onto M kernel-level threads

Good idea? Bad Idea?

# Question



Compare event-driven programming with multithreaded concurrency. Which is better in which circumstances, and why?