



CS 423

Operating System Design: OS support for Synchronization

Professor Adam Bates
Fall 2018

Goals for Today



- Learning Objectives:
 - Discuss OS support for Synchronization
- Announcements:
 - MP1 available on Compass2G. **Due February 19th!**
 - Next week — scheduling (skipping chapter 6 for now)



Reminder: Please put away devices at the start of class

Implementing Synchronization



- Take 1: using memory load/store
 - See too much milk solution/Peterson's algorithm
- Take 2: *(corrected from last class!)*

```
Lock::acquire() {  
    disableInterrupts();  
}
```

```
Lock::release() {  
    enableInterrupts();  
}
```

Above solution “works” on single processor...

Queueing Lock Implementation (1 Proc)



```
Lock::acquire() {
    disableInterrupts();
    if (value == BUSY) {
        waiting.add(myTCB);
        myTCB->state = WAITING;
        next = readyList.remove();
        switch(myTCB, next);
        myTCB->state = RUNNING;
    } else {
        value = BUSY;
    }
    enableInterrupts();
}
```

```
Lock::release() {
    disableInterrupts();
    if (!waiting.Empty()) {
        next = waiting.remove();
        next->state = READY;
        readyList.add(next);
    } else {
        value = FREE;
    }
    enableInterrupts();
}
```

Question



Why won't this work for multiprocessing?



- Read-modify-write instructions
 - Atomically read a value from memory, operate on it, and then write it back to memory
 - Intervening instructions prevented in hardware
- Examples
 - Test and set
 - Intel: xchgb, lock prefix
 - Compare and swap
- Any of these can be used for implementing locks and condition variables!

Spinlocks



- A spinlock is a lock where the processor waits in a loop for the lock to become free
 - Assumes lock will be held for a short time
 - Used to protect the CPU scheduler and to implement locks

```
Spinlock::acquire() {  
    while (testAndSet(&lockValue) == BUSY)  
        ;  
}
```

```
Spinlock::release() {  
    lockValue = FREE;  
    memorybarrier();  
}
```

How many spinlocks?



Neat. So how many spinlocks do we need?

How many spinlocks?



- Many data structures requiring synchronization!
 - Queue of waiting threads on lock X
 - Queue of waiting threads on lock Y
 - List of threads ready to run
- One spinlock per kernel?
 - Spinlock becomes bottleneck!
- Instead:
 - One spinlock per lock
 - One spinlock for the scheduler ready list
 - *Per-core ready list: one spinlock per core*

What thread is currently running?



- Thread scheduler needs to find the TCB of the currently running thread
 - To suspend and switch to a new thread
 - To check if the current thread holds a lock before acquiring or releasing it
- On a uniprocessor, easy: just use a global
- On a multiprocessor, various methods:
 - Compiler dedicates a register (e.g., r31 points to TCB running on the this CPU; each CPU has its own r31)
 - If hardware has a special per-processor register, use it
 - Fixed-size stacks: put a pointer to the TCB at the bottom of its stack
 - Find it by masking the current stack pointer



Lock implementation —

```
Lock::acquire() {  
    disableInterrupts();  
    spinLock.acquire();  
    if (value == BUSY) {  
        waiting.add(myTCB);  
        suspend(&spinlock);  
    } else {  
        value = BUSY;  
    }  
    spinLock.release();  
    enableInterrupts();  
}
```

```
Lock::release() {  
    disableInterrupts();  
    spinLock.acquire();  
    if (!waiting.Empty()) {  
        next = waiting.remove();  
        scheduler->makeReady(next);  
    } else {  
        value = FREE;  
    }  
    spinLock.release();  
    enableInterrupts();  
}
```



Scheduler implementation —

```
Sched::suspend(SpinLock *lock) {  
    TCB *next;  
  
    disableInterrupts();  
    schedSpinLock.acquire();  
    lock->release();  
    myTCB->state = WAITING;  
    next = readyList.remove();  
    thread_switch(myTCB, next);  
    myTCB->state = RUNNING;  
    schedSpinLock.release();  
    enableInterrupts();  
}  
  
Sched::makeReady(TCB *thread) {  
  
    disableInterrupts();  
    schedSpinLock.acquire();  
    readyList.add(thread);  
    thread->state = READY;  
    schedSpinLock.release();  
    enableInterrupts();  
}
```

Locks for user space??



- Kernel-managed threads
 - Manage data structures in kernel space
 - System calls to communicate w/ scheduler
- User-managed threads
 - Implement functionality in thread library
 - Can't disable interrupts, but can temporarily disable upcalls to avoid preemption in library scheduler, etc.

Locks in Linux



- Most locks are free most of the time. Linux implementation takes advantage of this fact!
- Fast path:
 - If lock is FREE, and no one is waiting, two instructions to acquire the lock
 - If no one is waiting, two instructions to release the lock
- Slow path
 - If lock is BUSY or someone is waiting, use multiproc impl.
- User-level locks also optimized:
 - Fast path: count is mapped to proc address space, no sys call needed when count is 0.
 - Slow path: system call to kernel, use kernel lock when waiting thread

Locks in Linux



Lock struct contains 3 (not two) states...

```
struct mutex {  
    /* 1: unlocked ;  
       0: locked;  
       negative : locked, possible waiters */  
    atomic_t count;  
    spinlock_t wait_lock;  
    struct list_head wait_list;  
};
```

Lock acquire code is a macro (to avoid proc call)...

```
lock decl (%eax)                // atomic decrement  
                                // %eax is pointer to count  
jns 1f                          // jump if not signed  
                                // (i.e., if value is now 0)  
call slowpath_acquire  
1: ...
```

Semaphores



- Semaphore has a non-negative integer value
 - P() atomically waits for value to become > 0 , then decrements
 - V() atomically increments value (waking up waiter if needed)
- Semaphores are like integers except:
 - Only operations are P and V
 - Operations are atomic
 - If value is 1, two P's will result in value 0 and one waiter
- Semaphores are useful for
 - Unlocked wait: interrupt handler, fork/join

Compare Implementations



Lock implementation —

```
Lock::acquire() {
    disableInterrupts();
    spinLock.acquire();
    if (value == BUSY) {
        waiting.add(myTCB);
        suspend(&spinlock);
    } else {
        value = BUSY;
    }
    spinLock.release();
    enableInterrupts();
}
```

```
Lock::release() {
    disableInterrupts();
    spinLock.acquire();
    if (!waiting.Empty()) {
        next = waiting.remove();
        scheduler->makeReady(next);
    } else {
        value = FREE;
    }
    spinLock.release();
    enableInterrupts();
}
```

Compare Implementations



Semaphore implementation —

```
Semaphore::P() {  
    disableInterrupts();  
    spinLock.acquire();  
    if (value == 0) {  
        waiting.add(myTCB);  
        suspend(&spinlock);  
    } else {  
        value--;  
    }  
    spinLock.release();  
    enableInterrupts();  
}
```

```
Semaphore::V() {  
    disableInterrupts();  
    spinLock.acquire();  
    if (!waiting.Empty()) {  
        next = waiting.remove();  
        scheduler->makeReady(next);  
    } else {  
        value++;  
    }  
    spinLock.release();  
    enableInterrupts();  
}
```

Semaphores Harmful?



- Semaphores conflate mutual the roles of locks and condition variables (mutual exclusion, shared data).
 - Simpler code verification: *prove every lock is eventually unlocked.*
- Semaphores have state!
 - What does value=3 mean? Programmer must carefully map object state to semaphore value.
 - CVs, in contrast, allows us to wait on arbitrary state. A better abstraction.
- However, semaphores have good uses, including...
 - Unlocked waits, e.g., interrupt handler that synchronizes communication between I/O device and waiting threads.

Semaphore Bounded Buffer



```
get() {  
    fullSlots.P();  
    mutex.P();  
    item = buf[front % MAX];  
    front++;  
    mutex.V();  
    emptySlots.V();  
    return item;  
}
```

```
put(item) {  
    emptySlots.P();  
    mutex.P();  
    buf[last % MAX] = item;  
    last++;  
    mutex.V();  
    fullSlots.V();  
}
```

Initially: front = last = 0; MAX is buffer capacity
mutex = 1; emptySlots = MAX; fullSlots = 0;



How can we implement Condition Variables using semaphores?

Take 1:

```
wait(lock) {  
    lock.release();  
    semaphore.P();  
    lock.acquire();  
}  
signal() {  
    semaphore.V();  
}
```

Problems?



How can we implement Condition Variables using semaphores?

Take 2:

```
wait(lock) {
    lock.release();
    semaphore.P();
    lock.acquire();
}
signal() {
    if (semaphore is not empty)
        semaphore.V();
}
```

Problems?



How can we implement Condition Variables using semaphores?

Take 2:

```
wait(lock) {
    semaphore = new Semaphore;
    queue.Append(semaphore);    // queue of waiting threads
    lock.release();
    semaphore.P();
    lock.acquire();
}
signal() {
    if (!queue.Empty()) {
        semaphore = queue.Remove();
        semaphore.V();    // wake up waiter
    }
}
```

Problems?

Implementing CVs w/ Semaphores



Implementation used for Microsoft Windows before native support was offered:

Take 4:

```
//Put thread on queue of waiting threads...
```

```
void CV::wait(Lock *lock){  
    semaphore = new Semaphore(0);  
    waitQueue. Append(semaphore)  
    lock.release();  
    semaphore.P();  
    lock.acquire();  
}
```

```
//Wake up one waiter if any.
```

```
void CV::signal() {  
    if(!waitQueue.isEmpty()) {  
        semaphore = queue.Remove();  
        semaphore.V();  
    }  
}
```