

# Goals for Today



- Learning Objective:
  - Define a taxonomy for virtualization architectures
- Announcements, etc:
  - C4 Submission for 3/9, 3/15 open (sorry for the delay)
  - Request for “Informal Early Feedback” forthcoming, probably on Monday
  - Midterm debrief forthcoming, probably on Wednesday



**Reminder:** Please put away devices at the start of class



# CS 423

## Operating System Design: Virtual Machines

Professor Adam Bates  
Spring 2017

# Virtual Machines



- What is a virtual machine?
  - Examples?
- Benefits?



- Creation of an isomorphism that maps a virtual guest system to a real host:
  - Maps guest state  $S$  to host state  $V(S)$
  - For any sequence of operations on the guest that changes guest state  $S1$  to  $S2$ , there is a sequence of operations on the host that maps state  $V(S1)$  to  $V(S2)$

# Important Interfaces



- Application programmer interface (API):
  - High-level language library such as `clib`
- Application binary interface (ABI):
  - User instructions (User ISA)
  - System calls
- Hardware-software interface:
  - Instruction set architecture (ISA)

# What's a machine?



- Machine is an entity that provides an interface
  - Language view:
    - Machine = Entity that provides the API
  - Process view:
    - Machine = Entity that provides the ABI
  - Operating system view:
    - Machine = Entity that provides the ISA

# What's a virtual machine?



- Virtual machine is an entity that emulates a guest interface on top of a host machine
  - Language view:
    - Virtual machine = Entity that emulates an API (e.g., JAVA) on top of another
    - Virtualizing software = compiler/interpreter
  - Process view:
    - Machine = Entity that emulates an ABI on top of another
    - Virtualizing software = runtime
  - Operating system view:
    - Machine = Entity that emulates an ISA
    - Virtualizing software = virtual machine monitor (VMM)

# Purpose of a VM



- Emulation
  - Create the illusion of having one type of machine on top of another
- Replication (/ Multiplexing)
  - Create the illusion of multiple independent smaller guest machines on top of one host machine (e.g., for security/isolation, or scalability/sharing)
- Optimization
  - Optimize a generic guest interface for one type of host



# Types of VMs



- Emulate (ISA/ABI/API) for purposes of (Emulation/Replication/Optimization) on top of (the same/different) one.

# Types of VMs



- Emulate (ISA/ABI/API) for purposes of (Emulation/Replication/Optimization) on top of (the same/different) one.
  - Process/language virtual machines (emulate ABI/API)
  - System virtual machines (emulate ISA)

# Types of VMs



- Emulate (ISA/ABI/API) for purposes of (Emulation/Replication/Optimization) on top of (the same/different) one.
  - Process/language virtual machines (emulate ABI/API)
  - System virtual machines (emulate ISA)

# Ex1: Multiprogramming



- Emulate what interface?
- For what purpose?
- On top of what?



- Emulate one ABI on top of another
  - Emulate a Intel IA-32 running Windows on top of PowerPC running MacOS (i.e., run a process compiled for IA-32/Windows on PowerPC/MacOS)
    - Interpreters: Pick one guest instruction at a time, update (simulated) host state using a set of host instructions
    - Binary translation: Do the translation in one step, not one line at a time. Run the translated binary

# Writing an Emulator



- Create a simulator data structure to represent:
  - Guest memory
    - Guest stack
    - Guest heap
  - Guest registers
- Inspect each binary instruction (machine instruction or system call)
  - Update the data structures to reflect the effect of the instruction

# Ex2: Binary Optimization



- Emulate one ABI on top of itself for purposes of optimization
  - Run the process binary, collect profiling data, then implement it more efficiently on top of the same machine/OS interface.

# Ex3: Language VMs



- Emulate one API on top of a set of different ABIs
  - Compile guest API to intermediate form (e.g., JAVA source to JAVA bytecode)
  - Interpret the bytecode on top of different host ABIs
- Examples:
  - JAVA
  - Microsoft Common Language Infrastructure (CLI), the foundation of .NET



# Types of VMs



- Emulate (ISA/ABI/API) for purposes of (Emulation/Replication/Optimization) on top of (the same/different) one.
  - Process/language virtual machines (emulate ABI/API)
  - System virtual machines (emulate ISA)

# Types of VMs



- Emulate (ISA/ABI/API) for purposes of (Emulation/Replication/Optimization) on top of (the same/different) one.
  - Process/language virtual machines (emulate ABI/API)
  - System virtual machines (emulate ISA)



- Implement VMM (ISA emulation) on bare hardware
  - Efficient
  - Must wipe out current operating system to install
  - Must support drivers for VMM
- Implement VMM on top of a host OS (Hosted VM)
  - Less efficient
  - Easy to install on top of host OS
  - Leverages host OS drivers



- Implement VMM (ISA emulation) on bare hardware
  - Efficient
  - Must wipe out current operating system to install
  - Must support drivers for VMM
- Implement VMM on top of a host OS (Hosted VM)
  - Less efficient
  - Easy to install on top of host OS
  - Leverages host OS drivers

**TYPE ONE  
HYPERVISOR**

**TYPE TWO  
HYPERVISOR**

# Whole System VMs



- Emulate one ISA on top of another
  - Typically runs on top of host OS (e.g., install Windows compiled for IA-32 on top of MacOS running on PowerPC)
  - Note: this is different from a process virtual machine that emulates the Windows interface and user IA-32 instructions on top of MacOS running on PowerPC



- Problem: Emulate guest ISA on host ISA



- Problem: Emulate guest ISA on host ISA
- Solution: Basic Interpretation

```
inst = code (PC)
opcode = extract_opcode (inst)
switch (opcode) {
    case opcode1 : call emulate_opcode1 ()
    case opcode2 : call emulate_opcode2 ()
    ...
}
```



- Problem: Emulate guest ISA on host ISA
- Solution: Basic Interpretation

```
new          inst = code (PC)
              opcode = extract_opcode (inst)
              routineCase = dispatch (opcode)
              jump routineCase
              ...
routineCase  call routine_address
              jump new
```



# Threaded Implementation



[ body of emulate\_opcode1 ]

inst = code (PC)

opcode = extract\_opcode (inst)

routine\_address = dispatch (opcode)

jump routine\_address

[ body of emulate\_opcode2]

inst = code (PC)

opcode = extract\_opcode (inst)

routine\_address = dispatch (opcode)

jump routine\_address