

Goals for Today



- Learning Objective:
 - Define a taxonomy for virtualization architectures
- Announcements, etc:
 - “Informal Early Feedback” at end of class today
 - Midterm debrief forthcoming *on Friday*
 - MP2 extension: **now due on March 23rd**



Reminder: Please put away devices at the start of class



CS 423

Operating System Design: Virtual Machines

Professor Adam Bates
Spring 2017

What's a virtual machine?



- Virtual machine is an entity that emulates a guest interface on top of a host machine
 - Language view:
 - Virtual machine = Entity that emulates an API (e.g., JAVA) on top of another
 - Virtualizing software = compiler/interpreter
 - Process view:
 - Machine = Entity that emulates an ABI on top of another
 - Virtualizing software = runtime
 - Operating system view:
 - Machine = Entity that emulates an ISA
 - Virtualizing software = virtual machine monitor (VMM)

Different views == who are we trying to fool??

Purpose of a VM



- Emulation
 - Create the illusion of having one type of machine on top of another
- Replication (/ Multiplexing)
 - Create the illusion of multiple independent smaller guest machines on top of one host machine (e.g., for security/isolation, or scalability/sharing)
- Optimization
 - Optimize a generic guest interface for one type of host

Types of VMs



- Emulate (ISA/ABI/API) for purposes of (Emulation/Replication/Optimization) on top of (the same/different) one.
 - Process/language virtual machines (emulate ABI/API)
 - System virtual machines (emulate ISA)



- Language VMs
 - Emulate same API as host (e.g., application profiling?)
 - Emulate different API than host (e.g., Java API)
- Process VMs
 - Emulate same ABI as host (e.g., multiprogramming)
 - Emulate different ABI than host (e.g., Java VM, MAME)
- System VMs
 - Emulate same ISA as host (e.g., KVM, VBox, Xen)
 - Emulate different ISA than host (e.g., MULTICS simulator)

Point of Clarification



- Emulation: General technique for performing any kind of virtualization (API/ABI/ISA)
- Not to be confused with *Emulator* in the colloquial sense (e.g., Video Game Emulator), which often refers to ABI emulation.

Writing an Emulator



- Problem: Emulate guest ISA on host ISA

Writing an Emulator



- Problem: Emulate guest ISA on host ISA
- Create a simulator data structure to represent:
 - Guest memory
 - Guest stack
 - Guest heap
 - Guest registers
- Inspect each binary instruction (machine instruction or system call)
 - Update the data structures to reflect the effect of the instruction



- Problem: Emulate guest ISA on host ISA
- Solution: Basic Interpretation, switch on opcode

```
inst = code (PC)
opcode = extract_opcode (inst)
switch (opcode) {
    case opcode1 : call emulate_opcode1 ()
    case opcode2 : call emulate_opcode2 ()
    ...
}
```



- Problem: Emulate guest ISA on host ISA
- Solution: Basic Interpretation

```
new          inst = code (PC)
              opcode = extract_opcode (inst)
              routineCase = dispatch (opcode)
              jump routineCase
              ...
routineCase  call routine_address
              jump new
```

Threaded Interpretation...



[body of emulate_opcode1]

inst = code (PC)

opcode = extract_opcode (inst)

routine_address = dispatch (opcode)

jump routine_address

[body of emulate_opcode2]

inst = code (PC)

opcode = extract_opcode (inst)

routine_address = dispatch (opcode)

jump routine_address

Note: Extracting Opcodes



- `extract_opcode (inst)`
 - Opcode may have options
 - Instruction must extract and combine several bit ranges in the machine word
 - Operands must also be extracted from other bit ranges
- Pre-decoding
 - Pre-extract the opcodes and operands for all instructions in program.
 - Put them on byte boundaries (intermediate code)
 - Must maintain two program counters. Why?

Note: Extracting Opcodes



lwz r1, 8(r2)

add r3, r3, r1

stw r3, 0(r4)

07		
1	2	08

08		
3	1	03

37		
3	4	00

Direct Threaded Impl.



- Replace opcode with address of emulating routine

Routine_address07		
1	2	08

Routine_address08		
3	1	03

Routine_address37		
3	4	00



- Emulation:
 - Guest code is traversed and instruction classes are mapped to routines that emulate them on the target architecture.
- Binary translation:
 - The entire program is translated into a binary of another architecture.
 - Each binary source instruction is emulated by some binary target instructions.

Challenges



- Can we really just read the source binary and translate it statically one instruction at a time to a target binary?
 - What are some difficulties?



- Code discovery and dynamic translation
 - How to tell whether something is code or data?
 - Consider a jump instruction: Is the part that follows it code or data?
- Code location problem
 - How to map source program counter to target program counter?
 - Can we do this without having a table as long as the program for instruction-by-instruction mapping?

Things to Notice



- You only need source-to-target program counter mapping for locations that are *targets of jumps*. Hence, only map those locations.
- You always know that something is an instruction (not data) in the source binary if the source program counter eventually ends up pointing to it.
- The problem is: You do not know targets of jumps (and what the program counter will end up pointing to) at static analysis time!
 - Why?



- Incremental Pre-decoding and Translation
 - As you execute a source binary block, translate it into a target binary block (this way you know you are translating valid instructions)
 - Whenever you jump:
 - If you jump to a new location: start a new target binary block, record the mapping between source program counter and target program counter in map table.
 - If you jump to a location already in the map table, get the target program counter from the table
 - Jumps must go through an emulation manager. Blocks are translated (the first time only) then executed directly thereafter



Informal Early Feedback