

# Goals for Today



- Learning Objective:
  - Learn about directory structures
  - Run through some disk performance exercises
- Announcements, etc:
  - C4: I removed 2 papers from your reading list
    - No longer need to read Multics Security Eval or SELinux
  - MP3 is out! Due **April 18th**.



**Reminder:** Please put away devices at the start of class



# CS 423

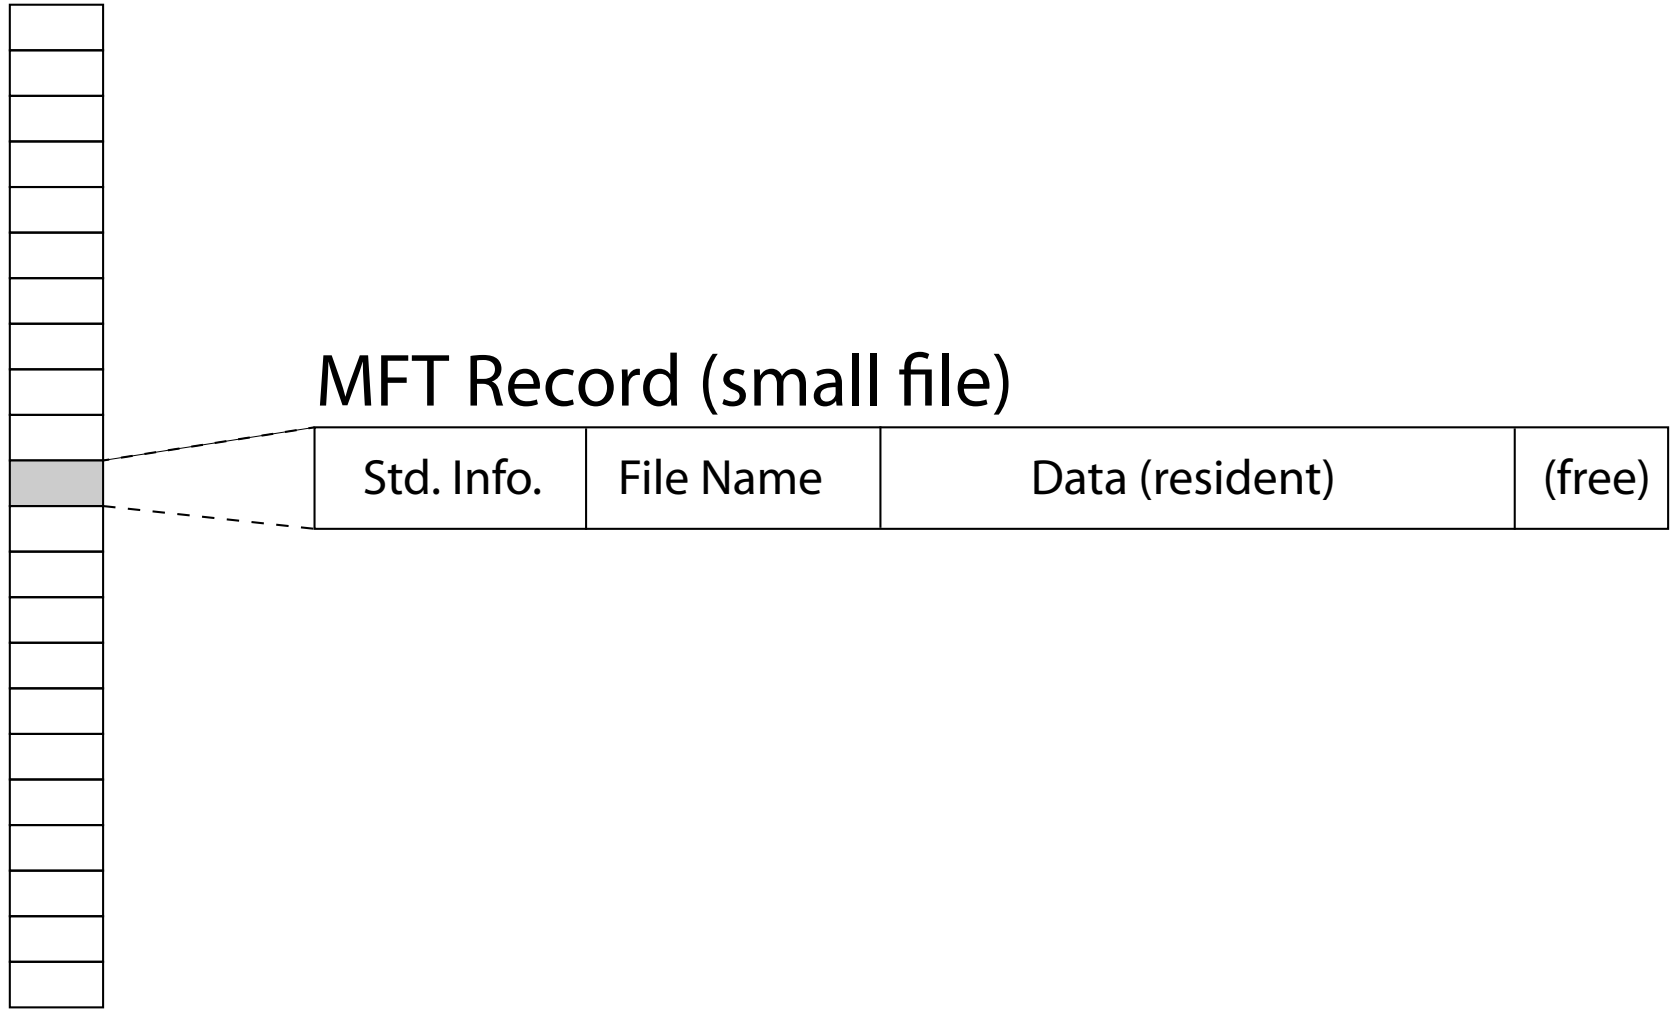
## Operating System Design: Directory Structures & Disk Performance

Professor Adam Bates  
Spring 2018

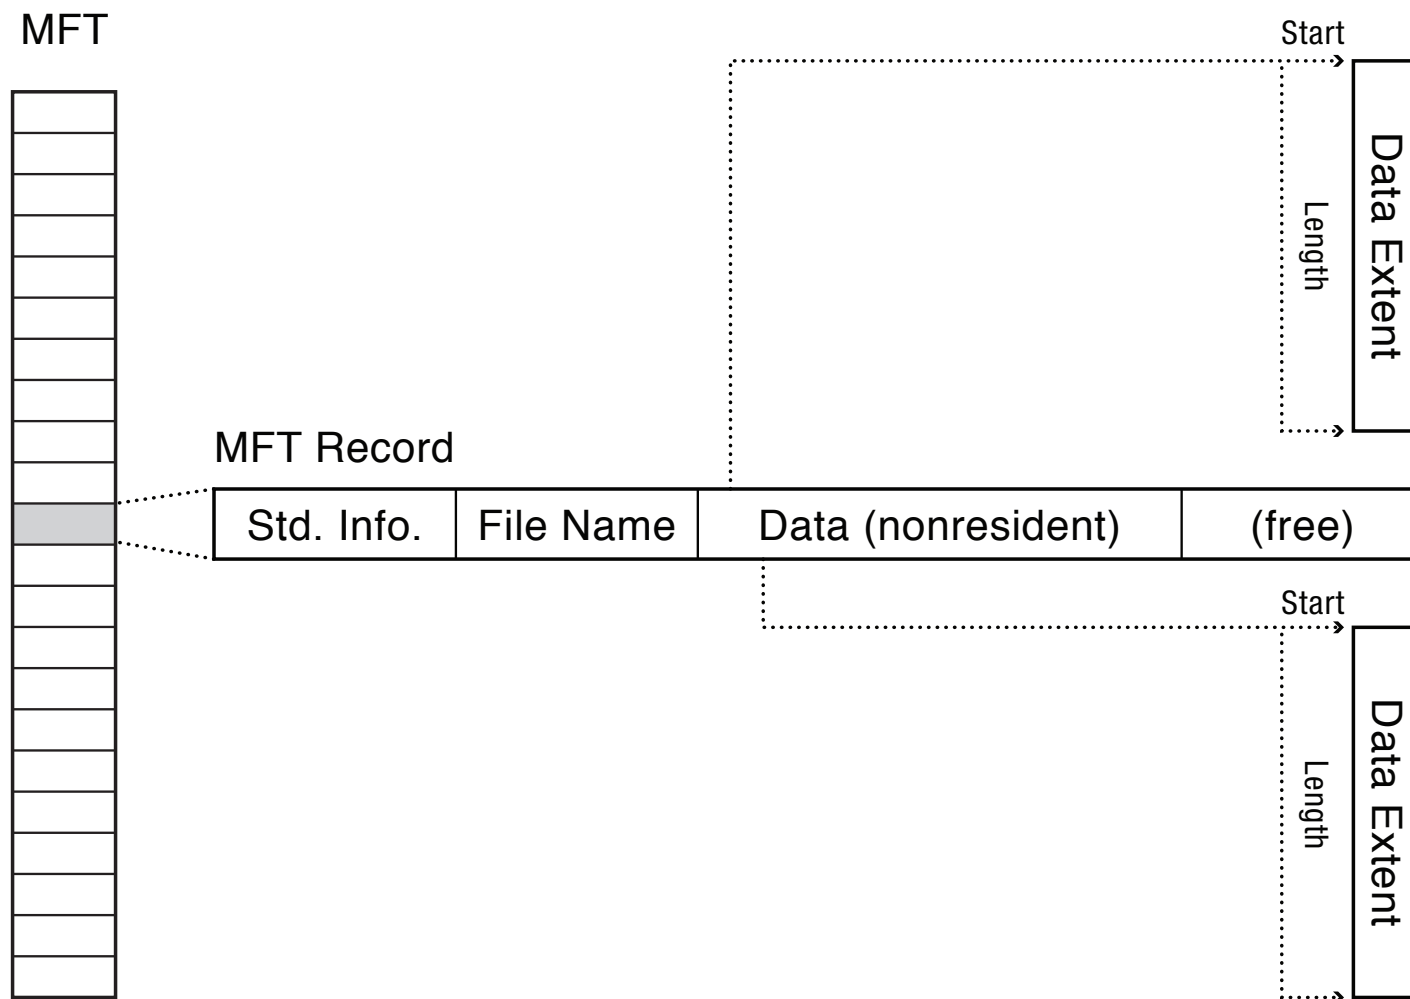


- Master File Table
  - Flexible 1KB storage for metadata and data
- Extents
  - Block pointers cover runs of blocks
  - Similar approach in linux (ext4)
  - File create can provide hint as to size of file
- Journalling for reliability

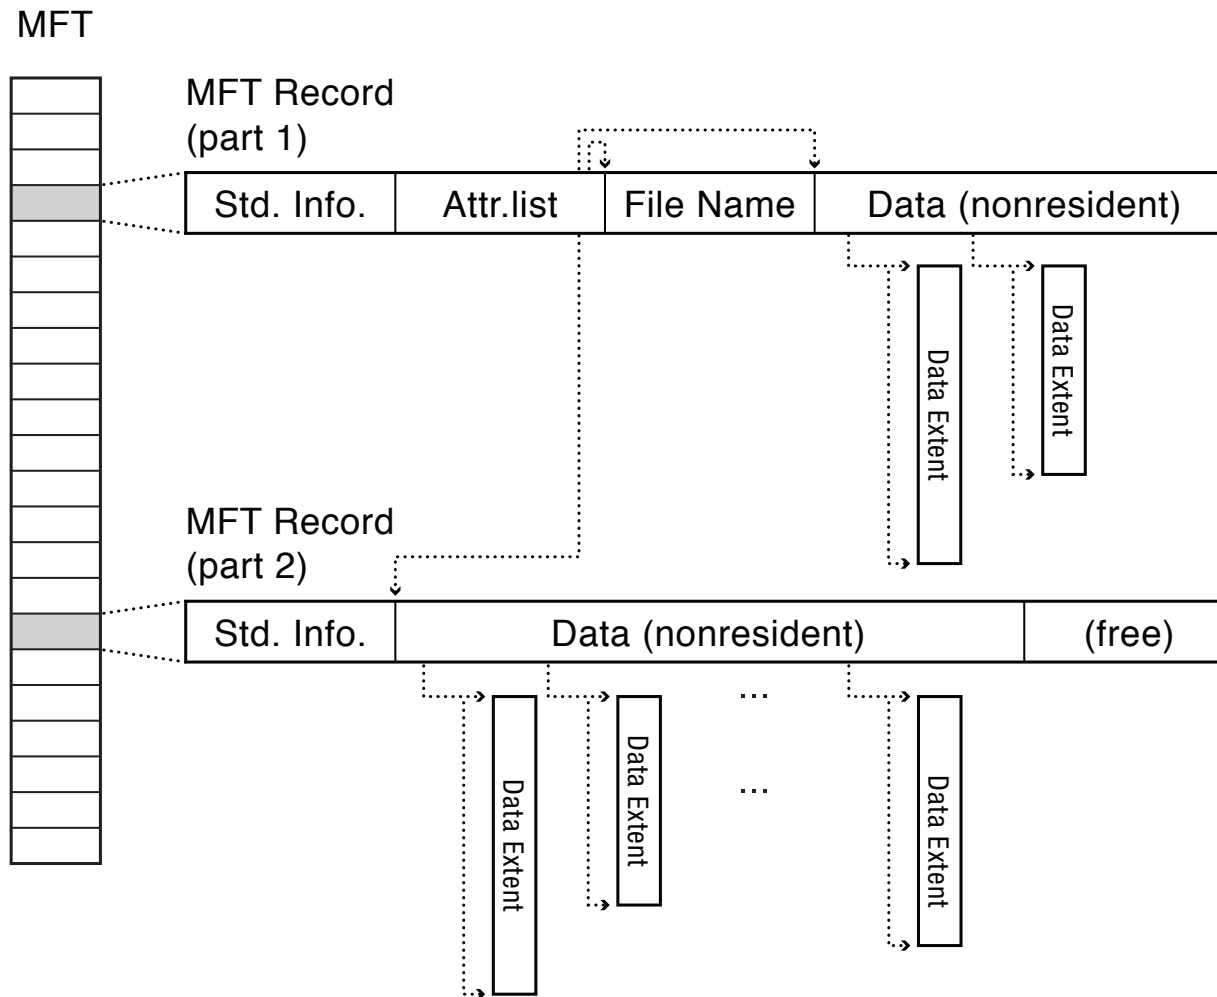
## Master File Table



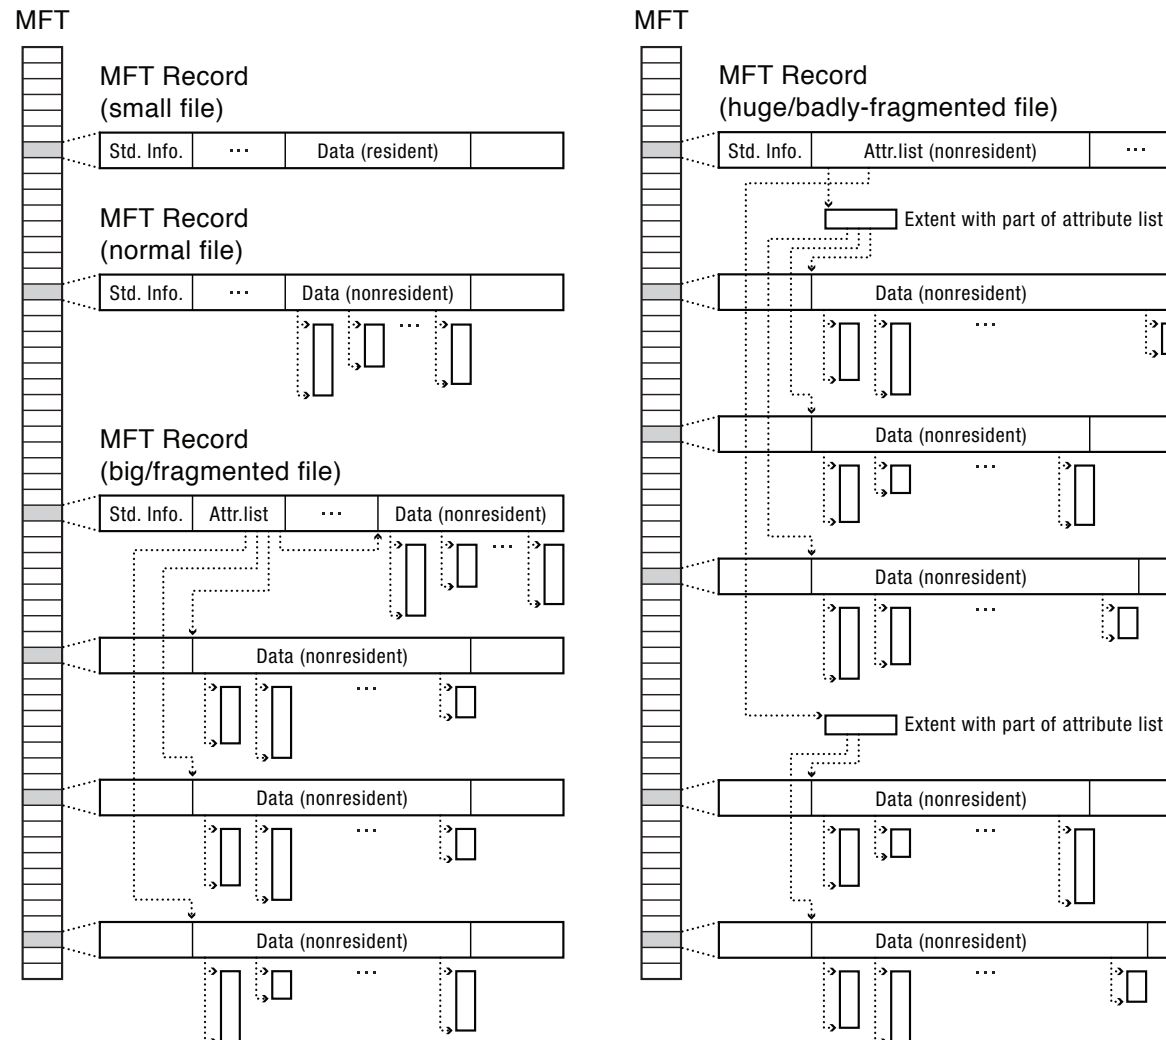
# NTFS



# NTFS Indirect Block



# NTFS



# Named Data






# Directory Structures



- maps symbolic names into logical file names
  - search
  - create file
  - list directory
  - backup, archival, file migration
- Directories are also files!



# Directory Internals







- Directories are also files! Special files
  - Denoted by “File Type” in inode
  - Hold access paths to the files in the file system
- They have data blocks (inodes) on disk
- Enables support for very large directories
- The partition superblock points to the inode of the root directory.

# Single-level Directory




## Directory

Name of File			
--------------------	--	--	--



## Files


# Tree-Structured Directories



# Directory Layout



- Represent directory as a list of files
- Linear search to find filename
- Suitable for small directories






# **What can we do to improve efficient in large directories?**

# B Trees



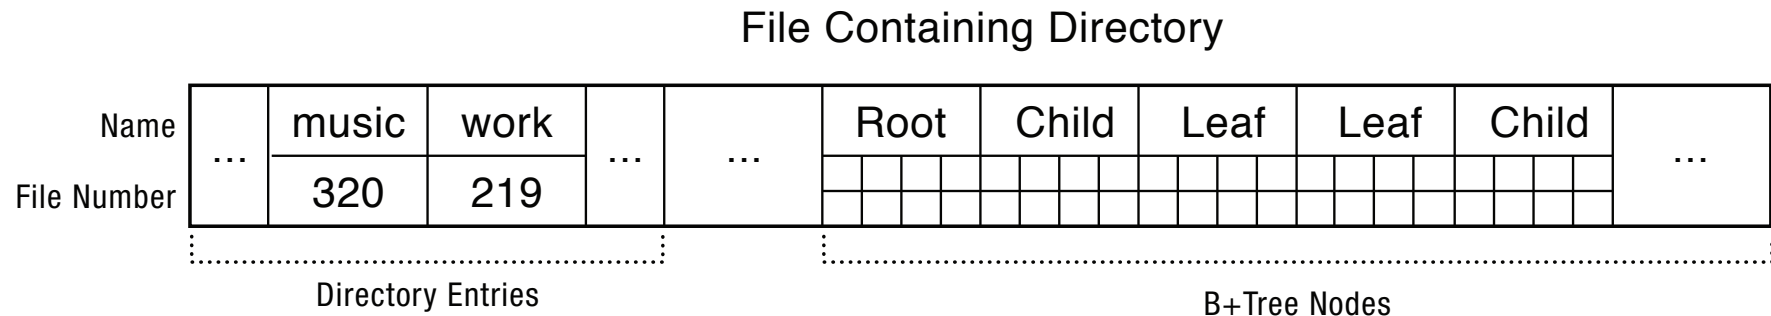
- Logarithmic search to find filename
- Suitable for large directories



# B Trees

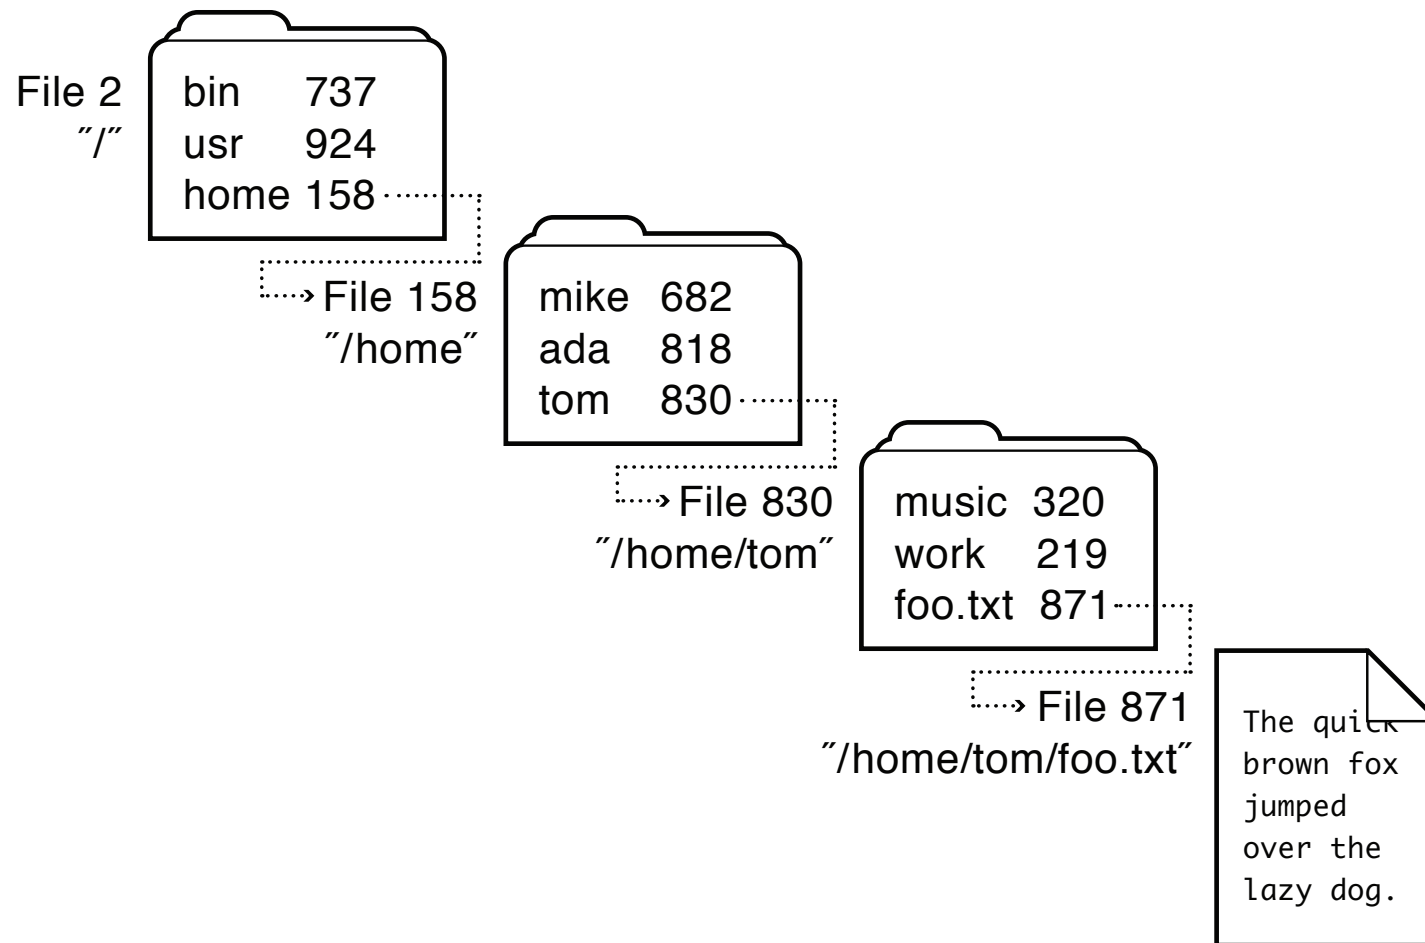


- Logarithmic search to find filename
- Suitable for large directories





# Recursive Filename Lookup



# Tree-Structured Directories




- arbitrary depth of directories
- leaf nodes are files
- interior nodes are directories
- path name lists nodes to traverse to find node
- use absolute paths from root
- use relative paths from current working directory pointer

# Tree-Structured Directories



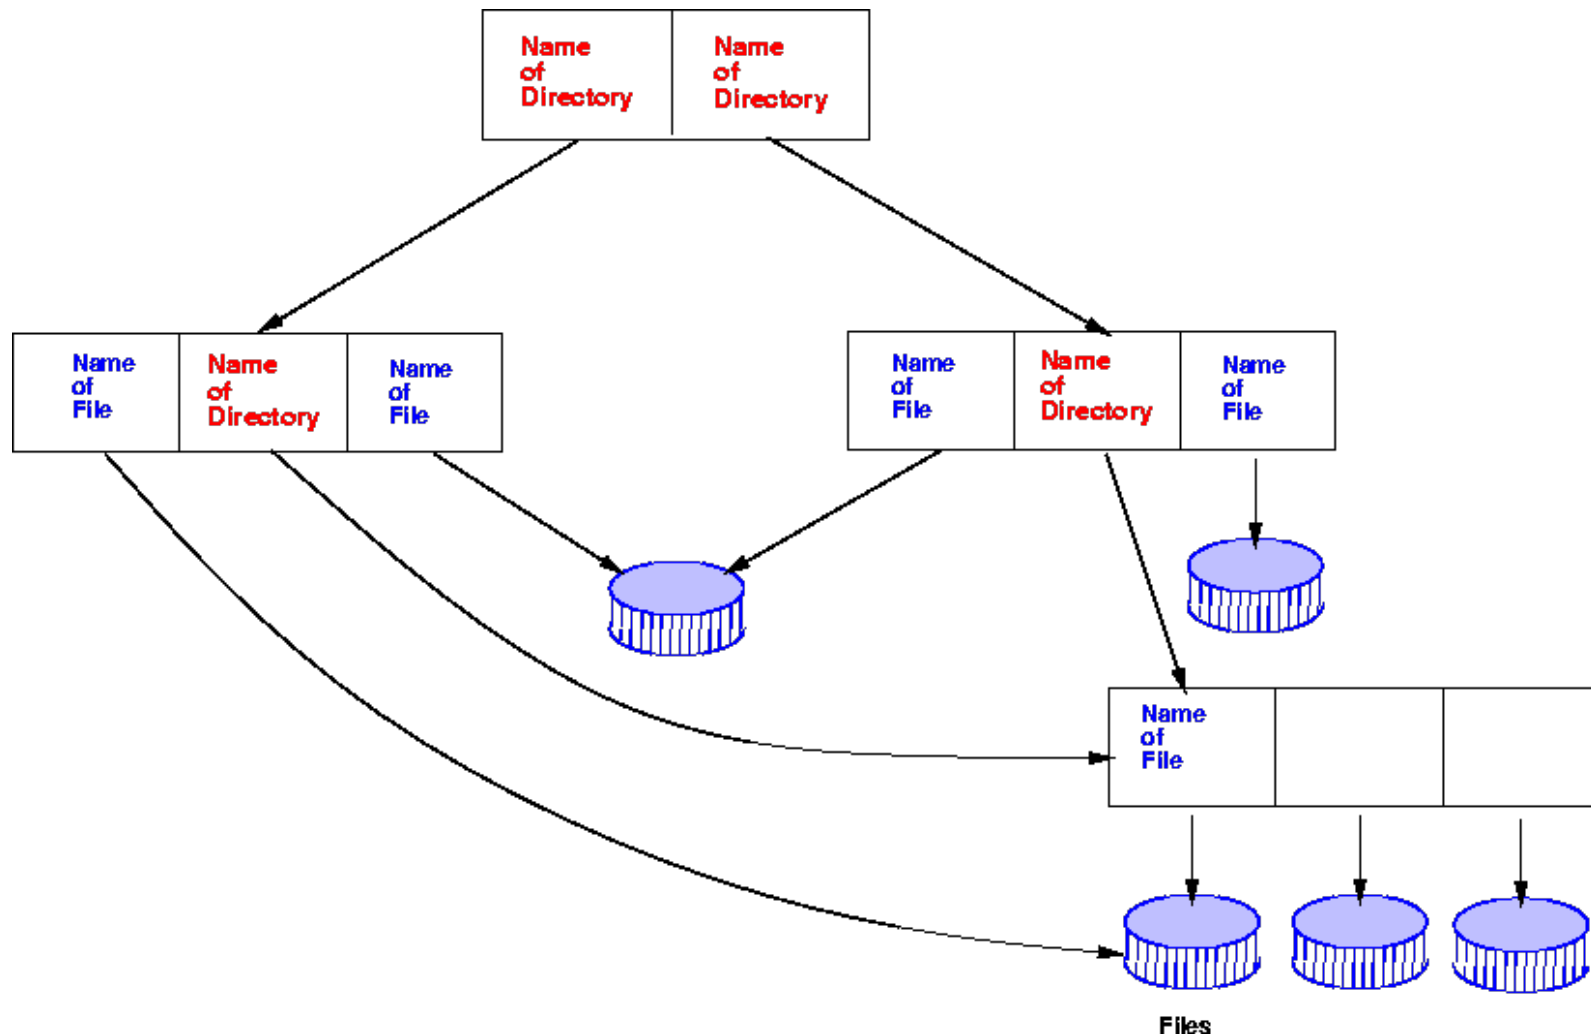
- We usually think of directories as trees...



# Acyclic Graph Structured Dir's




- But in practice they're actually acyclic graphs!



# Symbolic Links



- **Symbolic** links are different than regular links (often called **hard links**). Created with **ln -s**
- Can be thought of as a directory entry that points to the name of another file.
- Does not change link count for file
  - When original deleted, symbolic link remains
- They exist because:
  - Hard links don't work across file systems
  - Hard links only work for regular files, not directories





## **Some notes on disk performance...**

# Typical Modern Spec's



- Disk rotation speed: 5,000-15,000 RPM
- Number of sectors per track: 500-2000
- Number of tracks: 100,000-200,000
- Average seek latency: 2ms
- Block size: 0.5K-4K
- Multiple zones (layout is different for inner and outer tracks)
- Multiple bays (stacked drives)

# Estimated Sustained Average Transfer Rate




- Suppose that a disk drive spins at 7200 RPM (revolutions per minute), has a sector size of 512 bytes, and holds 160 sectors per track.
- What is sustained average transfer rate of this drive in megabytes per second?

A: Track.

B: Sector.

C: Sector of Track.

D: File







- Suppose that a disk drive spins at 7200 RPM (revolutions per minute), has a sector size of 512 bytes, and holds 160 sectors per track.
- What is sustained average transfer rate of this drive in megabytes per second?

Disk spins 120 times per second (7200 RPM/60)

Each spin transfers a track of 80 KB (160 sectors x 0.5K)

Sustained average transfer rate is  $120 \times 80 = 9.6 \text{ MB/s}$ .

# Average Performance of Random Access



- Suppose that a disk drive spins at 7200 RPM (revolutions per minute), has a sector size of 512 bytes, and holds 160 sectors per track.
- Average seek time for the drive is 8 milliseconds
- Estimate # of random sector I/Os per second that can be done and the effective average transfer rate for random-access of a sector?

**Disk spins 120 times per second**

**Average rotational cost is time to travel half track:  $1/120 * 50\% = 4.167\text{ms}$**

**Transfer time is 8ms to seek**

**+ 4.167 ms rotational latency**

**+ 0.052 ms (reading one sector takes  $0.0005\text{MB} / 9.6\text{MB}$ ).**

**= 12.219 ms**

**# of random sector access/second =  $1/0.012219 = 81.8$**

**Effective transferring rate:  $0.5 \text{ KB} * 81.8 = 0.0409 \text{ MB/s}$ .**



- Electronically Erasable Programmable Read Only Memory (EEPROM)
- Example specifications (NAND Flash):
  - Page size: 2KB (approx.)
  - Block size: 64 pages (128KB)
  - Device size: 16K blocks
- Random READ: 25 $\mu$ s, Sequential READ: 25ns
- WRITE performance
  - PROGRAM PAGE: 220 $\mu$ s, BLOCK ERASE: 1.5ms
- Endurance: 100,000 PROGRAM/ERASE cycles