

Goals for Today



- Learning Objective:
 - Learn how to achieve reliability + availability in storage!!
- Announcements, etc:
 - MP3 is out! Due **April 18th**.
 - MP3 Walkthrough on Monday, slides up later today



Reminder: Please put away devices at the start of class



CS 423

Operating System Design: Reliable Storage

Professor Adam Bates
Spring 2018

Storage is hard ; - (



*“In each cluster's first year, it's typical that 1,000 individual machine failures will occur; **thousands of hard drive failures will occur**; one power distribution unit will fail, bringing down 500 to 1,000 machines for about 6 hours; 20 racks will fail, each time causing 40 to 80 machines to vanish from the network; 5 racks will "go wonky," with half their network packets missing in action; and the cluster will have to be rewired once, affecting 5 percent of the machines at any given moment over a 2-day span, Dean said. And there's about a 50 percent chance that the cluster will overheat, taking down most of the servers in less than 5 minutes and taking 1 to 2 days to recover.”*

- Jeff Dean, Google Fellow (2008)

Storage Goals



- Storage reliability: data fetched is what you stored
 - Problem when machines randomly fail!
- Storage availability: data is there when you want it
 - Problem when disks randomly fail!
 - More disks \Rightarrow higher probability of some disk failing
 - Data available $\sim \text{Prob}(\text{disk working})^k$
 - If failures are independent and data is spread across k disks
 - For large k , probability system works $\rightarrow 0$

File System Reliability



- What can happen if disk loses power or machine software crashes?
 - Some operations in progress may complete
 - Some operations in progress may be lost
 - Overwrite of a block may only partially complete
- File systems need durability (as a minimum!)
 - Data previously stored can be retrieved (maybe after some recovery step), regardless of failure

Storage Reliability Problem



- Single logical file operation can involve updates to multiple physical disk blocks
 - inode, indirect block, data block, bitmap, ...
- At a physical level, operations complete one at a time
 - Want concurrent operations for performance
- How do we guarantee consistency regardless of when crash occurs?

Transaction Concept



- A transaction is a grouping of low-level operations that are related to a single logical operation
- Transactions are atomic — operations appear to happen as a group, or not at all (at logical level)
 - At physical level of course, only a single disk/flash write is atomic
- Transactions are durable — operations that complete stay completed
 - Future failures do not corrupt previously stored data
- (In-Progress) Transactions are isolated — other transactions cannot see the results of earlier transactions until they are committed
- Transactions exhibit consistency — sequential memory model



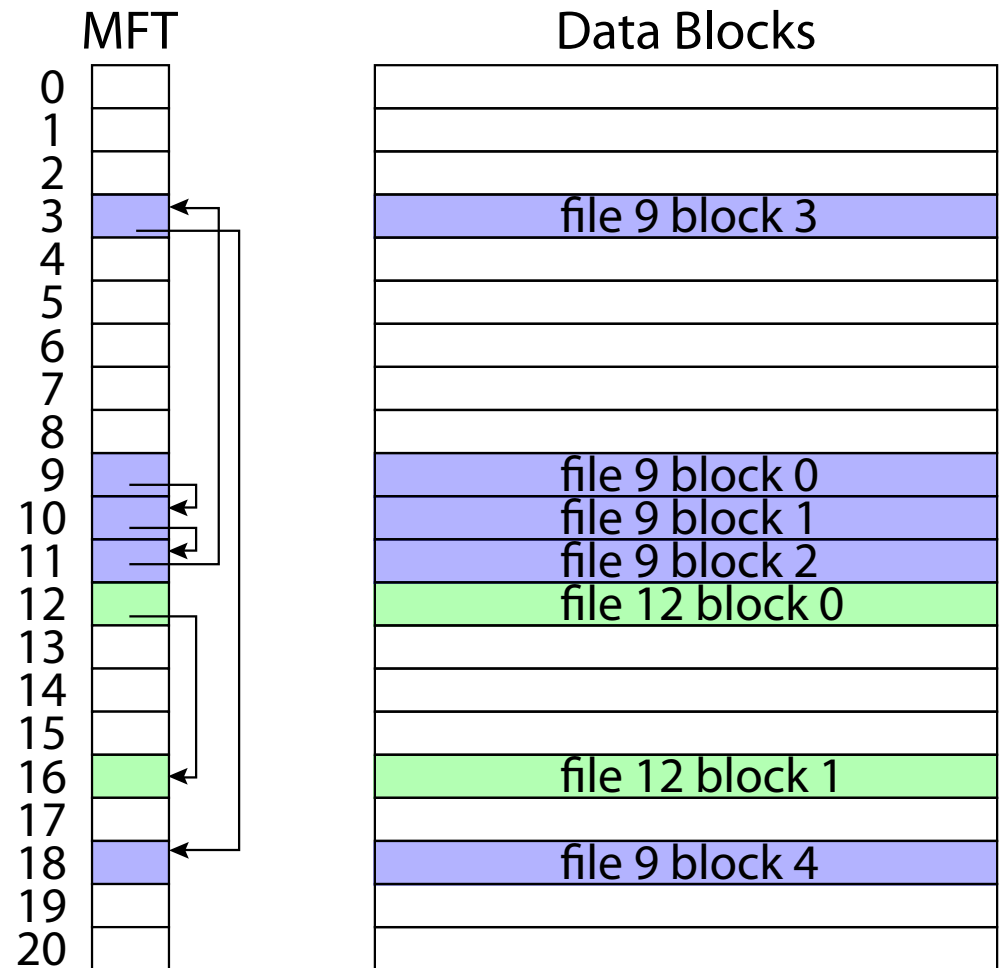
- Sequence operations in a specific order
 - Careful design to allow sequence to be interrupted safely
- Post-crash recovery
 - Read data structures to see if there were any operations in progress
 - Clean up/finish as needed
- Approach taken in FAT, FFS (fsck), and many app-level recovery schemes (e.g., Word)

Reliability Attempt #1: Careful Ordering



FAT: Append Data to File

- Add data block
- Add pointer to data block
- Update file tail to point to new MFT entry
- Update access time at head of file



Reliability Attempt #1: Careful Ordering

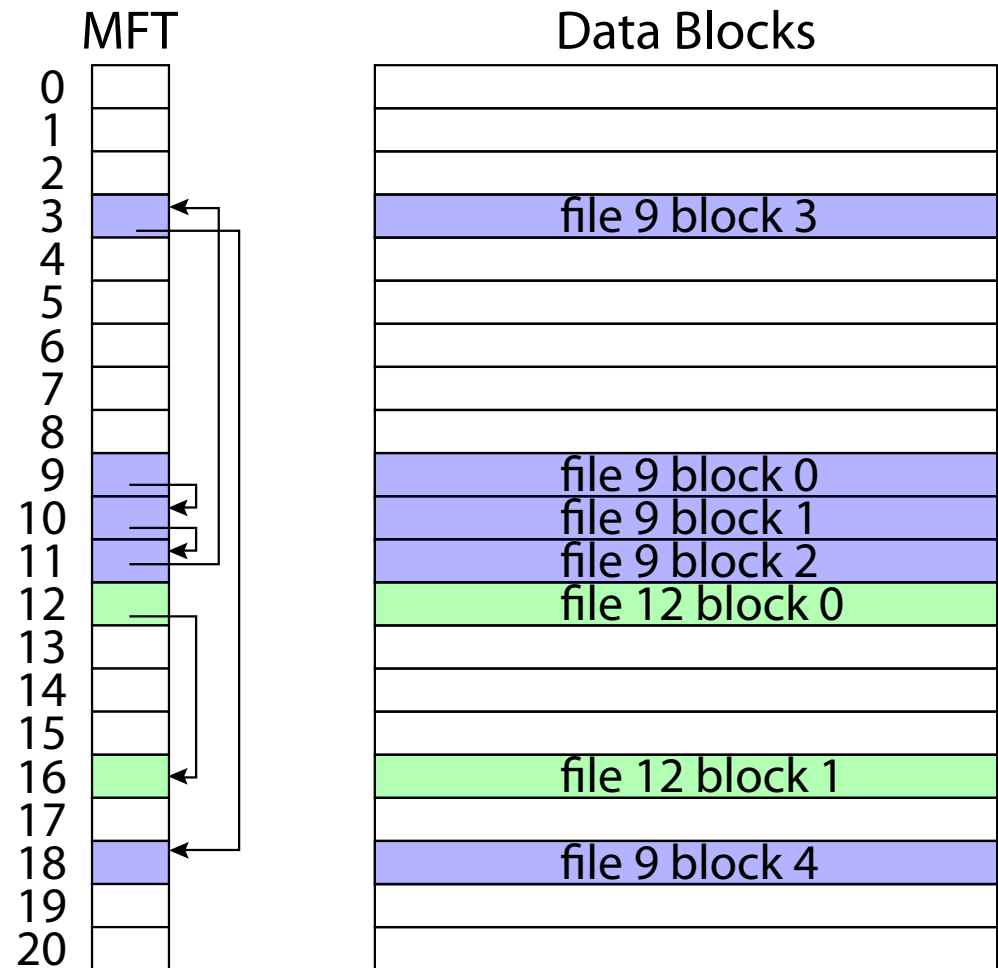


FAT: Append Data to File

- Add data block
- Add pointer to data block
- Update file tail to point to new MFT entry
- Update access time at head of file

Recovery

- Scan MFT
- If entry is unlinked, delete data block
- If access time is incorrect, update

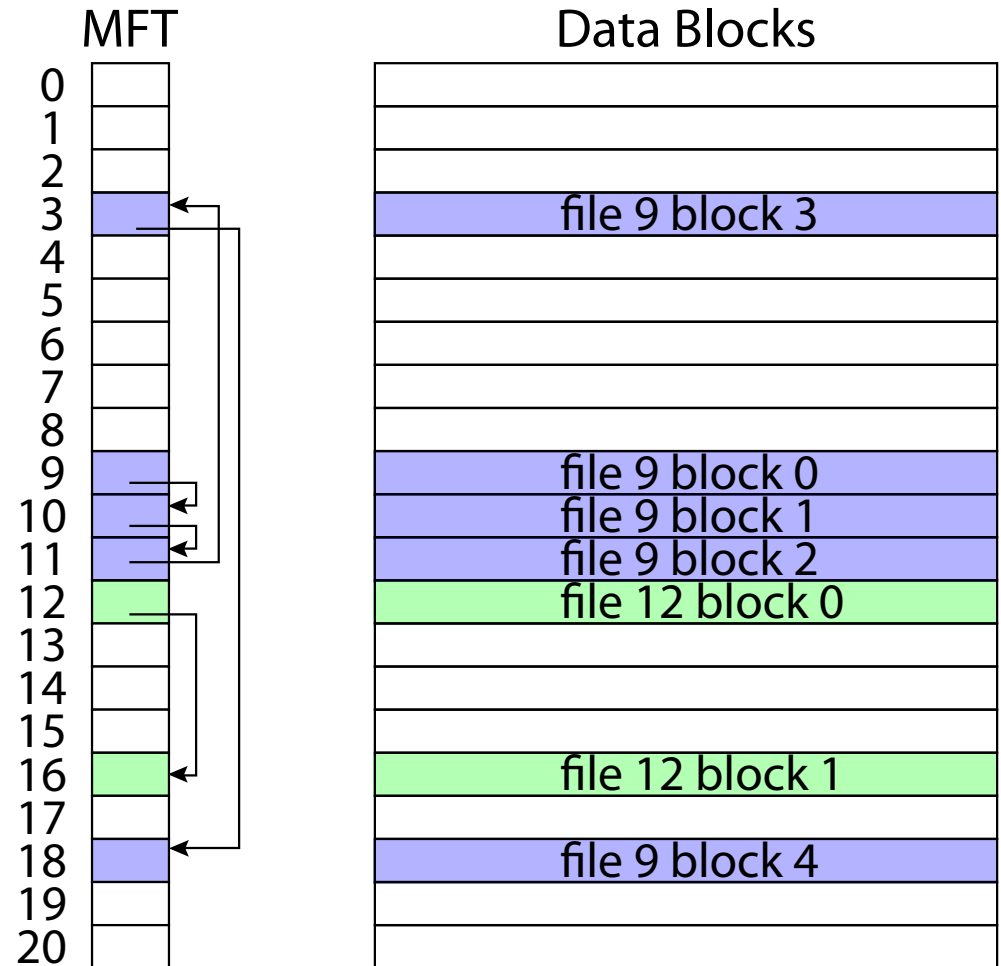


Reliability Attempt #1: Careful Ordering



FAT: Create New File

- Allocate data block
- Update MFT entry to point to data block
- Update directory with
file name -> file number
- What if directory spans multiple disk blocks?
- Update modify time for directory



Reliability Attempt #1: Careful Ordering




FAT: Create New File

- Allocate data block
- Update MFT entry to point to data block
- Update directory with
file name -> file number
- What if directory spans multiple disk blocks?
- Update modify time for directory

Recovery

- Scan MFT
- If any unlinked files (not in any directory), delete
- Scan directories for missing update times




Reliability Attempt #1: Careful Ordering




FFS: Create New File

- Allocate data block
- Write data block
- Allocate inode
- Write inode block
- Update bitmap of free blocks
- Update directory with
file name -> file number
- Update modify time for directory

Inode Array




Inode



Triple
Indirect
Blocks

Double
Indirect
Blocks



Reliability Attempt #1: Careful Ordering



FFS: Create New File


- Allocate data block
- Write data block
- Allocate inode
- Write inode block
- Update bitmap of free blocks
- Update directory with
file name -> file number
- Update modify time for directory

Recovery


- Scan inode table
- If any unlinked files (not in any directory), delete
- Compare free block bitmap against inode trees
- Scan directories for missing update/access times

Recovery time is proportional to size of disk!

Inode Array




Inode



Triple
Indirect
Blocks

Double
Indirect
Blocks




Reliability Attempt #1: Careful Ordering




FFS: Move a File

- Remove filename from old directory
- Add filename to new directory

Inode Array




Inode



Triple
Indirect
Blocks

Double
Indirect
Blocks



Reliability Attempt #1: Careful Ordering




FFS: Move a File

- Remove filename from old directory
- Add filename to new directory


Recovery

- Scan all directories to determine set of live files
- Consider files with valid inodes and not in any directory
 - New file being created?
 - File move?
 - File deletion?

Inode Array




Inode



Triple
Indirect
Blocks

Double
Indirect
Blocks



Reliability Attempt #1: Careful Ordering

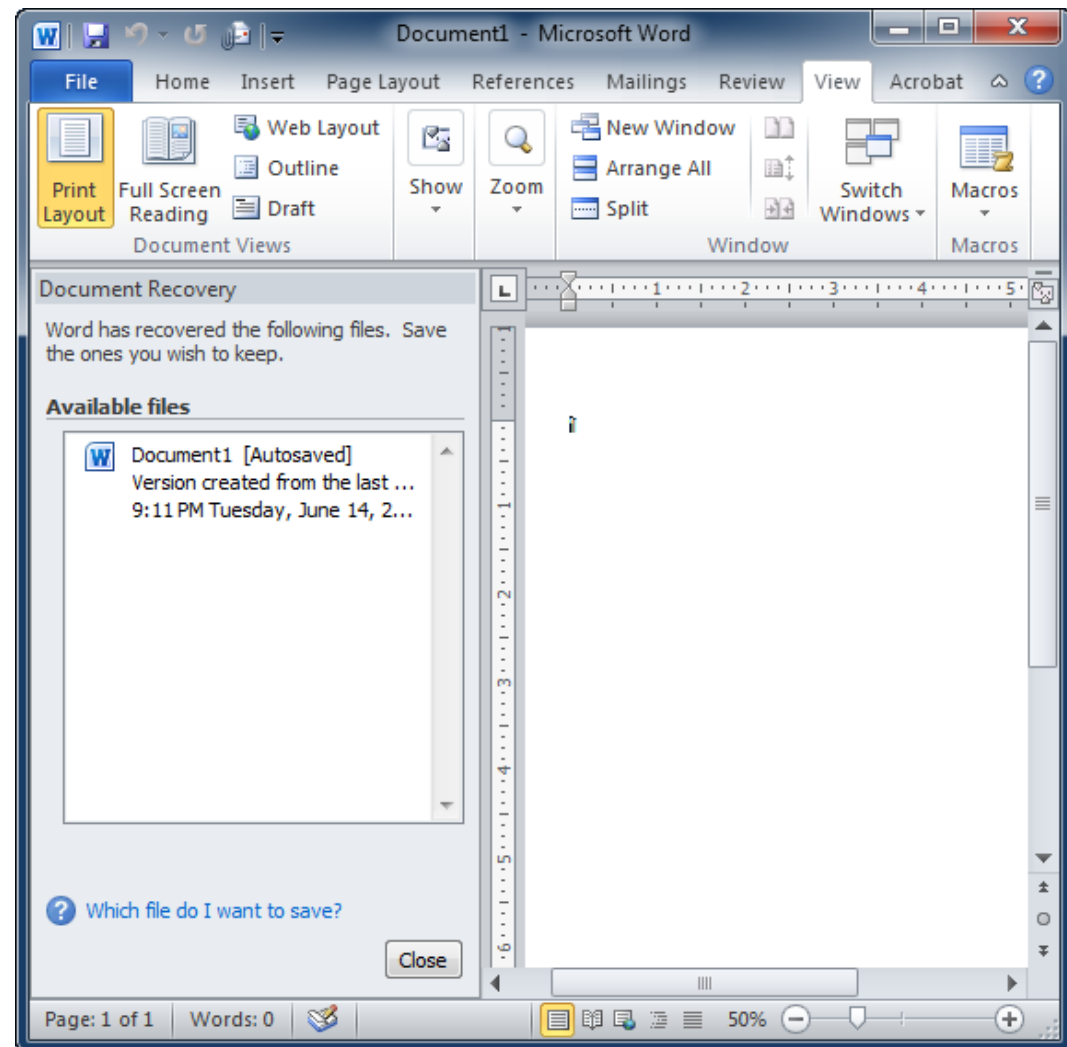


Application Level

- Write name of each open file to app folder
- Write changes to backup file
- Rename backup file to be file (atomic operation provided by file system)
- Delete list in app folder on clean shutdown

Recovery

- On startup, see if any files were left open
- If so, look for backup file
- If so, ask user to compare versions





FFS: Move and Grep

- Observation — careful ordering is not a panacea...

Process A moves file from x to y

```
mv x/file y/
```

Process B greps across x and y

```
grep x/* y/*
```

- Will Process B always see the contents of the file?



Pros

- Works with minimal support from the disk drive
- Works for most multi-step operations

Cons

- Can require time-consuming recovery after a failure
- Difficult to reduce every operation to a safely-interruptible sequence of writes
- Difficult to achieve consistency when multiple operations occur concurrently (e.g., FFS grep)

Reliability Attempt #2: Copy-on-Write

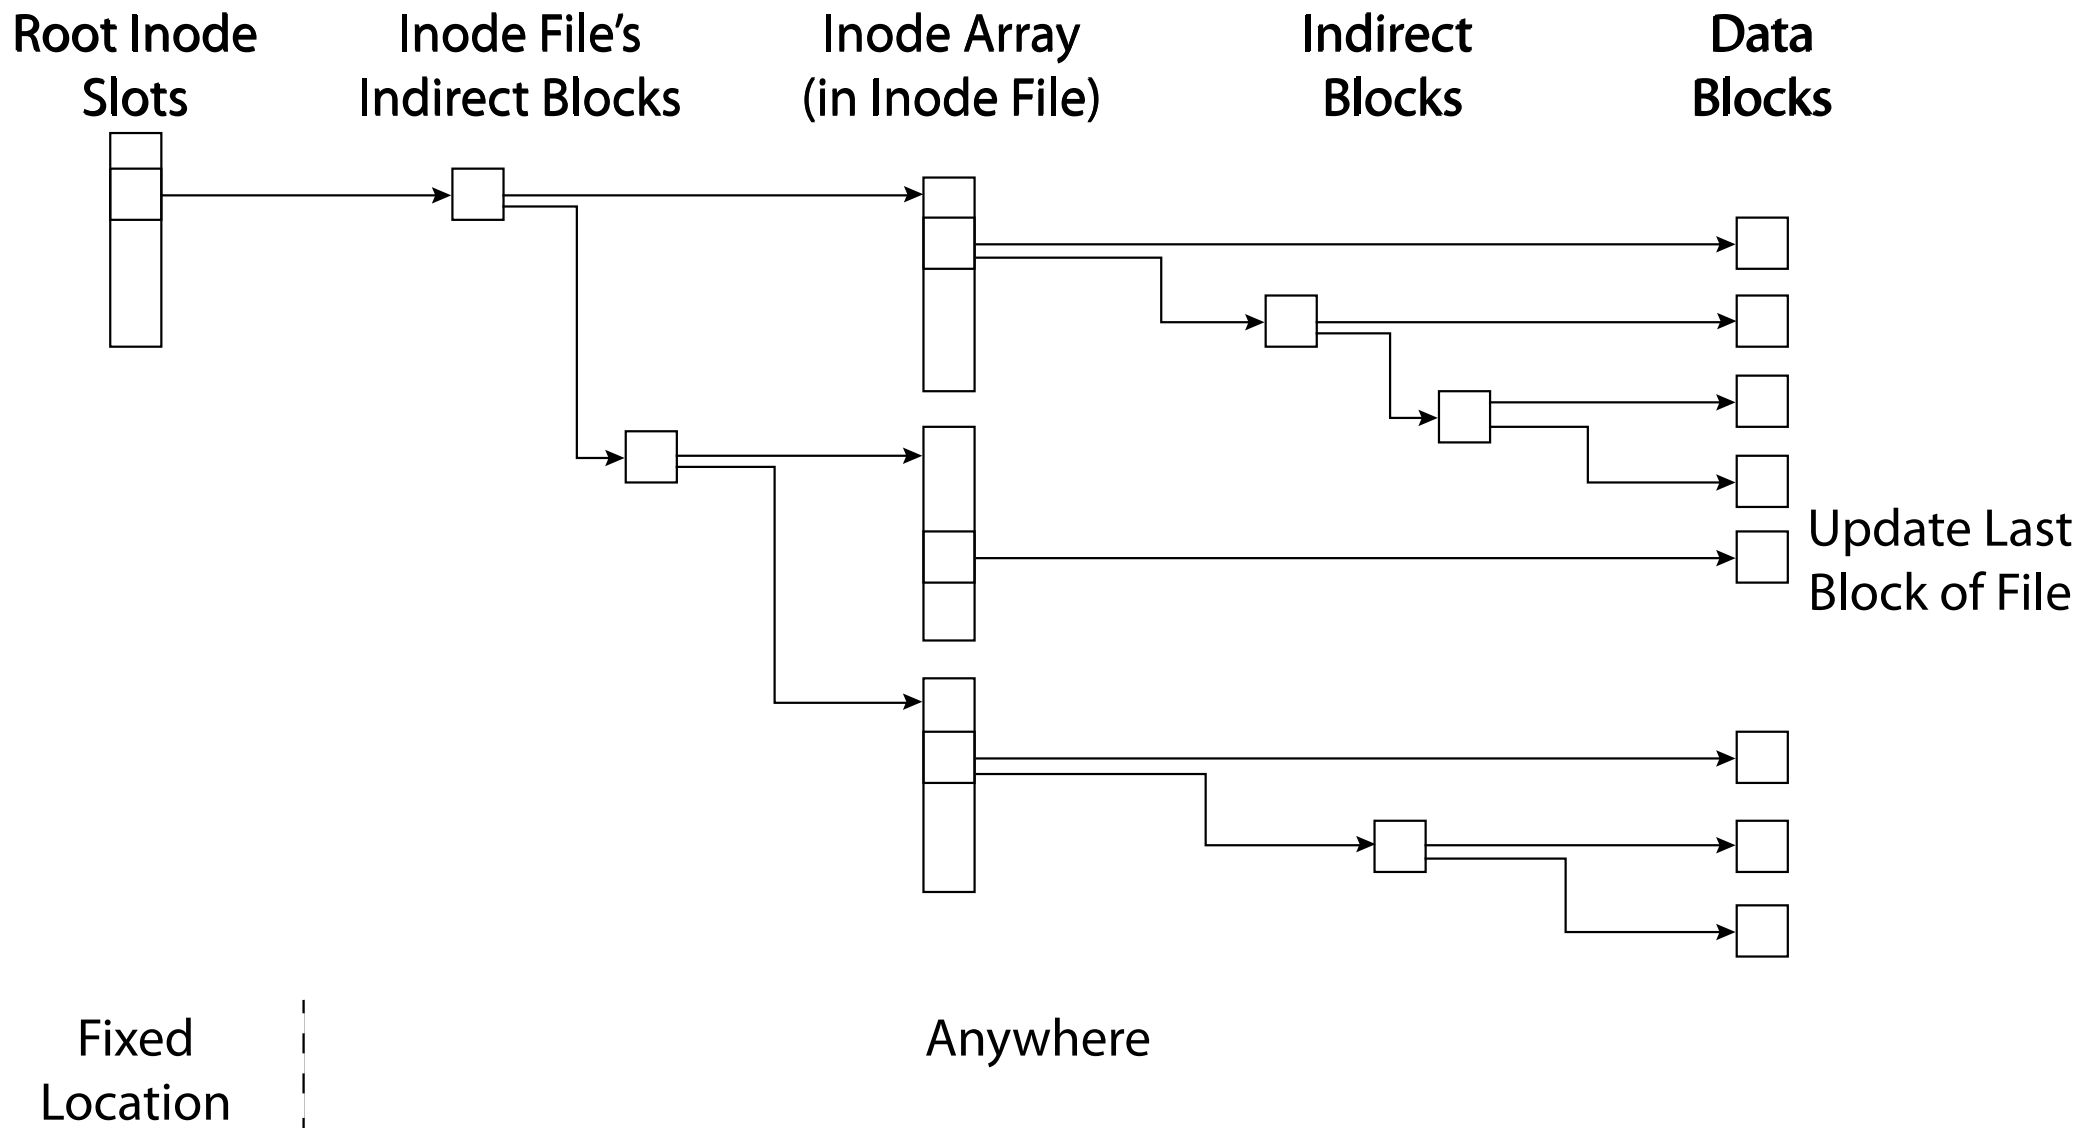


- To update file system, write a new version of the file system containing the update
 - Never update in place
 - Reuse existing unchanged disk blocks
- Seems expensive! But...
 - Updates can be batched
 - Almost all disk writes can occur in parallel
- Approach taken in network file server appliances (WAFL, ZFS)

Reliability Attempt #2: Copy-on-Write



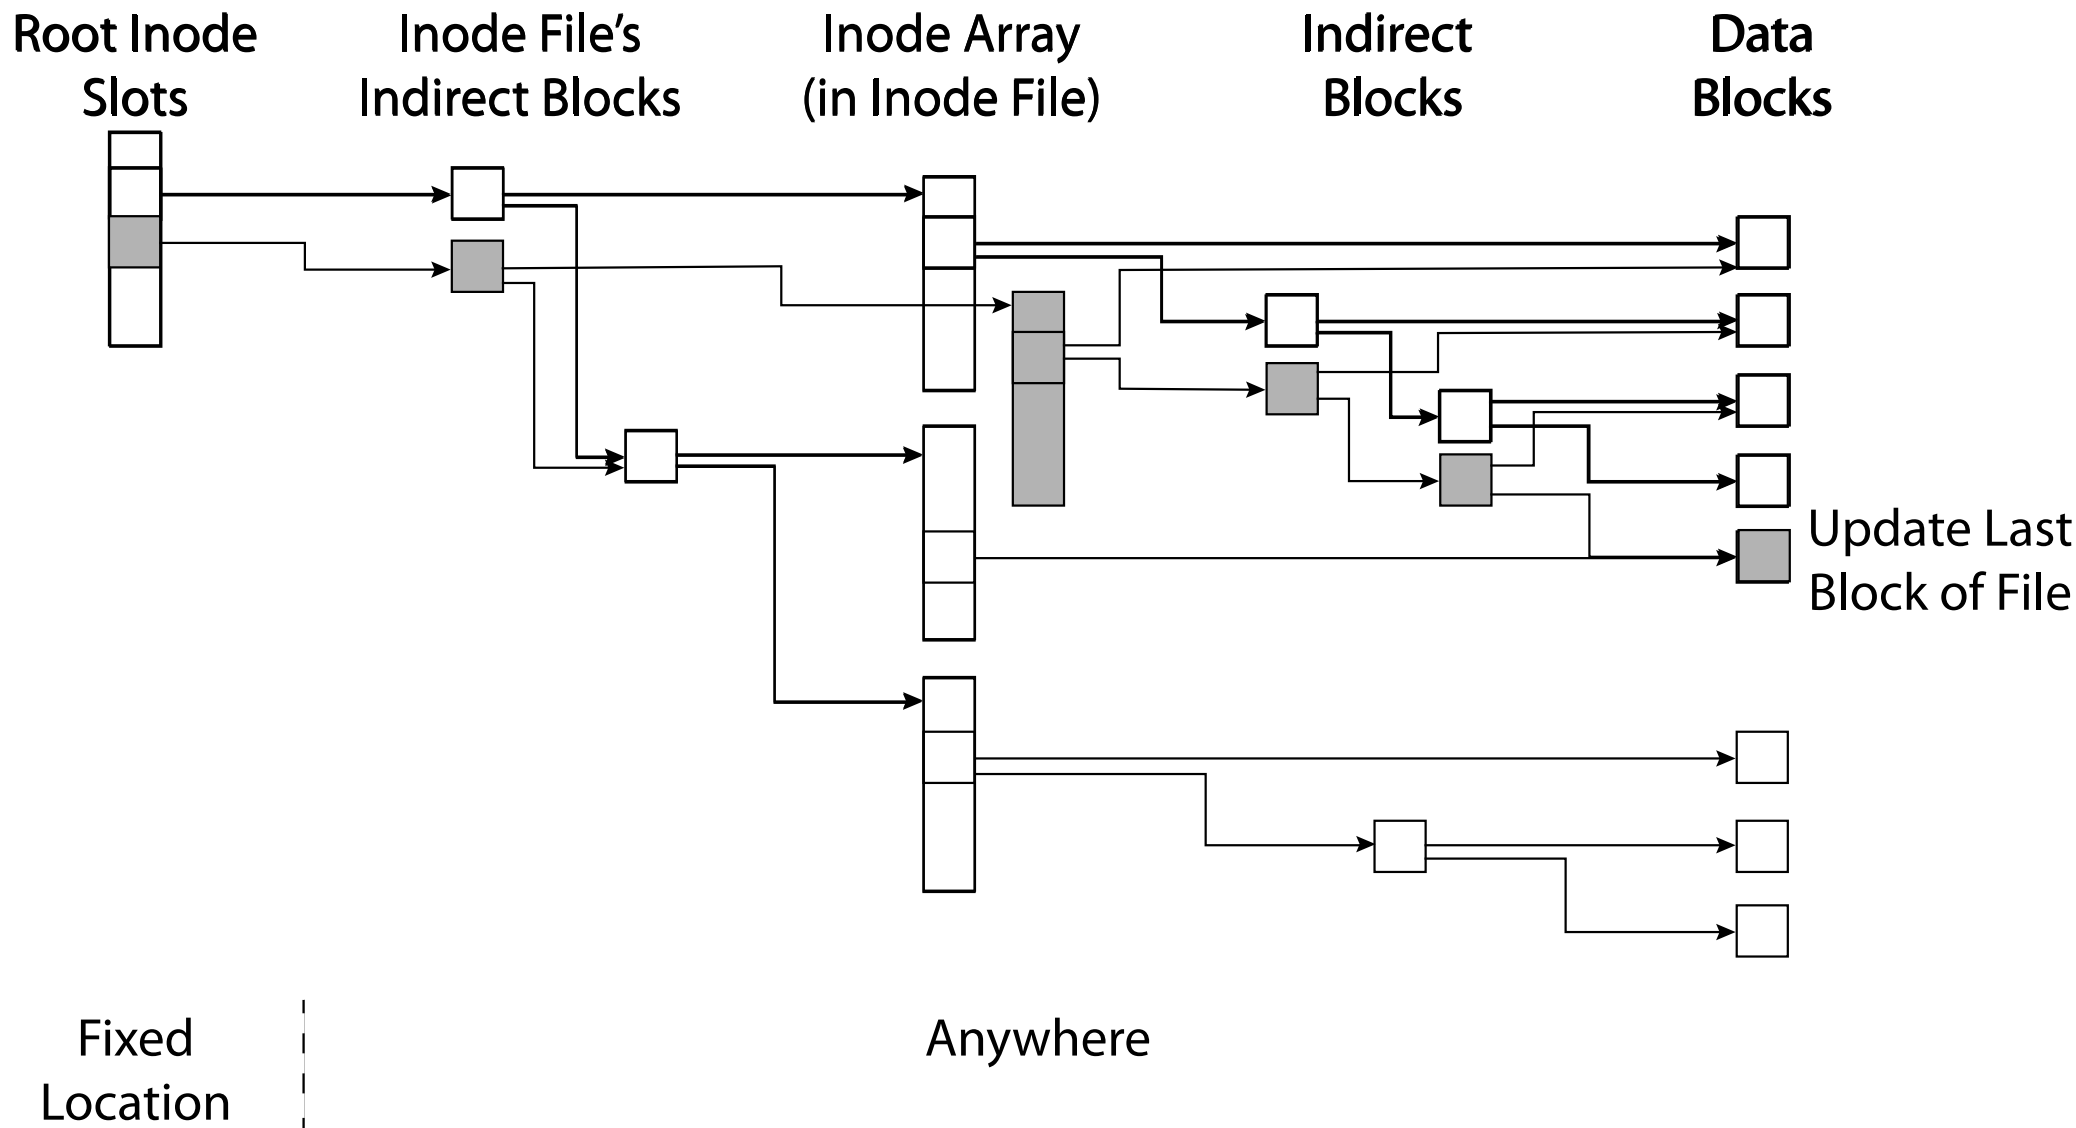
Copy on Write (Write Anywhere File Layout)



Reliability Attempt #2: Copy-on-Write




Copy on Write (Write Anywhere File Layout)



Reliability Attempt #2: Copy-on-Write



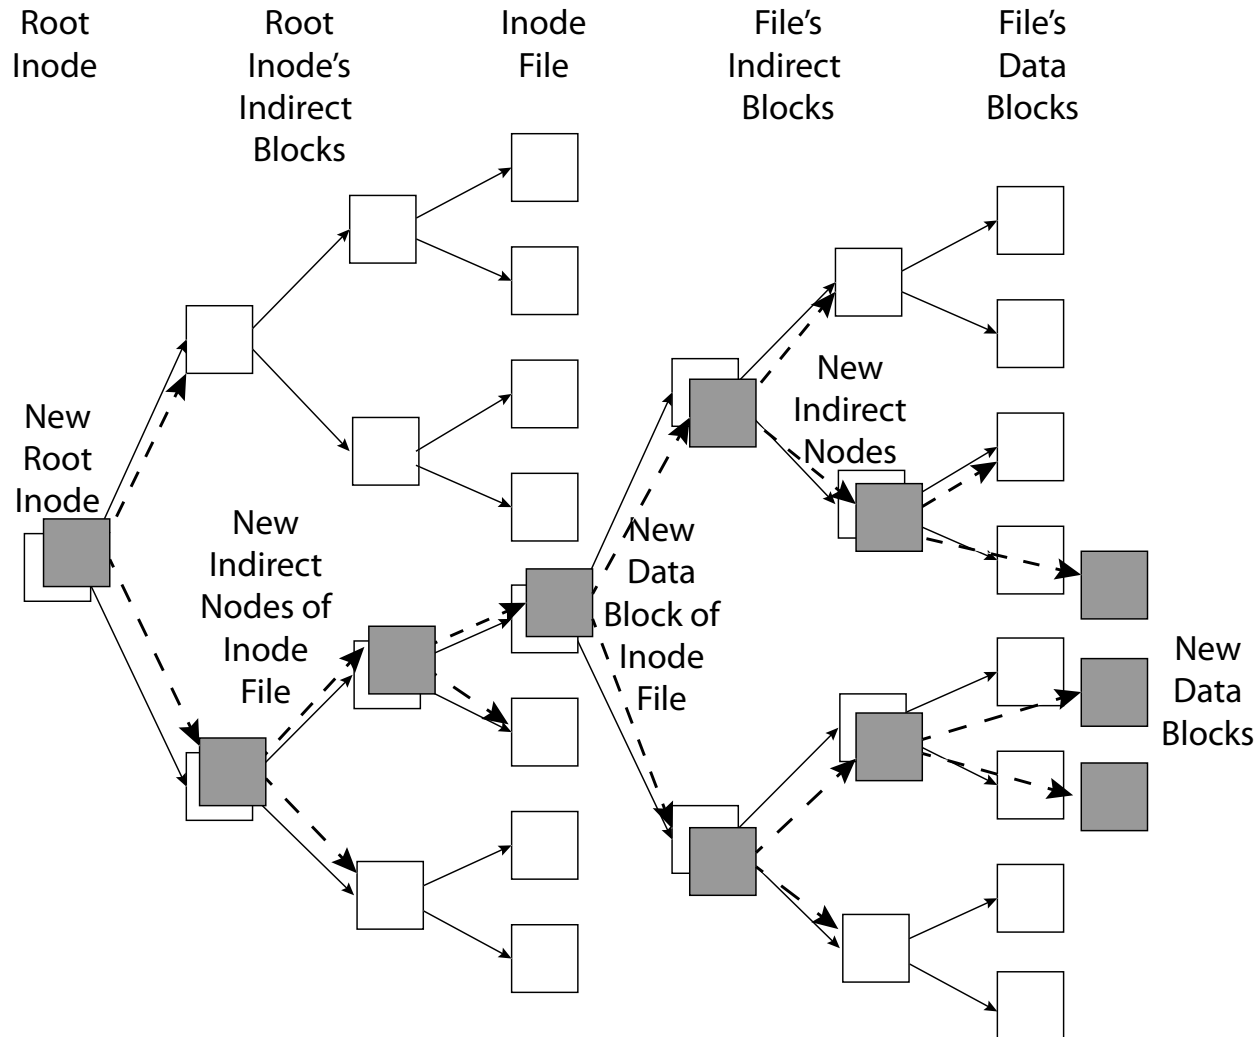
Batch Updates



Reliability Attempt #2: Copy-on-Write



Batch Update




Reliability Attempt #2: Copy-on-Write



FFS Updates


(updates are in-place)



Reliability Attempt #2: Copy-on-Write



Write Anywhere File Layout (WAFL) Updates (Uses Copy-on-Write)





Garbage Collection

- For write efficiency, want contiguous sequences of free blocks
 - Spread across all block groups
 - Updates leave dead blocks scattered
- For read efficiency, want data read together to be in the same block group
 - Write anywhere leaves related data scattered
- Solution? Background coalescing of live/dead blocks



Pros

- Correct behavior regardless of failures
- Fast recovery (root block array)
- High throughput (best if updates are batched)

Cons

- Potential for high latency
- Small changes require many writes
- Garbage collection essential for performance

Logging File Systems



- Instead of modifying data structures on disk directly, write changes to a journal/log
 - Intention list: set of changes we intend to make
 - Log/Journal is append-only
- Once changes are on log, safe to apply changes to data structures on disk
 - Recovery can read log to see what changes were intended
- Once changes are copied, safe to remove log

Redo Logging



- Prepare
 - Write all changes (in transaction) to log
- Commit
 - Single disk write to make transaction durable
- Redo
 - Copy changes to disk
- Garbage collection
 - Reclaim space in log
- Recovery
 - Read log
 - Redo any operations for committed transactions
 - Garbage collect log

Redo Logging



Before transaction start

Cache

Tom = \$200

Mike = \$100

Nonvolatile
Storage

Tom = \$200

Mike = \$100

Log:

Redo Logging



After Updates are Logged

Cache

Tom = \$100

Mike = \$200

Nonvolatile
Storage

Tom = \$200

Mike = \$100

Log: Tom = \$100 Mike = \$200

Redo Logging



After commit logged

Cache

Tom = \$100

Mike = \$200

Nonvolatile
Storage

Tom = \$100

Mike = \$200

Log: Tom = \$100 Mike = \$200 COMMIT

Redo Logging



After copy back

Cache

Tom = \$100

Mike = \$200

Nonvolatile
Storage

Tom = \$100

Mike = \$200

Log:

Redo Logging



After garbage collection

Cache

Tom = \$100

Mike = \$200

Nonvolatile
Storage

Tom = \$100

Mike = \$200

Log:

Redo Logging



Questions

- What happens if machine crashes?
 - Before transaction start
 - After transaction start, before operations are logged
 - After operations are logged, before commit
 - After commit, before write back
 - After write back before garbage collection
- What happens if machine crashes during recovery?

Redo Logging



Performance

- Log written sequentially
 - Often kept in flash storage
- Asynchronous write back
 - Any order as long as all changes are logged before commit, and all write backs occur after commit
- Can process multiple transactions
 - Transaction ID in each log entry
 - Transaction completed iff its commit record is in log

Transaction Isolation



- What if grep starts after changes are logged but before they are committed?

Process A moves file from x to y
`mv x/file y/`

Process B greps across x and y
`grep x/* y/*`

Transaction Isolation



- What if grep starts after changes are logged but before they are committed?

Process A moves file from x to y
`mv x/file y/`

Process B greps across x and y
`grep x/* y/*`

- Two Phase Locking: Release locks only AFTER transaction commit.
 - Prevents a process from seeing results of a transaction that might not commit!

Process A moves file from x to y
`Lock x, y`
`mv x/file y/`
`Commit & Release x, y`

Process B greps across x and y
`Lock x, y`
`grep x/* y/*`
`Release x, y`

Serializability



- With two phase locking and redo logging, transactions appear to occur in a sequential order (serializability)
 - Either: grep then move or move then grep
- Other implementations can also provide serializability
 - e.g., Optimistic concurrency control: abort any transaction that would conflict with serializability

Storage Availability



- Storage reliability: data fetched is what you stored
 - Transactions, redo logging, etc.
- Storage availability: data is there when you want it
 - More disks \Rightarrow higher probability of some disk failing
 - Data available $\sim \text{Prob}(\text{disk working})^k$
 - If failures are independent and data is spread across k disks
 - For large k , probability system works $\rightarrow 0$

RAID



“Redundant Array of Inexpensive Disks”

- Multiple disk drives provide reliability via **redundancy**.
- **Speeds up access times** even beyond sequential.
- Increases the **mean time to failure**






- RAID
 - multiple disks work cooperatively
 - Improve reliability by storing redundant data
 - **Striping (RAID 0)** improves performance with disk **striping** (use a group of disks as one storage unit)
 - **Mirroring (RAID 1)** keeps duplicate of each disk
 - Striped mirrors (**RAID 1+0**) or mirrored stripes (**RAID 0+1**) provides high performance and high reliability
 - **Block interleaved parity (RAID 4, 5, 6)** uses much less redundancy

RAID Level 0




- Level 0 is nonredundant disk array
- Files are striped across disks, no redundant info
- High read throughput
- Best write throughput (no redundant info to write)
- Any disk failure results in data loss



RAID Level I



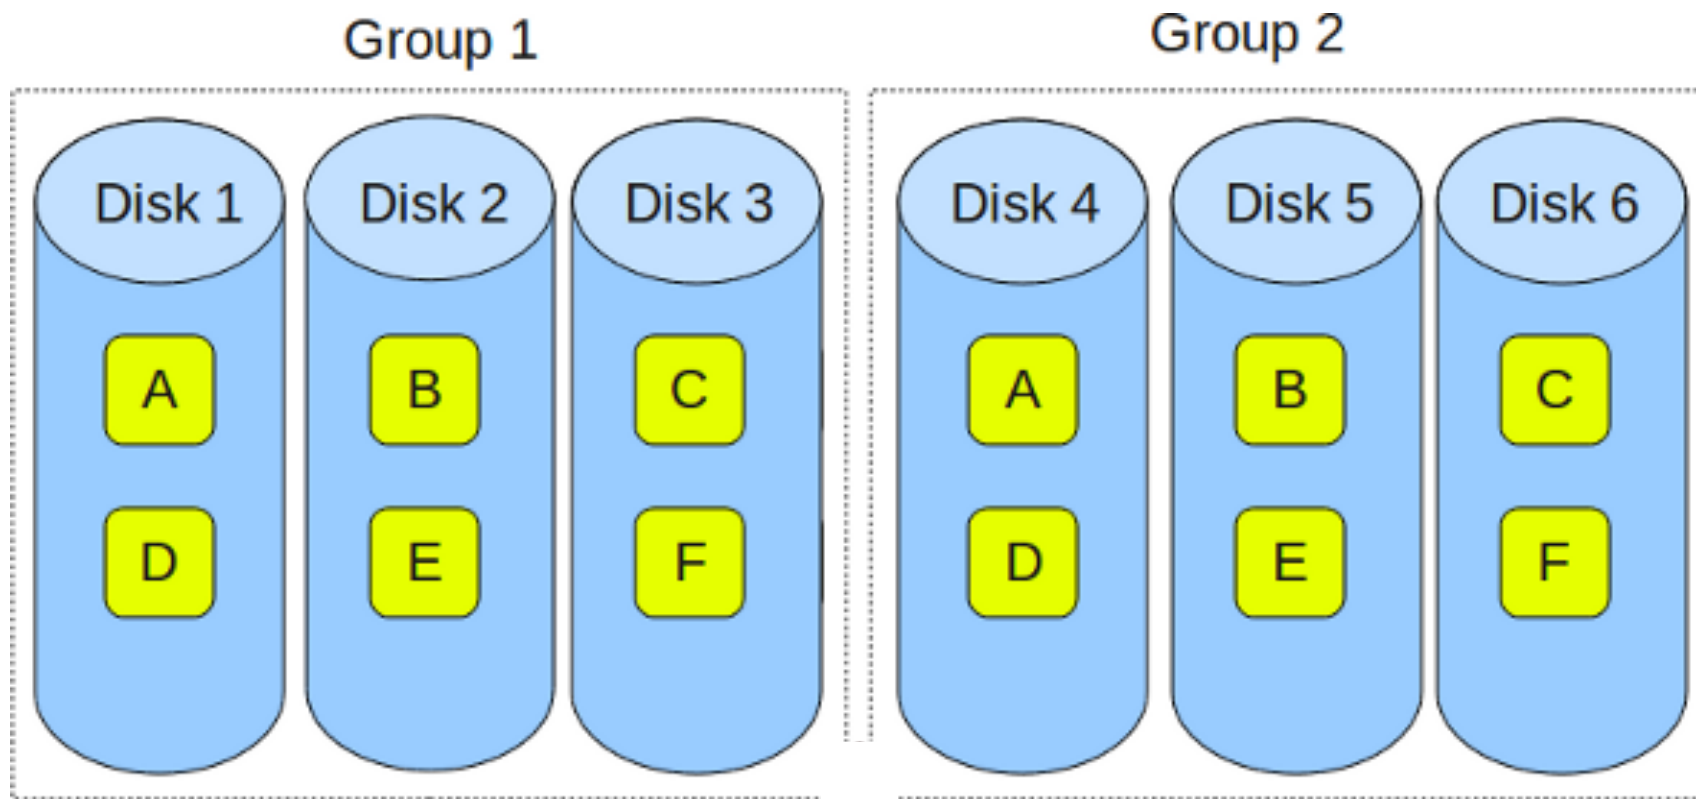
- Mirrored Disks
- Data is written to two places
 - On failure, just use surviving disk (easy to rebuild)
- On read, choose fastest to read
 - Write performance is same as single drive, read performance is 2x better
- Expensive (high space overhead)



RAID Level 0+1



- Stripe on a set of disks
- Then mirror of data blocks is striped on the second set.

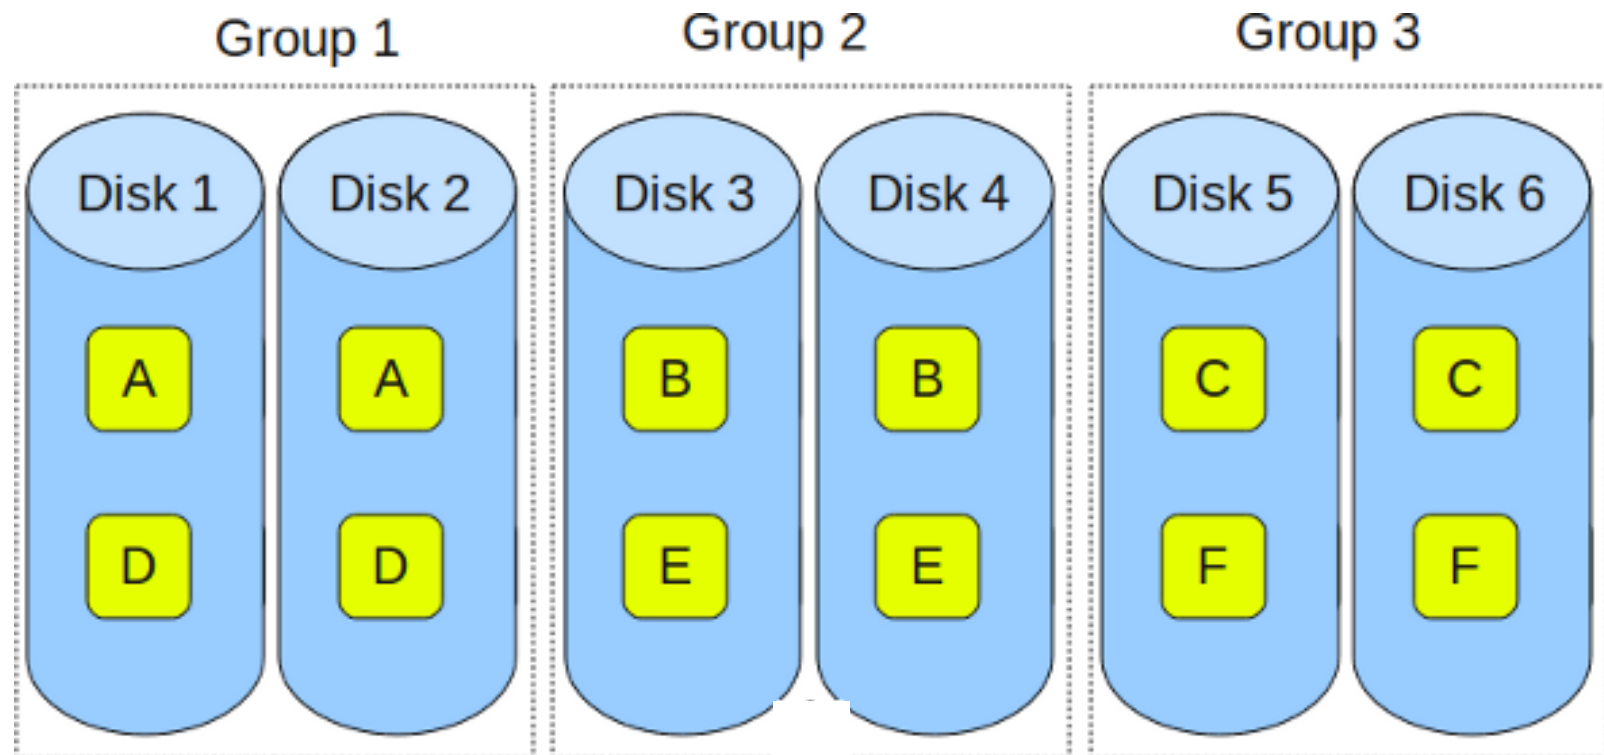


RAID 01 – Blocks Striped. (and Blocks Mirrored)

RAID Level 1+0



- Pair mirrors first.
- Then stripe on a set of paired mirrors



RAID 10 – Blocks Mirrored. (and Blocks Striped)