

Goals for Today



- Learning Objective:
 - Understand how energy usage informs OS design
- Announcements, etc:
 - MP4 due **May 7th**
 - Deadline provides more time than necessary; wanted to give you flexibility
 - HW1 available! Due **May 4th**
 - Just an “appetizer” for the final exam
 - Multiple attempts allowed, but first attempt is graded



Reminder: Please put away devices at the start of class



CS 423

Operating System Design: Energy + Power Considerations

Professor Adam Bates
Spring 2018

Why care about energy?



Low-end Computers:

- Resource-constrained battery-operated devices (laptops, phones, wireless sensors, ...)
- Processor speed grows faster than battery capacity: energy becomes a bottleneck

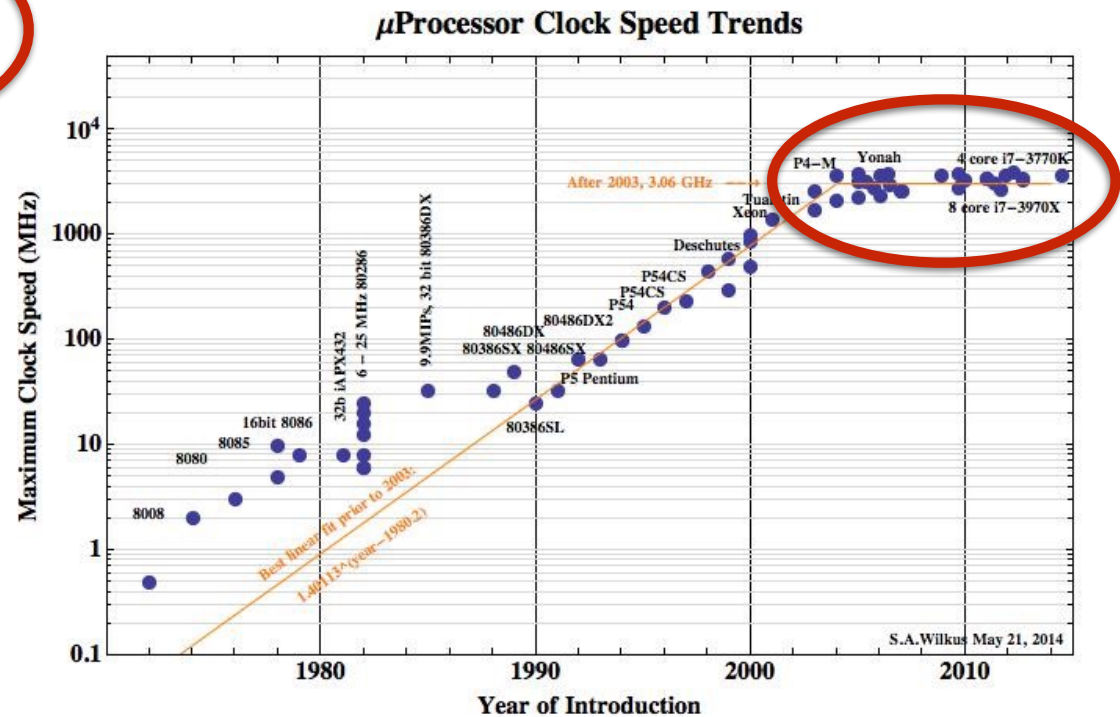
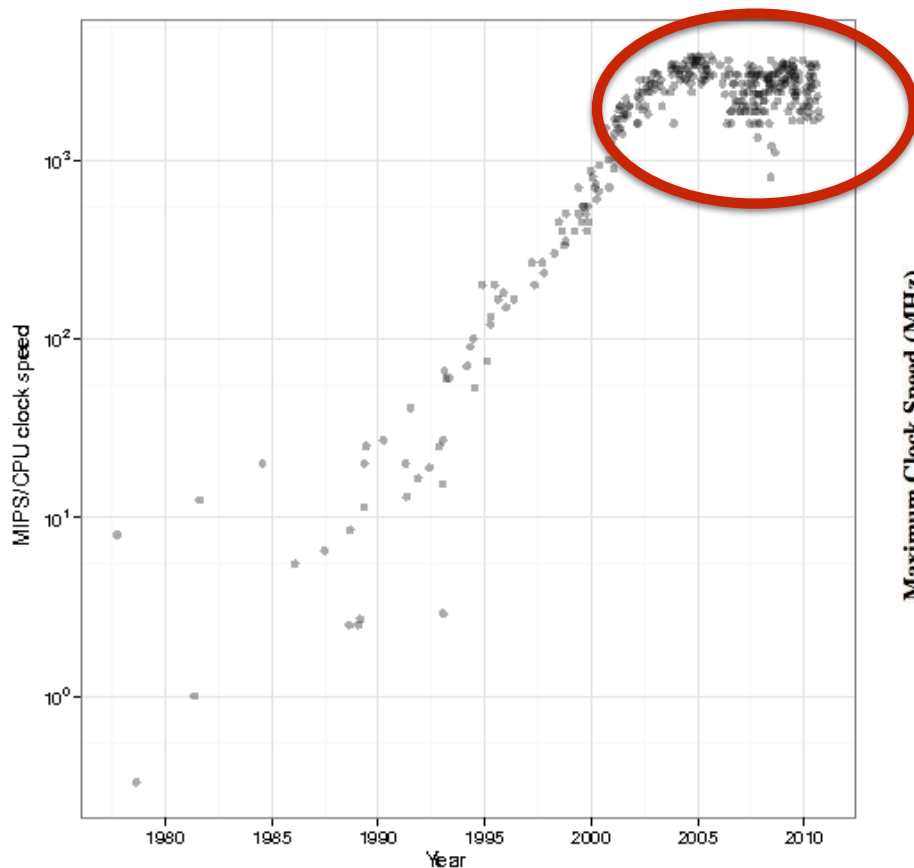
High-end Computers:

- Cost of energy is increasing: The energy bill is the second highest operational expense of data centers (Google, HP, IBM, ...)

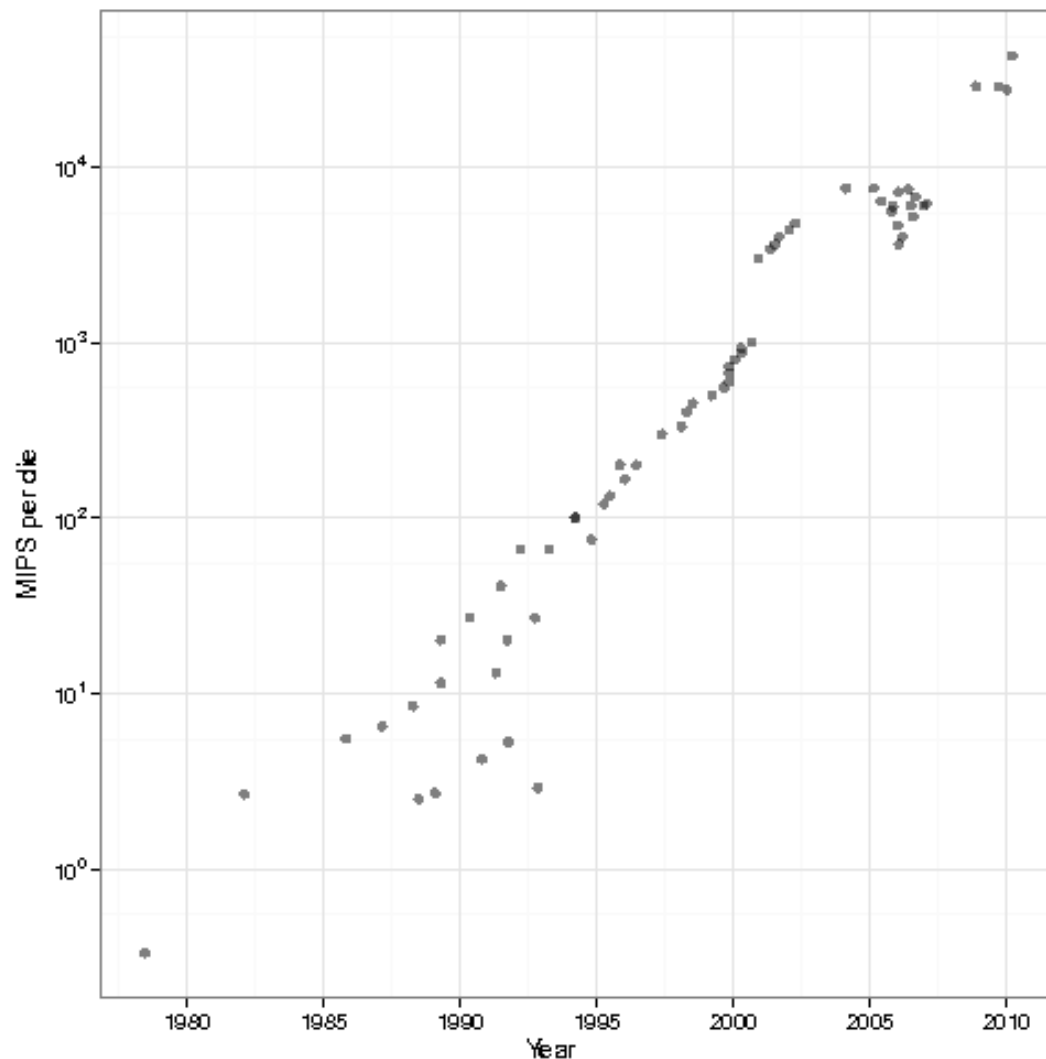
Microprocessor Clock Speed



- Moore's Law (1980-2005)
- Question: Why did the speed curve level off in 2005?



Computational Power (per Die)



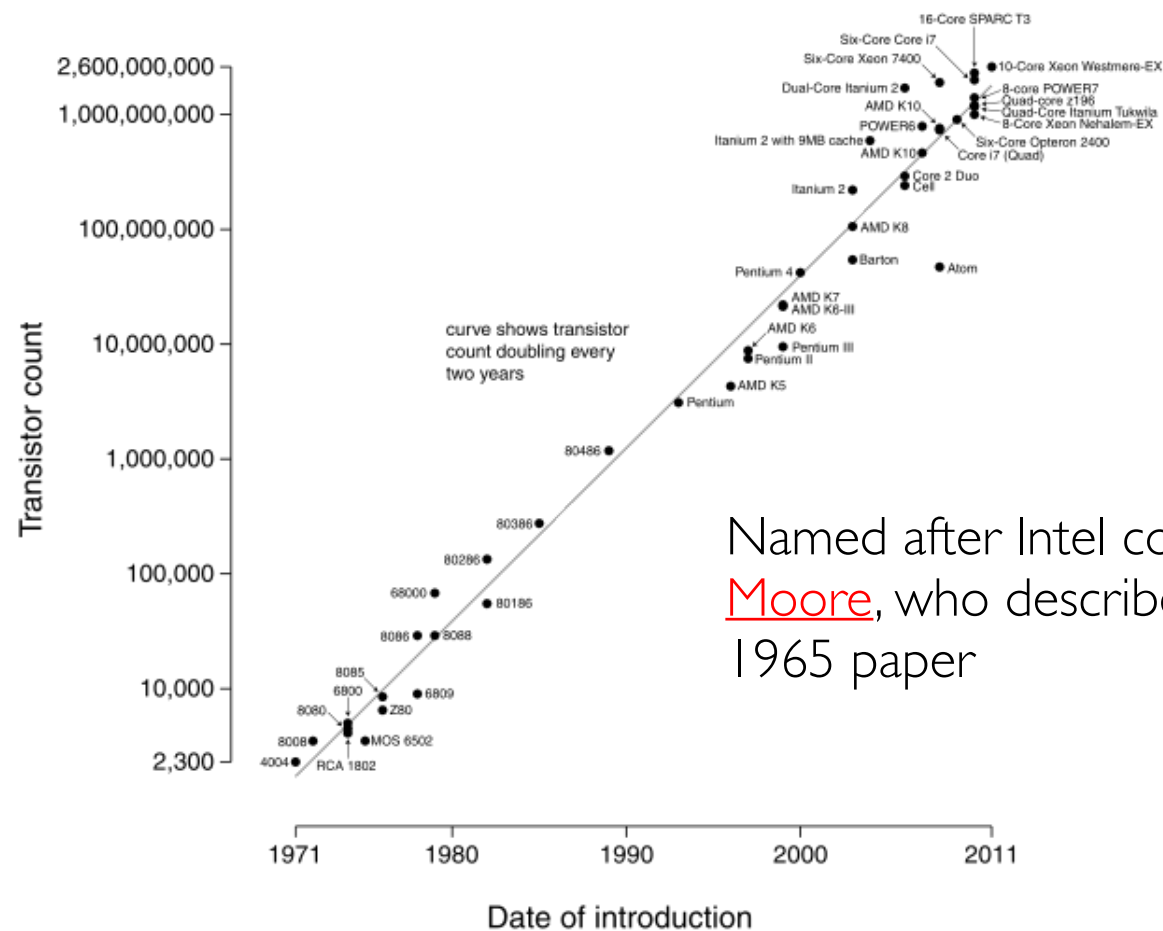
- Note the exponential rise in power consumption
- Question: how come it does not saturate?

Moore's Law



Moore's Law: Transistor count doubles every two years

Microprocessor Transistor Counts 1971-2011 & Moore's Law



Named after Intel co-founder [Gordon E. Moore](#), who described the trend in his 1965 paper



Part of UEFI since 2013:

- Exposes different power saving states in a platform-independent manner
- The standard was originally developed by Intel, Microsoft, and Toshiba (in 1996), then later joined by HP, and Phoenix.
- The latest version is "Revision 6" published in April 2015.

ACPI Global States



- **G0**: working
- **G1**: Sleeping and hibernation (several degrees available)
- **G2**., Soft Off: almost the same as G3 Mechanical Off, except that the power supply still supplies power, at a minimum, to the power button to allow wakeup. A full reboot is required.
- **G3**, Mechanical Off: The computer's power has been totally removed via a mechanical switch (as on the rear of a PSU).

ACPI Global States



- **G0**: working
- **G1**: Sleeping and hibernation (several degrees available)
- **G2**., Soft Off: almost the same as G3 Mechanical Off, except that the power supply still supplies power, at a minimum, to the power button to allow wakeup. A full reboot is required.
- **G3**, Mechanical Off: The computer's power has been totally removed via a mechanical switch (as on the rear of a PSU).



C-States:

- **C0**: is the operating state.
- **C1** (often known as Halt): is a state where the processor is not executing instructions, but can return to an executing state instantaneously. All ACPI-conformant processors must support this power state.
- **C2** (often known as Stop-Clock): is a state where the processor maintains all software-visible state, but may take longer to wake up. This processor state is optional.
- **C3** (often known as Sleep) is a state where the processor does not need to keep its cache, but maintains other state. This processor state is optional.



C-States:

- **C0**: is the operating state.
- **C1** (often known as Halt): is a state where the processor is not executing instructions, but can return to an executing state instantaneously. All ACPI-conformant processors must support this power state.
- **C2** (often known as Stop-Clock): is a state where the processor maintains all software-visible state, but may take longer to wake up. This processor state is optional.
- **C3** (often known as Sleep) is a state where the processor does not need to keep its cache, but maintains other state. This processor state is optional.



P-States:

- **P0** max power and frequency
- **P1** less than P0, voltage/frequency scaled
- **P2** less than P1, voltage/frequency scaled
- ...
- **P_n** less than P(_n-1), voltage/frequency scaled



- Terminology
 - R : Power spent on computation
 - V : Processor voltage
 - f : Processor clock frequency
 - R_0 : Leakage power
- Power spent on computation is:
 - $R = k_v V^2 f + R_0$
where k_v is a constant

Energy of Computation



- Power spent on computation is:
 - $R = k_v V^2 f + R_0$
- Consider a piece of computation of length C clock cycles and a processor operating at frequency f
- The execution time is $t = C/f$
- Energy spent is:
 - $E = R t = (k_v V^2 f + R_0)(C/f)$



- Power spent on computation is:
 - $R = k_v V^2 f + R_0$
- Energy spent is:
 - $E = R t = (k_v V^2 f + R_0)(C/f)$
- Question:
 - Does it make sense to operate the processor at a reduced speed to save energy? Why or why not?



Is reducing processor frequency good or bad?

- Does it make sense to operate the processor at a reduced speed to save energy? Why or why not?

Possible Answer:

$$E = R t = (k_v V^2 f + R_0)(C/f) = k_v V^2 C + R_0 C/f$$

- Conclusion: E is minimum when f is maximum.
 - Operate at top speed
- Is this really true? What are the underlying assumptions?



Reducing voltage and frequency:

- In reality, processor voltage can be decreased if clock frequency is decreased
 - Voltage and frequency can be decreased roughly proportionally.
 - In this case (where $V \sim f$):

$$R = k_f f^3 + R_0$$

$$E = (k_f f^3 + R_0)(C/f) = k_f f^2 C + R_0 C/f$$



Reducing voltage and frequency:

- In reality, processor voltage can be decreased if clock frequency is decreased
 - Voltage and frequency can be decreased roughly proportionally.
 - In this case (where $V \sim f$):

$$R = k_f f^3 + R_0$$

$$E = (k_f f^3 + R_0)(C/f) = k_f f^2 C + R_0 C/f$$

- *Question: Does reducing frequency (and voltage) increase or decrease total energy spend on a task?*



- There exists a minimum frequency below which no energy savings are achieved

$$E = k_f f^2 C + R_0 C / f$$

$$dE/df = 2k_f f C - R_0 C / f^2 = 0$$

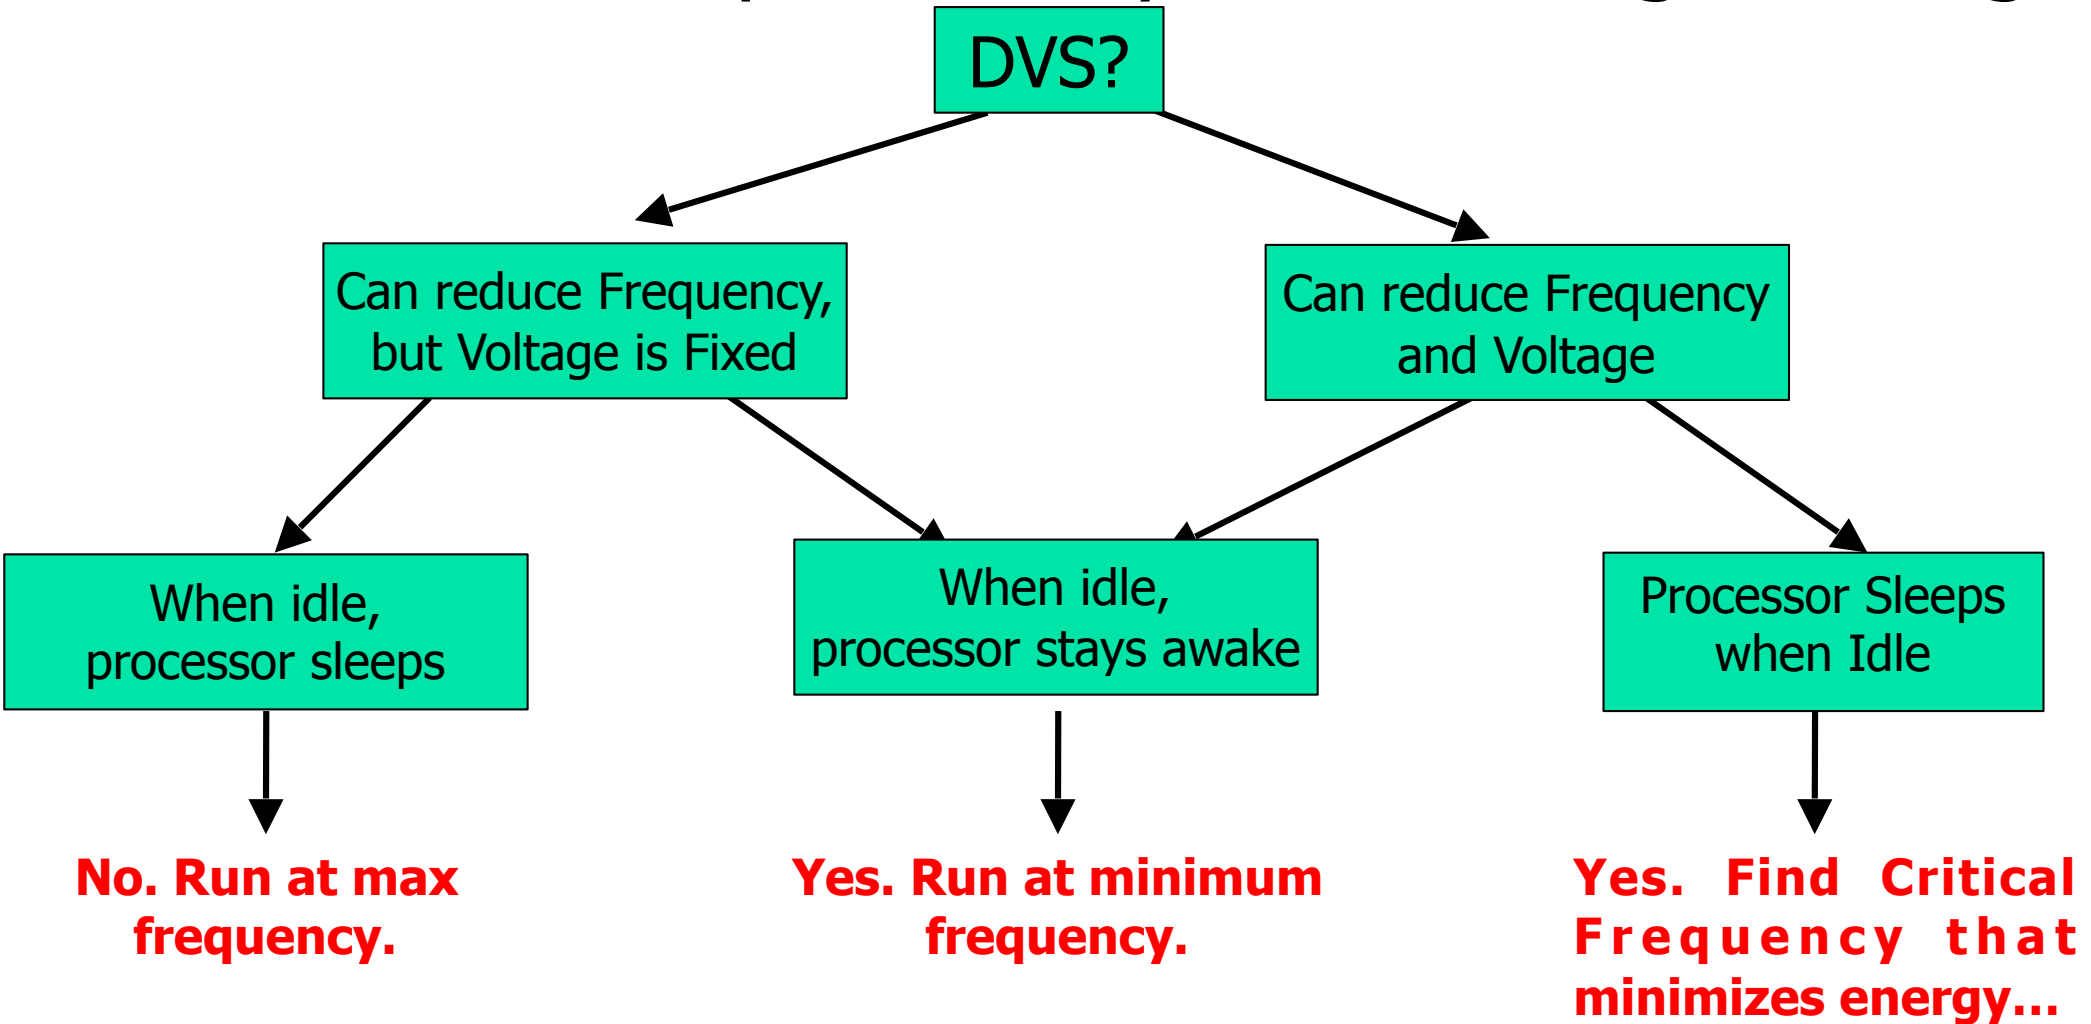
$$f = \sqrt[3]{\frac{R_0}{2k_f}}$$



- Linux defines multiple DVS modes (called CPUfreq “governors”):
 - Performance (highest frequency)
 - Powersave (lowest frequency)
 - Userspace (“root” user controls frequency)
 - OnDemand (adaptively change frequency depending on load)



When should we perform dynamic voltage scaling?



Accounting for Off-chip



- In the preceding discussion, we assumed that task execution time at frequency f is C/f , where C is the total cycles needed
- In reality some cycles are lost waiting for memory access and I/O (Off-chip cycles).
 - Let the number of CPU cycles used be C_{cpu} and the time spent off-chip be $C_{off-chip}$
 - Execution time at frequency f is given by

$$C_{cpu}/f + C_{off-chip}$$

The Cost of Wakeup



- Turning Processor off...
- Energy expended on wakeup, E_{wake}
- To sleep or not to sleep?

The Cost of Wakeup



- Turning Processor off...
- Energy expended on wakeup, E_{wake}
- To sleep or not to sleep?
 - Not to sleep (for time t):

$$E_{no-sleep} = (k_v V^2 f + R_0) t$$

- To sleep (for time t) then wake up:

$$E_{sleep} = P_{sleep} t + E_{wake}$$

The Cost of Wakeup



- Turning Processor off...
- Energy expended on wakeup, E_{wake}

- To sleep or not to sleep?

- Not to sleep (for time t):

$$E_{no-sleep} = (k_v V^2 f + R_0) t$$

- To sleep (for time t) then wake up:

$$E_{sleep} = P_{sleep} t + E_{wake}$$

- To save energy by sleeping: $E_{sleep} < E_{no-sleep}$

$$t > \frac{E_{wake}}{k_v V^2 f + R_0 - P_{sleep}}$$

The Cost of Wakeup



- Turning Processor off...
- Energy expended on wakeup, E_{wake}

- To sleep or not to sleep?

- Not to sleep (for time t):

$$E_{no-sleep} = (k_v V^2 f + R_0) t$$

- To sleep (for time t) then wake up:

$$E_{sleep} = P_{sleep} t + E_{wake}$$

- To save energy by sleeping: $E_{sleep} < E_{no-sleep}$

$$t > \frac{E_{wake}}{k_v V^2 f + R_0 - P_{sleep}}$$

**Minimum sleep
interval**



- DPM refers to turning devices off (or putting them in deep sleep modes)
- Device wakeup has a cost that imposes a minimum sleep interval (a breakeven time)
- DPM must maximize power savings due to sleep while maintaining schedulability



**How does dynamic
power management
affect scheduling?**

The Problem with work-conserving scheduling:

Task 1 ($C=2$, $P=12$)



Task 2 ($C=1$, $P=16$)



The Problem with work-conserving scheduling:

Task 1 ($C=2$, $P=12$)



Task 2 ($C=1$, $P=16$)



Minimum sleep period

The Problem with work-conserving scheduling:

Task 1 ($C=2, P=12$)



Task 2 ($C=1, P=16$)



No opportunity to sleep! : - (



Minimum sleep period

The Problem with work-conserving scheduling:

Task 1 ($C=2$, $P=12$)



Task 2 ($C=1$, $P=16$)



Solution: Must batch!

Minimum sleep period



- From the perspective of minimizing energy, is it always a good idea to use up all processors?

How many proc to use?



- Consider using one processor at frequency f versus two at frequency $f/2$
- Case 1: Total power for one processor
 - $k_f f^3 + R_0$
- Case 2: Total power for two processors
 - $2 \{k_f (f/2)^3 + R_0\} = k_f f^3 / 4 + 2 R_0$

How many proc to use?



- Consider using one processor at frequency f versus two at frequency $f/2$
- Case 1: Total power for one processor
 - $k_f f^3 + R_0$
- Case 2: Total power for two processors
 - $2 \{k_f (f/2)^3 + R_0\} = k_f f^3 / 4 + 2 R_0$
- The general case: n processors
 - $n \{k_f (f/n)^3 + R_0\} = k_f f^3 / n^2 + n R_0$

How many proc to use?



- The general case: n processors
 - $Power = n \{k_f (f/n)^3 + R_0\} = k_f f^3 / n^2 + n R_0$
 - $dPower/dn = -2 k_f f^3 / n^3 + R_0 = 0$

$$n = \sqrt[3]{\frac{2k_f f^3}{R_0}}$$

- What if n is not an integer?