

Goals for Today



- Learning Objective:
 - Present final exam details + review content
- Announcements, etc:
 - MP3 Soft Extension: Submit by **TODAY** for -10pts
 - MP4 due **May 7th**
 - Deadline provides more time than necessary; wanted to give you flexibility
 - vSphere console was temporarily down, is back up now
 - Review Homework will be posted later today.



Reminder: Please put away devices at the start of class



CS 423

Operating System Design: Final Exam Overview

Professor Adam Bates
Spring 2018

Final Exam Details



- May 4th, 1:30pm - 3:30pm
- You will have 2 hours
- Scantron Multiple choice
- 30-40 Questions
 - Questions per minute will be less than Midterm
- **Openbook:** Textbooks, paper notes, printed sheets allowed. No electronic devices permitted (or necessary)!
- **Content:** All lecture and text material covered after the midterm content (i.e., starting with Virtualization)



Final Exam Content



- Virtualization (Emulation, Binary Translation...)
- File Systems (Disk Scheduling, Directories, Reliability...)
- Security (Access control, Encryption, Attacks, Reference monitors)
- Guest Lectures (Hardware Attacks, Process VMs)
- Remaining Special Topics (Energy, Linux Audit)
- *Exam questions will not be explicitly cumulative, but I can't guarantee that content from before the midterm won't come up in some fashion.*

Virtualization



- Key Concepts:
 - Different purposes for virtualization
 - Different virtualization layers
 - Emulation versus Binary Translation
 - Dynamic Binary Translation Challenges + Optimizations
 - Challenges of Process VMs
 - e.g., Emulating Target Architecture
 - Interpretation/Emulation versus Translation

What's a virtual machine?



- Virtual machine is an entity that emulates a guest interface on top of a host machine
 - Language view:
 - Virtual machine = Entity that emulates an API (e.g., JAVA) on top of another
 - Virtualizing software = compiler/interpreter
 - Process view:
 - Machine = Entity that emulates an ABI on top of another
 - Virtualizing software = runtime
 - Operating system view:
 - Machine = Entity that emulates an ISA
 - Virtualizing software = virtual machine monitor (VMM)

Different views == who are we trying to fool??

Purpose of a VM



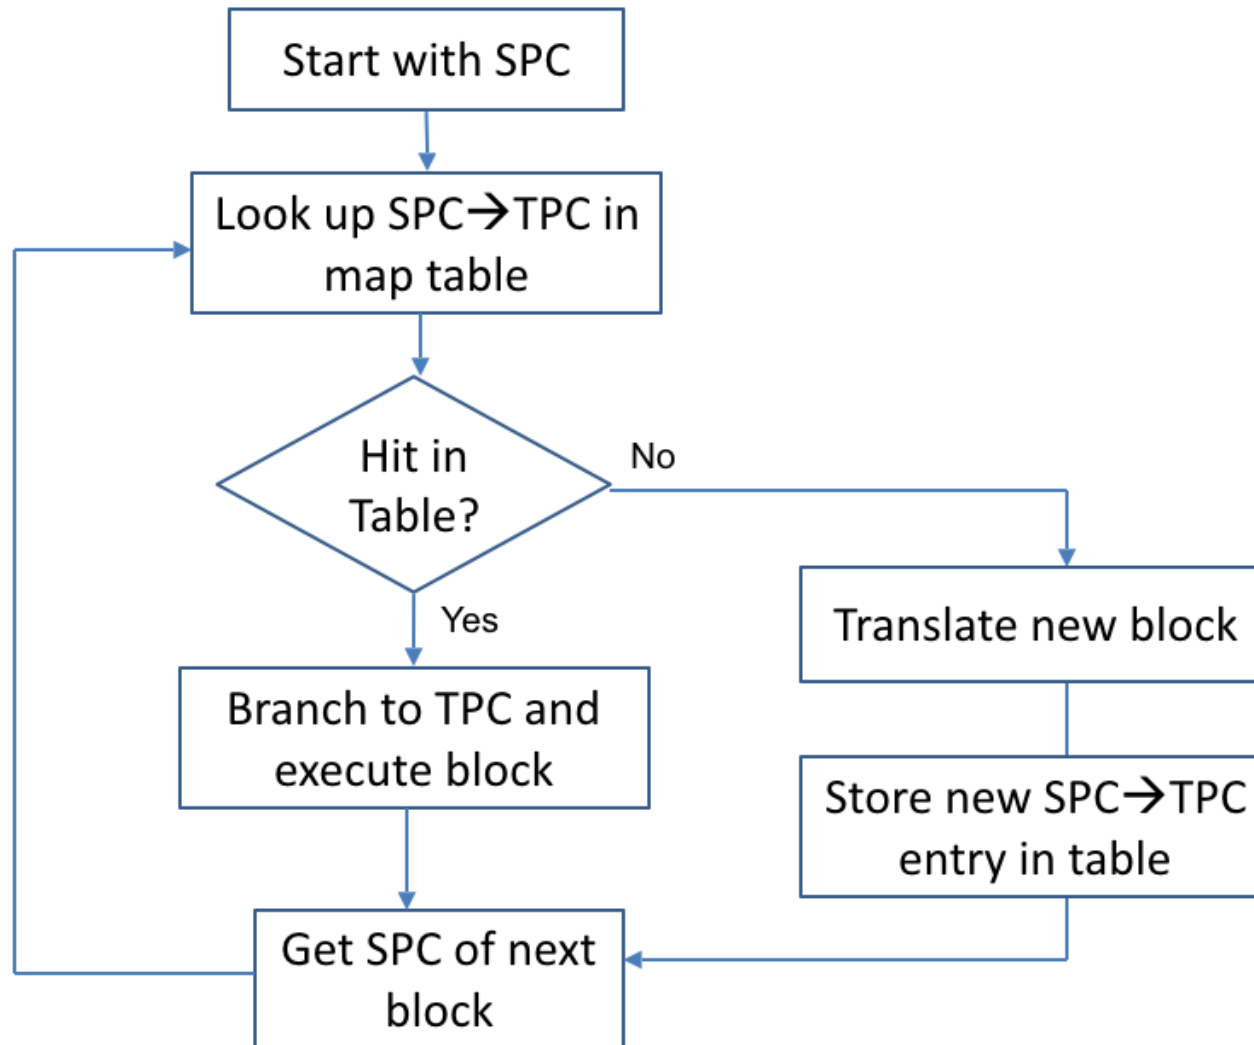
- Emulation
 - Create the illusion of having one type of machine on top of another
- Replication (/ Multiplexing)
 - Create the illusion of multiple independent smaller guest machines on top of one host machine (e.g., for security/isolation, or scalability/sharing)
- Optimization
 - Optimize a generic guest interface for one type of host

Writing an Emulator



- Problem: Emulate guest ISA on host ISA
- Create a simulator data structure to represent:
 - Guest memory
 - Guest stack
 - Guest heap
 - Guest registers
- Inspect each binary instruction (machine instruction or system call)
 - Update the data structures to reflect the effect of the instruction

Dynamic Binary Translation



Instruction Emulation



- Interpretation versus binary translation?

- Interpretation:


- no startup overhead
- High overhead per instruction

- Binary translation:

- High startup overhead
- Low overhead per instruction

- Can we combine the best of both worlds?

- Small program: Do interpretation
- Large program: Do binary translation





- Key Concepts:
 - Disk Scheduling
 - Concepts + Modern Implementations
 - Data Layout on Disk
 - File Allocation Strategies
 - Concepts + Modern Implementations
 - Locality
 - Directory Structures
 - Representing Large Directories
 - Reliability
 - Transaction Concept + Implementations
 - RAID

Disk Scheduling




- Which disk request is serviced first?
 - FCFS
 - Shortest seek time first
 - Elevator (SCAN)
 - C-SCAN (Circular SCAN)

A: Track.

B: Sector.

C: Sector of Track.

D: File



Disk Scheduling Decision — Given a series of access requests, on which track should the disk arm be placed next to maximize fairness, throughput, etc?

Linux I/O Schedulers



- What disk (I/O) schedulers are available in Linux?

```
$ cat /sys/block/sda/queue/scheduler  
noop [deadline] cfq
```

^ scheduler enabled on our VMs

- As of Linux 2.6.10, it is possible to change the IO scheduler for a given block device on the fly!
- How to enable a specific scheduler?

```
$ echo SCHEDNAME > /sys/block/DEV/queue/scheduler
```

- SCHEDNAME = Desired I/O scheduler
- DEV = device name (e.g., hda)

Disk Layout for a FS



Disk layout in a typical file system:

| | | | |
|---------------|----------------|-----------------------------------|------------------|
| Boot block | Super block | File metadata (i-node in Unix) | File data blocks |
|---------------|----------------|-----------------------------------|------------------|

- Superblock defines a file system
 - size of the file system
 - size of the file descriptor area
 - free list pointer, or pointer to bitmap
 - location of the file descriptor of the root directory
 - other meta-data such as permission and various times
- For reliability, replicate the superblock

Contiguous Allocation




- Request in advance for the size of the file
- Search bit map or linked list to locate a space
- File header
 - first sector in file
 - number of sectors
- Pros
 - Fast sequential access
 - Easy random access
- Cons
 - External fragmentation
 - Hard to grow files

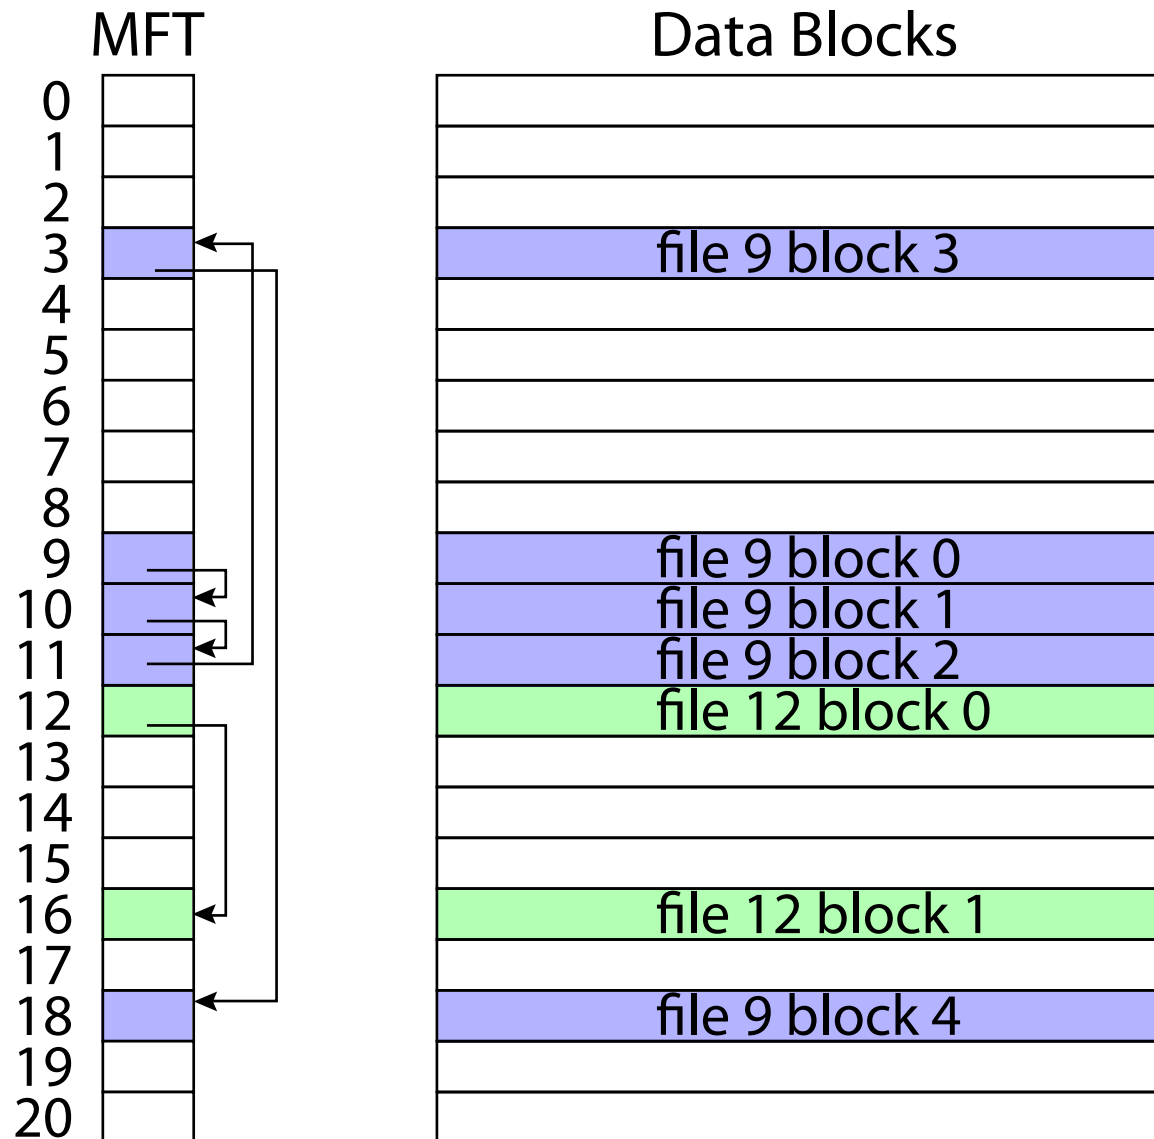
Linked Files




- File header points to 1st block on disk
- Each block points to next
- Pros
 - Can grow files dynamically
 - Free list is similar to a file
- Cons
 - random access: horrible
 - unreliable: losing a block means losing the rest



MS File Allocation Table (FAT)



Berkeley FFS / UNIX FS

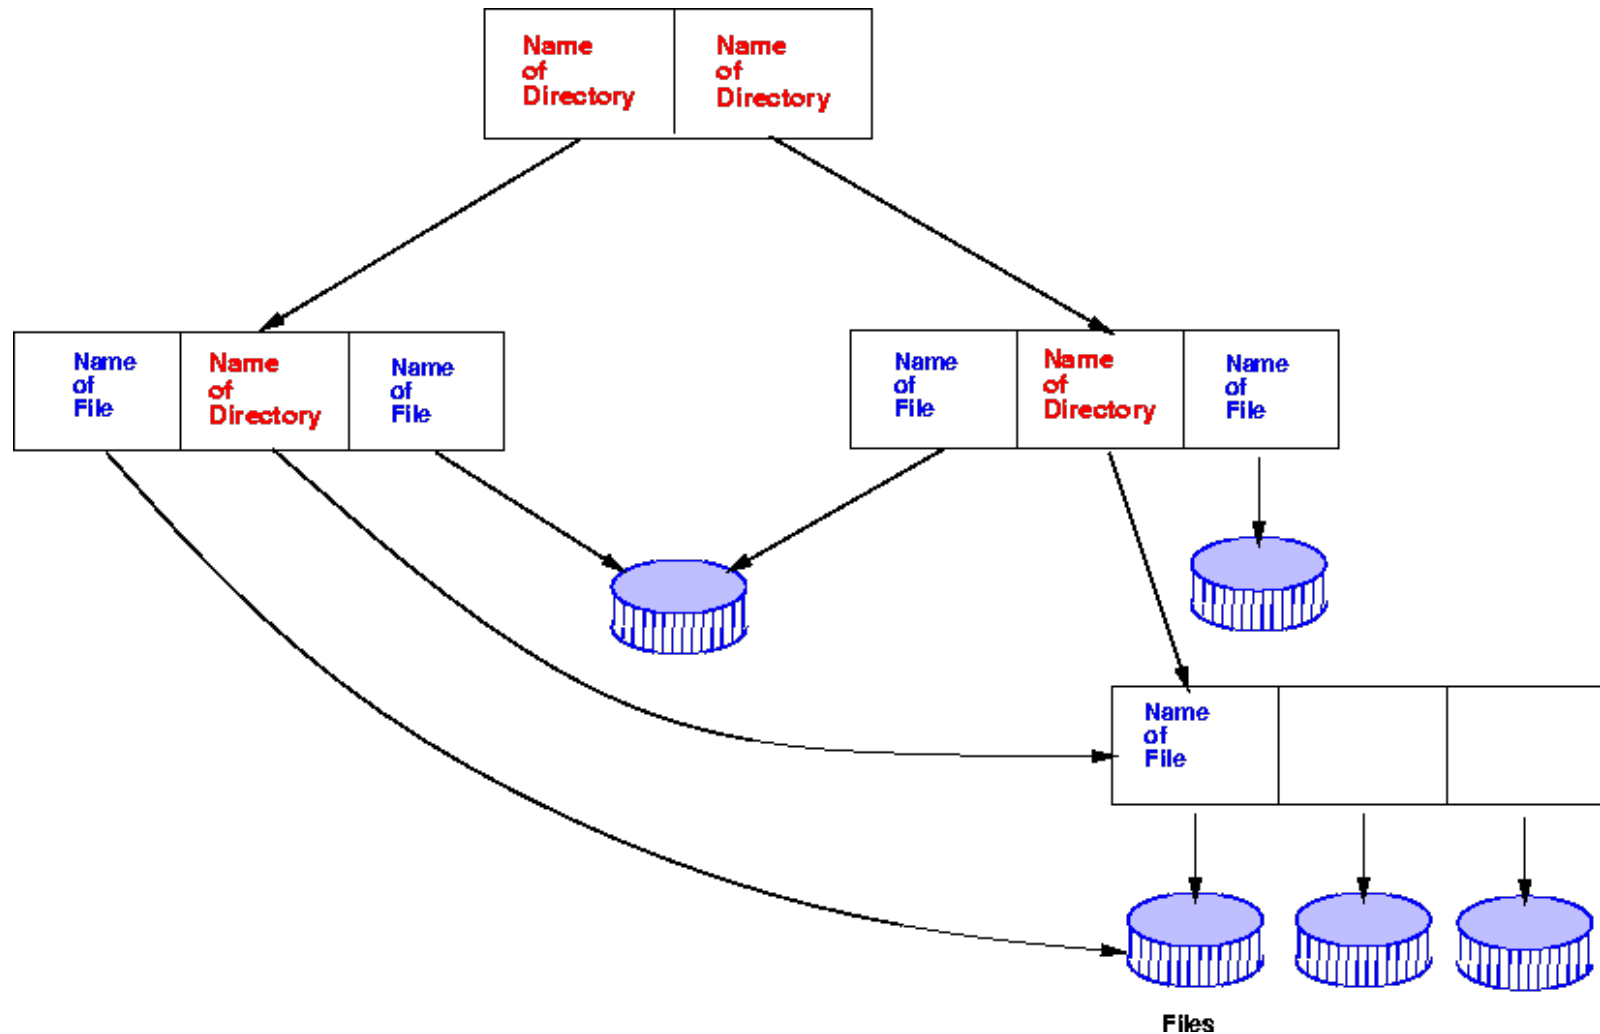


Berkeley FFS Locality



- How does FFS provide locality?
- Block group allocation
 - Block group is a set of nearby cylinders
 - Files in same directory located in same group
 - Subdirectories located in different block groups
- inode table spread throughout disk
 - inodes, bitmap near file blocks
- First fit allocation
 - Property: Small files may be a little fragmented, but large files will be contiguous


Acyclic Graph Structured Dir's



Directory Layout



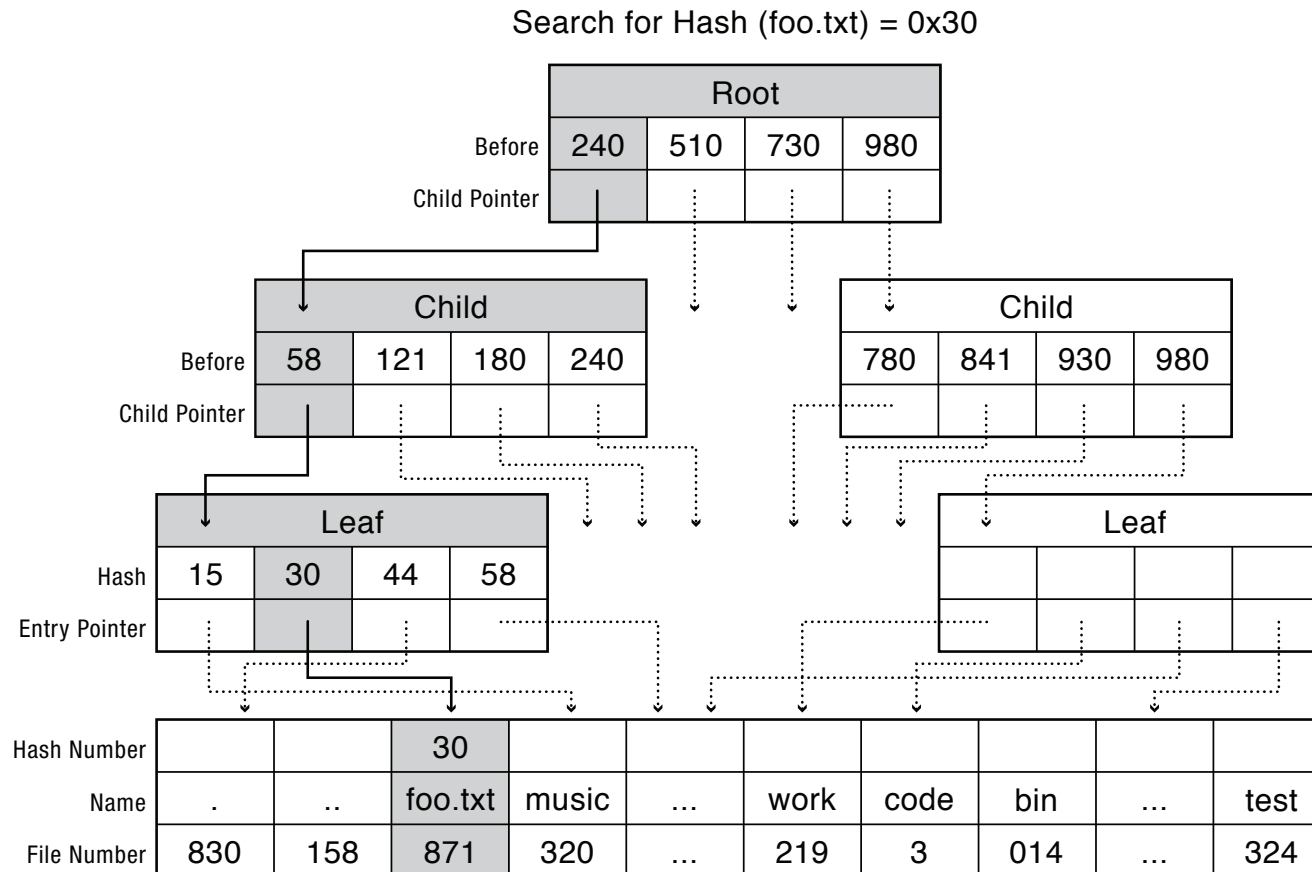
- Represent directory as a list of files
- Linear search to find filename
- Suitable for small directories



B Trees



- Logarithmic search to find filename
- Suitable for large directories



Transaction Concept



- A transaction is a grouping of low-level operations that are related to a single logical operation
- Transactions are atomic — operations appear to happen as a group, or not at all (at logical level)
 - At physical level of course, only a single disk/flash write is atomic
- Transactions are durable — operations that complete stay completed
 - Future failures do not corrupt previously stored data
- (In-Progress) Transactions are isolated — other transactions cannot see the results of earlier transactions until they are committed
- Transactions exhibit consistency — sequential memory model



Pros

- Works with minimal support from the disk drive
- Works for most multi-step operations

Cons

- Can require time-consuming recovery after a failure
- Difficult to reduce every operation to a safely-interruptible sequence of writes
- Difficult to achieve consistency when multiple operations occur concurrently (e.g., FFS grep)



Pros

- Correct behavior regardless of failures
- Fast recovery (root block array)
- High throughput (best if updates are batched)

Cons

- Potential for high latency
- Small changes require many writes
- Garbage collection essential for performance

Logging File Systems



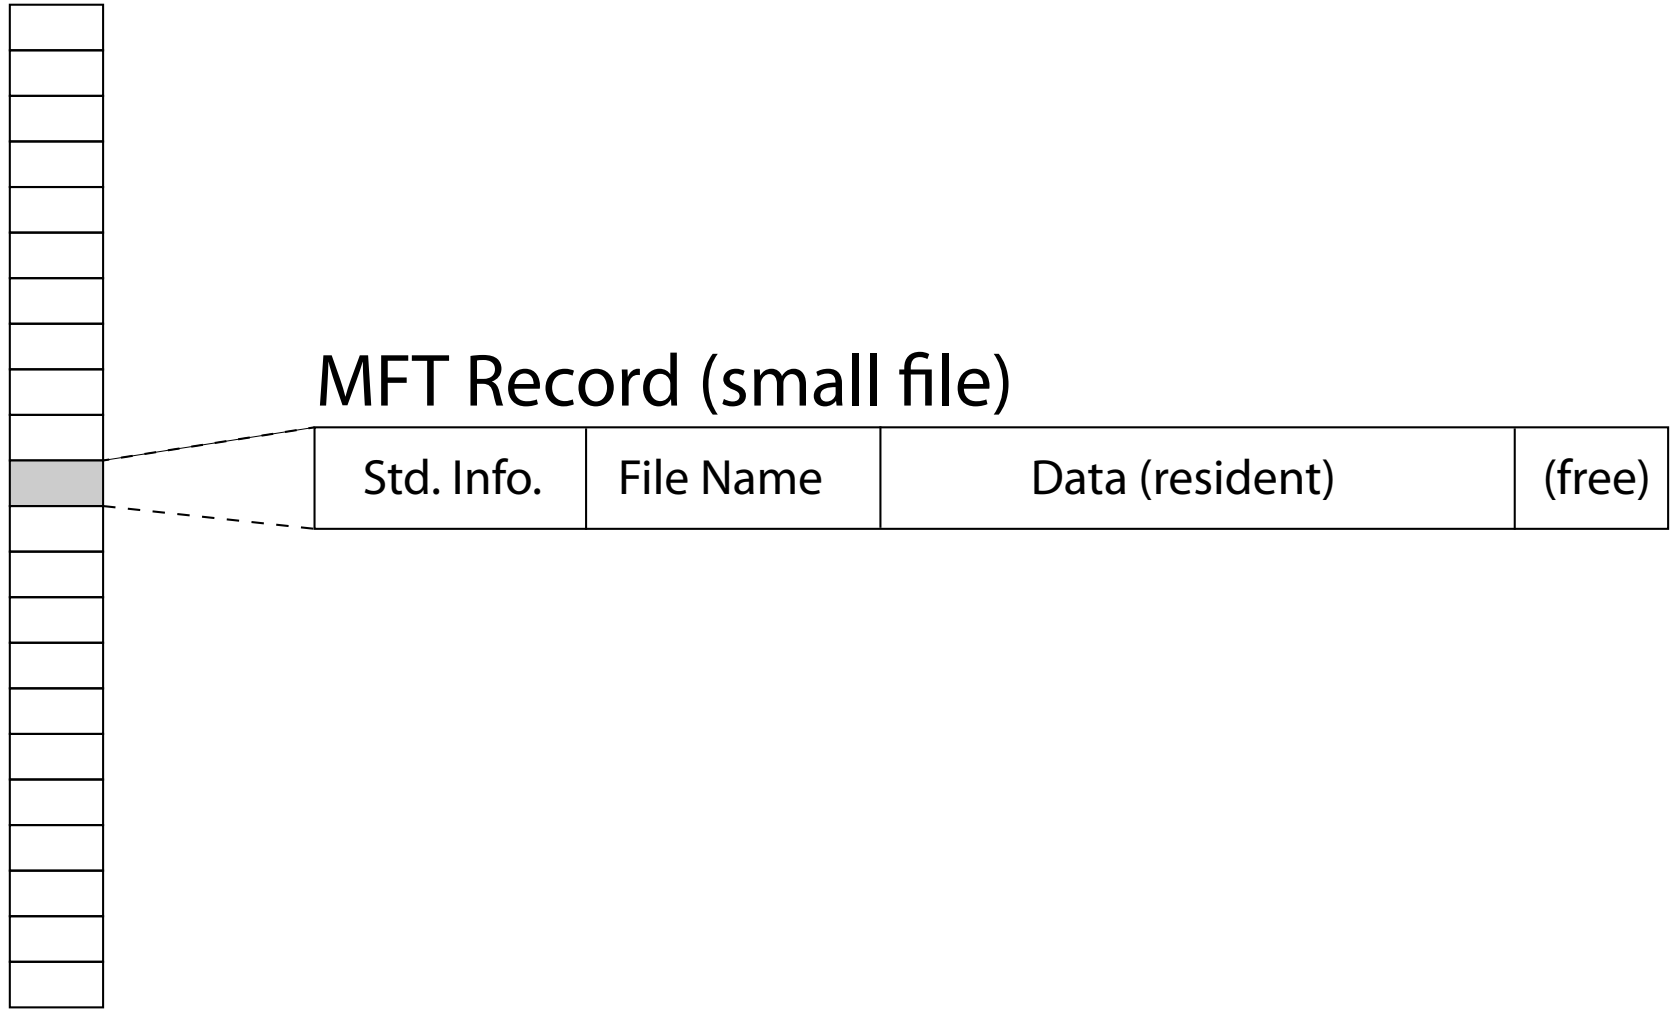
- Instead of modifying data structures on disk directly, write changes to a journal/log
 - Intention list: set of changes we intend to make
 - Log/Journal is append-only
- Once changes are on log, safe to apply changes to data structures on disk
 - Recovery can read log to see what changes were intended
- Once changes are copied, safe to remove log

What about NTFS?

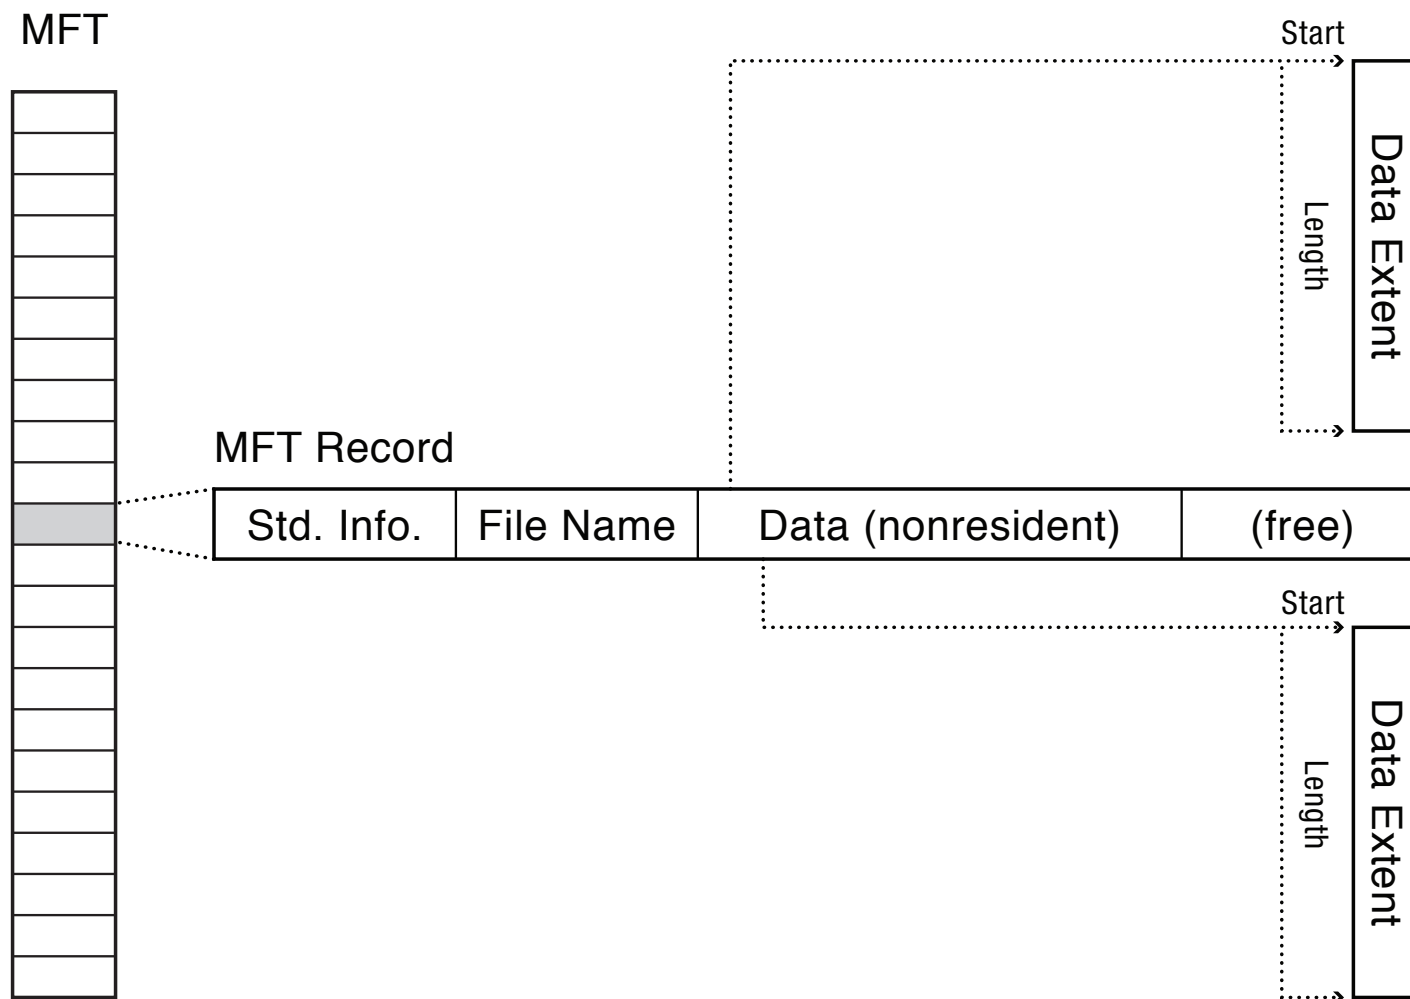


- Improved Metadata support
 - Flexible 1KB storage for metadata and data
- Scalability Features
 - MFT is optimized for 4KB resident data
 - Extents: a middle ground between contiguous and non-contiguous allocation.
 - Block pointers cover runs of blocks
 - Similar approach in linux (ext4)
- NTFS uses journalling for reliability

Master File Table



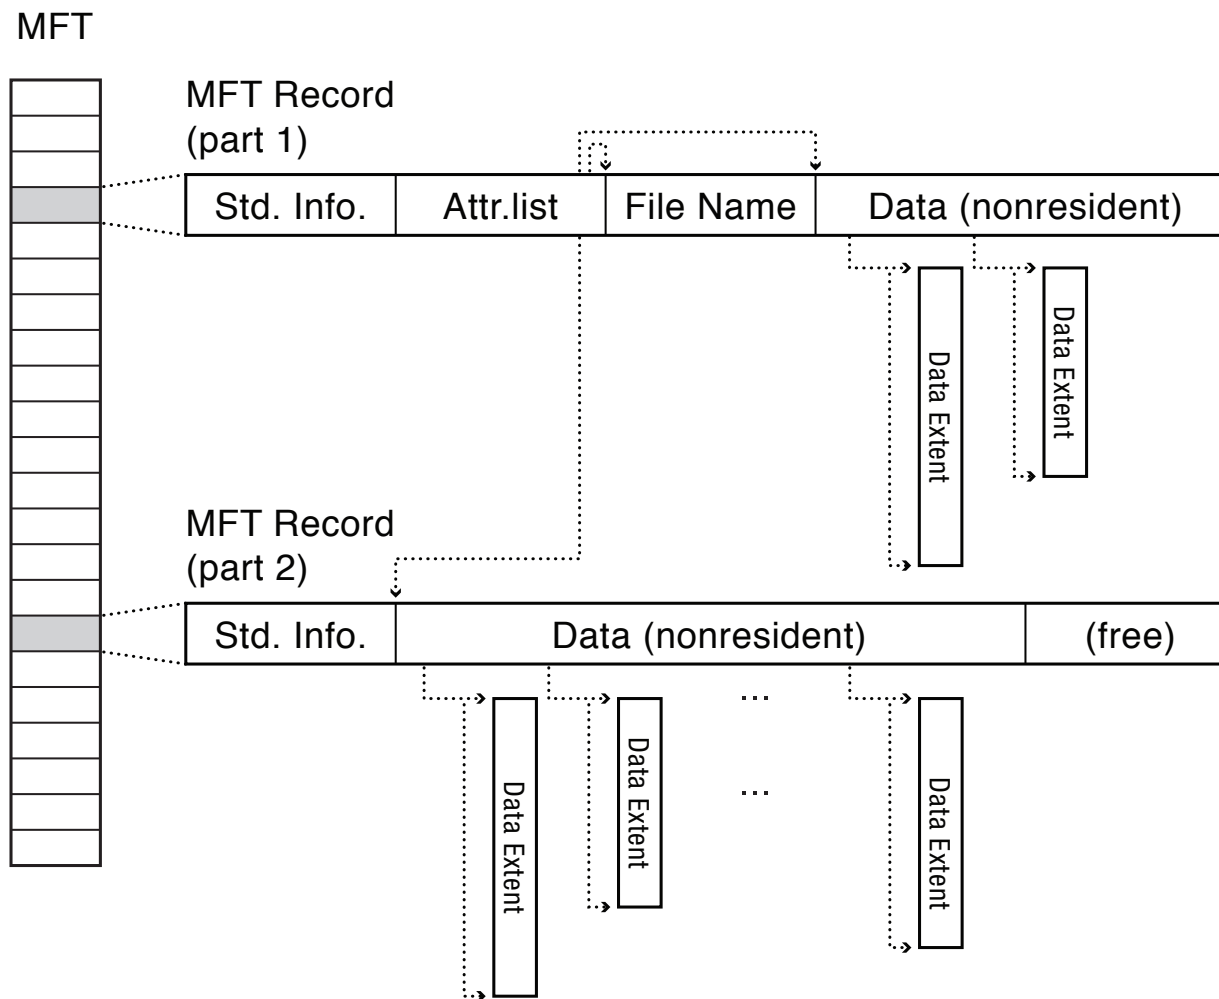
NTFS



NTFS Indirect Block



What if file is too large to fit all extent pointers in one data cluster?






- RAID
 - multiple disks work cooperatively
 - Improve reliability by storing redundant data
 - **Striping (RAID 0)** improves performance with disk **striping** (use a group of disks as one storage unit)
 - **Mirroring (RAID 1)** keeps duplicate of each disk
 - Striped mirrors (**RAID 1+0**) or mirrored stripes (**RAID 0+1**) provides high performance and high reliability
 - **Block interleaved parity (RAID 4, 5, 6)** uses much less redundancy

RAID Level 0




- Level 0 is nonredundant disk array
- Files are striped across disks, no redundant info
- High read throughput
- Best write throughput (no redundant info to write)
- Any disk failure results in data loss



RAID Level 1



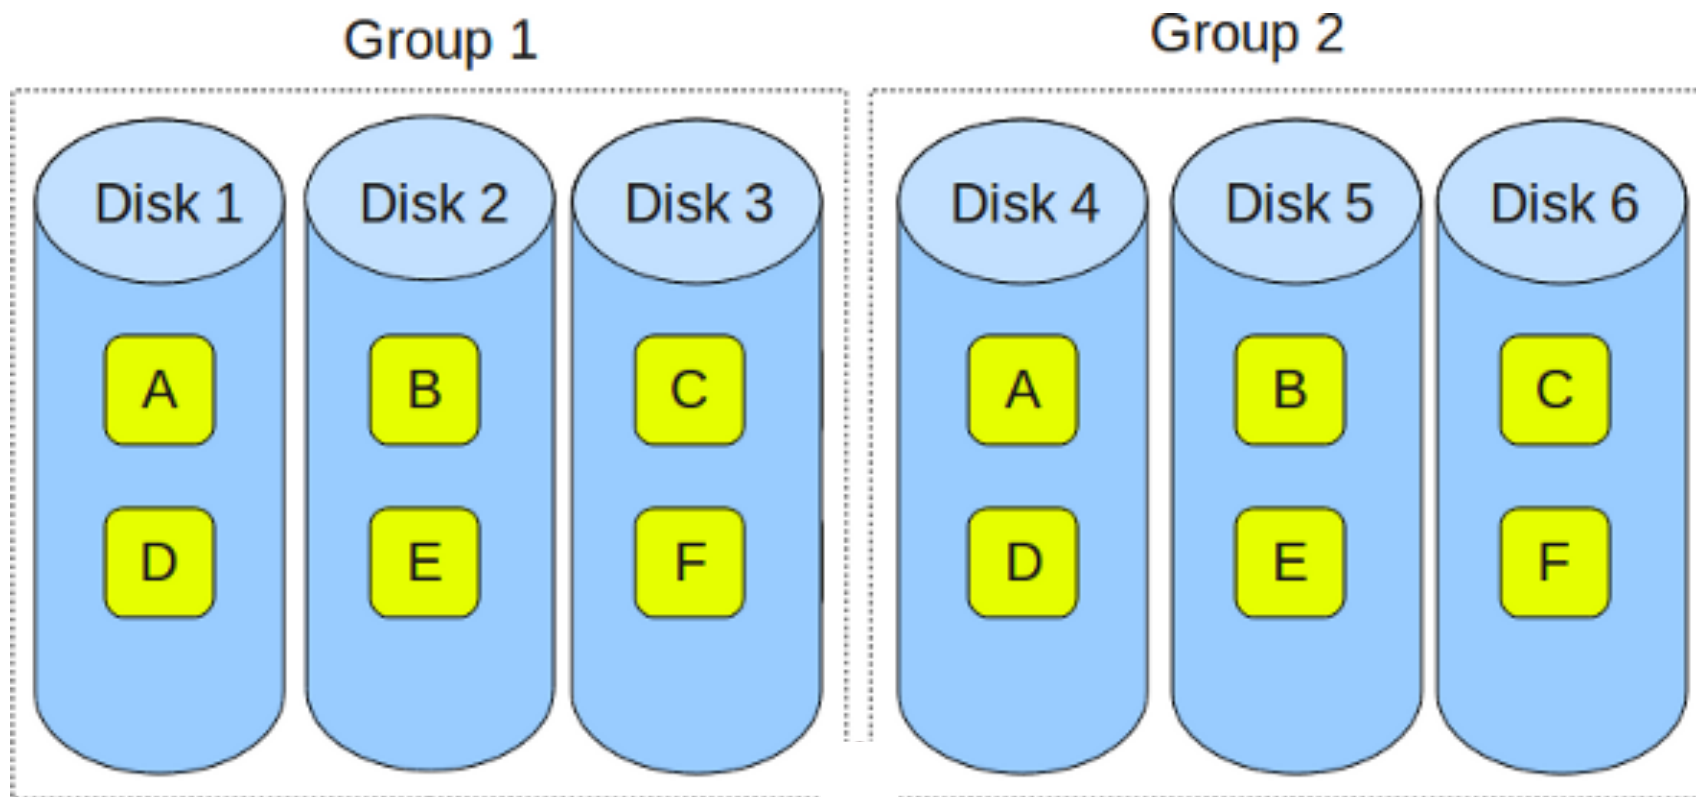
- Mirrored Disks
- Data is written to two places
 - On failure, just use surviving disk (easy to rebuild)
- On read, choose fastest to read
 - Write performance is same as single drive, read performance is 2x better
- Expensive (high space overhead)



RAID Level 0+1



- Stripe on a set of disks
- Then mirror of data blocks is striped on the second set.




RAID 01 – Blocks Striped. (and Blocks Mirrored)

RAID Level 1+0



- Pair mirrors first.
- Then stripe on a set of paired mirrors



RAID 10 – Blocks Mirrored. (and Blocks Striped)




- Key Concepts:
 - Least Privilege
 - Encryption — have a High-level / Block Box comprehension. You won't need to prove RSA.
 - Authentication + Passwords
 - Why do secure systems (epically) fail?
 - Cache Side-Channels (Fletcher lecture)
 - Access Control
 - e.g., DAC, Capabilities, Bell-LaPadula
 - Cryptography versus Access Control
 - Reference Monitors, LSM, SELinux

How to study for security?



- No corresponding chapter in textbook!?!?
- We won't be straying far from the lectures/slides
- Google + Wikipedia concepts if you need further clarification (just evaluate credibility of sources).




Principle of Least Privilege




- Grant each principal the least permission possible for them to do their assigned work
 - Minimize code running inside kernel
 - Minimize code running as sysadmin
- Practical challenge: hard to know
 - what permissions are needed in advance
 - what permissions should be granted
 - Ex: to smartphone apps
 - Ex: to servers

Symmetric Key (DES, IDEA)



- Single key (symmetric) is shared between parties, kept secret from everyone else
 - Ciphertext = $(M)^K$; Plaintext = $M = ((M)^K)^K$
 - if K kept secret, then both parties know M is authentic and secret

Public Key (RSA, PGP)




Keys come in pairs: public and private

- $M = ((M)^{K\text{-public}})^{K\text{-private}}$
- Ensures secrecy: can only be read by receiver

2-Factor Authentication




- Can be difficult for people to remember encryption keys and passwords
- Instead, store K -private inside a chip
 - use challenge-response to authenticate smartcard
 - Use PIN to prove user has smartcard



Ex1: Tenex Password Vuln



- Observation: Programs have **a lot** of control over how their virtual memory works.
- Attack #1: Trap-To-User Bit Exploit



Trap-To-User: Alert me if this 2nd page is accessed!

- Attack #2: Exploit timing side-channel

Processing time for password check was proportional to the number of correct characters at the front of the attacker's guess.

Reflections on Trusting Trust



- Thompson's Takeaway: You can't fully trust code that you didn't write yourself!
- Presented as a thought experiment during Thompson's Turing Award Lecture. Didn't really happen... we think??
- Hard to re-secure a machine after penetration. How do you know you've removed all the backdoors?
- It's hard to detect that a machine has been penetrated
- Any system with bugs is vulnerable
 - and all systems have bugs

Discretionary Access Control (DAC)



Access Mask defines permissions for User, Group, and Other

```
C:\WINDOWS\system32\cmd.exe - ssh lootah@polaris.cse.psu.edu
drwx----- 2 lootah gcse 2048 Mar 2 2005 .ssh
-rw-r--r-- 1 lootah gcse 0 Mar 8 22:21 ssl.txt
drwxr-xr-x 3 lootah gcse 2048 Mar 24 11:27 .subversion
-rw-r--r-- 1 lootah gcse 77 Mar 23 16:23 .sversionrc
drwx----- 3 lootah gcse 2048 Mar 28 17:49 .thumbnails
drwx----- 7 lootah gcse 2048 Jun 20 14:57 .Trash
-rw----- 1 lootah gcse 5828 Mar 1 2005 .TAuthority
-rw-r--r-- 1 lootah gcse 5884 Jul 3 1999 .twmrc
-rw----- 1 lootah gcse 737 Mar 8 17:36 .viminfo
drwxr-xr-x 2 lootah gcse 2048 Jun 30 19:21 www
-rw----- 1 lootah gcse 62 Apr 2 17:54 .xauthkKIS4n
-rw----- 1 lootah gcse 64 Apr 2 18:06 .xauthNphamR
-rw----- 1 lootah gcse 5455 Aug 7 14:52 .xauthority
-rw----- 1 lootah gcse 62 Apr 2 18:23 .xauthSH9Uuv
-rw----- 1 lootah gcse 62 Apr 2 18:24 .xauthHzecq
-rw----- 1 lootah gcse 62 Apr 2 18:20 .xauthHJSG51
-rw----- 1 lootah gcse 2103 Sep 1 10:55 .xdefaults
-rw----- 1 lootah gcse 70712 Aug 25 2004 .xftcache
-rw-r--r-- 1 lootah gcse 2433 Jul 3 1999 .xresources
-rw-r--r-- 1 lootah gcse 10496 May 5 17:58 .xscreensaver
lrwxrwxrwx 1 lootah gcse 9 Jun 11 06:09 .xsession -> .Xsession
-rwxr-xr-x 1 lootah gcse 4495 Mar 4 2005 .Xsession
-rw-r--r-- 1 lootah gcse 255 Feb 1 2005 .xsession-errors
drwx----- 2 lootah gcse 2048 Sep 22 2004 .xpics
[lootah@polaris ~]$
```

`chmod u=rwx,g=rx,o=r myfile`
`chmod 754 myfile` <- Same thing

4 stands for "read",
2 stands for "write",
1 stands for "execute", and
0 stands for "no permission."

Problems?




- What might go wrong with DAC or Capabilities?
 - Security is left to the discretion of subjects
 - Impossible to guarantee security of system
 - Security of system changes over time.
- Solution?
 - Mandatory Access Control: Operating system constrains the ability of subjects (even owners) to perform operations on objects according to a system-wide *security policy*.

Bell-LaPadula Model



- A multi-level security model that provides strong confidentiality guarantees.
- Formalizes Classified Information
- State machine (Lattice) specifies permissible actions





- Designed by the NSA
- A more flexible solution than MLS
- SELinux Policies are comprised of 3 components:
 - Labeling State defines security contexts for every file (object) and user (subject).
 - Protection State defines the permitted $\langle \text{subject}, \text{object}, \text{operation} \rangle$ tuples.
 - Transition State permits ways for subjects and objects to move between security contexts.
- Enforcement mechanism designed to satisfy reference monitor concept

LSM Architecture



- Linux Kernel modified in 5 ways:
 - Opaque security fields added to certain kernel data structures
 - Security hook function calls inserted at various points with the kernel code
 - A generic security system call added
 - Function to allow modules to register and unregister as security modules
 - Move capabilities logic into an optional security module

Reference Monitor



Cool. But how do we implement these models in an operating system?

