



CS 423

Operating System Design: MP4 Walkthrough

Mohammad Noureddine
Spring 2018

Goals for Today



- Learning Objective:
 - Understand Linux Security Modules
 - Go through implementation details of MP4
- Announcements, etc:
 - MP3 Soft Extension
 - No office hour next week!
 - Make up office hour: Monday 04/30 at 3:30 pm and Piazza



Reminder: Please put away devices at the start of class

Preliminaries



- Please **do NOT revert to snapshots** taken prior to the migration of the VMs
- **Take stable snapshots** before starting this MP
- Your security module will **affect kernel boot**
 - Work incrementally
 - Start with empty functions, add logic in small doses
- Follow MP **submission instructions** carefully

Goals of this MP



- Understand Linux Security Modules
- Understand basic concepts behind Mandatory Access Control (MAC)
- Understand and use filesystem extended attributes
- Add custom kernel configuration parameters and boot parameters
- Obtain least privilege policy for `/usr/bin/passwd`

Linux Security Modules



- Came out of a presentation that the NSA did in 2001
 - Security Enhanced Linux (SELinux)
- Kernel provided support for Discretionary Access Control
 - Did not provide framework for different security models w/o changes to core kernel code
- Linux Security Modules (LSM) proposed as a solution
 - Not to be fooled by the term “module”
 - LSMs are **NOT** loadable at runtime

Example LSMs

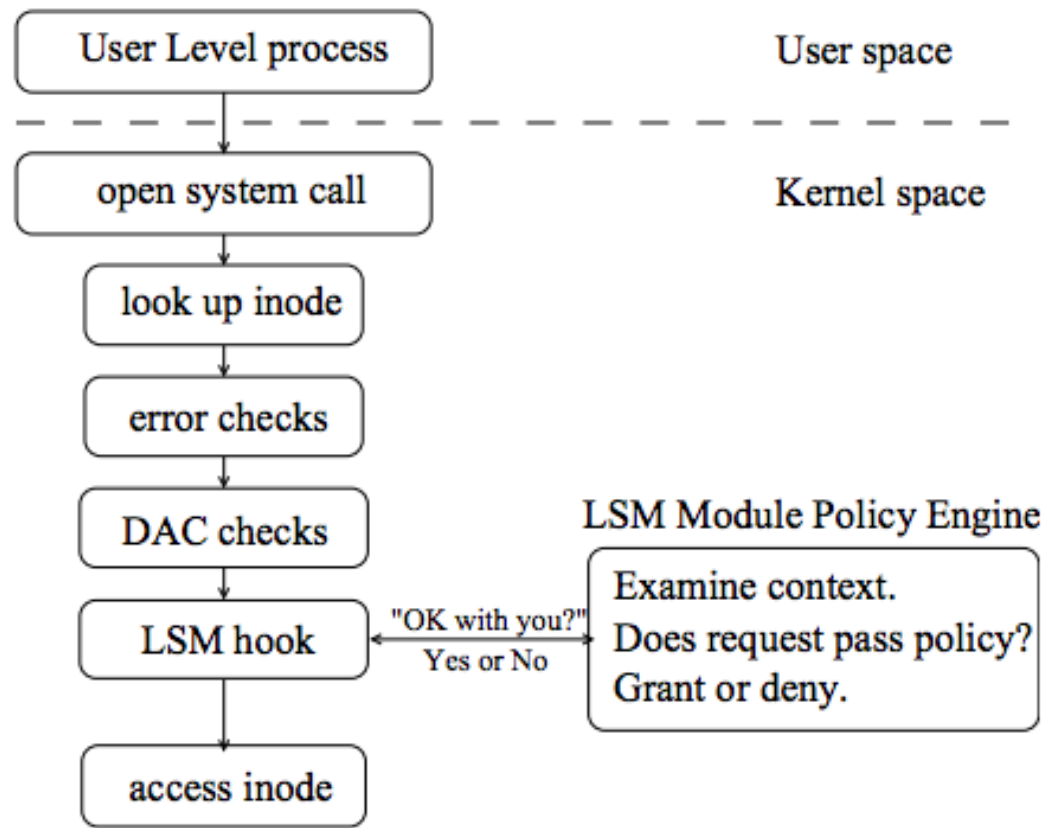


- Some of the LSMs approved in the current kernel
 - AppArmor
 - SELinux
 - Smack
 - TOMOYO Linux
 - Yama
- Must be configured at build-time and at boot time

How Do LSMs Work?



- Hooks inserted throughout important functionalities of the kernel



Question



- In which context does the LSM run?

Question



- Q: In which context does the LSM run?
- A: In the kernel context just before the kernel fulfills a request

```
union security_list_options {
    int (*binder_set_context_mgr)(struct task_struct *mgr);
    int (*binder_transaction)(struct task_struct *from,
                              struct task_struct *to);
    int (*binder_transfer_binder)(struct task_struct *from,
                                  struct task_struct *to);
    int (*binder_transfer_file)(struct task_struct *from,
                                struct task_struct *to,
                                struct file *file);

    int (*ptrace_access_check)(struct task_struct *child,
                              unsigned int mode);
    int (*ptrace_traceme)(struct task_struct *parent);
    int (*capget)(struct task_struct *target, kernel_cap_t *effective,
                  kernel_cap_t *inheritable, kernel_cap_t *permitted);
    int (*capset)(struct cred *new, const struct cred *old,
                  const kernel_cap_t *effective,
                  const kernel_cap_t *inheritable,
                  const kernel_cap_t *permitted);
    int (*capable)(const struct cred *cred, struct user_namespace *ns,
```

Major and Minor LSM



- Major LSM
 - Only one major LSM can run in the system
 - Examples: SELinux, Smack, etc.
 - Can access **subjective security blobs**
 - Data structures reserved only for the use of major LSMs
- Minor LSM
 - Can be stacked
 - Does not have access to the security blobs
 - Examples: YAMA

Security Blobs?



- Reserved fields in various kernel data structures
 - `task_struct`, `inode`, `sk_buff`, `file`, `linux_binprm`
- Controlled by the major security module running
- `struct cred` is the security context of a thread
 - `task->cred->security` is the tasks's subjective security blob
 - A task can only modify its own credentials
 - No need for locks in this case!
 - Need rcu read locks to access another tasks's credentials



- Q: What is Mandatory Access Control anyway?



- Q: What is Mandatory Access Control anyway?
- Access rights are based on regulations defined by a central authority
- Strictly enforced by the kernel
- Label objects by sensitivity
 - Unclassified, confidential, secret, top secret
- Label users (subjects) by clearance
 - Grant access based on combination of subject and object labels

Labeling our System



- We will developed a major security module
- To keep things simple, we will focus on tasks that carry the label **target**
- We will focus on only labeling **inodes**
 - We can use the security blobs
 - Alternatively, we will use **extended filesystem attributes**
- How do we label our tasks then?
 - We will use the **inode label of the binary** that is used to launch the process

FS Extended Attributes



- Provides custom file attributes that are not interpreted by the file system
- Attributes under the prefix `security` will be used for interpretation by an LSM
- We will be using `security.mp4` in our implementation
- e.g.,
 - `setfattr -n security.mp4 -v target target_binary`
 - `setfattr -n <prefix>.<suffix> -v <value> <file>`
 - `getfattr -d -m - <file>`

MP4 Challenges



- Label management
 - How to assign and maintain labels
 - How to transfer labels from inodes to tasks
- Access control
 - Who gets to access what
 - Enforce MAC policy
- Kernel configuration
 - `Kconfig` environment
 - Change boot parameters

Step 1: Compilation



- Customize kernel configuration using the Kconfig environment
- First add custom config option to `security/mp4/Kconfig`

```
1 config SECURITY_MP4_LSM
2     bool "CS423 machine problem 4 support"
3     depends on NET
4     depends on SECURITY
5     select NETLABEL
6     select SECURITY_NETWORK
7     default n
8     help
9         This selects the cs423 machine problem 4 security lsm to be
10        compiled with the kernel.
11        If you are unsure how to answer this question, answer N.
```

Step 1: Compilation



- Now when you run `make oldconfig`, `make` will ask you whether to enable
 - `CONFIG_SECURITY_MP4_LSM`
- You can also use it for static compiler macros in your code. e.g.

```
#ifdef CONFIG_SECURITY_MP4_LSM
void do_something(void) { printf("MP4 active\n"); }
#else
void do_something(void) { }
#endif
```

Step 1: Compilation



- You can also use `make menuconfig` to see your config option visually

```
[*] SHA1 hash of loaded profiles
[*] Yama support
[*] CS423 machine problem 4 support
[*] Integrity subsystem
```

- You might want to turn `DEBUG_INFO` off to speed up the generation of the `.deb` files

Step 1: Compilation



- After the first compilation, you do not need to recompile the entire kernel again
- Reminder: `make clean` removes all of the object files and will cause the entire kernel to be recompiled
- For incremental builds, use: `make -j<num_proc>`
- To produce `.deb` files again:
 - `make bindeb-pkg LOCALVERSION=...`

Step 1: Boot params



- Next we want to enable the mp4 module as the major security module in our system
- The kernel accepts the key-value pair `security=<module>` as part of its boot parameters
- **Update** `/etc/default/grub`:

```
GRUB_CMDLINE_LINUX_DEFAULT="security=mp4"
```

- **Don't forget to update grub!**

Step 2: Implementation



- We will implement our module in three steps:
 1. Register the module and enable it as the only major security module ([Provided to you at no cost in mp4.c](#))
 2. Implement the labels initialization and management
 3. Implement the mandatory access control logic
- We provide you with helper functions in [mp4_given.h](#)
- Use `pr_info`, `pr_err`, `pr_debug`, `pr_warn` macros
 - `#define pr_fmt(fmt) "cs423_mp4: " fmt`

Step 2.1: Startup



- We provide you with the startup code to get your started
- We will implement the following security hooks:

```
static struct security_hook_list mp4_hooks[] = {  
  
    LSM_HOOK_INIT(inode_init_security, mp4_inode_init_security),  
    LSM_HOOK_INIT(inode_permission, mp4_inode_permission),  
  
    LSM_HOOK_INIT(bprm_set_creds, mp4_bprm_set_creds),  
  
    LSM_HOOK_INIT(cred_alloc_blank, mp4_cred_alloc_blank),  
    LSM_HOOK_INIT(cred_free, mp4_cred_free),  
    LSM_HOOK_INIT(cred_prepare, mp4_cred_prepare)  
};
```

Step 2.2: Label Semantics



```
/* mp4 labels along with their semantics */
#define MP4_NO_ACCESS 0      /* may not be accessed by target,
                               * but may by everyone other */

#define MP4_READ_OBJ 1      /* object may be read by anyone */

#define MP4_READ_WRITE 2    /* object may read/written/appended by the target,
                               * but can only be read by others */

#define MP4_WRITE_OBJ 3     /* object may be written/appended by the target,
                               * but not read, and only read by others */

#define MP4_EXEC_OBJ 4      /* object may be read and executed by all */

/* NOTE: FOR DIRECTORIES, ONLY CHECK ACCESS FOR THE TARGET SID, ALL OTHER NON
* TARGET PROCESSES SHOULD DEFAULT TO THE LINUX REGULAR ACCESS CONTROL
*/
#define MP4_READ_DIR 5      /* for directories that can be read/exec/access
                               * by all */

#define MP4_RW_DIR 6        /* for directory that may be modified by the target
                               * program */
```


Step 2.2: Label Map



```
if (strcmp(cred_ctx, "read-only") == 0)
    return MP4_READ_OBJ;
else if (strcmp(cred_ctx, "read-write") == 0)
    return MP4_READ_WRITE;
else if (strcmp(cred_ctx, "exec") == 0)
    return MP4_EXEC_OBJ;
else if (strcmp(cred_ctx, "target") == 0)
    return MP4_TARGET_SID;
else if (strcmp(cred_ctx, "dir") == 0)
    return MP4_READ_DIR;
else if (strcmp(cred_ctx, "dir-write") == 0)
    return MP4_RW_DIR;
else
    return MP4_NO_ACCESS;
```

Step 2.2: Label Mgmt



- We are interested in three operations:
 1. Allocate/free/copy subject security blobs
 2. When a process starts, check the inode of the binary that launches it.
 - If it is labeled with target, mark `task_struct` as target
 - `mp4_bprm_set_creds`
 3. Assign read-write label to inodes created by the target application
 - `mp4_inode_init_security`

Step 2.2: Attributes



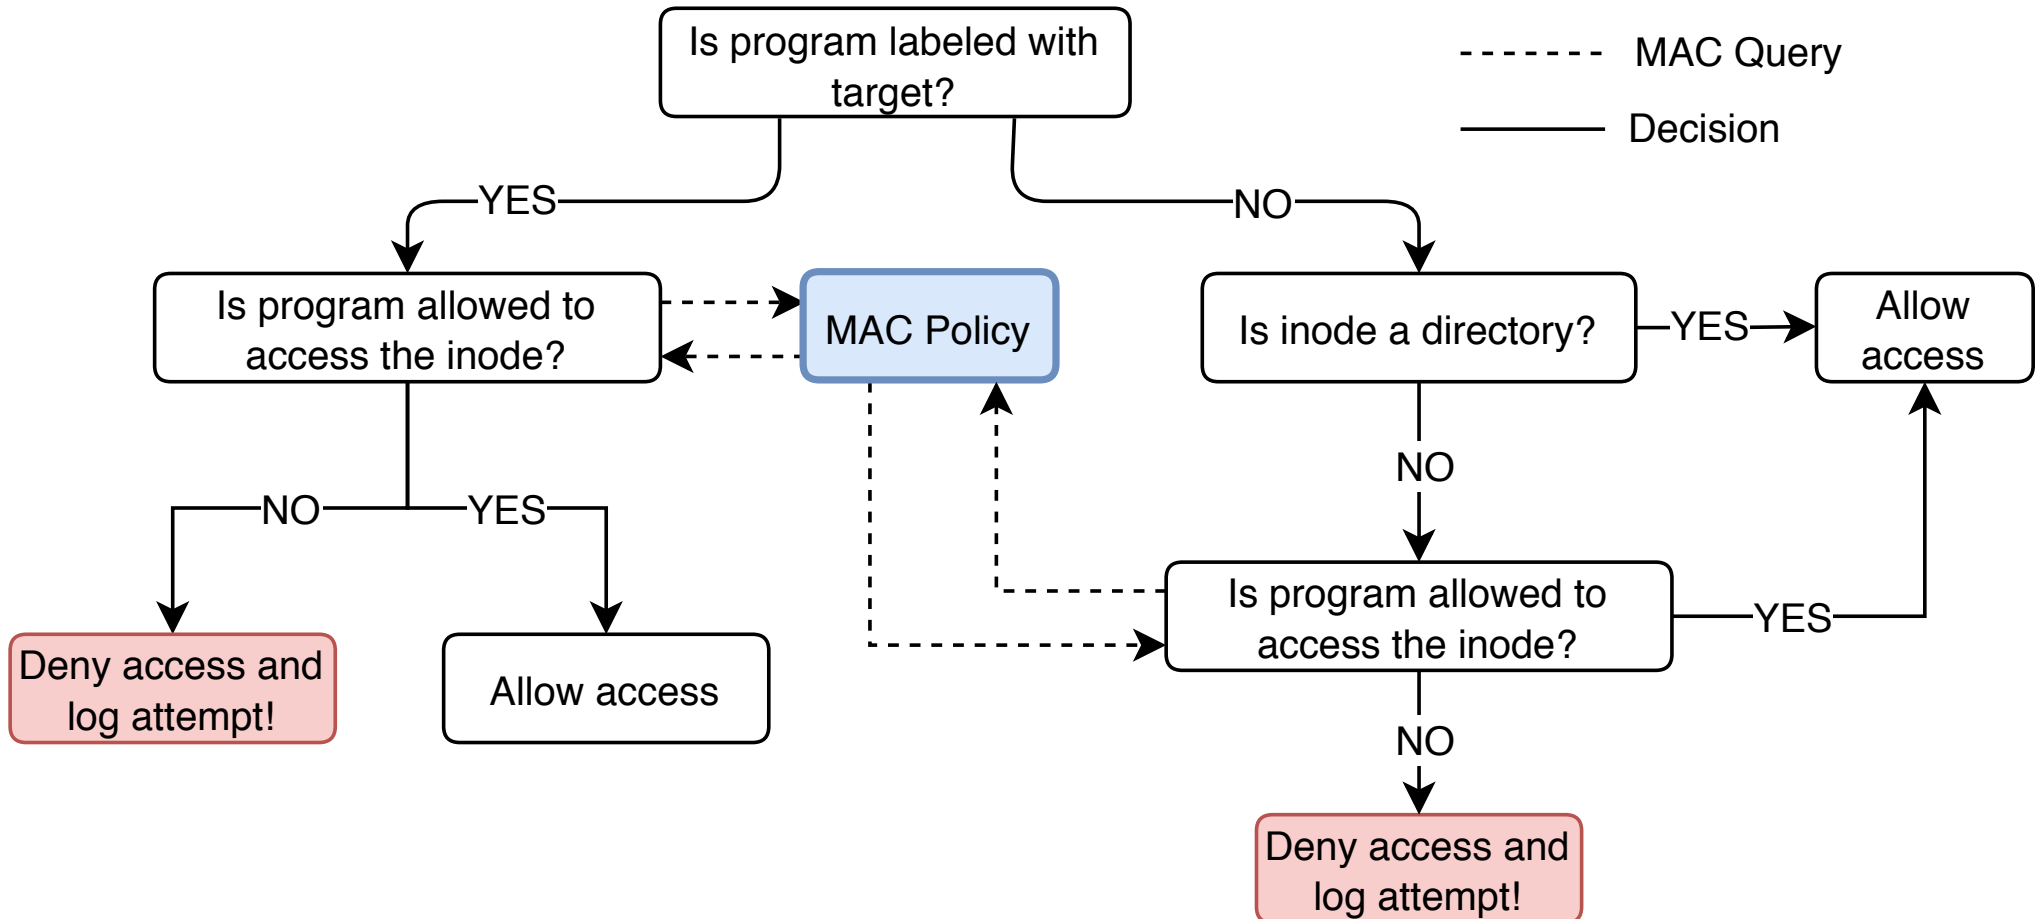
- How do we obtain an inode's extended attributes?
- Few hints:
 - Given an `struct inode *`, we can ask for its `struct dentry *`
 - You can query some kernel functions if there is something you need to know
 - This is important if you don't know how much memory to allocate
 - Watch for the `ERANGE` errno
 - It is **very important** to put back a dentry after you use it
 - `dput(dentry);`

Step 2.3: Implement AC



- Translate label semantics into code
 - `mp4_inode_permission`
- Operation masks are in `linux/fs.h`
- Obtain current task's subject blob using `current_cred()`
- To speed things up during boot, we want to skip certain directories
 - Obtain inode's path (hint: use `dentry`!)
 - Call `mp4_should_skip_path` from `mp4_given.h`

Step 2.3: Implement AC



Step 2.3: Implement AC



- You **MUST** log attempts that are denied access
- To minimize the chances of bricking your machine:
 - Always take a snapshot that takes you back to stable state
 - Implement AC logic, but always return access granted and print appropriate messages
 - Check you messages, if all is according to plan, update your code to return appropriate values
 - Test your return codes

Step 3: Testing



- Test your security module on simple functions
 - vim, cat, etc.
 - avoid operation critical programs (ls, cd, bash, etc.)
 - **Note**, to grant read access `/home/netid/file.txt`,
 - **must** have access to all three of `/home`, `/home/netid/`, and `/home/netid/file.txt`
- Always restore you system state to a place where all labels are removed before you reboot

Step 3: Testing



- Suggested method of testing:
 - Create two scripts: `mp4_test.perm` and `mp4_test.perm.unload`
 - source first script to load, source the other to unload

- In `mp4_test.perm`:

```
setfattr -n security.mp4 -v target /usr/bin/cat
...
setfattr -n security.mp4 -v read-only /home/netid/file.txt
```

- In `mp4_test.perm.unload`, undo everything before reboot:

```
setfattr -x security.mp4 /usr/bin/cat
...
setfattr -x security.mp4 /home/netid/file.txt
```


Final Step: Obtain Policy



- Goal is to obtain least privilege policy for the program `/usr/bin/passwd`
- **DO NOT TRY TO CHANGE THE PASSWORD FOR YOUR NETID ACCOUNT**
- Create dummy user account and change its password
- Use **strace** to obtain the set of files and access requests that `passwd` uses
- Generate `passwd.perm` and `passwd.perm.unload` based on the outcome
- Test your module's enforcement of the policy!