



Main Memory

Yajin Zhou (<http://yajin.org>)

Zhejiang University



Review

- Deadlock problem: Four conditions for deadlock
- System model
- Resource allocation graph
 - No cycle -> no deadlock, cycle -> possible deadlock
- Handling deadlocks
 - deadlock prevention
 - Break four conditions, ordering of resources
 - deadlock avoidance
 - Extra information, safe state
 - Single instance: resource allocation graph -> claim edge
 - Multiple instances: banker's algorithm
- deadlock detection
 - Wait-for graph, similar with banker's algorithm



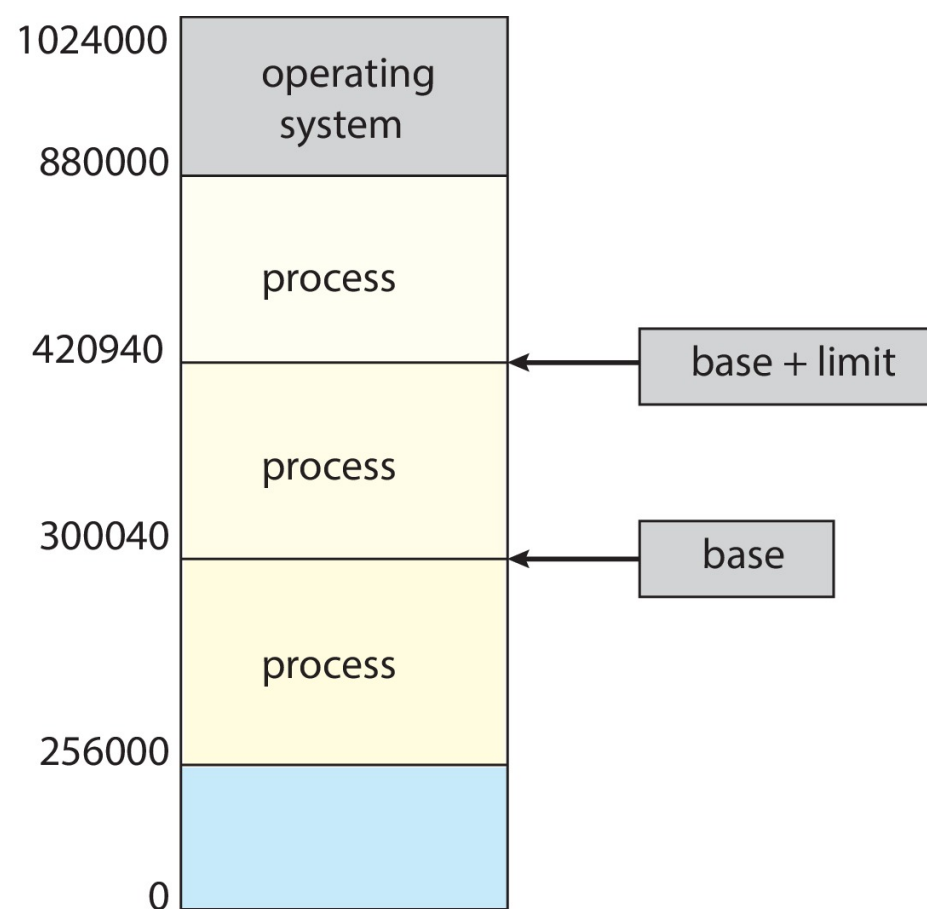
Background

- Program must be brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are only storage CPU can access directly
- Memory unit only sees a stream of:
 - addresses + read requests, or
 - address + data and write requests
- Register access is done in one CPU clock (or less)
- Main memory can take many cycles, causing a **stall**
- **Cache** sits between main memory and CPU registers
- **Protection** of memory required to ensure correct operation
 - hardware vs software



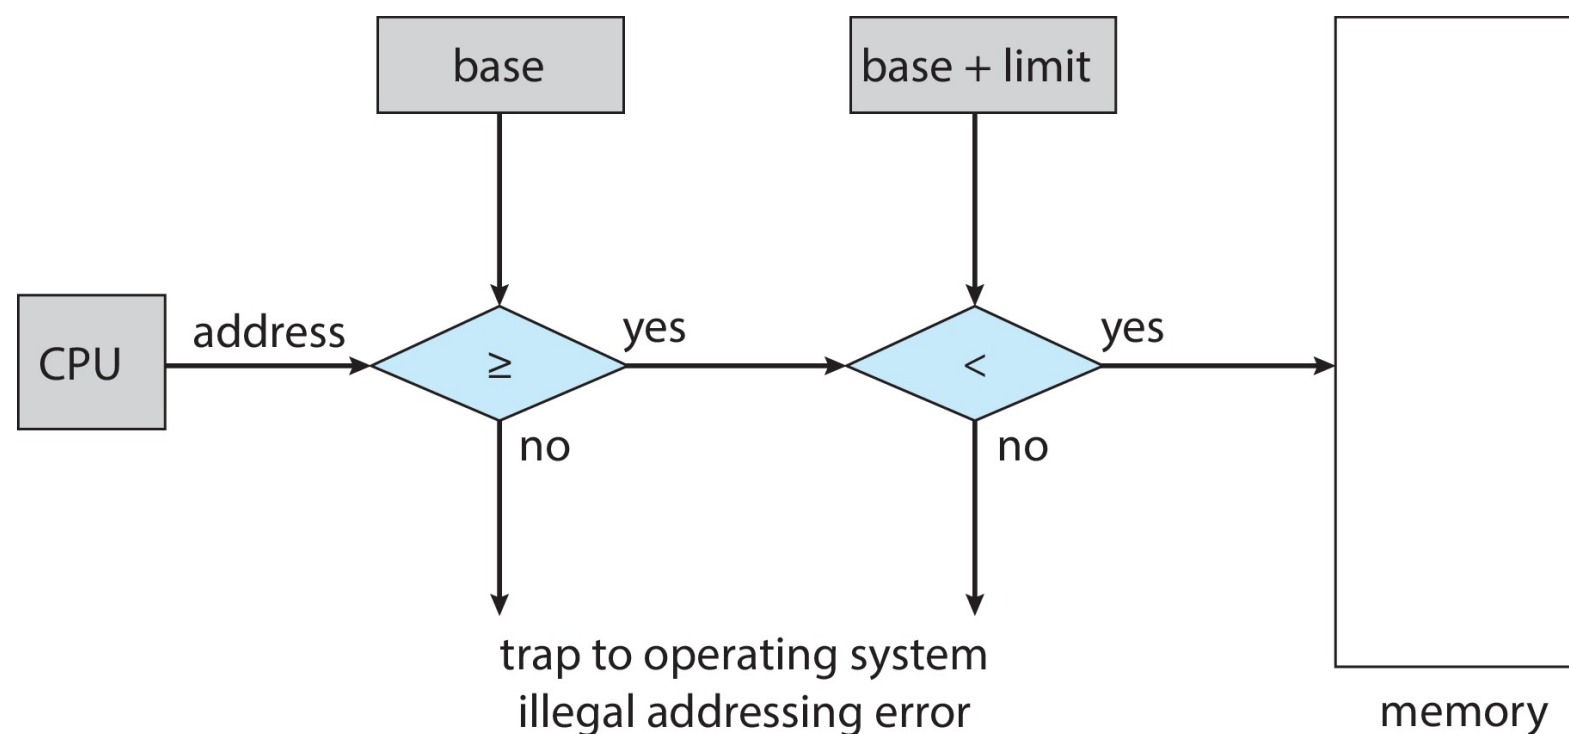
Protection

- Need to ensure that a process can access only those addresses in its address space.
- We can provide this protection by using a pair of **base** and **limit** registers define the logical address space of a process



Hardware Address Protection

- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user
- the instructions to loading the base and limit registers are privileged





Address Binding

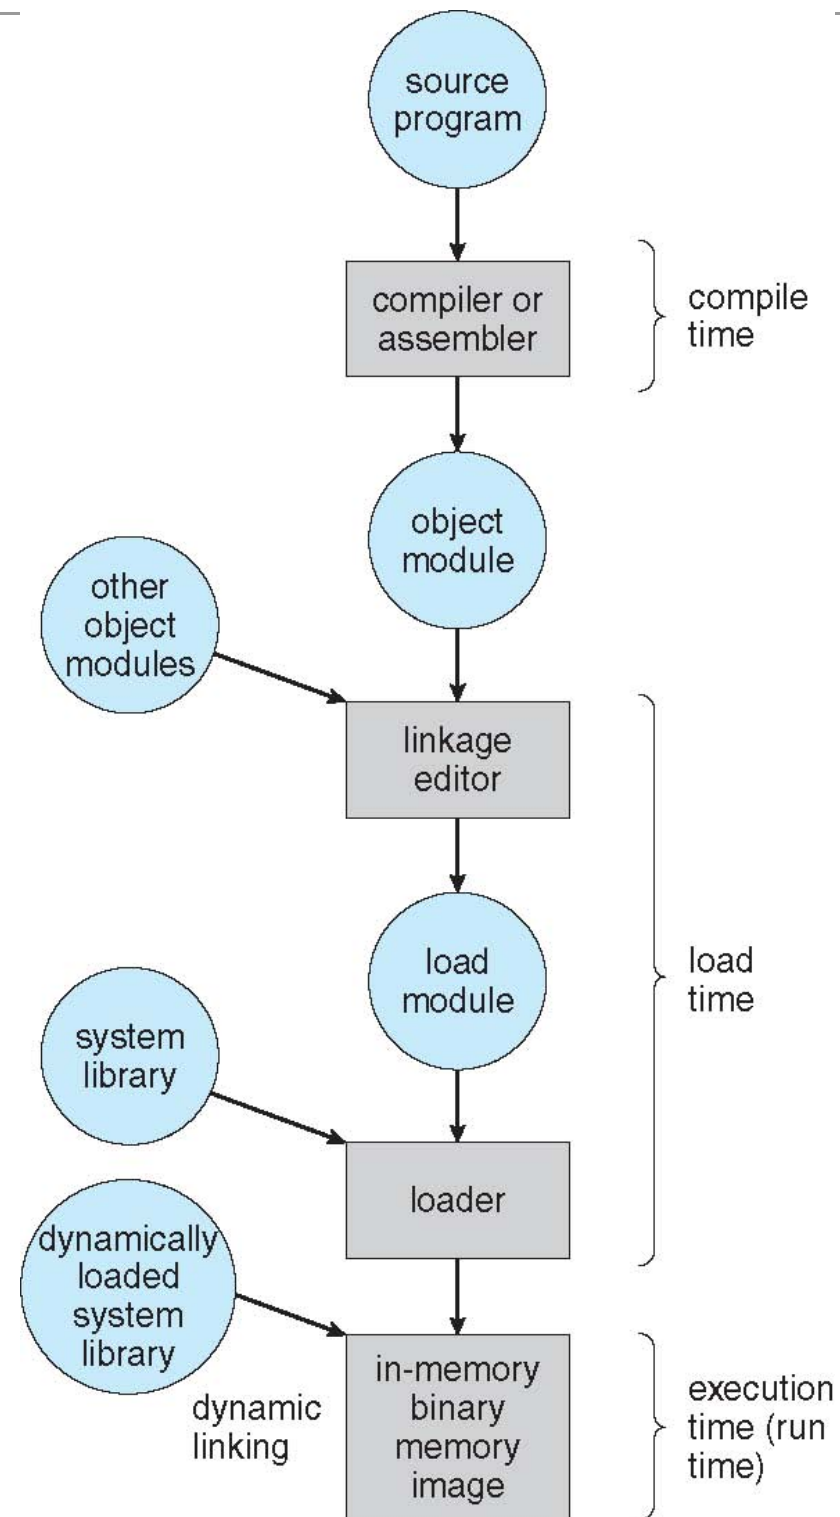
- Inconvenient to have first user process address always at 0 (**why?**)
 - (shared) libraries, *nil* pointer dereference detection...
- Addresses are represented in different ways at different stages of a program's life
 - **source code** addresses are usually **symbolic** (e.g., temp)
 - **compiler** binds symbols to **relocatable addresses**
 - e.g., “14 bytes from beginning of this module”
 - **linker** (or loader) binds relocatable addresses to **absolute addresses**
 - e.g., 0x0e74014
 - Each binding maps one address space to another



Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages
 - **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
 - **Load time:** Must generate relocatable code if memory location is not known at compile time
 - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - Need **hardware support** for address maps (e.g., base and limit registers)

Multi-step Processing of a User Program






Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a **separate physical address** space is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as virtual address
 - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program

Memory-Management Unit (MMU)

- Hardware device that at run time maps virtual to physical address



- Many methods possible, covered in the rest of this chapter




Memory-Management Unit

- Consider simple scheme. which is a generalization of the base-register scheme.
- The base register now called **relocation register**
- The value in the relocation register is added to every address generated by a user process at the time it is sent to memory
- The user program deals with logical addresses; it never sees the real physical addresses
 - Execution-time binding occurs when reference is made to location in memory
 - Logical address bound to physical addresses



Memory-Management Unit

- Consider simple scheme. which is a generalization of the base-register scheme.
- The base register now called **relocation register**
- The value in the relocation register is added to every address generated by a user process at the time it is sent to memory
- Logical address 0 to max, physical address: R, R+max. What if these two sizes differ?





Dynamic Loading

- The entire program does need to be in memory to execute
- Routine is not loaded until it is called
- Better memory-space utilization; unused routine is never loaded
- All routines kept on disk in relocatable load format
- Advantage
 - Routine is loaded only when it is needed -> Useful when large amounts of code are needed to handle infrequently occurring cases
- No special support from the operating system is required
 - Implemented through program design
 - OS can help by providing libraries to implement dynamic loading



Dynamic Linking

- **Static linking** – system libraries and program code combined by the loader into the binary program image
- Dynamic linking –linking postponed until **execution time**
 - Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
 - Stub replaces itself with the address of the routine, and executes the routine
- Dynamic linking is particularly useful for libraries
 - System also known as **shared libraries**
 - Consider applicability to patching system libraries
 - Versioning may be needed
 - What will happen without dynamic linking?
- Help from OS: share libraries between processes



Contiguous Allocation

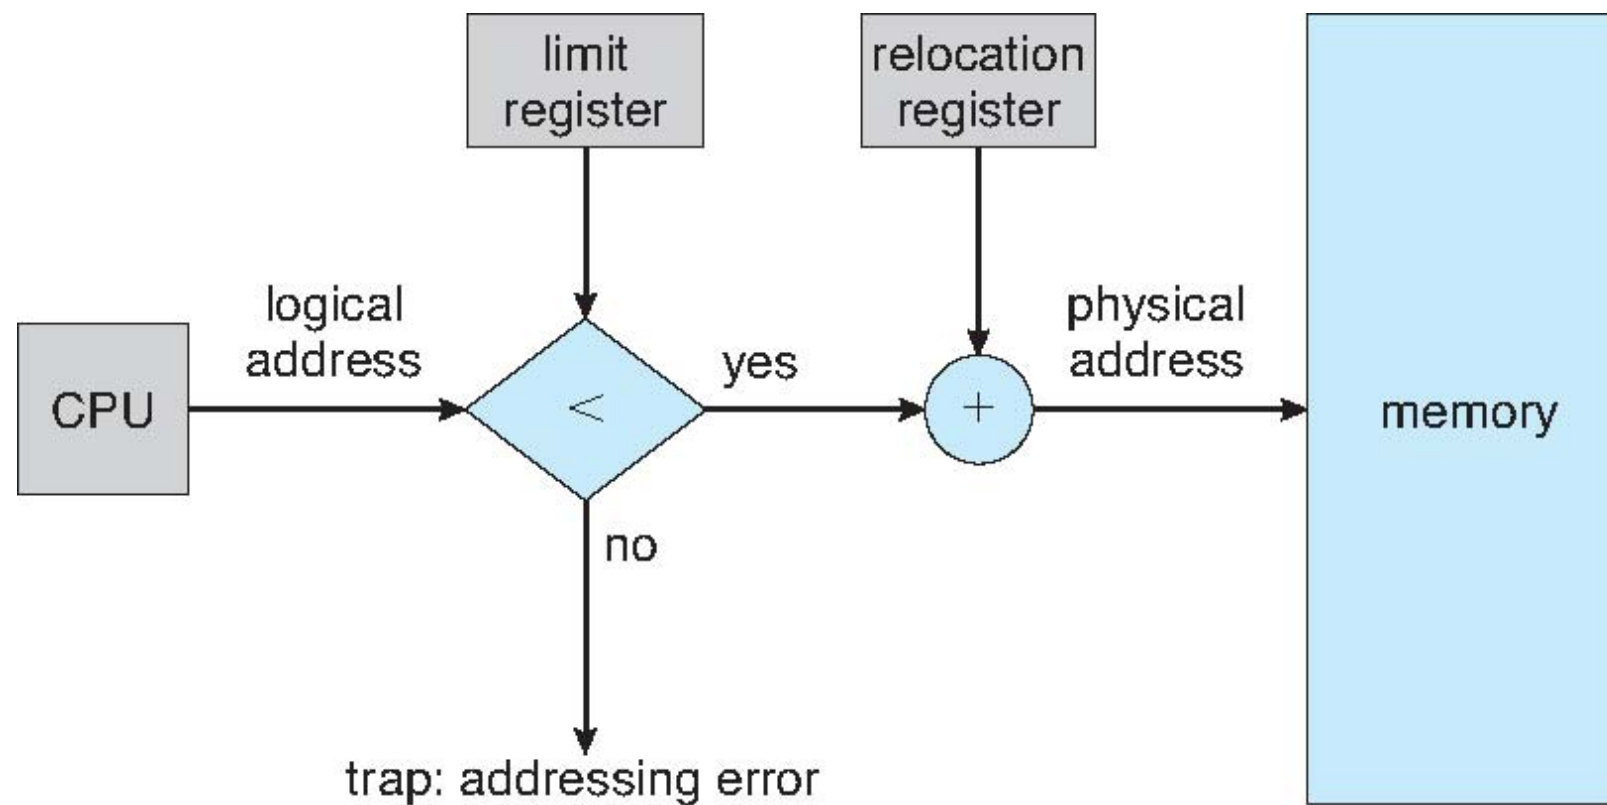
- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually into two **partitions**:
 - Resident operating system, usually held in low memory with interrupt vector
 - User processes then held in high memory
 - Each process contained in single contiguous section of memory



Contiguous Allocation: protection

- Relocation registers used to protect user processes from each other, and from changing operating-system code and data
 - Base register contains value of smallest physical address
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address dynamically
 - Can then allow actions such as kernel code being **transient** and kernel changing size

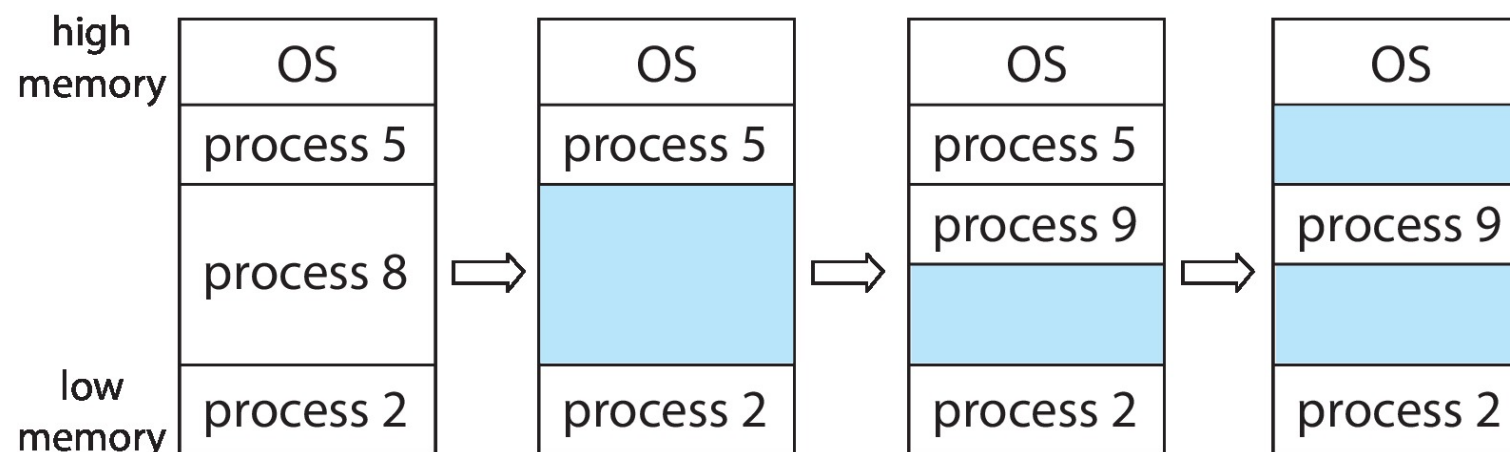
Hardware Support for Relocation and Limit Registers





Memory Allocation: Variable Partition

- Multiple-partition allocation
- Degree of multiprogramming limited by number of partitions
- **Variable-partition** sizes for efficiency (sized to a given process' needs)
- **Hole** – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Process exiting frees its partition, adjacent free partitions combined
- Operating system maintains information about:
 - a) allocated partitions b) free partitions (hole)





Memory Allocation

- How to satisfy a request of size n from a list of free memory blocks?
 - **first-fit:** allocate from the first block that is big enough
 - **best-fit:** allocate from the smallest block that is big enough
 - must search entire list, unless ordered by size
 - produces the smallest leftover hole
 - **worst-fit:** allocate from the largest hole
 - must also search entire list
 - produces the largest leftover hole
- **Fragmentation** is big problem for all three methods
 - first-fit and best-fit usually perform better than worst-fit



Fragmentation

- **External fragmentation**
 - unusable memory between allocated memory blocks
 - **total amount** of free memory space **is larger than a request**
 - the request cannot be fulfilled because the free memory is not **contiguous**
 - external fragmentation can be reduced by **compaction**
 - shuffle memory contents to place all free memory in one large block
 - program needs to be **relocatable** at runtime
 - Performance overhead, timing to do this operation
 - Another solution: **paging**
 - 50-percent rule: N allocated blocks, $0.5N$ will be lost due to fragmentation. $1/3$ is unusable!



Fragmentation

- **Internal fragmentation**
 - memory allocated may be larger than the requested size
 - this size difference is memory *internal to a partition*, but not being used
 - Example: free space 18464 bytes, request 18462 bytes
- Sophisticated algorithms are designed to avoid fragmentation
 - none of the first-/best-/worst-fit can be considered sophisticated



Review

- Address binding
- Logical address vs physical address
- MMU: address translation + protection
- Memory allocation
 - Contiguous allocation: first, best, worst fit
- Fragmentation: external/internal



Paging

- Physical address space of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available
 - Avoids **external fragmentation** -> **avoid for compacting**
 - Avoids problem of varying sized memory chunks
- Basic methods
 - Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes
 - Divide logical memory into blocks of same size called **pages**
 - Keep track of all free frames
 - To run a program of size **N** pages, need to find **N** free frames and load program
 - Set up a **page table** to translate logical to physical addresses
 - Backing store likewise split into pages
 - Still have Internal fragmentation



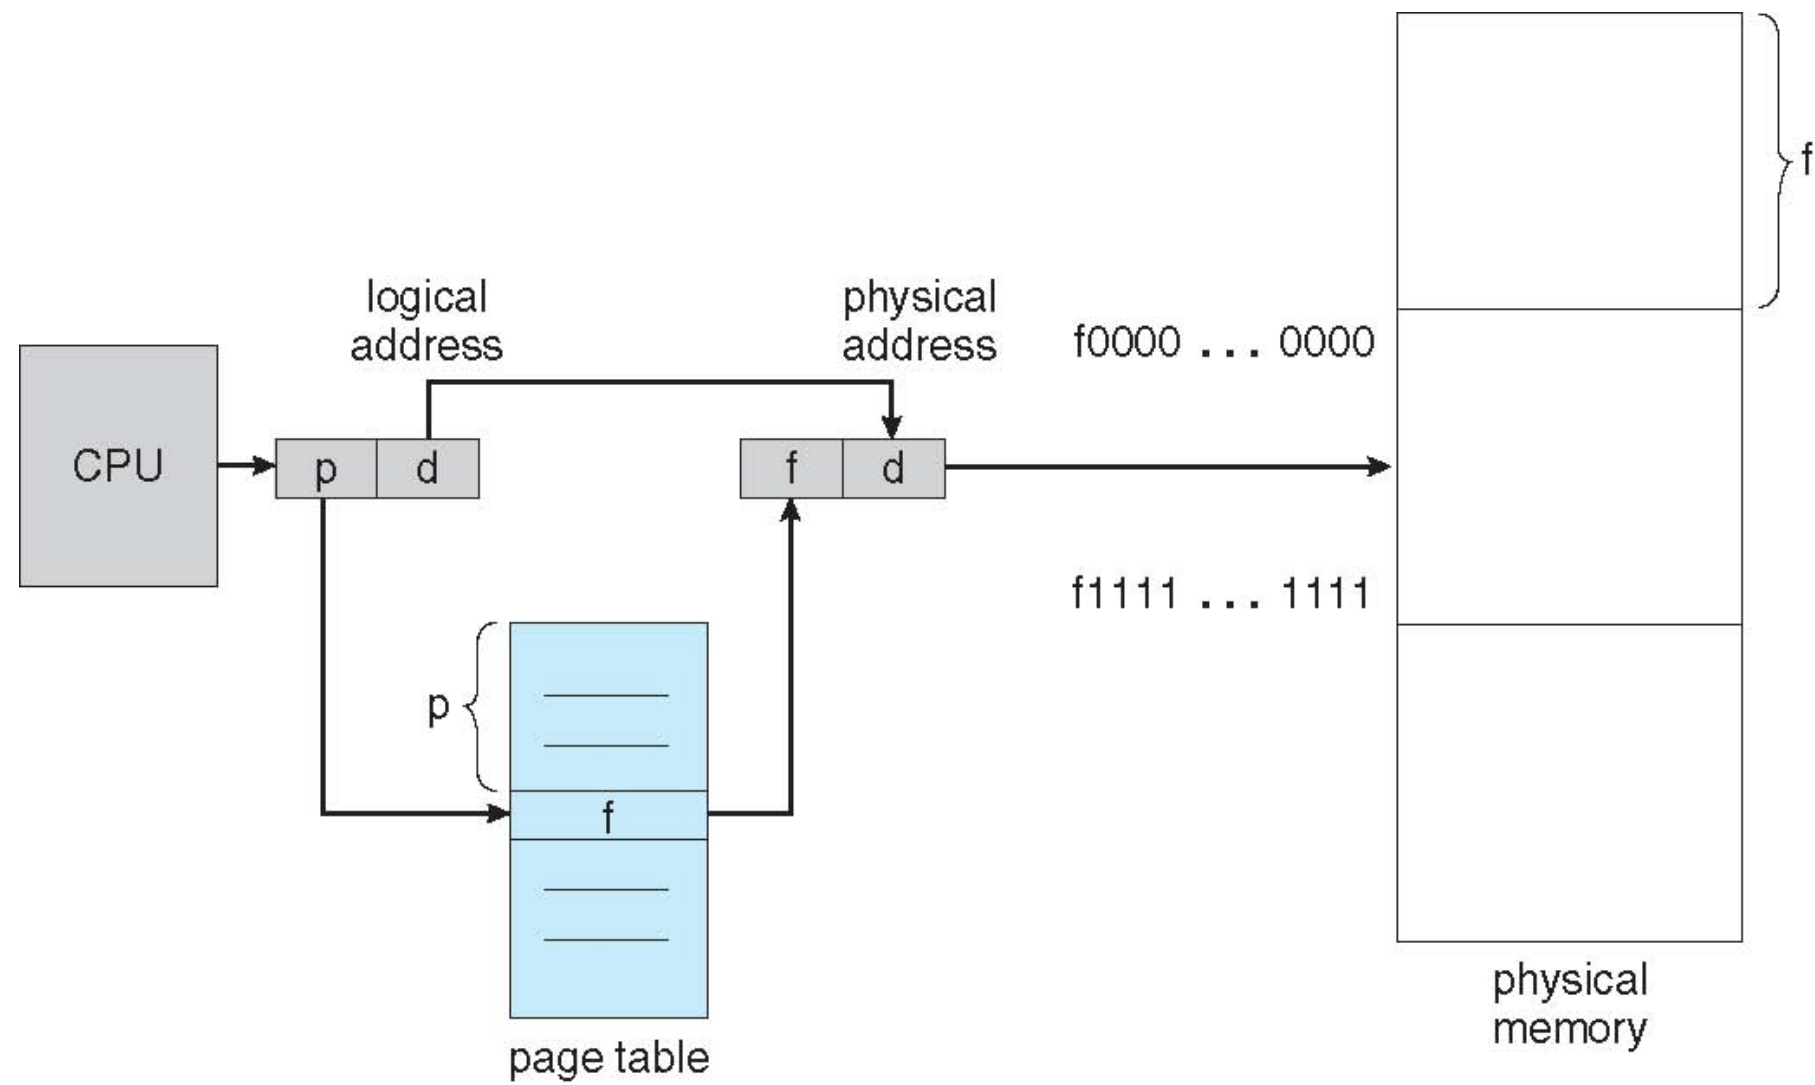
Paging: Address Translation

- A logical address is divided into:
 - **page number** (p)
 - used as an index into a page table
 - page table entry contains the corresponding **physical** frame number
 - **page offset** (d)
 - offset within the page/frame
 - combined with frame number to get the physical address


page number	page offset
p	d
$m - n \text{ bits}$	$n \text{ bits}$

m bit logical address space, n bit page size

Paging Hardware



Paging Example




logical
memory

0	1
1	4
2	3
3	7


page table

frame
number



physical
memory

Paging Example II



$m = 4$ and $n = 2$ 32-byte memory and 4-byte pages



Paging: Internal Fragmentation


- Paging has **no external fragmentation**, but **internal fragmentation**
 - e.g., page size: 2,048, program size: 72,766 (35 pages + 1,086 bytes)
 - internal fragmentation: $2,048 - 1,086 = 962$
 - **worst** case internal fragmentation: **1 frame – 1 byte**
 - **average** internal fragmentation: $1 / 2$ frame size
- Small frame sizes more desirable than large frame size?
 - memory becomes larger, and page table takes memory
 - page sizes actually grow over time
 - 4KB → 2MB → 4MB → 1GB → 2GB



Frame Table


- OS is managing physical memory, it should be aware of the allocation details of physicals
 - Which frame is free, and how many frames have been allocated
 - ...
- One entry for each physical frame
- the allocated frame belongs to which process

Free Frames



(a)

Before allocation



(b)

new-process page table

0	14
1	13
2	18
3	20

After allocation

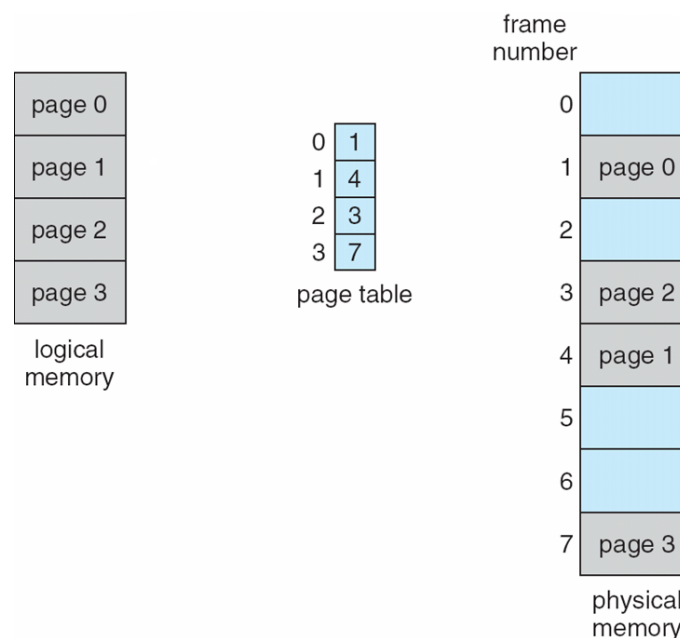


Hardware Support: Simplest Case

- Page table is in a set of dedicated registers
 - Advantages: very efficient - access to register is fast
 - Disadvantages: the table size is very small, and the context switch need to save and restore these registers

Hardware Support: Alternative Way

- One big page table maps logical address to physical address
 - the page table should be **kept in main memory**
 - **page-table base register (PTBR)** points to the page table
 - **does PTBR contain physical or logical address?**
 - **page-table length register (PTLR)** indicates the size of the page table
- Every data/instruction access requires **two memory accesses**
 - one for the page table and one for the data / instruction
 - CPU can cache the translation to avoid one memory access (**TLB**)





TLB

- TLB (**translation look-aside buffer**) caches the address translation
 - if page number is in the TLB, no need to access the page table
 - if page number is not in the TLB, need to replace one TLB entry
 - TLB usually use a fast-lookup hardware cache called **associative memory**
 - TLB is usually small, 64 to 1024 entries
- Use with page table
 - TLB contains a few page table entries
 - Check whether page number is in TLB
 - If -> frame number is available and used
 - If not -> **TLB miss**. access page table and then fetch into TLB
 - TLB flush: TLB entries are full
 - TLB wire down: TLB entries should not be flushed



TLB

- TLB and context switch
 - Each process has its own page table
 - switching process needs to switch page table
 - **TLB must be consistent with page table**
 - Option I: Flush TLB at every context switch, **or**,
 - Option II: Tag TLB entries with **address-space identifier** (ASID) that uniquely identifies a process
 - some TLB entries can be **shared** by processes, and fixed in the TLB
 - e.g., TLB entries for the kernel
- TLB and operating system
 - MIPS: OS should deal with TLB miss exception
 - X86: TLB miss is handled by hardware

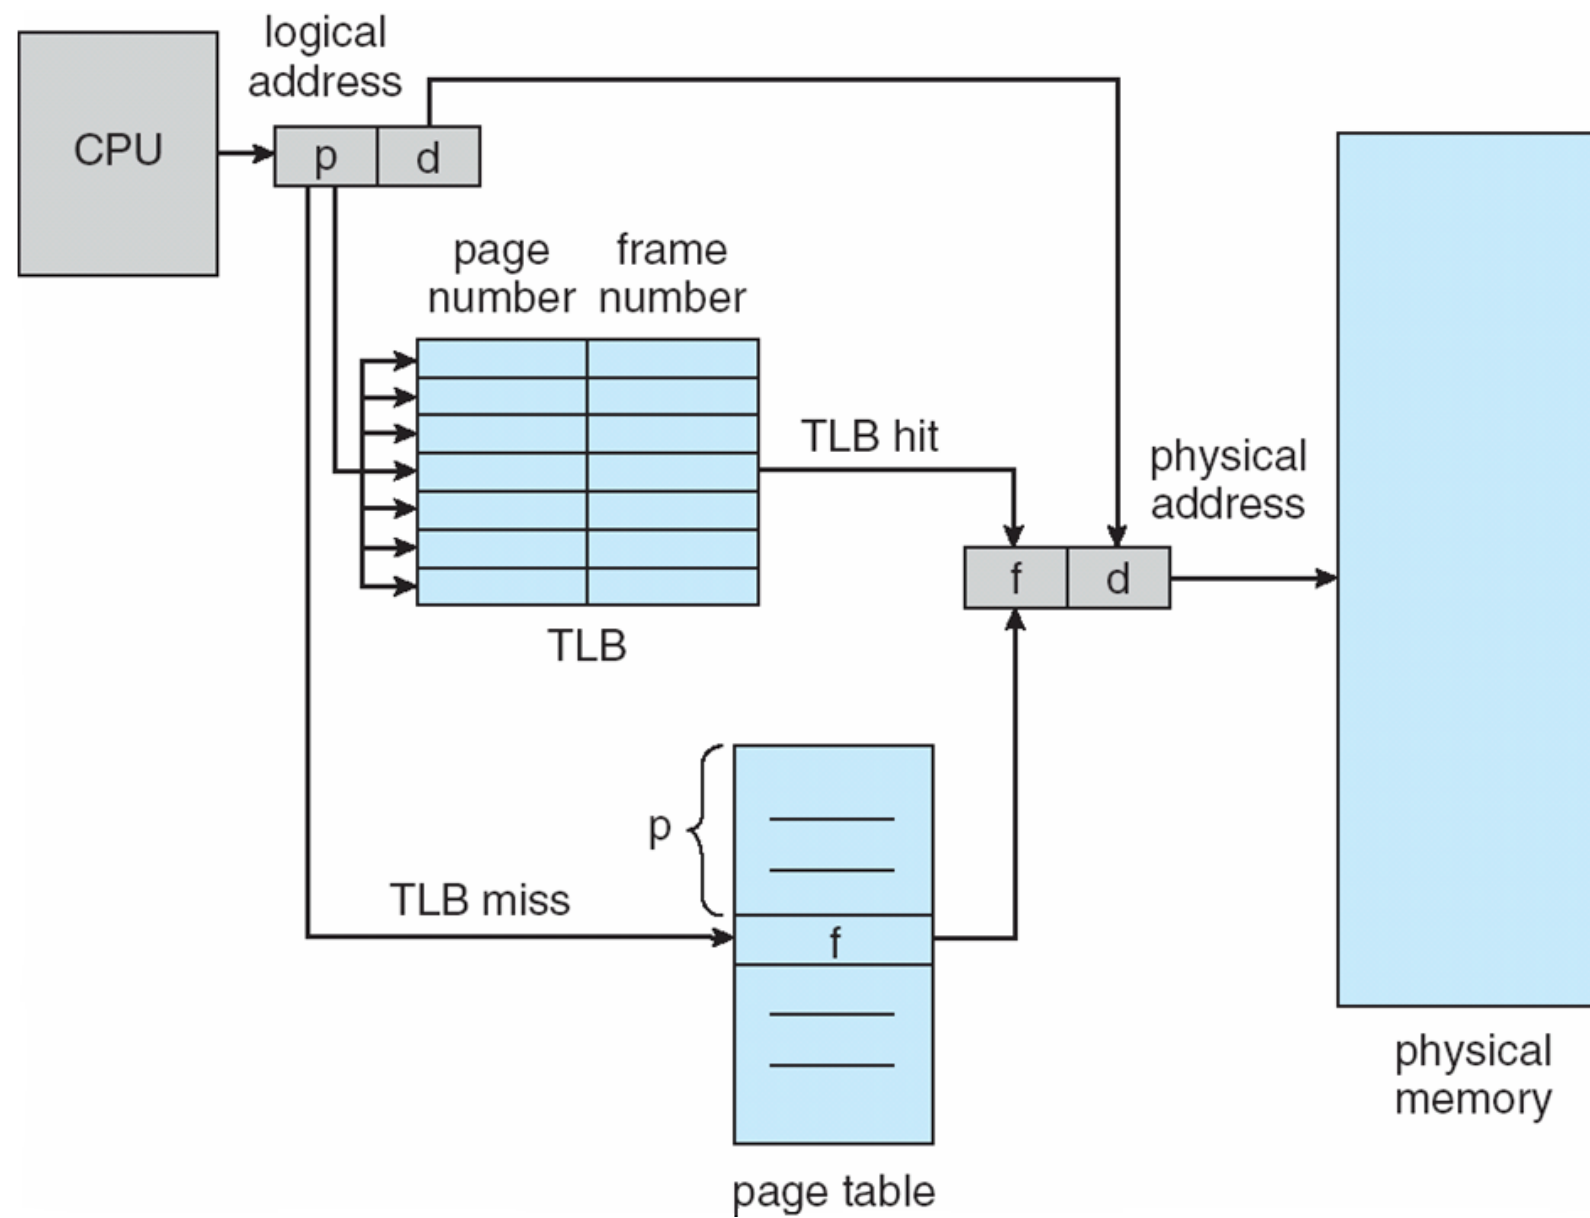


Associative Memory

- Associative memory: memory that supports **parallel search**
- Associative memory is not addressed by “addresses”, but **contents**
 - if p is in associative memory’s key, return frame# (value) directly
 - think of hash tables

Page #	Frame #
1	7
2	12
3	15
4	31

Paging Hardware With TLB





Effective Access Time

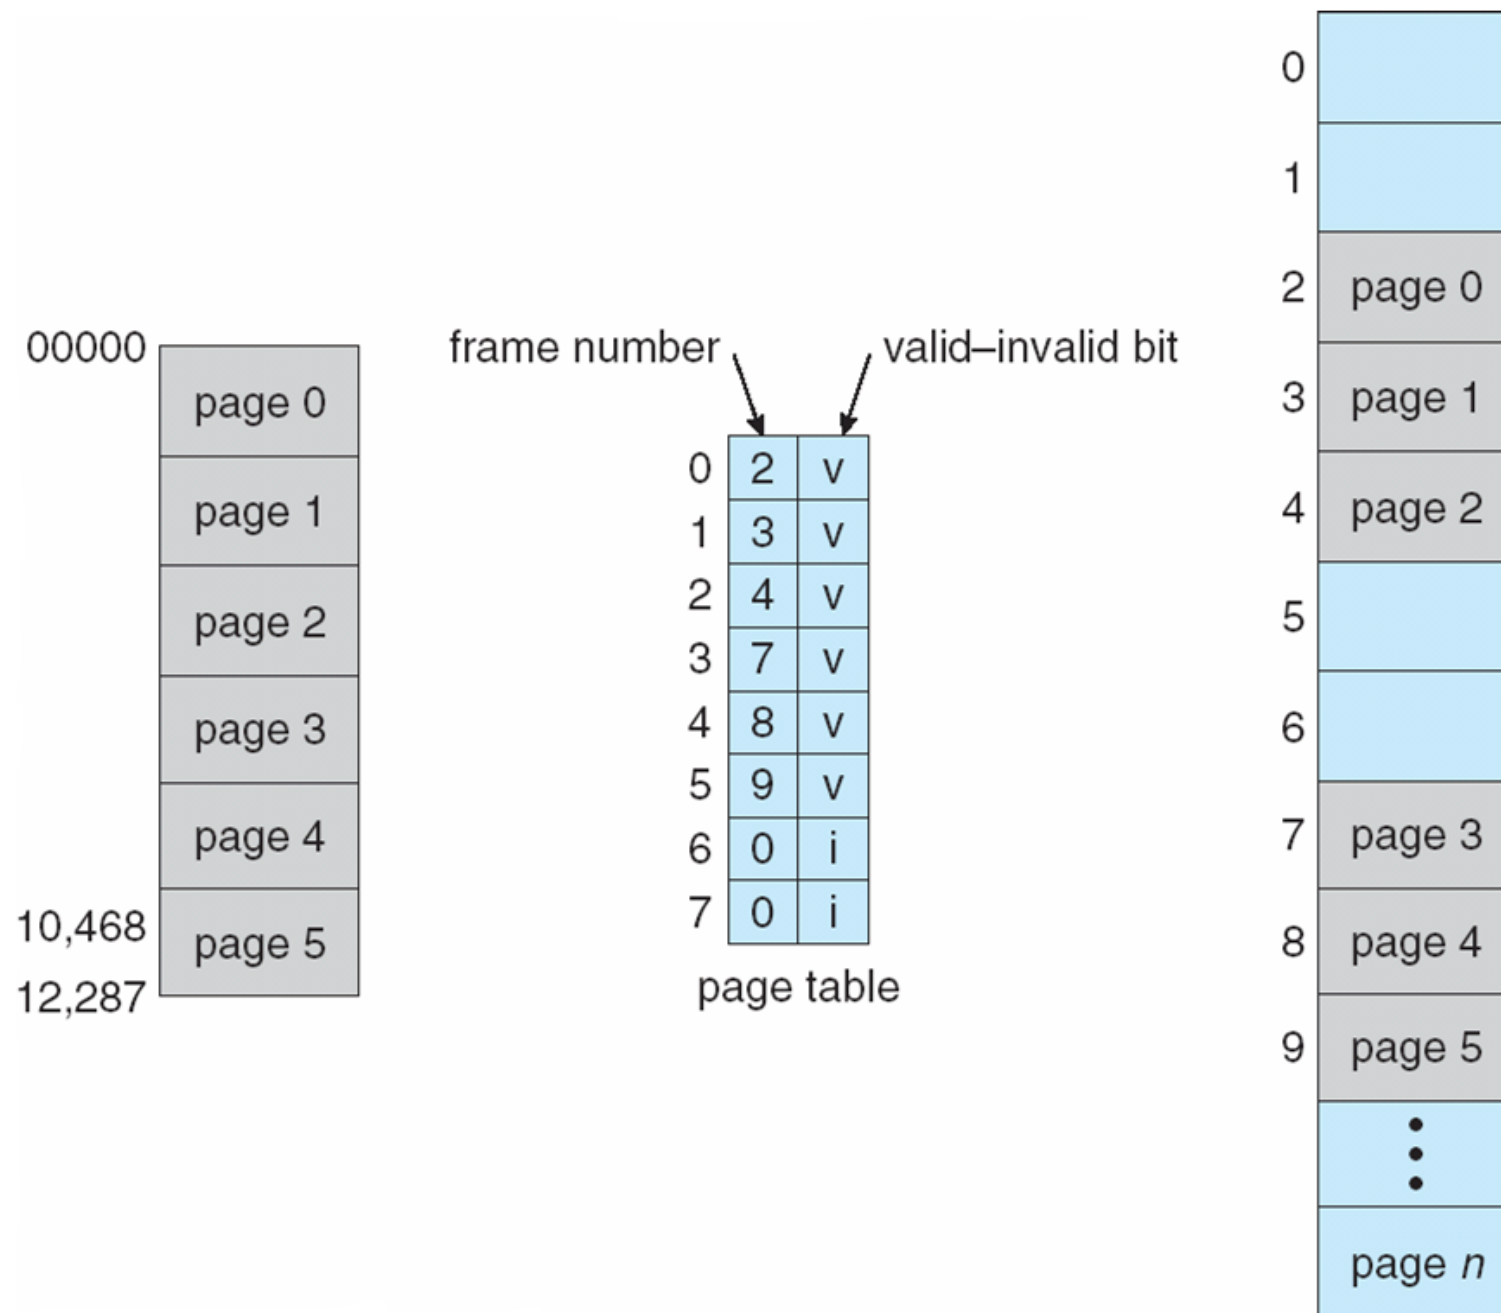
- **Hit ratio** – percentage of times that a page number is found in the TLB
- An 80% hit ratio means that we find the desired page number in the TLB 80% of the time.
- Suppose that 10 nanoseconds to access memory.
 - If we find the desired page in TLB then a mapped-memory access take 10 ns
 - Otherwise we need **two memory access** so it is 20 ns: page table + memory access
- Effective Access Time (EAT)
- $$\text{EAT} = 0.80 \times 10 + 0.20 \times 20 = 12 \text{ nanoseconds}$$
- implying 20% slowdown in access time
- Consider a more realistic hit ratio of 99%,
- $$\text{EAT} = 0.99 \times 10 + 0.01 \times 20 = 10.1\text{ns}$$
- implying only 1% slowdown in access time.



Memory Protection

- Accomplished by protection bits with each frame
- Each page table entry has a **present** (aka. valid) bit
 - present: the page has a valid physical frame, thus can be accessed
- Each page table entry contains some protection bits
 - **kernel/user, read/write, execution?, kernel-execution?**
 - why do we need them?
- Any violations of memory protection result in a trap to the kernel

Memory Protection





Memory Protection (more)

- **NX bit**
 - segregate areas of memory for use by either storage of processor instructions (code) or for storage of data
 - Intel: XD(execute disable), AMD: EVP (enhanced virus protection), ARM: XN (execute never)
- PXN: Privileged Execute Never (intel: SMEP)
 - A Permission fault is generated if the processor is executing at **PL1(kernel)** and attempts to execute an instruction fetched from the corresponding memory region when this PXN bit is 1 (usually user space memory)

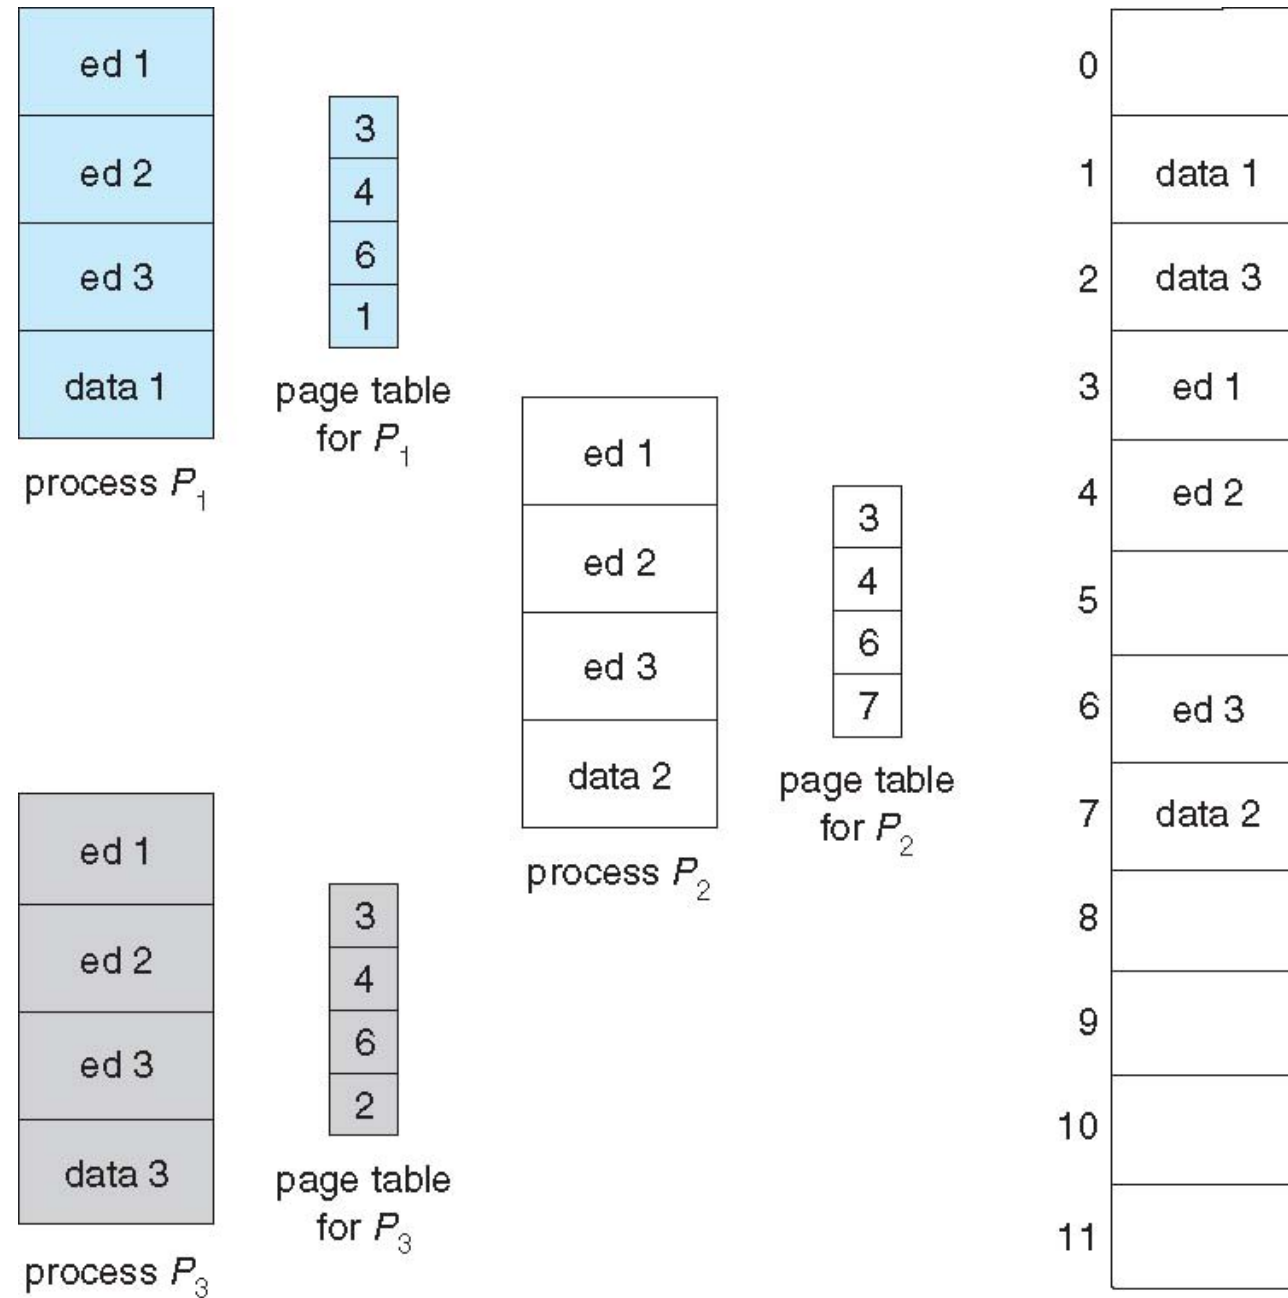


Page Sharing

- Paging allows to share memory between processes
 - e.g., one copy of **code** shared by **all processes of the same program**
 - text editors, compilers, browser..
 - shared memory can be used for **inter-process communication**
 - shared libraries
- Reentrant code: non-self-modifying code: never changes between execution
- Each process can, of course, have its private code and data



Page Sharing





Structure of Page Table


- One-level page table can consume lots of memory for page table
 - e.g., 32-bit logical address space and 4KB page size
 - page table would have 1 million entries ($2^{32} / 2^{12}$)
 - if each entry is 4 bytes → 4 MB of memory for page table alone
 - each process requires its own page table
 - page table must be **physically contiguous**
- To reduce memory consumption of page tables:
 - **hierarchical page table**
 - **hashed page table**
 - **inverted page table**



Hierarchical Page Tables

- Break up the logical address space into **multiple-level** of page tables
 - e.g., two-level page table
 - first-level page table contains the frame# for second-level page tables
 - “page” the page table
- Why hierarchical page table can save memory for page table?


Two-Level Page Table






Two-Level Paging

- A logical address is divided into:
 - a **page directory number** (first level page table)
 - a **page table number** (2nd level page table)
 - a **page offset**
- Example: 2-level paging in 32-bit Intel CPUs
 - 32-bit address space, 4KB page size
 - 10-bit page directory number, 10-bit page table number
 - each page table entry is 4 bytes, one frame contains 1024 entries (2^{10})



Address-Translation Scheme







64-bit Logical Address Space

- 64-bit logical address space requires more levels of paging
 - two-level paging is not sufficient for 64-bit logical address space
 - if page size is 4 KB (2^{12}), outer page table has 2^{42} entries, inner page tables have 2^{10} 4-byte entries
 - one solution is to add more levels of page tables
 - e.g., three levels of paging: 1st level page table is 2^{34} bytes in size
 - and possibly 4 memory accesses to get to one physical memory location
- usually **not support full 64-bit virtual address space**
 - AMD-64 supports 48-bit
 - canonical form: 48 through 63 of valid virtual address must be copies of bit 47



64-bit Logical Address Space

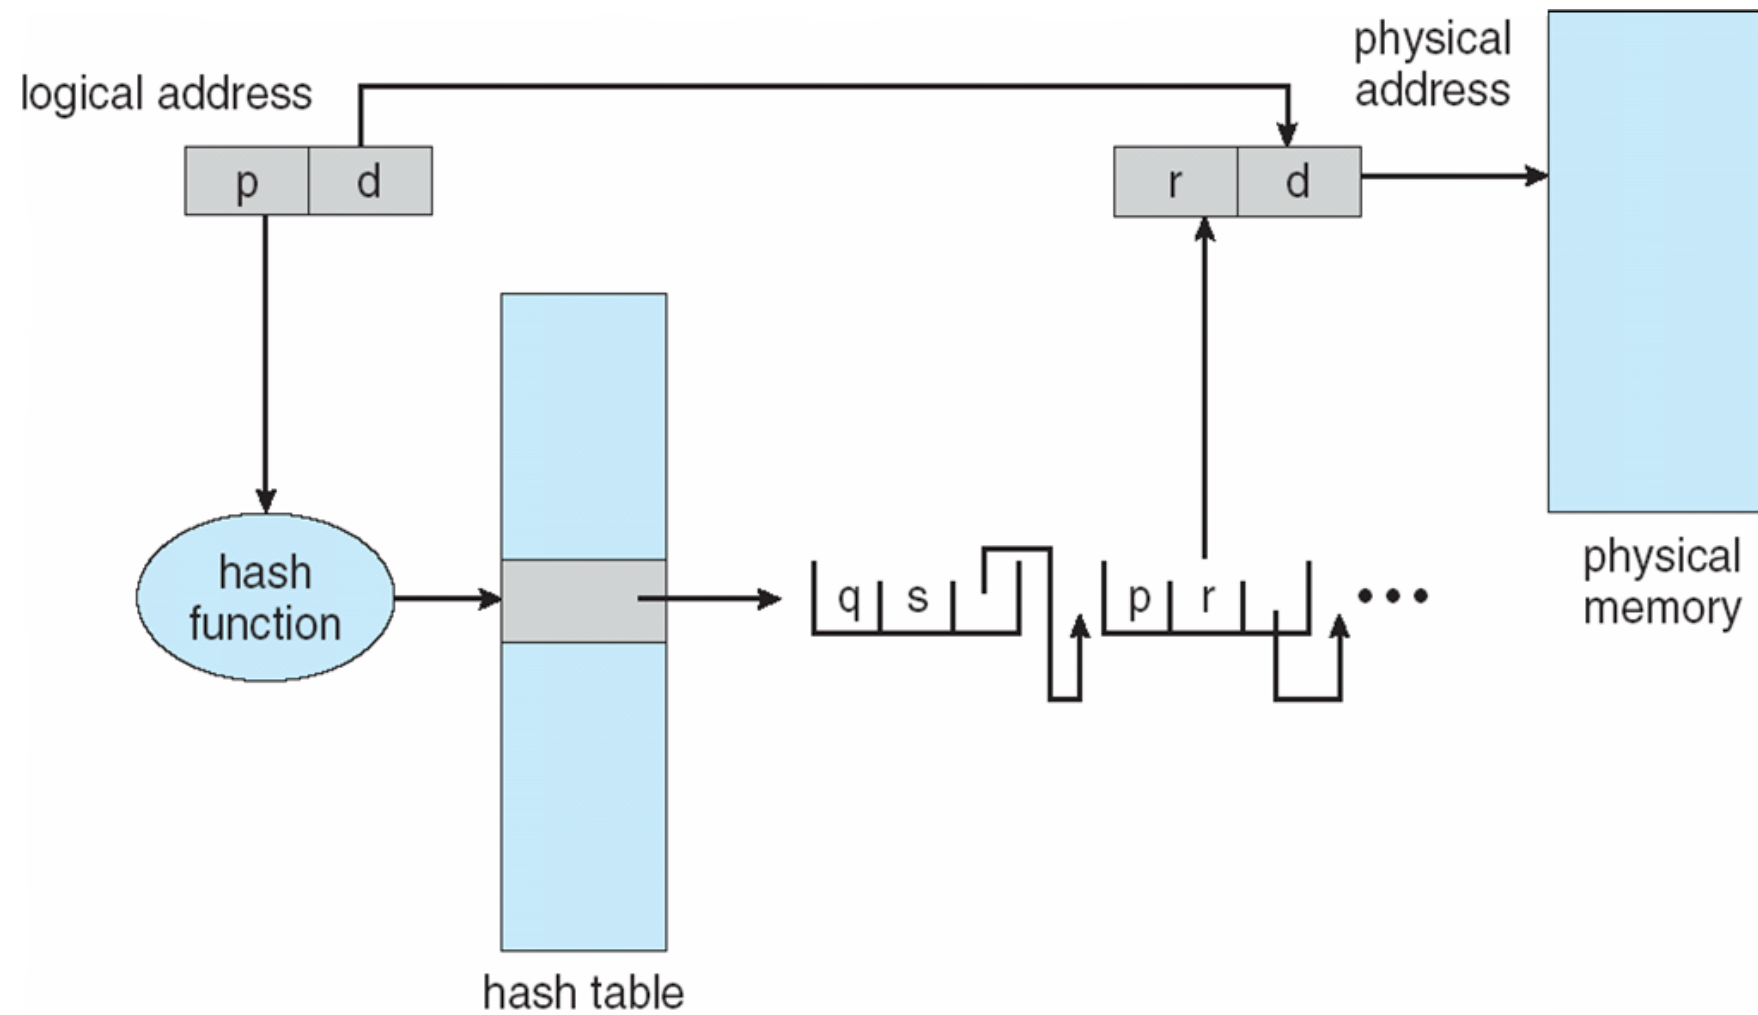




Hashed Page Tables

- In hashed page table, virtual page# is hashed into a frame#
 - the page table contains a chain of elements hashing to the same location
 - each element contains: **page#**, **frame#**, and a **pointer to the next element**
 - virtual page numbers are compared in this chain searching for a match
 - if a match is found, the corresponding frame# is returned
- Hashed page table is common in address spaces > 32 bits
- **Clustered page tables**
 - Each entry refers to several pages

Hashed Page Table

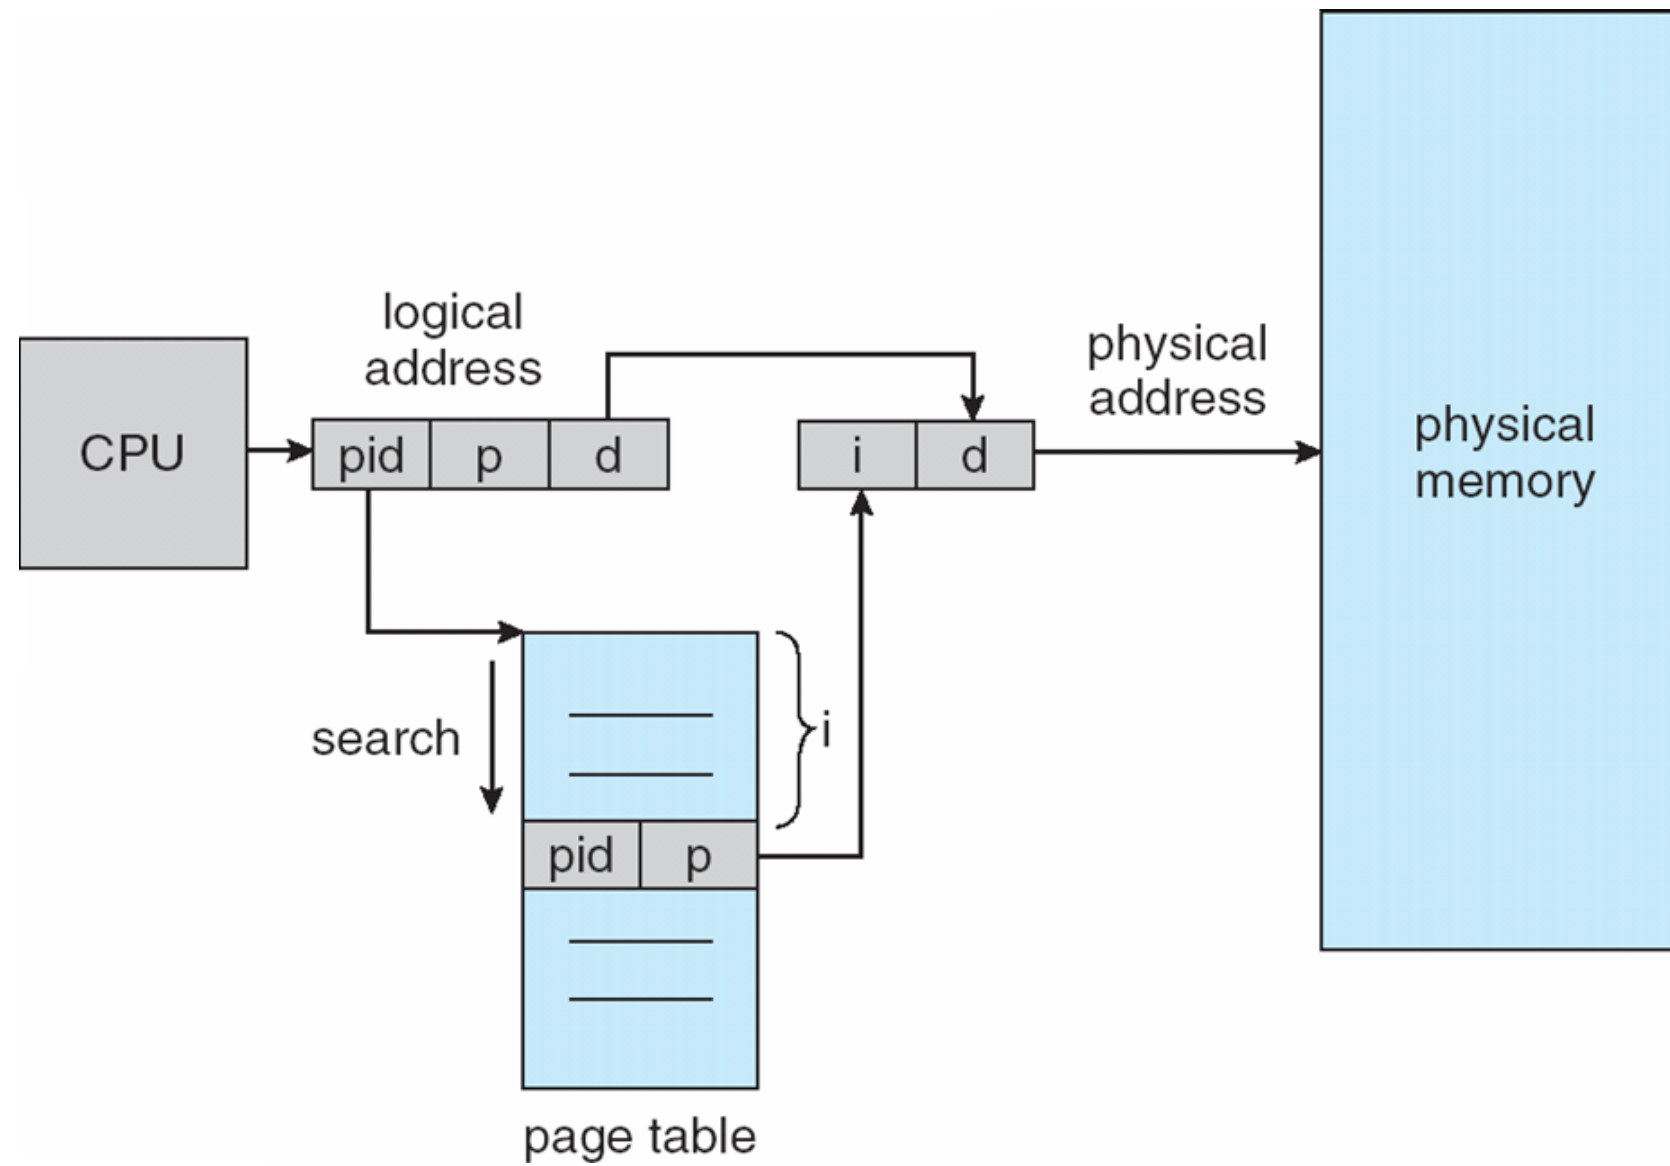




Inverted Page Table

- Inverted page table tracks allocation of physical frame to a process
 - one entry for each physical frame → fixed amount of memory for page table
 - each entry has the **process id** and the **page#** (virtual address)
- Sounds like a brilliant idea?
 - to translate a virtual address, it is necessary to search the (**whole**) page table
 - can use TLB to accelerate access, TLB miss could be very expensive
 - how to implement shared memory?
 - a physical frame can only be mapped into one process!
 - Because one physical memory page cannot have multiple virtual page entry!

Inverted Page Table






Swapping

- **Swapping** extends physical memory with backing disks
 - a process can be swapped temporarily out of memory to a backing store
 - backing store is usually a (fast) disk
 - the process will be brought back into memory for continued execution
 - does the process need to be swapped back in to same physical address?
- Swapping is usually only initiated under memory pressure
- Context switch time can become very high due to swapping
 - if the next process to be run is not in memory, need to swap it in
 - disk I/O has high latency

Swapping



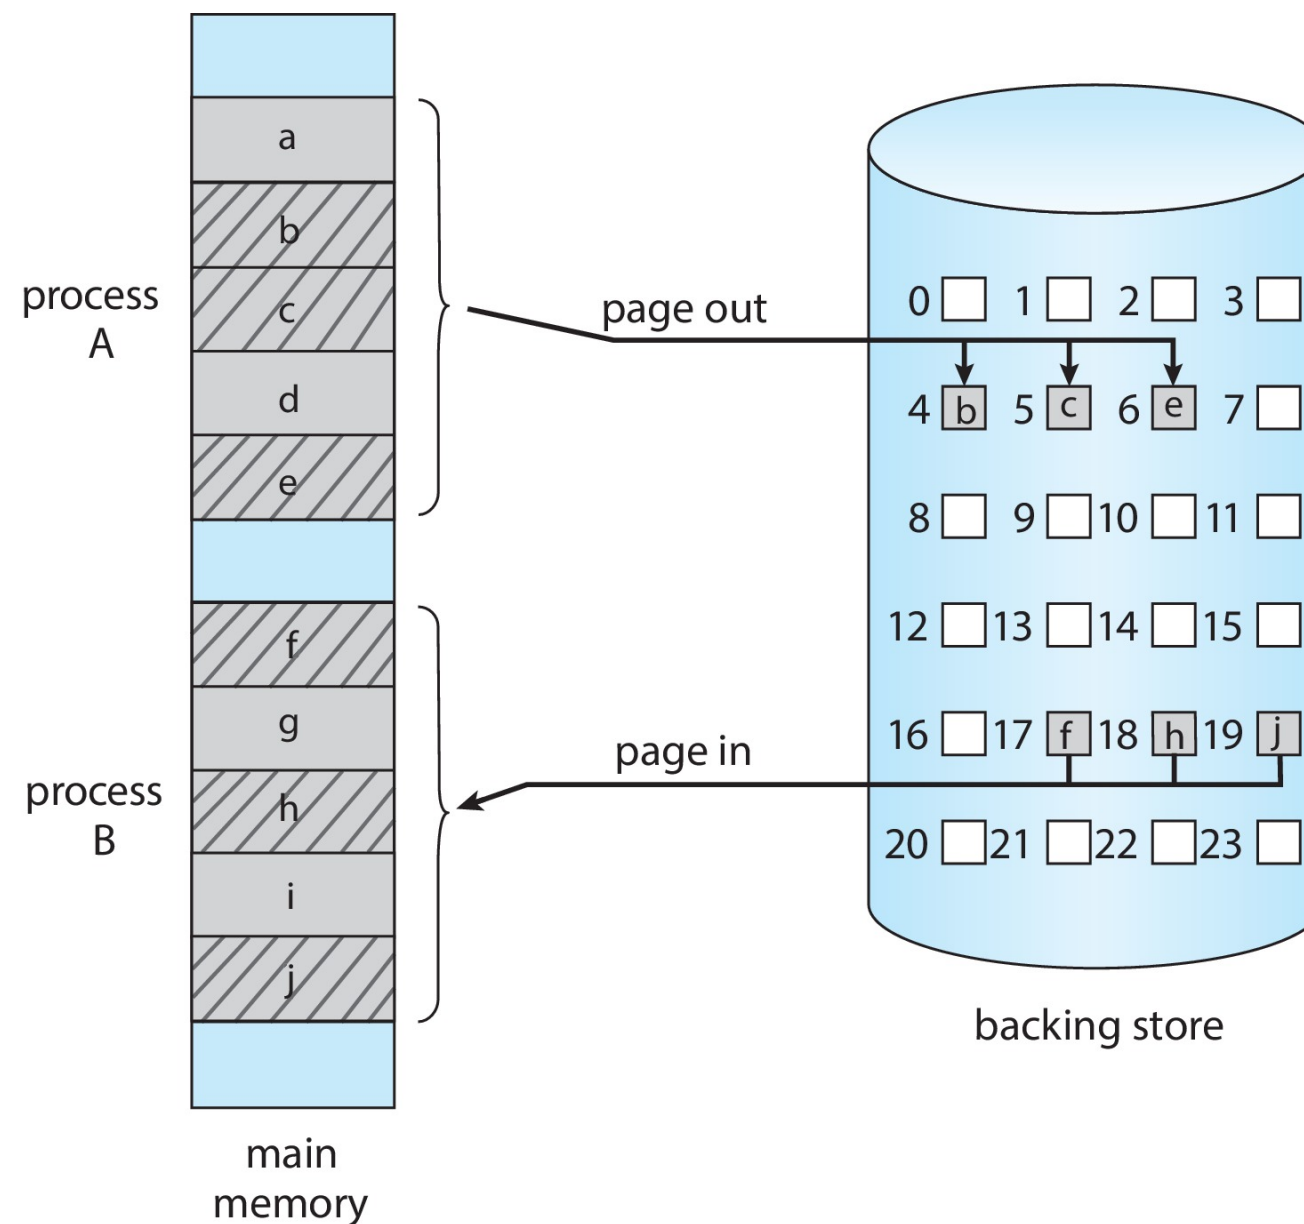


Context Switch Time including Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be **very high**
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - Swap out time of 2,000 **ms**
 - Plus swap in of same sized process
 - Total context switch swapping component time of 4000ms (4 seconds)
- Can reduce if reduce size of memory swapped – by knowing how much memory really being used

Swapping with Paging

- Swap pages instead of entire process





Swapping on Mobile Systems

- Not typically supported
 - Flash memory based
 - Small amount of space
 - Limited number of write cycles
 - Poor throughput between flash memory and CPU on mobile platform
- Instead use other methods to free memory if low
 - iOS **asks** apps to voluntarily relinquish allocated memory
 - Read-only data thrown out and reloaded from flash if needed
 - Failure to free can result in termination
 - Android terminates apps if low free memory, but first writes **application state** to flash for fast restart
 - Both OSes support paging as discussed below



Example: The Intel 32 and 64-bit Architectures

- Dominant industry chips
- Pentium CPUs are 32-bit and called IA-32 architecture
- Current Intel CPUs are 64-bit and called IA-64 architecture
- Many variations in the chips, cover the main ideas here




Example: The Intel IA-32 Architecture

- Supports both segmentation and segmentation with paging
 - Each segment can be 4 GB
 - Up to 16 K segments per process
 - Divided into two partitions
 - First partition of up to 8 K segments are private to process (kept in **local descriptor table (LDT)**)
 - Second partition of up to 8K segments shared among all processes (kept in **global descriptor table (GDT)**)



Example: The Intel IA-32 Architecture (Cont.)

- CPU generates logical address
 - **Selector** given to segmentation unit
 - Which produces linear addresses



- *s*-> segment number, *g*-> local/global, *p*->protection
- Segment selector is stored in the segment registers: CS, DS etc.
- GDTR, LDTR -> base address of the descriptor table
- descriptor: base, limit and other bits





Example: The Intel IA-32 Architecture

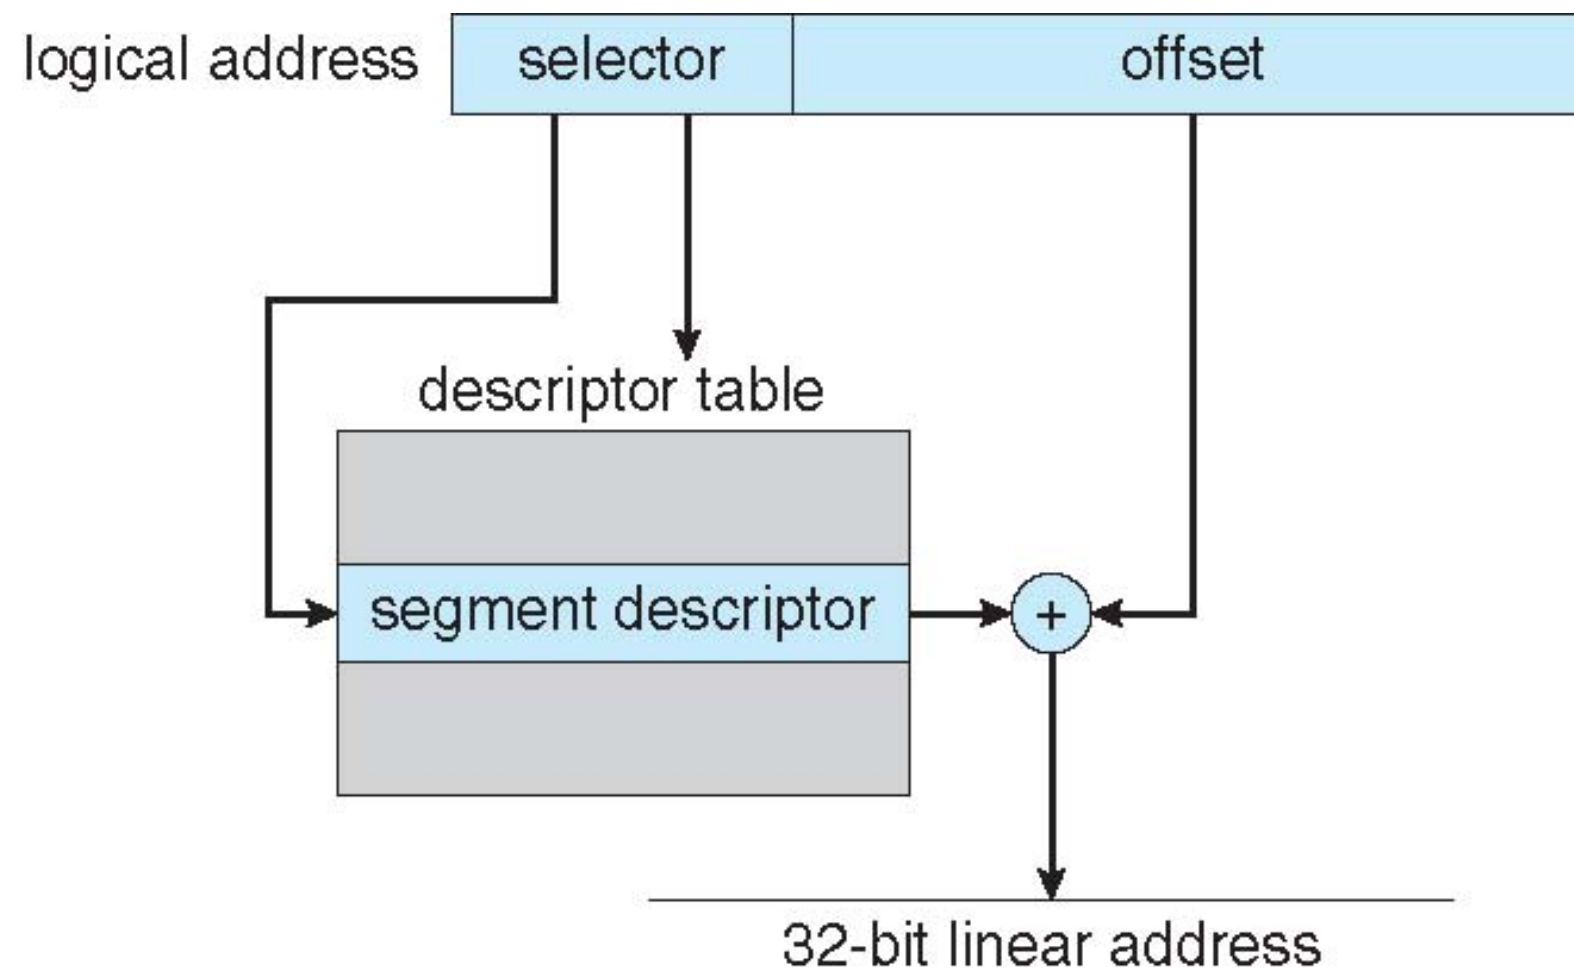
- Linear address given to paging unit
 - Which generates physical address in main memory
 - Paging units form equivalent of MMU
 - Pages sizes can be 4 KB or 4 MB




Logical to Physical Address Translation in IA-32



Intel IA-32 Segmentation



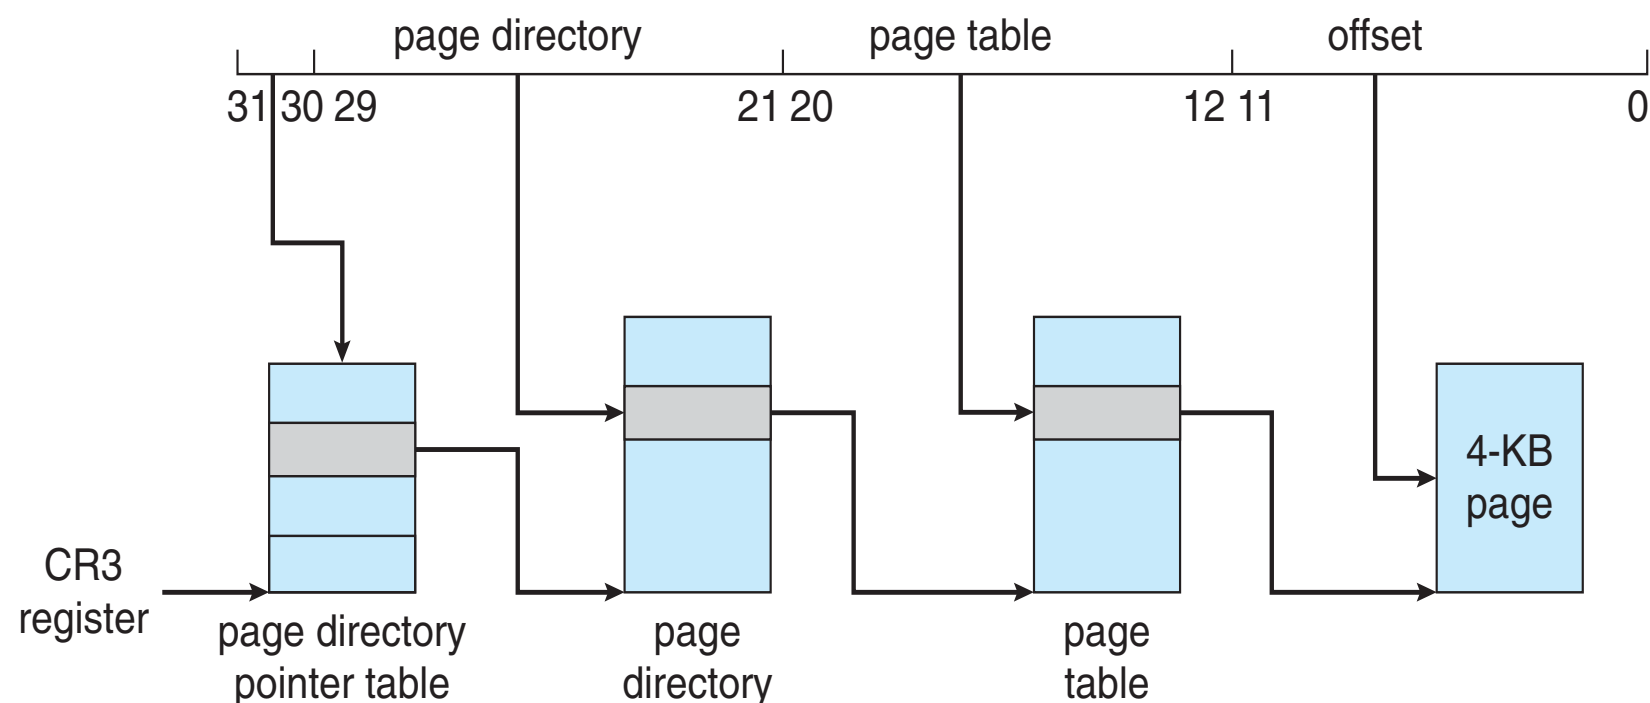
Intel IA-32 Paging Architecture





Intel IA-32 Page Address Extensions

- 32-bit address limits led Intel to create **page address extension (PAE)**, allowing 32-bit apps access to more than 4GB of memory space
 - Paging went to a 3-level scheme
 - Top two bits refer to a **page directory pointer table**
 - Page-directory and page-table entries moved to 64-bits in size
 - Net effect is increasing address space to 36 bits – 64GB of physical memory

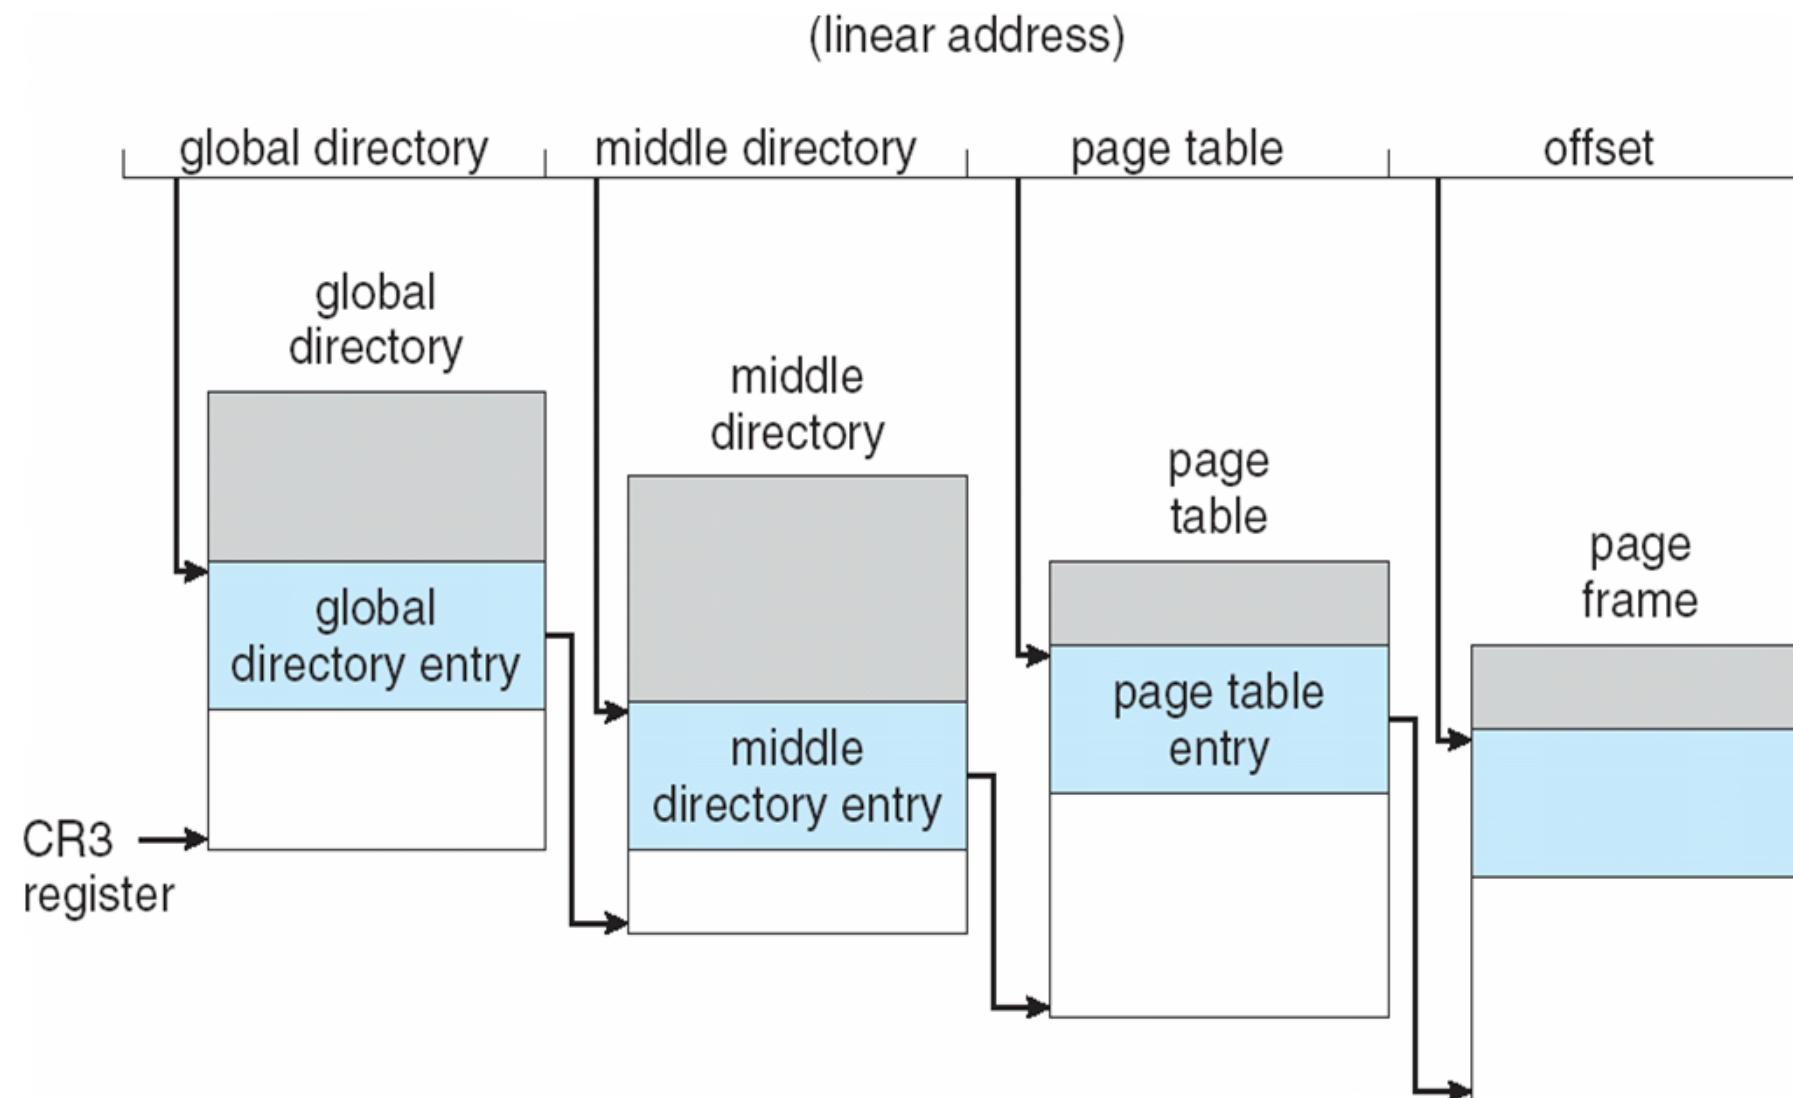




Linux Support for Intel Pentium

- Linux uses only 6 segments
 - kernel code, kernel data, user code, user data
 - task-state segment (TSS), default LDT segment
- Linux only uses two of four possible modes
 - kernel: ring 0, user space: ring 3
- Uses a generic four-level paging for 32-bit and 64-bit systems
 - for two-level paging, middle and upper directories are omitted
 - older kernels have three-level generic paging


Three-level Paging in Linux





Intel x86-64

- Current generation Intel x86 architecture
- 64 bits is ginormous (> 16 exabytes)
- In practice only implement 48 bit addressing
 - Page sizes of 4 KB, 2 MB, 1 GB
 - **Four levels** of paging hierarchy
- Can also use PAE so virtual addresses are 48 bits and physical addresses are 52 bits



HW9 is out